I. Project Context & Success Criteria

1.0 About Project

Rosotravel's pretend to create new category in Activity & Experience

Rosotravel becomes:

• Decision making platform

• adaptive
• human
• real
• dynamic
• transparent
• expert-driven

It is the only platform where:

• content adapts
• the experience is personalized
• the last tour is visible
• guides are front and center
• itineraries reflect user intent
• pages evolve over time
• the system learns continuously
• travelers feel understood

Short description

Scope 2 — Programmatic Decision Platform, SEO & Website Architecture

This scope defines the full transformation of Rosotravel.com from a traditional tour-operator website into a logic-driven, programmatic decision platform for tours and activities. The core of the system is not "AI as a brand promise", but a governed architecture that combines curated product selection, destination hierarchy, SEO logic, data structures, content rules, schema, internal linking, CMS workflows, LLM-readable feeds, migration safeguards, payment/partner extensions, and continuous optimization.

AI is used as an enabling layer for drafting, enrichment, validation, snippets, FAQ, metadata and scalable content operations. The strategic value remains human and decision-first: Rosotravel shows fewer but better options, explains why they were selected, proves trust through real signals, and helps travelers choose faster with higher confidence.

This scope covers the new website structure, page templates, navigation, product decision layers, RAW ingestion, entity/database model, keyword ownership, token libraries, JSON-LD/schema builders, relation graph, internal linking engine, CMS/marketing panels, publishing layer, SEO governance, migration logic, LLM/API outputs, performance requirements, payment processor, partner programs, reviews loop, data layer and CRM foundations.

In short, this is the blueprint for rebuilding Rosotravel.com as a scalable, human-curated and programmatically operated experience platform — designed to grow organic traffic, improve conversion, support LLM visibility, protect SEO safety, and reduce dependency on OTAs.

What this scope covers at a high level

1. **Platform positioning and product concept**
   Rosotravel is repositioned from a traditional tour operator into a decision-first experience platform: curated, adaptive, human, transparent, expert-driven and designed to reduce choice overload.
2. **Complete website architecture and UX logic**
   The scope defines the future structure of the website: homepage, country, region, city, things-to-do, attraction, product, near-page, blog, expert, bundle and explore/map templates, including navigation, breadcrumbs, page hierarchy, URL architecture, internal linking and user journeys.
3. **Decision Platform product presentation**
   The scope defines how products are selected and shown through curated shortlists, Top Picks, decision rails, “Our picks”, Top Tickets, bundles, Magic Finder, Browse Curated Catalog and controlled “more options” layers instead of an endless OTA-style catalog.
4. **Dynamic Experience Engine and audience-based rendering**
   The scope defines how the interface adapts presentation for different traveler profiles, audiences, intents and constraints by selecting prewritten variants, reordered curated blocks, audience-aware rails, bundles and recommendations without creating uncontrolled indexable URL variants.
5. **Human-Face, Proof and Trust Layer**
   The scope defines components such as Rosotravel Live, Last Tour Snapshot, Trip DNA, guide/expert visibility, verified reviews, reputation strips, trust badges, “Why this pick”, “Shortlisted X of Y”, and other proof-led UI elements that make the platform feel real, human and trustworthy.
6. **Technical data architecture and entity model**
   The scope defines the underlying data architecture: destination entities, product entities, RAW objects, versioning, hashes, external IDs, product-to-city/attraction relationships, hierarchical entity model, content fields, product classification, portfolio logic, and batch lifecycle.
7. **AI, tokens, schema and relation systems**
   The scope defines the global token library, copy tokens, SEO tokens, JSON-LD builders, schema mapping by page type, schema relations, hierarchy signals, internal linking relations, nearby/related blocks, entity relations and machine-readable structured data outputs.

8. **Programmatic SEO and keyword ownership system**
   The scope defines SEMrush/GSC keyword retrieval, keyword clustering, embeddings, Qdrant collections, keyword ownership locking, anti-cannibalization, canonical logic, mixed-intent handling, similarity thresholds, near-page rules and controlled post-publish SEO expansion.
9. **AI content, drafting, validation and publishing pipeline**
   The scope defines how RAW/enriched data becomes AI-drafted content, metadata, FAQs, alt text, schema, snippets, internal linking suggestions and LLM-ready outputs, including validation gates, CMS locks, Copyscape/N-gram checks, human QA sampling and publishing hard gates.
10. **CMS, marketing panel and operational tooling**
   The scope defines new CMS and admin capabilities: Marketing Panel, Product Ops, SEO Control Center, Keyword Ownership panel, Near-Page Control Panel, Redirect/Deprecation panel, Media Library, Optimization Inbox, alerts, roles, permissions, audit logs and manual optimization tools.
11. **LLM, API and machine-to-machine layer**
   The scope defines AI sitemap, AI summary feeds, inventory feeds, public API documentation, `/ai-query`, MCP/read-only LLM tools, LLM visibility feeds, agent-readable tour search, availability links, bundles, structured feeds and bot/agent booking support.
12. **Migration, performance, security and NFR**
   The scope defines legacy URL handling, one-time migration, staging/testing environment, SEO migration safeguards, publication caps, crawl budget/backpressure, DevOps monitoring, caching, backup, security, RBAC, signed requests, audit logs and performance requirements.
13. **Commerce, payment, partners and distribution extensions**
   The scope also includes selected business extensions such as a new payment processor, affiliate/travel-agent programs, coupon codes, deeplinks, partner assets, content/revenue-share programs and marketing product feeds for dynamic ads, affiliates and partners.
14. **Data layer, CRM, reviews and growth loop foundations**
   The scope includes the foundation for the marketing loop: data layer, tracking, review collection, verified reviews, CRM/review flows, performance monitoring, optimization loops and future audience-to-conversion growth architecture.

General Rules:

01. LLM Readability and easy booking for agent

02. AI first

➢ All pages, content, metadata, keywords, near-pages, snippets, and images should be generated automatically by AI using predefined prompts, rules, and templates.
➢ AI-generated outputs must follow the keyword ownership system (no cannibalization) and fully respect the SEO architecture.
➢ All schemas, alt texts, internal linking, and LLM snippets are produced automatically by AI.

- ➢ The system must regenerate content automatically when data changes (Diff Engine) or when performance drops (CTR, impressions, bounce).

03. Human Optimized

- ➢ Humans do not write content from scratch — they only optimize, correct, or improve AI drafts where needed
- ➢ Any manually edited section becomes locked and cannot be overwritten by AI unless explicitly re-enabled in CMS.
- ➢ Human QA is based on sampling (3–5%) to validate tone, accuracy, UX, and keyword fit before publishing a batch
- ➢ Expert-written content applies only to editorial formats: blog posts, expert insights, and brand storytelling

04. Project short description

The system begins with a **predefined, authoritative list of destinations** (countries and regions only). These destinations are uploaded manually, normalized into a single canonical English form, deduplicated, and mapped into a strict parent–child hierarchy. No SEO content is generated at this stage.

Products (internal and external API feeds) are then ingested as **RAW objects**, each stored with a version hash. While products are added, they are linked to destinations using external identifiers (e.g. Viator destinationId or during adding product you choose a city which is connected). A destination becomes **active** once at least one product is linked. Active destinations are automatically published with a minimal landing page containing only an H1 and a product listing. Inactive destinations are never indexed, published, or included in sitemaps.

For active destinations, a marketing administrator can manually trigger **"Run SEO Auto"**. This one-time (per version) action retrieves a limited set of **destination-level keywords** from SEMrush. No embeddings, clustering, or manual keyword selection are used at the destination level. The system automatically selects one primary and up to two secondary keywords and prepares destination-level snippets, meta data, hero image, alt text and FAQs according to the country/region page fields. Regeneration creates a new AI version but does not affect indexation state.

All internal and external products enter the system as RAW objects with version numbers. Internal products receive a one-time AI rewrite and then become manual-only; external products allow AI rewrites unless a human edits a specific field, which becomes permanently manual. A differentiation engine compares RAW versions to detect changes and selectively ddeupdates only the necessary sections (e.g., prices vs. descriptions).

Keywords are retrieved from SEMrush and GSC, embedded with OpenAI, clustered in LangChain, and permanently assigned to products, attractions, things-to-do pages, immediate near-pages, or blog topics. Before any keyword ownership assignment, the system must aggregate all keywords for a destination (city/region) across products, attractions and historical data, and only then distribute keywords according to page-type ownership rules. Keyword ownership locks prevent cannibalization across the entire domain. Only unused keywords are allocated to new pages.

AI drafts full content and schema for each entity (products, attractions, cities, regions, countries, categories, expert pages, and immediate near-pages), generating titles, descriptions, highlights, initial FAQ placeholders, metadata, alt texts, structured data, and LLM snippets. Validation runs through GPT consistency checks, Copyscape duplication scans, and Screaming Frog N-gram detection, followed by human sampling for quality control. After QC, the system retrieves SEMrush “Questions” and GSC real-user queries, embeds them, deduplicates them via Qdrant similarity, assigns ownership (to prevent FAQ cannibalization), and generates final FAQ content via OpenAI before human QA review.

During QA, a factual data enrichment module pulls Google Places API (opening hours, geo, website, phone, ratings) and Wikidata API (historical facts, entity type, identifiers, fallback coordinates) to ensure accuracy and to populate FAQ + schema + attraction global map.

Approved batches are published with final URLs, multilingual versions, JSON-LD schema, sitemaps, AI sitemaps, and LLM-ready ai-summary.json feeds.

After publishing, all live pages are re-embedded into Qdrant so future clustering, content generation, and optimization have accurate semantic context. The global attractions dataset integrates Google Places coordinates and Wikidata entity IDs, powering both the /explore/global-map/ interface and the /explore/ai-dataset/ feed used by ChatGPT, Perplexity, and other LLMs.

The system continuously monitors performance (GSC + GA4), flags pages for re-optimization, and executes AI-driven improvements for low-CTR or low-impression URLs. Exactly five weeks after publishing, the system generates dynamic near-pages using real emerging queries, embedding similarity, and SERP gaps. Informational-intent blog topics are generated separately, written manually by experts to boost E-E-A-T.

The cycle repeats continuously—expanding coverage, refining keyword architecture, improving rankings, and increasing conversions with every batch.

- ➢ **This pipeline runs end-to-end on every batch.**

01. All products (internal and external API feeds) enter as RAW objects, are stored, hashed, and assigned a version number. Each RAW feed forms a batch grouped by country or region.
02. Internal products receive a one-time AI rewrite, then become locked for manual-only editing, while external products allow AI rewriting unless a human edits a field.
03. The Diff Engine compares RAW vX with RAW vX+1 and only reprocesses changed sections; operational changes update instantly, while text changes re-enter AI unless the section is locked.
04. All images are optimized via Cloudinary (WebP, CDN, lazy-load), while missing attraction/city/region visuals may be generated by Midjourney.
05. SEMrush and GSC retrieve keyword datasets (navigational, transactional, commercial, informational, mix of 2), which are embedded and stored in Qdrant for clustering and duplication checks.

06. LangChain performs semantic clustering, assigns funnel stages, and allocates keywords to products, attractions, things-to-do pages, immediate near-pages and blog topics.
07. A factual data enrichment module retrieves structured operational and encyclopedic data from Google Places API (opening hours, address, coordinates, website, phone, ratings, user counts) and Wikidata API (entity type, inception, historical facts, sameAs identifiers, fallback coordinates). This data enriches schema, supports FAQ accuracy, and powers the global-map dataset.
08. Keyword ownership is locked per entity; once assigned, keywords cannot be reused by other pages to prevent cannibalization.
09. OpenAI generates full content drafts including meta tags, H1/H2/H3s, full descriptions, highlights, itinerary, initial FAQ placeholders, alt texts, snippets and structured schema JSON-LD.
10. GPT validation checks accuracy, logic, naturalness, and SEO compliance, while Copyscape and N-gram scanning verify originality and template diversity.
11. After automated QC, the system retrieves FAQ-intent keywords (SEMrush Questions + GSC query patterns), embeds and clusters them, removes duplicates via Qdrant similarity, assigns FAQ ownership, and generates final FAQ content using OpenAI.
12. Human QA sampling reviews 3–5 percent of each batch; if the sample passes, the full batch proceeds, otherwise prompts or keyword rules are adjusted before regeneration.
13. Publishing generates final URLs and outputs all page types (products, attractions, things-to-do, category pages, expert pages, immediate near-pages), updates sitemaps, AI-sitemap and AI dataset feeds.
    - ➢ The final published content is embedded again into Qdrant including isPartOf, hasPart, nearby and related relations.
14. OpenAI translation performs multilingual translation with semantic preservation, while maintaining keyword locks where relevant.
15. Performance monitoring gathers GSC and GA4 metrics (impressions, CTR, bounce, conversions), feeding LangChain’s weekly classification into OK, needs optimization or A/B testing queues.
    - ➢ AI-driven optimization improves meta titles, descriptions or intros while respecting locked sections.
    - ➢ Exactly 5 weeks after publication, dynamic near-pages are generated based on real GSC queries and embedding similarity; missing topics spawn new drafts automatically.
    - ➢ Blog topic pools are generated from informational keywords before publishing, and blog posts are always written manually by experts.
16. The improvement loop repeats continuously, using data to refine future batches and structural decisions.

1.1 Goals

01. Change structure whole structure of the website, logics, internal linking, hierarchy, new templates of landing pages
02. Programmatic SEO and AI content generation for hundreds of thousands pages
03. Architecture adjusted to LLM needs with proper data structure with JSON-LD i Schema.org
04. AI agent smoothly understand and can book products on rosotravel for customer
05. Portal can personalise content and portfolio for customer without questionaries (understand intentions)
06. Create modules allowing marketing team to optimize selected landing pages
07. Design proper front-end to maximalize ROI of above changes

Expected Results until December 2026 y:

➢ **Minimum (pessimistic): ≥300,000 organic sessions / month, CR ≥ 1,0%, SEO ROI ≥ 300%**

Organic sessions are considered as Google and also LLM referrals or AI Agent bookings.

1. 2 Business outcomes

01. 100% of selected new/changed tours published within **24–48h** (median < 24h).
02. SEO visibility growth via hundreds of thousands of **long-tail** pages and structured data.
03. Keywords with all intentions are automatically matched with landing pages based on Semrush API, generating pages and all the elements
04. Having marketing module for SEO optimization
05. Automatically detect low quality score landing pages
06. Pages consumable by **LLMs** (ChatGPT/Gemini/Perplexity) through AI-oriented feeds.
07. Stable costs via **incremental reprocessing** (only changed sections).
08. Create content comply with rules **E-E-A-T**

Website is growing automatically along with new products and new semantic real searches creating landing pages based on Semrush API

1.3 Company concept

01. Rosotravel positioning switch from tour operator to "The world's most recommended experience expert" (Decision Platform)

Where you are today
Rosotravel is primarily recognized as a local tour operator, delivering proprietary tours and experiences across 100+ cities.
What this gives us:
- Authenticity: real local delivery, not a marketplace listing.
- Local knowledge: we understand the flow of a day, the pain points, the meeting logistics, the "what actually matters".
- Trust built through delivery: when we run a tour ourselves, guests feel the difference.

What holds the current model back (at scale):
- Discovery ceiling: organic growth is limited when the brand is perceived mainly as "another operator", not a global expert source.
- Marketplace dependency: demand and visibility are partially outsourced to GYG/Viator-type platforms.
- Choice overwhelm elsewhere: travelers see thousands of similar listings and lose confidence; a great operator still gets buried in noise.
- Inconsistent perception across destinations: even when we deliver premium experiences, the "why choose Rosotravel" story is not always visible at the decision moment.

Where we want to be by end of 2026
Rosotravel becomes the world's most recommended experience expert by doing one thing better than anyone else:
turning overwhelming choice into a confident decision — and then delivering a day that feels complete.

This is not "more tours".
It is "better decisions + better fulfilment" at global scale.

How we win (what we compete on)
We do not compete with GetYourGuide/Viator on volume or price wars.
We compete on:
1) Curation (decision quality)
- We do not list everything.
- We shortlist the best options that truly differ (no near-duplicate "same sightseeing" spam).
- We compare like-with-like, so the "winner" is meaningful, not accidental.

2) Trust and proof (not promises)
- Clear standards, consistent quality floors, verified reviews.
- Human expertise presented in a scalable way (experts, guides, signatures, live proof blocks).
- "Why these picks" + "what we skipped" trade-offs so the user understands the decision.

3) Discovery effectiveness (finding the right page for the right intent)
- The site becomes the place users (and search engines / LLMs) trust for "the best way to do X in Y".
- Each intent maps to a canonical page type (city / things-to-do / attraction / near page / bundle) to avoid SEO chaos and cannibalization.

4) Fulfilment (a day that actually works)
- Realistic pacing, clean logistics, fewer surprises.
- The traveler finishes the day with the feeling: “that made sense” — not “we rushed through a checklist”.

Identity (public-facing; premium, human-first)
- “Outstanding memories — chosen by people who live the experience.”
- “Feel the place, not the planning.”
- “We don’t list everything — we help you choose the experience that makes your day work.”
- “Chosen with standards. Delivered with care.”

Market Position (Strategic Map)

Competitive Positioning Table

Dimension	**GetYourGuide / Viator**	**Rosotravel (target)**	**Classic local companies**
Offer	Thousands of similar listings	Decision-first curation + curated experiences	Limited
Trust / Quality	Low control (aggregation model)	Consistent standards + proof-led trust	Variable
Personalization	Limited	Traveler-profile decisions (Audience + Intent + Constraints)	Manual
Price / Margin	Low (price wars)	Mid–high (value justified by trust & fulfillment)	High but unscalable
Technology	Marketplace distribution	Scalable discovery + governed SEO + LLM visibility + operational proof	None
Brand perception	Commercial marketplace	“Trusted global local expert” (world’s most recommended)	Niche operator

What “trusted global local expert” means operationally
- Global: consistent standards and decision surfaces across destinations.
- Local: grounded in real delivery, real experts, and verified traveler outcomes.
- Expert: we explain decisions and trade-offs, not just show listings.
- Recommended: the brand is chosen because it reduces risk and improves the quality of the day.

02. Conception from bullet points

Conception from bullet points (Decision Platform mechanism — readable, sequential)

Goal of this pipeline

Build a scalable "decision engine" that:

- ingests global inventory (internal + partners),

- selects a curated Set per destination (Model C),

- publishes decision-first pages (City / Things-to-do / Attraction / Near pages / Bundles),

- enforces keyword ownership to avoid cannibalization,

- produces machine-readable outputs for search engines and LLMs,

- and runs a controlled post-publish optimization loop.

0) Predefined destination upload (authoritative list)

What happens:

- We upload an authoritative list of destinations (countries + regions first; cities may be added later depending on migration plan).

- Destinations are stored as canonical entities with stable IDs and strict parent-child relations.

Outcome:

- A deterministic destination graph that will power routing and templates.

Tools/modules:

- Destination registry (DB)

- Slug normalization + canonical routing rules (code)

1) Destination normalization & hierarchy mapping (no SEO content yet)

What happens:

- Deduplicate and normalize destination names into a canonical form (English SSOT + localized fields later).

- Map hierarchy: Country → Region → City.

- No SEO content is generated at this stage.

Outcome:

- Clean routing surfaces without indexable duplicates.

Tools/modules:

- Slug normalization engine (code)

- Validation scripts (module)

2) Product inventory ingestion (internal + external API feeds) into RAW

What happens:

- Ingest all products (internal + partner feeds) into one RAW inventory store.

- Each product is stored as a single record with provider payload retained.

Outcome:

- “RAW All Products” becomes the operational source-of-truth inventory.

Tools/modules:

- Provider ingestion (APIs)

- RAW storage (DB)

- Graph sync monitor (module)

3) Destination ↔ product linking (external_id / destinationId + primary validation)

What happens:

- Link each product to destinations using provider references (e.g., Viator destinationId) and internal mapping rules.

- Validate “primary city” assignment (must reflect where the tour actually starts).

Outcome:

- Correct RAW_count(city) and correct destination scope for selection and publishing.

Tools/modules:

- Destination mapping logic (code)

- Validation via meeting/pickup mapping (implementation confirmed in discovery)

4) RAW data versioning & Diff Engine (what changed and why)

What happens:

- Every RAW refresh produces hashes and versions.

- Diff Engine detects operational vs textual changes.

Outcome:

- We reprocess only what changed, and we keep auditability.

Tools/modules:

- Unified version domains + hashing (code)

- Data Diff Engine (module)

5) Portfolio governance (RAW → Pool → Set) under Model C

What happens:

- RAW = everything ingested (never shown by default).

- Pool = eligible candidates (auto add/remove by thresholds).

- Set = curated membership used for main stake pages (manual confirmation; stable baseline).

Hard policy:

- Main stake uses Set-only (curated), with INSTANT-only gating for Set eligibility.

Outcome:

- A controlled curated portfolio per destination that can scale without becoming an OTA catalog.

Tools/modules:

- Product Ops UI (panel)

- Eligibility engine + alerts (code/module)

6) Product classification (Phase C) (turn inventory into “decision-ready” candidates)

What happens:

- Apply category groups (01 tours+activities / 02 tickets+workshops / 03 transfers) for thresholds.

- Classify ticket-only vs tour-like to protect “Top Tickets” blocks.

- Build archetype core/variant groups for like-with-like competition and dedup.

- Apply audience/intent/constraint/value tier tags contract (hard vs soft).

- Apply always-on language advantage rule within archetype winner selection.

Outcome:

- Deterministic winners per archetype and clean page composition without duplicates.

Tools/modules:

- Classification engine (code)

- Similarity checks support (Qdrant where needed)

- QA validation scripts (module)

7) Image ingestion & optimization (media pipeline)

What happens:

- Ingest provider images and store them in RAW media library.

- Optimize via Cloudinary (WebP, responsive, lazy-load).

- Optional future: generate missing destination/POI assets via a generator (not required for MVP).

Outcome:

- Fast pages + consistent assets + indexable alt text generation path.

Tools/modules:

- Cloudinary pipeline (module)

- Media library UI (module)

8) Keyword retrieval (for non-destination pages) + keyword ownership locking

What happens:

- Retrieve keyword datasets from SEMrush and validate with GSC signals.

- Allocate keywords to canonical page types (product / attraction / things-to-do / category / near pages / blog topics).

Hard rule:

- Keyword ownership locks prevent cannibalization across the domain.

Outcome:

- Each intent maps to one canonical page type; scale without duplicate pages.

Tools/modules:

- SEMrush API (keyword retrieval)

- GSC API (validation)

- Keyword_map + ownership locks (DB)

- Qdrant similarity support for collision checks

9) Factual data enrichment (trust + schema + POI graph)

What happens:

- Enrich attractions and destinations with structured factual data.

- Use factual sources to power visitor info, schema and map datasets.

Outcome:

- Higher accuracy, stronger trust signals, better LLM comprehension.

Tools/modules:

- Google Places API

- Wikidata API

- Enrichment validator (module)

10) Drafting (content + snippets + schema + alt text) + validation

What happens:

- Generate drafts for page bodies, metadata, structured data (JSON-LD), and image alt text.

- Snippets are not generated live; they are prewritten variants stored in CMS and selected by DEE.

- Validate drafts for compliance and uniqueness.

Outcome:

- Scalable production of decision-focused pages and consistent structured data.

Tools/modules:

- Drafting layer (LLM prompts/templates)

- Schema generator (JSON-LD)

- Validation: GPT rules + Copyscape + N-gram checks (modules)

11) FAQ pipeline (retrieval → dedup → ownership → drafting)

What happens:

- Retrieve FAQ-intent queries from SEMrush Questions + GSC queries.

- Deduplicate via similarity thresholds.

- Assign ownership so the same FAQ topic does not duplicate across pages.

- Draft final FAQs and attach FAQPage schema.

Outcome:

- Real query-driven FAQ coverage without FAQ cannibalization.

Tools/modules:

- SEMrush Questions + GSC queries (APIs)

- Qdrant similarity for dedup

- FAQ ownership store (DB)

- FAQ drafting (LLM)

12) Human QA sampling + publishing

What happens:

- Human QA sampling (3–5%) validates accuracy, tone, UX, and keyword fit.

- Publish pages: product, attraction, things-to-do, category/near pages, city/region/country, expert, bundles.

Outcome:

- Controlled launch quality at scale.

Tools/modules:

- QA dashboard (module)

- Publishing queue + caps (module)

13) Post-publish embedding + feeds + sitemaps

What happens:

- Embed published content for future similarity/optimization work.

- Update sitemap + AI sitemap + ai-summary feed.

Outcome:

- Better internal clustering, better LLM-readable outputs, better crawling.

Tools/modules:

- Qdrant embeddings store

- Sitemap builder

- AI feed generator

14) Multilingual translation + hreflang

What happens:

- Translate content to supported site languages.

- Generate hreflang and self-canonical correctly.

Important:

- Site language (hreflang) is separate from “tour language” filters (runtime).

Outcome:

- Multi-language growth without canonical mistakes.

Tools/modules:

- Translation pipeline

- Hreflang generator (module)

15) Dynamic filters engine (non-indexable)

What happens:

- Runtime filters exist for UX but do not create indexable parameter pages.

Outcome:

- No SEO duplicates; clean crawl budget.

Tools/modules:

- Filter routing engine (module)

16) Global attractions map dataset generation (POI graph)

What happens:

- Build a POI graph dataset for the global map and machine-readable feeds.

Outcome:

- Visual exploration + structured discovery surface for LLMs.

Tools/modules:

- POI graph builder (code)

- Map dataset feed generator (module)

17) Monitoring + optimization loop

What happens:

- Monitor performance (GSC + GA4).

- Create Optimization Inbox jobs (allowed job types only).

- Run controlled post-optimization (CTR/impressions).

Outcome:

- Continuous improvement without breaking ownership locks.

Tools/modules:

- GSC + GA4 ingestion

- Optimization Inbox + publishing caps

18) Dynamic near-pages creation (gated; after observation window)

What happens:

- Generate new decision near-pages from emerging queries and SERP gaps only after a gated window (e.g., 5 weeks).

Outcome:

- Expand long-tail coverage safely, based on real demand.

Tools/modules:

- Near-page generator (governed)

- Similarity + ownership validators

19) Blog topic generation (expert-written only)

What happens:

- Generate topic pool from informational keywords.

- Articles are written manually by experts to strengthen authority and brand storytelling.

Outcome:

- E-E-A-T growth without turning the brand into generic programmatic content.

Tools/modules:

- Topic pool generator (SEMrush/GSC signals)

- Editorial workflow (UI)

Optional manual trigger (separate from the main pipeline)

“Run SEO Auto” (destination-level keywords + FAQs, no clustering)

Use case:

- quick baseline for destination pages without full clustering.

Constraints:

- must not break ownership rules or indexation rules.

Tools/modules:

- SEMrush API (light retrieval)

- Destination-level CMS drafting + FAQ drafting

03. General Rules:

a. LLM Readability and easy booking for agents

Rosotravel pages and feeds must be understandable not only for humans, but also for machines (LLMs, partner tools, travel agent workflows).
This requires:
- clear entity structure (what we sell, where, when, and in which variants),
- stable IDs and canonical URLs,
- structured data and machine-readable feeds,
- and a deterministic checkout/booking pathway (no reliance on scraping HTML).

b. Automation-driven production (prompts, rules, templates)

The system must be able to produce and maintain the majority of site content at scale using governed templates and rules.
This includes:
- page content drafts (for all core templates),
- metadata (meta title, meta description, H1/H2 structure),
- prewritten snippet variants (selected by DEE rules; not generated live),
- structured data (JSON-LD),
- image processing (optimized assets + alt text),
- internal linking outputs (based on the routing graph and ownership constraints),
- and controlled near-page creation (only under governance rules and caps).

c. Governance rule (non-negotiable)

All automated outputs must respect:
- keyword ownership locking (no cannibalization),
- canonical/indexation rules,
- and the deterministic routing contract.

No “helpful” automated rewrite may create a duplicate intent page, competing slug, or conflicting canonical.

d. Regeneration triggers (controlled, not chaotic)

The system must regenerate only when it is justified and safe:

A) Data-driven change (Diff Engine)
- When RAW data changes (pricing, availability structure, logistics, inclusions/exclusions, itinerary shape),
the Diff Engine determines what changed and triggers only the necessary downstream updates.
B) Performance-driven optimization (Optimization Inbox)
- When performance signals indicate underperformance (CTR, impressions, bounce, conversion),
the system creates optimization jobs in the Optimization Inbox.
These jobs are governed:
- they cannot change keyword ownership,
- they cannot change canonical identity,
- and they are executed under publishing caps and review rules.

e. Human-governed optimization (quality standard)

Humans are responsible for quality control and brand tone, not for writing everything from scratch.
Human responsibilities:
- review and improve drafts where needed,
- correct factual issues and UX clarity,
- and lock critical sections that must remain stable.

f. Locking rule (non-negotiable)

Any manually edited section becomes locked:
- it cannot be overwritten by automated regeneration
- unless explicitly re-enabled in CMS by an authorized role.

g. Human QA policy (sampling at scale)

Human QA is performed by sampling (3–5%) per batch run to validate:
- tone and brand alignment (Experiencing + Fulfilment + Reliability),
- factual accuracy where applicable,
- UX clarity (“decision-first, not catalog-first”),
- and keyword fit without cannibalization.
If the sample fails:
- prompts/rules must be adjusted,
- and the affected subset is regenerated before publishing continues.

h. Expert-written content (editorial only)

Human experts write content only where authenticity and authority are the primary value:
- blog posts (informational guides),

- expert insights,
- brand storytelling (mission, values, “how we work”).

This editorial layer strengthens E-E-A-T without turning the entire site into manual production.

i. Resulting operating model

- Automation produces scale and consistency under governance.
- Humans provide the final quality standard and brand integrity.
- The system remains deterministic, auditable, and safe for SEO, partners, and LLM discovery.

a. Project short description

The system begins with a **predefined, authoritative list of destinations** (countries and regions only). These destinations are uploaded manually, normalized into a single canonical English form, deduplicated, and mapped into a strict parent–child hierarchy. No SEO content is generated at this stage.

Products (internal and external API feeds) are then ingested as **RAW objects**, each stored with a version hash. While products are added, they are linked to destinations using external identifiers (e.g. Viator destinationId or during adding product you choose a city which is connected). A destination becomes **active** once at least one product is linked. Active destinations are automatically published with a minimal landing page containing only an H1 and a product listing. Inactive destinations are never indexed, published, or included in sitemaps.

For active destinations, a marketing administrator can manually trigger **“Run SEO Auto”**. This one-time (per version) action retrieves a limited set of **destination-level keywords** from SEMrush. No embeddings, clustering, or manual keyword selection are used at the destination level. The system automatically selects one primary and up to two secondary keywords and prepares destination-level snippets, meta data, hero image, alt text and FAQs according to the country/region page fields. Regeneration creates a new AI version but does not affect indexation state.

All internal and external products enter the system as RAW objects with version numbers. Internal products receive a one-time AI rewrite and then become manual-only; external products allow AI rewrites unless a human edits a specific field, which becomes permanently manual. A differentiation engine compares RAW versions to detect changes and selectively updates only the necessary sections (e.g., prices vs. descriptions).

Keywords are retrieved from SEMrush and GSC, embedded with OpenAI, clustered in LangChain, and permanently assigned to products, attractions, things-to-do pages, immediate near-pages, or blog topics. Before any keyword ownership assignment, the system must aggregate all keywords for a destination (city/region) across products, attractions and historical data, and only then distribute keywords according to page-type ownership rules. Keyword ownership locks prevent cannibalization across the entire domain. Only unused keywords are allocated to new pages.

AI drafts full content and schema for each entity (products, attractions, cities, regions, countries, categories, expert pages, and immediate near-pages), generating titles, descriptions, highlights, initial FAQ placeholders, metadata, alt texts, structured data, and LLM snippets. Validation runs through GPT consistency checks, Copyscape duplication scans, and Screaming Frog N-gram detection, followed by human sampling for quality control. After QC, the system retrieves SEMrush “Questions” and GSC real-user queries, embeds them, deduplicates them via Qdrant similarity, assigns ownership (to prevent FAQ cannibalization), and generates final FAQ content via OpenAI before human QA review.

During QA, a factual data enrichment module pulls Google Places API (opening hours, geo, website, phone, ratings) and Wikidata API (historical facts, entity type, identifiers, fallback coordinates) to ensure accuracy and to populate FAQ + schema + attraction global map.

Approved batches are published with final URLs, multilingual versions, JSON-LD schema, sitemaps, AI sitemaps, and LLM-ready ai-summary.json feeds.

After publishing, all live pages are re-embedded into Qdrant so future clustering, content generation, and optimization have accurate semantic context. The global attractions dataset integrates Google Places coordinates and Wikidata entity IDs, powering both the /explore/global-map/ interface and the /explore/ai-dataset/ feed used by ChatGPT, Perplexity, and other LLMs.

The system continuously monitors performance (GSC + GA4), flags pages for re-optimization, and executes AI-driven improvements for low-CTR or low-impression URLs. Exactly five weeks after publishing, the system generates dynamic near-pages using real emerging queries, embedding similarity, and SERP gaps. Informational-intent blog topics are generated separately, written manually by experts to boost E-E-A-T.

The cycle repeats continuously—expanding coverage, refining keyword architecture, improving rankings, and increasing conversions with every batch.

- ➢ **This pipeline runs end-to-end on every batch.**

17. All products (internal and external API feeds) enter as RAW objects, are stored, hashed, and assigned a version number. Each RAW feed forms a batch grouped by country or region.
18. Internal products receive a one-time AI rewrite, then become locked for manual-only editing, while external products allow AI rewriting unless a human edits a field.
19. The Diff Engine compares RAW vX with RAW vX+1 and only reprocesses changed sections; operational changes update instantly, while text changes re-enter AI unless the section is locked.
20. All images are optimized via Cloudinary (WebP, CDN, lazy-load), while missing attraction/city/region visuals may be generated by Midjourney.
21. SEMrush and GSC retrieve keyword datasets (navigational, transactional, commercial, informational, mix of 2), which are embedded and stored in Qdrant for clustering and duplication checks.
22. LangChain performs semantic clustering, assigns funnel stages, and allocates keywords to products, attractions, things-to-do pages, immediate near-pages and blog topics.

23. A factual data enrichment module retrieves structured operational and encyclopedic data from Google Places API (opening hours, address, coordinates, website, phone, ratings, user counts) and Wikidata API (entity type, inception, historical facts, sameAs identifiers, fallback coordinates). This data enriches schema, supports FAQ accuracy, and powers the global-map dataset.
24. Keyword ownership is locked per entity; once assigned, keywords cannot be reused by other pages to prevent cannibalization.
25. OpenAI generates full content drafts including meta tags, H1/H2/H3s, full descriptions, highlights, itinerary, initial FAQ placeholders, alt texts, snippets and structured schema JSON-LD.
26. GPT validation checks accuracy, logic, naturalness, and SEO compliance, while Copyscape and N-gram scanning verify originality and template diversity.
27. After automated QC, the system retrieves FAQ-intent keywords (SEMrush Questions + GSC query patterns), embeds and clusters them, removes duplicates via Qdrant similarity, assigns FAQ ownership, and generates final FAQ content using OpenAI.
28. Human QA sampling reviews 3–5 percent of each batch; if the sample passes, the full batch proceeds, otherwise prompts or keyword rules are adjusted before regeneration.
29. Publishing generates final URLs and outputs all page types (products, attractions, things-to-do, category pages, expert pages, immediate near-pages), updates sitemaps, AI-sitemap and AI dataset feeds.
    - ➢ The final published content is embedded again into Qdrant including isPartOf, hasPart, nearby and related relations.
30. OpenAI translation performs multilingual translation with semantic preservation, while maintaining keyword locks where relevant.
31. Performance monitoring gathers GSC and GA4 metrics (impressions, CTR, bounce, conversions), feeding LangChain’s weekly classification into OK, needs optimization or A/B testing queues.
    - ➢ AI-driven optimization improves meta titles, descriptions or intros while respecting locked sections.
    - ➢ Exactly 5 weeks after publication, dynamic near-pages are generated based on real GSC queries and embedding similarity; missing topics spawn new drafts automatically.
    - ➢ Blog topic pools are generated from informational keywords before publishing, and blog posts are always written manually by experts.
32. The improvement loop repeats continuously, using data to refine future batches and structural decisions.

1.4 Rosotravel Marketing USP (Decision Platform — “the world’s most recommended experience expert”)

Rosotravel is building a new category in travel:

a decision-first experience platform where travelers see fewer options, but choose with higher confidence — because every recommendation is curated by experts, meaningfully differentiated, and proven by real experience signals.

This USP is built on two layers (in this order):

Expert Decision Platform (Portfolio Curation)

Rosotravel does not operate like a catalog.
We act like a global experience expert: we scan large inventories and shortlist only the options that are worth a traveler's time.

What "expert curation" means in practice
- We select a curated portfolio per destination from a much larger RAW inventory (internal + partners).
- We compare experiences "like-with-like" (archetypes), so winners are meaningful.
- We avoid near-duplicate "same sightseeing" options that create noise.
- We enforce quality standards (e.g., rating floors and eligibility gates) so the default experience remains premium.
- We present clear decision outputs: Top Picks, decision rails, "2-of-many" attraction picks, and curated 1/2/3-day bundles.
- We explain decisions with short, human-readable reasoning:
why these picks, what we skipped, and what trade-offs exist.

What users should feel
- "They already did the research for me."
- "These options are different for a reason."
- "I can choose faster and feel safer."

Dynamic Experience Engine (DEE) (Audience-driven presentation)

DEE is not the brand headline.
DEE is the system layer that adapts what the expert-curated platform shows to a traveler's audience, intent, and constraints.

DEE does not generate text live on the fly.
Instead:
- content variants (snippets and microcopy) are prepared in advance,
- and the system deterministically selects the best variant and best curated blocks for the current profile.

What DEE adapts (examples)
- prewritten snippet variants (selected, not generated live)
- ordering and emphasis of curated blocks (Top Picks, rails, bundles)
- segment-aware bundle substitutions (only when quality allows; never by lowering standards)
- FAQ presentation and ownership-based coverage
- safe "More Options" flows (Phase B) when a traveler wants something specific

Outcome
- The expert platform defines the shortlist.
- DEE ensures the shortlist feels “made for me” without becoming a catalog.

Proof Layer (Human-Face / Live Experience)
A decision platform must prove it is real.
Rosotravel adds a proof layer that makes experiences feel human and verifiable before booking.

Human-Face and Live Experience Layer (Proof Layer)

This layer complements DEE by showing the human truth behind experiences.
Core concept
Rosotravel reveals the “real life” behind tours and activities:
- who guided the most recent tours,
- what guests felt (themes, not generic claims),
- how the experience plays out in real conditions,
- and proof signals that reduce uncertainty before booking.

Instead of anonymous, timeless reviews only, Rosotravel adds “fresh proof”:
“What happened recently, with real people, in real conditions.”

Why this matters
- Trust increases when signals are fresh, specific, and consistent.
- Human context restores authenticity in a space dominated by generic marketplace listings.
- It supports the decision platform promise: fewer options, stronger confidence.

Core proof components (examples)
A) Last Tour Snapshot (Rosotravel Live)
Displayed on product pages and selectively on attraction/city pages.
Can show:
- recent departure window (date range)
- guide/expert attribution (when applicable)
- group size signal (range)
- “Guests loved” themes (from verified data)
- recent media strip when available (under moderation rules)

B) Trip DNA (Experience Profile)
A compact 4–5 dimension profile that supports fast decisions:
- pace / walking intensity
- crowd level
- learning vs sightseeing balance
- comfort/logistics level
- “best for” cues

C) City Proof Block (“This week in {{City}}”)
A small operational proof block:

- tours run (rolling window)
- travelers hosted (rolling window)
- satisfaction/rating signal
- trending experts/guides (when applicable)

Decision Platform Messaging (Must-have, editable)
A curated platform must explain itself.
Therefore, every core template must include small, editable messaging components that clarify:
- Why you see a shortlist (not a catalog)
- Why these picks were chosen (2–4 bullets)
- What we skipped / trade-offs (2–4 bullets)
- What to do if you want something specific (“More options” / Magic Finder)
All these messages must be editable in CMS and must not create indexable URL variants.

Brand promise (external messaging)
Short versions:
- “Travel that feels personal.”
- “Choose faster. Enjoy more.”
- “Curated decisions. Real experiences.”
- “See what’s real. Choose what fits.”

Long version:
“Rosotravel helps you choose the right experience by shortlisting the best options, explaining the trade-offs, and proving the reality behind each tour with trusted signals — so your day works, and the memory lasts.”

What users should remember
- Rosotravel shows fewer options, but each option is meaningfully different.
- The shortlist is curated by experts, not by a catalog algorithm.
- DEE makes the experience feel personalized by selecting the right curated blocks and prewritten variants.
- The platform feels human and reliable: real experts, real signals, fewer surprises.

1.5 General Project description & rules

01. Internal Marketing Principles (Rules for Consistent Execution)

These rules ensure Rosotravel is perceived as a **decision platform (expert-led)**, not a catalog.

1) Every decision surface must have an expert anchor

Every key page (City, Things-to-do, Attraction, Bundle, Near page) must include at least one **Decision Trust Component**:

- Why you see a shortlist (1–2 lines)
- Why these picks (2–4 bullets)
- Trade-offs / What we skipped (2–4 bullets)
- What to do if you want something specific (More Options / Magic Finder)

All decision messaging must be **CMS-editable**.

2) DEE adapts presentation, not core facts

DEE does **not** generate text in real time.
DEE only:

- selects prewritten snippet variants,
- reorders curated blocks,
- chooses which curated set to surface for the inferred traveler profile.

All snippet variants are pre-generated and stored in CMS.

3) Every product page must communicate relevance

Each product page must clearly show:

- who it is best for (traveler cues),
- what you get (spec-style clarity),
- logistics truth (meeting/pickup),
- trust proof signals (verified reviews, recent operational indicators where applicable).

4) Every City and Attraction page must show recent real-world proof (where volume allows)

- City pages should include "This week in {{City}}" proof block (rolling window + fallback).
- Attraction pages should include compact proof signals ("recent tours at this attraction" where volume allows).

All metrics must be computed from **verified operational sources** (not manual claims).

5) Adaptiveness must never break governance

Any rewrite or optimization must:

- respect keyword ownership locks,
- respect canonical/indexing rules,
- avoid creating duplicate intent pages.

Optimization jobs must be controlled via the **Optimization Inbox** (allowed job types only).

6) Humans define the standard; automation scales it

Humans do not write everything from scratch.
Humans set:

- the decision model rules,
- the tone and standards,
- the approval/locking policy.

Automation executes at scale under those rules.

Product Model and Content Architecture — Core Logic (“Three USP Layers”)

USP #1 — Expert Decision Platform (Portfolio Curation)

Rosotravel curates a destination portfolio from a much larger inventory.
The platform selects like-with-like winners and avoids near-duplicate “same sightseeing” noise.

Outputs include:

- City Page Top Picks, Rails, Top Tickets blocks
- Attraction “Our 2 picks”
- Bundles (1/2/3 days)
- Governed near pages (decision pages for long-tail intent)

USP #2 — Dynamic Experience Engine (DEE) (Audience-driven presentation)

DEE adapts what the expert platform shows to the traveler profile:

- selects prewritten snippet variants,
- reorders curated blocks,
- chooses which curated rails/bundles to surface.

DEE does **not** create new indexable URL variants.

USP #3 — Human Face / Live Proof Layer

Rosotravel adds proof-led components:

- Rosotravel Live (operational proof where available)
- Trip DNA (experience profile)
- experts/guides presence and attribution
- verified reviews and freshness framing

This builds trust in a curated shortlist model.

Three Content Layers (Inventory → Curation → Packages)

Layer A — Verified Experiences (Rosotravel Originals)

- Owned/operated experiences across multiple cities (maintained in Ops dashboard).
- Full control of quality, logistics, photos, descriptions, and guides.
- Tone: crafted and delivered by real local experts.
- Policy: Originals are strategically preferred but still must meet quality/fit rules (no auto-winners without fit).

Layer B — Curated Partner Collection (Set-driven)

- Partner inventory is ingested into RAW and filtered into Pool and Set under Model C.

- B2C main stake uses Set only (premium, decisive).
- Wider browsing exists only in Phase B “Browse Curated Catalog” and remains band-limited.

Hard rule:

- Partner items shown in curated surfaces must pass quality floors and dedup rules.

Layer C — Expert Bundles (1/2/3-day itineraries)

- Every city has 3 baseline bundles (1/2/3 days), profile-aware substitution if quality allows.
- Every bundle must include Decision Proof:
    - Why these picks (2–4 bullets)
    - Trade-offs / What we skipped (2–4 bullets)
- Major cities may be manually edited; long-tail cities follow templates and governed generation.

Bundles are decision products, not “just a list”.

UX / SEO / Operations — Translating It Into Site Architecture

Layer	Goal	Human Role	System Role (tools/modules)
I. Structural SEO & Routing	canonical URLs, indexation safety, schema eligibility	define rules + governance	sitemap builder, hreflang generator, slug normalization, schema generator
II. Programmatic	scalable pages (city/ttd/attraction/near/product)	define tone + QA standard	drafting templates, validation scripts,

Page Production			uniqueness checks, image pipeline
III. Decision Platform Outputs (Model C)	winners, rails, bundles, dedup	Product Ops governance	portfolio builder, archetype/dedup engine, eligibility rules, batch manager
IV. Proof & Trust Layer	“proof not promises”	approve messaging tokens	Rosotravel Live metrics, review/UGC modules, Trip DNA components
V. LLM Visibility & Feeds	machine-readable discovery	editorial oversight	AI sitemap, ai-summary feed, API feed, structured data sync
VI. Media Layer	authentic visual presence	curate/approve assets	Cloudinary optimization, media library, transcript/summary tools

---

Brand Narrative & Tone (aligned to mission + values)

Mission (internal anchor)

We create outstanding memories that travelers will cherish for a lifetime.

Vision (direction)

To become the world’s most recommended experience expert.

Values in messaging

- Experiencing: real moments, real places, real stories.

- Fulfilment: the day feels complete and worth it.
- Reliability: clear confirmations and smooth delivery.
- Dedication: we stay involved and solve issues fast.
- TrailSetting: we set standards and keep improving.

Tone & visuals

- Voice: warm, knowledgeable, trustworthy (“local expert meets travel companion”).
- Visuals: human, natural light, real context; avoid stock “catalog feel”.
- Trust signals: Expert Pick, Verified, quality standard tooltips (CMS-editable).

What we must avoid in external messaging

- Do not lead with “powered by AI”.
- Do not claim “real-time rewriting”.
- Do not overpromise “everything for everyone”.

Lead with: **expert decisions + proof + fulfilment**.

Operating model (automation-driven, human-standard)

Automation builds the infrastructure

- drafts content and schema from templates,
- maintains consistency via diff + governed regeneration,
- powers feeds for SEO and LLM discovery.

Humans enforce the standard

- approve and lock critical sections,
- review sampling (3–5%) per batch,

- refine prompts/rules when samples fail,
- keep brand tone premium and truthful.

Editorial is human-written

- blog posts,
- expert insights,
- brand storytelling pages.

This strengthens authority without turning the whole site into manual production.

Outcome

Scalable growth without losing authenticity:

- decision-first pages users trust,
- proof-led signals that reduce risk,
- experiences that deliver memorable fulfilment.

➢ II. Website structure & User Experience Specification

REF: Migration of current data (products, countries, cities and others/ At the end of the document there is a step about the one time migration process.

Inventory:

a. AI prompts library
➢ Link
b. Pages summary:
01. Homepage
02. Country Page
03. Region Page
04. City Page
05. Things to do Pages (city level)
06. Attraction Page

07. Product Page
08. Category / Near Pages
09. Blog Pages
10. Packages / Expert Bundles
11. Explore / AI feed

2.1 Navigation & menus (header/footer/breadcrumb)

J. Internal Linking, Navigation Graph & Menu Logic

J.01 Purpose of This Layer (Full Definition)

The Internal Linking & Navigation layer ensures that:

1. Users can **move naturally** from country → region → city → attraction → product → near-page → blog, and back.

2. Search engines and LLMs see a **clear, consistent site graph** that matches the SEO hierarchy and schema relations.

3. Internal links are generated **automatically wherever possible**, with only minimal configuration needed in the CMS.

4. Changes in the semantic/SEO layer (keywords, ownership, embeddings) automatically influence **related links and “nearby” blocks**, without breaking UX.

5. The **header, footer, breadcrumbs and contextual modules** are all consistent and driven by the same navigation graph.

This layer sits “on top” of the content and semantic layer and translates them into **real navigational paths, menus, and internal links**.

J.02 Navigation Graph Model (Logical Overview)

The **Navigation Graph** is the combination of:

1. **Hierarchical edges** (location tree)

    - Country → Region → City → Attraction → Product → Near-pages

    - Country → Region → City → Things-to-do category → Products

    - City → Experience categories → Products
2. **Horizontal edges** (topical & geo relationships)
    - Attraction ↔ Nearby attraction
    - Product ↔ Related product
    - Product ↔ Package
    - City ↔ Nearby city
    - Article ↔ Attraction / City / Product (via about / primaryTopic / subjectOf)
3. **Menu nodes** (Header & Footer)
    - Destinations, Experiences, Travel Guides, Experts, Explore Map, Sign In
    - About, For Businesses, Support, Legal, Sitemap, AI Sitemap, API feed
4. **Entry points**
    - Homepage
    - Search / Start Planning
    - Explore Map

The graph is implemented partly as **relational data (PostgreSQL / Laravel models)** and partly as **semantic data (Qdrant embeddings + relations)**. The internal linking engine uses both:

- Relational: for “official” hierarchy & menu.
- Semantic: for “related / recommended / nearby” blocks.

J.03 Node Types and Link Types (Navigation Entities)

Node types (top-level entities):

- Country
- Region

- City
- Attraction (POI)
- Tour / Product
- Package / Bundle
- Things-to-do category (in city)
- Near-page (long-tail / guide / comparison)
- Travel guide article (country / city / topic)
- Blog article
- Expert profile
- Creator profile
- Explore node (map POI)
- System pages (About, Contact, Help Center, B2B etc.)

Core link types (must match schema + semantic layer):

1. Parent / Child
   - Country → Region → City → Attraction → Product
   - Stored in DB as `parent_id` / `parent_type` and represented in schema as `isPartOf` / `hasPart`.
2. Nearby (geo)
   - Attraction → Nearby attractions
   - City → Nearby cities
   - Computed via geo bounding box & poi_vectors.
3. Related (semantic)
   - Product ↔ Product
   - Product ↔ Package

    - Product ↔ Near-page
    - Attraction ↔ Travel guide / Blog
    - Based on Qdrant pages_vectors similarity.
4. Category membership
    - City → Experience categories (Day trips, Family-friendly, Night tours, etc.)
    - Category → Products
5. Editorial links
    - Manually curated internal links (e.g. from travel guide to specific products), stored as explicit relationships in DB and **never overridden by AI**.

J.04 Breadcrumb Logic (Consistent Across All Page Types)

Breadcrumb pattern

For location-based entities:

- Country page:
  `Home > Destinations > {{Country_Name}}`
- Region page:
  `Home > Destinations > {{Country_Name}} > {{Region_Name}}`
- City page:
  `Home > Destinations > {{Country_Name}} > {{Region_Name}} > {{City_Name}}`
- Attraction page:
  `Home > Destinations > {{Country_Name}} > {{Region_Name}} > {{City_Name}} > {{Attraction_Name}}`
- Product page:
  `Home > Destinations > {{Country_Name}} > {{Region_Name}} > {{City_Name}} > {{Tour_Name}}`
- Things-to-do (city) page:
  `Home > Destinations > {{Country_Name}} > {{Region_Name}} >`

`{{City_Name}} > Things to do`

- Package page:
  `Home > Destinations > {{Country_Name}} > {{Region_Name}} > {{City_Name}} > {{Package_Name}}`

For non-location:

- Blog article:
  `Home > Travel Guides > {{Topic / City / Country}} > {{Article Title}}`
- Expert profile:
  `Home > Experts > {{Country}} > {{City}} > {{Expert_Name}}`
- B2B page:
  `Home > For Businesses > {{Subpage}}`

CMS implementation:

- Each entity already has: `country_id`, `region_id`, `city_id`, `attraction_id` (where applicable).
- A generic **Breadcrumb Builder** takes the entity, walks up parent relations and prints breadcrumbs + JSON-LD `BreadcrumbList`.
- Marketer does **not** build breadcrumbs manually

J.05 Header Navigation Logic (Global Entry Points)

Header (sales elements):

- Destinations
- Experiences
- Travel Guides
- Experts
- Explore Map
- Sign In

Destinations menu

- Clicking “Destinations” → `/destinations/` hub.
- Hover/mega-menu (optional later):
    - Featured countries (top revenue / trending)
    - Quick links: “Top cities in Europe”, “Top cities in USA”

Experiences menu

- “Experiences” → `/experiences/` hub (lists experience categories).
- Sub-links (within hub): Day trips, Family-friendly, Food & Wine, Museum tours, etc.

Travel Guides menu

- “Travel Guides” → `/travel-guides/` hub.
- Sub-structure internally:
    - Guides by continent → country → city → topics.

Experts menu

- “Experts” → `/experts/` hub.
- Split into:
    - Rosotravel experts (core team)
    - Local experts by country/city

Explore Map

- “Explore Map” → `/explore/global-map/`
- Primary use: visual entry into attractions + cities.

Sign In

- `/my-account/login` (or `/login`), same for all locales, then user sees localized content based on language.

Important: Header links **never change** by language; only the labels are translated. URLs remain the same structure with language prefix.

J.06 Footer Navigation Logic

Footer sections (as per earlier spec):

1. **About Rosotravel**
    - About Rosotravel
    - Who We Are
    - Mission & Values
    - How We Work
    - Sustainability
    - Careers
    - Press & Media
2. **Explore**
    - Destinations
    - Experiences
    - Travel Guides
    - Experts
    - Creators
    - Explore Map
3. **For Businesses**
    - Travel Agents / DMC

- Affiliate Program
- Content Creator Program

4. **Support**
   - Contact Us
   - Help Center
   - Manage Booking
   - Refund Policy
   - Safety Standards
5. **Legal & Technical**
   - Terms & Conditions
   - Privacy Policy
   - Cookies
   - Sitemap
   - AI Sitemap
   - API Feed

Internal logic:

- Footer links are **static navigation nodes** (configured in CMS menu builder).
- They do **not depend** on semantic layer.
- Language-specific versions use same path with language prefix and translated labels.

2.2 UI flow (front end) -& Flow/journeys

UI Flow is how users think and navigate (Home → Country → Region → City → Things to do → Category → Tour), whereas the URL is a simplified SEO-friendly representation (/destinations/italy/rome/things-to-do/), with the full structure still preserved in schema and breadcrumbs.

01. General Hierarchy data structure (backend)
a. Full User Flow Diagram

Status
Header and footer UI structure is accepted as correct for Model C.
Required work is routing + SEO safety + governance, not redesign.

A) Hard routing rules (prevent raw-catalog leakage)
All navigation links in header/footer must route to curated surfaces only:
- City Page: /{lang}/{country}/{city}/
- Things-to-do: /{lang}/{country}/{city}/things-to-do/
- City category pages (curated): /{lang}/{country}/{city}/{category}/
- Attraction/POI pages: /{lang}/{country}/{city}/{attraction}/
- Near pages (programmatic decision pages): /{lang}/{country}/{city}/{keyword-slug}/
- Bundle pages: /{lang}/{country}/{city}/itineraries/{duration}-day/
- Guides: /{lang}/travel-guides/... (per guide taxonomy)
- Experts: /{lang}/experts/... (governed)
- Map: /explore/global-map/ (or /explore/world-map/ per naming)

Hard rule:
No menu item may link to filter-parameter URLs that create duplicate indexable pages.

B) Experiences mega-menu — Quick links must land on decision pages
Quick links like:
- “Best for first-timers”
- “With kids”
- “Skip-the-line”
must route to one of:
- curated category pages, OR
- governed near pages (decision-first SEO pages),
never to parameter-based filtered listing URLs.

C) Search routing (deterministic)
Search results must route deterministically by entity type:
- City → City Page
- Attraction/POI → Attraction Page
- Category intent → City category page or near page (if keyword_map exists)
- Product → Product Page
- Bundle/itinerary intent → Bundle page (if exists) else Things-to-do

D) Experts pages governance
If “Local experts in {City}” pages are generated programmatically:
- apply caps and thin-page prevention rules (SEO governance),
- require minimum content completeness per expert page,
- avoid generating for cities without sufficient expert coverage.

E) Footer technical links (good practice)

Footer “Legal & Technical” links are correct and recommended:
- Sitemap
- AI Sitemap
- API Feed
Ensure these URLs follow governance rules and do not expose restricted/internal data.

F) Breadcrumb logic (SEO + UX)
Breadcrumbs must follow canonical hierarchy:
Home → Country → Region (optional) → City → (Things-to-do / Category / Attraction) → Product/Bundle
Breadcrumbs must always use canonical URLs (no segment variants, no parameter URLs).

01. Header UI flow

Skopiuj kod

```scss
                  ┌──────────────────────────┐
                  │          HOMEPAGE          │
                  └─────────────┬────────────┘
                                │
         ┌──────────────────────┼──────────────────────┐
         │                      │                      │
   ▼ Destinations          ▼ Experiences         ▼ Travel Guides
(country → region → city →  (activity → region →   (city guides,
attraction → tours → bundles)   city → tours)       blog, tips)
         │                      │                      │
         │                      │                      │
   ▼ Product Page          ▼ Product Page          ▼ City Guide
         │                      │                      │
         └──────────────────────┼──────────────────────┘
                                │
                         ▼ Start Planning
                   (search bar → autocomplete
               → routes to best destination/category)

         ┌──────────────────────┼──────────────────────┐
         │                      │                      │
     ▼ Experts             ▼ Global Map          ▼ Support Pages
(profiles → country → city) (interactive POI) (About, Contact, B2B, etc.)
```

➢ The header is the sales engine, routing every user toward the optimal funnel geographic (Destinations), activity-based (Experiences), informational (Guides), trust-driven (Experts), visual (Attraction/Global Map), or direct high-intent search (Start Planning).

Destinations

Central hub for geographic exploration.
Leads users into the hierarchical travel flow:

Country → Region → City → Things to do → Attractions → Product / Bundle.
Supports SEO through strong geo-intention routing ("things to do in Rome", "Paris attractions").

Experiences

Activity-driven discovery layer.
Helps users who search by **type of activity** (day trips, food tours, skip-the-line, family-friendly).
Converts "I don't know the city yet" users into funnel-ready segments.

Travel Guides

Long-form content and SEO landing pages for inspiration and E-E-A-T.
Includes city guides, country guides, seasonal guides, itineraries.
High-intent informational traffic source pitching Rosotravel tours contextually.

Experts

Human-face layer.
Profiles for global and local experts, guides, creators — building trust, authenticity, and LLM visibility.
High E-E-A-T impact + users value local insight.

Global Attractions Map

Interactive visual exploration tool.
Functions as **a universal entry point** for users without a fixed destination.
Also acts as the primary **LLM Feed View** (structured data accessible to AI models).

Start Planning / Search Trips

The main CTA and core conversion entry point.
Intent detector → routes users to the right landing page (City, Attraction, Experience, Product).
Acts like GYG/Booking hybrid search + future foundation for AI trip planner.

Sign In

Account access for bookings, tickets, preferences.
Also supports future personalization (recent searches, saved trips, expert recommendations).

02. Footer UI flow

```bash
                        ┌─────────────────────────┐
                        │        ANY PAGE         │
                        │ (user scrolls to footer)│
                        └────────────┬────────────┘
                                     │
      ┌───────────────┬──────────────┼─────────────────────┬───────────────┐
      │               │              │                     │               │
 ▼ About Rosotravel  ▼ Explore      ▼ For Businesses      ▼ Support       ▼ Legal &
(Brand / Story / E-E-A-T) (SEO discovery) (B2B growth)   (Help / Contact) (Complianc
      │               │              │                     │               │
      │               │              │                     │               │
┌─────┴───────┐ ┌─────┴───────┐ ┌────┴──────────┐  ┌──────┴──────┐  ┌─────┴─────
│ /about/     │ │ /destinations/│ │ /b2b/travel- │  │ /contact/   │  │ /legal/te
│ /about/who-we-│ │ /experiences/ │ │ agents/      │  │ /help-center/│ │ /legal/pr
│ are/        │ │ /travel-guides│ │ /b2b/affiliate/│ │ /manage-    │  │ vacy/
│ /about/mission│ │ /experts/   │ │ /b2b/creators/ │ │ booking/    │  │ /cookies/
│ -values/    │ │ /creators/    │ │              │  │ /refund-    │  │ /sitemap/
│ /about/how-we-│ │ /explore/   │ │              │  │ policy/     │  │ /ai-sitem
│ work/       │ │ global-map/   │ │              │  │ /safety-    │  │ /api-feed
│ /about/sustain│ │             │ │              │  │ standards/  │  │
└─────────────┘ └─────────────┘ └──────────────┘  └─────────────┘  └──────────
      │                 │                │                  │               │
      ▼                 ▼                ▼                  ▼               ▼
Brand trust, story  → user returns to  Partner sign-up / RFP → Help / contact → Legal reassurance
+ E-E-A-T signals     exploration       forms, onboarding      ticket solving   API docs, sitema
```

Interpretation:

- ➢ **About Rosotravel** – builds E-E-A-T, brand story, human touch.
- ➢ **Explore** – from the footer the user returns to “exploration” (Destinations, Experiences, Guides, Experts, Map).
- ➢ **For Businesses** – special flow for agents, affiliates, creators.
- ➢ **Support** – all support paths (FAQ, contact, manage booking).
- ➢ **Legal & Technical** – compliance + technical (Sitemap, AI Sitemap, API Feed).

➢

03. Header Detailed Diagram

`Home → Destinations → Country → Region → City → Things to do → Category/Attraction(POI) → Product`

sql Skopiuj kod

```sql
┌──────────────────────────────────────────────────────────────────────┐
│  Rosotravel (logo)   Destinations   Experiences   Travel Guides   Experts  │
│                                                                     ▼      │
│                      Global Attractions Map         Start Planning         │
│                                              [ Search trips □□□□ ]           │
│                                                              Sign In       │
└──────────────────────────────────────────────────────────────────────┘
```

➢ **Behavior**:
- ○ Destinations – opens a mega-menu:
  - ■ column 1: “Popular countries” (Italy, UK, USA, Thailand…)
  - ■ column 2: “Top cities” (Rome, London, Dubai…)
  - ■ column 3: “Inspiration” (City guides, Top attractions, Things to do pages).
- ○ Experiences – a mega-menu with:
  - ■ categories: Day trips, Walking tours, Food & Wine, Family-friendly, Nightlife…
  - ■ quick links: “Best for first-timers”, “With kids”, “Skip-the-line”.
- ○ Travel Guides – links to:
  - ■ City Guides, Country Guides, Themes (Family, Winter trips, Romantic, etc.).
- ○ Experts:
  - ■ Meet our experts” (global),
  - ■ “Local experts in Rome”, “Local experts in London”, etc.
- ○ Global Attractions Map – leads to /explore/global-map/.
- ○ Start Planning (search bar):
  - ■ placeholder: “Where do you want to go?”
  - ■ possible fields: [Destination] [Dates (optional)] [Travel style (optional)]
  - ■ results → routing to: City, Things-to-do in city, Experiences (if category-like query), or specific tours.

➢ Header flow, footer flow, core user journeys.

➢ Breadcrumb logic.

Footer Detailed Diagram

mathematica Skopiuj kod

```
┌──────────────────────────────────────────────────────────────────────────────┐
| About Rosotravel           Explore                   For Businesses         |
|  - About Rosotravel         - Destinations             - Travel Agents / DMC  |
|  - Who We Are               - Experiences              - Affiliate Program    |
|  - Mission & Values         - Travel Guides            - Content Creator Program|
|  - How We Work              - Experts
|  - Sustainability           - Creators                 Support                |
|  - Careers                  - Explore Map              - Contact Us           |
|  - Press & Media                                       - Help Center          |
|                                                        - Manage Booking       |
| Legal & Technical                                      - Refund Policy        |
|  - Terms & Conditions                                  - Safety Standards     |
|  - Privacy Policy                                                             |
|  - Cookies
|  - Sitemap      AI Sitemap  API Feed                                          |
└──────────────────────────────────────────────────────────────────────────────┘
```

About Rosotravel

- About Rosotravel
- Who We Are
- Mission & Values
- How We Work
- Sustainability
- Careers
- Press & Media

Explore

- Destinations
- Experiences
- Travel Guides
- Experts
- Creators
- World Attraction Map (change name of explore map from photo)

For Businesses

- Travel Agents / DMC
- Affiliate Program
- Content Creator Program

Support

- Contact Us
- Help Center

- Manage Booking
- Refund Policy
- Safety Standards

Legal & Technical

➢ Terms & Conditions
➢ Privacy Policy
➢ Cookies
➢ Sitemap
➢ AI Sitemap
➢ API Feed

02. Core User Journeys (3 main scenarios)

a. SEO user (“things to do in Rome”)

Google Search →

lands on: /destinations/italy/lazio/rome/things-to-do/

├─ CTA: top categories (Day trips, Family-friendly, Night tours)

├─ Widgets: “Top attractions in Rome” → /destinations/italy/lazio/rome/attractions/colosseum/

└─ Tours carousel: /destinations/italy/lazio/rome/tours/{tour-slug}/

Content:

- CTA: top categories
- Widgets: “Top attractions in Rome” → attraction pages
- Tours carousel → product pages

b. Direct navigation (“I want family-friendly experiences”)

Homepage →

Header → Experiences →

/experiences/

→ /experiences/family-friendly/

→ filter by {country} + {city}

→ /experiences/family-friendly/italy/rome/

→ product: /destinations/italy/lazio/rome/tours/colosseum-family-tour/

c. Trust journey (About → Experts → Product)

Any page → scroll to footer →

About Rosotravel → /about/

→ Experts → /experts/

→ choose country /experts/italy/

→ city /experts/italy/rome/

→ expert profile /experts/italy/rome/maria-rossi/

→ “My favorite tours in Rome” → product pages

2.3 Page templates

Front-End

- Render JSON-LD in `<script type="application/ld+json">` per page.
- Ensure canonical URLs and hreflang pairs are reflected in schema when relevant.
- Defer loading where possible (but keep inline for crawl).

01. GLOBAL UI / TRUST COMPONENTS — Decision Platform Messaging & Trust Layer (SSOT) — v0 (Discovery-Driven)

Status
This section is an initial concept and an architectural requirement.
All placement rules, UI variants, and proof policies must be finalized during the 3-week discovery with engineering + SEO + Product/Commercial.
Do not treat this as final UI spec. Treat it as a mandatory workstream.

Purpose
Rosotravel is a decision platform, not a catalog.
Users must understand:
- why they see shortlists (not thousands of items),
- what standards we apply,
- what badges mean (with proof),
- where “more options” live and why they are separated,
- and why Rosotravel picks are trustworthy.

This layer must be present across ALL templates and flows:
- City Page
- Things-to-do Page
- Attraction Page
- Category Pages and Near-pages
- Product Page
- Phase B overlays (B2C + B2B/Travel Agent)
- Checkout attribution messaging (where relevant)

Core requirement:
Every label, tooltip, microcopy, CTA caption, and explanation text defined here MUST be editable in CMS (SSOT), versioned, and role-governed.

A) Messaging Objects (SSOT “Copy Tokens”)

Define a global “Decision Messaging Library” in CMS with stable IDs.
Each token includes:
- token_id (stable)
- default_text
- optional variants by audience profile
- optional variants by template type
- optional variants by context (main stake vs Phase B vs Travel Agent)
- tooltip_text (optional)
- governance: owner role + approval requirement
- last_updated_at + version

Examples of required token families:
- WHY_SHORTLIST (1–2 line explanation)
- SHORTLISTED_X_OF_Y (computed string template)
- QUALITY_STANDARD_TOOLTIP (computed thresholds; no hardcoded numbers)
- ON_YOUR_REQUEST_FRAMING (Phase B only)
- STILL_NOT_SURE_SAFETY_NET (CTA strip)
- DECISION_PLATFORM_TAGLINE (compact trust strip)
- BADGE_* tooltips (Guaranteed Access, Curated by founders, Verified, Best Value, Premium, etc.)
- SEE_ALL_OPTIONS_IN_SEGMENT (CTA)
- WHY_THIS_PICK (mini explanation line per card/rail)

Hard rule:
No template may hardcode these texts directly in code.
All must be editable via CMS tokens.

B) Mandatory Components (Concept v0)

B1) “Why you see few picks” (Decision Narrative)
A compact explanation shown near curated sections.
Goal: prevent the “you have low inventory” interpretation.

B2) “Shortlisted X-of-Y” proof line (computed)
Displayed per rail/section:
- X = items shown
- Y = eligible candidate count after gates + dedup in that rail/archetype context
Must be computed and logged (Explainability must show how Y was computed).

B3) “Quality Standard” tooltip (computed)
Explains:
- default standard (e.g., rating floors),
- why “on your request” exists,
- and that “on your request” is separated.
Numbers must be computed from Product Ops configuration (no hardcoded values).

B4) “Still not sure?” safety net CTA strip
A compact safety net displayed in key templates:
- Plan my trip (Trip Planner)
- Ask our expert
- Contact support (optional)

C) Badges with Explanations (SSOT + Proof Policies)

Badges must:
- have explicit meaning text (tooltip)
- have explicit proof policy (when allowed)
- be audited (why badge was shown for this item)

C1) Guaranteed Access (policy required)
Meaning:
Rosotravel guarantees entry conditions under defined operational rules.
Requirement:
Badge may be shown only if the product meets the Guaranteed Access proof policy.
Proof policy must be defined during discovery (data sources + enforcement).

C2) Curated by founders (policy required)
Meaning:
Founders manually curated/approved this pick.
Requirement:
May be shown only when explicit founder approval exists (not default).

C3) Best Value / Premium (policy required)
Meaning:
Best Value = strongest value-for-money within archetype.
Premium = higher comfort/quality execution within archetype.
Requirement:
Labels must align with value tier logic and must not be manually applied without audit.

C4) Verified booking / verified review (policy required)
Meaning:
Review is tied to a verified booking.
Requirement:
Badge shown only when verification is available.

D) Placement Rules (Discovery Workstream)

This section does not finalize exact placements yet.
Discovery must define a “Placement Matrix” for each template and key UI zone.

Discovery output must specify, for each template:
- which components appear
- where they appear (zones)
- when they appear (conditions)
- which copy token IDs are used
- whether it is always visible or collapsible
- mobile behavior

Templates that must be covered:
- City Page (Top Picks, Rails, Top attraction tours, Phase B entry)
- Things-to-do Page (Rails + feeds + Phase B entry)
- Attraction Page (Our 2 picks + ticket-first block + POI pack entry)
- Category Pages / Near-pages (decision explanation + shortlist proof)
- Product Page (why recommended + badges + operator trust cues)
- Phase B overlay (B2C + Travel Agent)
- Checkout surfaces (attribution messaging where required)

E) Product Page decision messaging (explicit requirement)

The Product Page must include decision messaging components in multiple places:
- near the primary CTA (why this is recommended; badges with meaning)
- near key specs (instant confirmation, cancellation, meeting/pickup)
- near reviews (verified review badge)
- optionally: “How we curate” short strip (one line)

Exact placement is defined during discovery via the Placement Matrix.

F) Phase B separation messaging (explicit requirement)

Phase B must always distinguish:
- curated “standard” picks
vs
- “on your request” (below standard / popular-but-lower-rated / non-instant)
This separation requires:
- clear framing copy (CMS token)
- visual separation (UI spec)
- zero mixing rules

G) Governance and Approval (must be implemented)

All messaging tokens and badge tooltips must be:
- editable in CMS
- versioned
- auditable (who changed what)
- role-governed (SEO Lead approval for messaging changes; Admin override logged)

H) Discovery Deliverables (must be produced)

Discovery must produce:
1) Placement Matrix (template × zone × component × conditions)
2) Badge Proof Policies (Guaranteed Access, Curated by founders, Verified, etc.)
3) Copy Token Library (token IDs + default texts + variants)
4) Explainability hooks (how X-of-Y and badge proofs are computed and displayed)
5) A/B testing plan (optional) and safe rollout plan

End of section (v0). Finalization required in discovery.

Messaging & Trust Layer (Discovery-Driven) - Matrix Decision Platform

Status
This is a v0 placement matrix to guide discovery.
Exact placement, conditions, and proof policies must be finalized during the 3-week discovery (Engineering + SEO + Product Ops).
All copy must be CMS-editable (SSOT tokens), versioned, and role-governed.

==
0) CMS Copy Token Library (v0) — token IDs to create (all editable)
==

Core decision narrative
- TOKEN_DECISION_TAGLINE_SHORT
Example: “Curated picks. Instant confirmations. Local experts.”
- TOKEN_WHY_SHORTLIST_1L

Example: “We shortlist the best options so you don’t need to compare hundreds of similar tours.”
- TOKEN_SHORTLISTED_X_OF_Y
Template: “Shortlisted {X} of {Y} eligible options.”
- TOKEN_QUALITY_STANDARD_TOOLTIP
Template: “Standard picks meet our quality floor. ‘On your request’ shows additional options when you ask for a specific POI/topic.”
- TOKEN_STILL_NOT_SURE_STRIP
Example: “Still not sure? Plan your trip or ask our expert.”

Phase B separation
- TOKEN_ON_YOUR_REQUEST_FRAMING
Example: “Popular attractions can be overcrowded and sometimes lower-rated. We recommend our handpicked picks above. If you specifically want this, here are options on your request.”
- TOKEN_PHASEB_TAB_MAGIC_FINDER
- TOKEN_PHASEB_TAB_BROWSE_CATALOG
- TOKEN_PHASEB_EXPAND_REGION
Example: “Expand to nearby destinations (region)”

Badges (each includes tooltip token)
- TOKEN_BADGE_GUARANTEED_ACCESS
- TOKEN_BADGE_GUARANTEED_ACCESS_TOOLTIP (proof-policy required)
- TOKEN_BADGE_CURATED_BY_FOUNDERS
- TOKEN_BADGE_CURATED_BY_FOUNDERS_TOOLTIP (explicit founder approval required)
- TOKEN_BADGE_BEST_VALUE
- TOKEN_BADGE_BEST_VALUE_TOOLTIP
- TOKEN_BADGE_PREMIUM
- TOKEN_BADGE_PREMIUM_TOOLTIP
- TOKEN_BADGE_VERIFIED_REVIEW
- TOKEN_BADGE_VERIFIED_REVIEW_TOOLTIP

Card-level explainability
- TOKEN_WHY_THIS_PICK_SHORT
Template: “Why this pick: {reason_1}, {reason_2}.”
- TOKEN_SEE_ALL_OPTIONS_IN_SEGMENT
Example: “See all options in this segment”
- TOKEN_CHOOSE_BY_INTENT
Example: “Choose by intent”

Travel Agent / Partners
- TOKEN_AGENT_DISCOUNT_STRIP
Template: “Your agent rates: Originals {X}% • Partners {Y}%”
- TOKEN_AGENT_BELOW_STANDARD_WARNING
Example: “Below Rosotravel standard — use only when the customer requests a specific POI/ticket.”
- TOKEN_AGENT_EXPORT_CTA (optional)

Trip Planner / Ask Expert
- TOKEN_TRIP_PLANNER_TEASER_1L
Example: “Not sure what fits your trip? Get a 1–3 day plan by email.”
- TOKEN_TRIP_PLANNER_CTA
Example: “Plan my trip”
- TOKEN_ASK_EXPERT_CTA
Example: “Ask our expert”

Ticket-first / POI page
- TOKEN_TICKET_FIRST_EXPLAINER
Example: “Ticket-first: we show the cleanest ticket/city-card options first. Guided tours appear only as upgrades.”

Live Snapshot
- TOKEN_LIVE_SNAPSHOT_TITLE_CITY
Example: “This week in {City} with Rosotravel”
- TOKEN_LIVE_SNAPSHOT_TITLE_POI
Example: “Live at {Attraction}”
- TOKEN_LIVE_SNAPSHOT_FALLBACK
Example: “Last month with Rosotravel in {City}”

Governance (meta)
- All tokens are CMS-editable with versioning + audit log.
- SEO Lead approval required for edits to: decision narrative, quality standard tooltip, badge tooltips.
- Product Ops Lead approval required for edits to: agent warnings/discount strip texts (optional).
- Admin can override with audit logging.

==
1) Placement Matrix — City Page (/{lang}/{country}/{city}/)
==

Zone definitions (City Page)
Z1 Hero Zone
Z2 Decision Stack (Bundles → Top Picks → Private CTAs → Rails → Top attraction tours → Top tickets)
Z3 Phase B Entry (Show more options button)
Z4 Trust Bridge (immediately after Phase B entry)
Z5 SEO/Content Zone (tips, neighborhoods, guides, reviews, FAQ, comments, planner)

Rows (City Page)

1.1 Z1 Hero Zone — Trust strip
- Component: Decision Platform Tagline strip
- Purpose: establish “decision platform” identity quickly
- Display: always

- Data: none
- Tokens: TOKEN_DECISION_TAGLINE_SHORT
- CMS-editable: yes
- TBD: final UI style (strip vs microtext)

1.2 Z2 Rails header — “Why you see few picks”
- Component: 1–2 line explanation near Rails
- Purpose: prevent “low inventory” perception; explain curated model
- Display: always (City Page)
- Data: none
- Tokens: TOKEN_WHY_SHORTLIST_1L
- CMS-editable: yes
- TBD: exact placement (above rails vs inside each rail header)

1.3 Z2 Rail header — “Shortlisted X-of-Y”
- Component: computed proof line per rail
- Purpose: show “we chose 2 of 20” proof
- Display: per rail when Y >= 6 (recommended) else optional
- Data: X (shown count), Y (eligible_count_after_gates_dedup for that rail context)
- Tokens: TOKEN_SHORTLISTED_X_OF_Y
- CMS-editable: yes (template text)
- TBD (critical): definition of Y per rail, logging, and explainability linkage

1.4 Z2 Rail header — Quality Standard tooltip
- Component: tooltip/icon explaining standard vs on-your-request
- Display: always near Rails and Phase B entry
- Data: thresholds from Product Ops config (no hardcoded numbers)
- Tokens: TOKEN_QUALITY_STANDARD_TOOLTIP
- CMS-editable: yes
- TBD: exact numbers shown or kept abstract (recommended: abstract + link “learn more”)

1.5 Z2 Product cards — Badge tooltips (when applicable)
- Component: badges: Guaranteed Access / Best Value / Premium / Verified Review / Curated by founders
- Display: only when proof policy satisfied
- Data: badge eligibility proofs (policy-defined)
- Tokens: TOKEN_BADGE_* + TOKEN_BADGE_*_TOOLTIP
- CMS-editable: yes
- TBD: proof policy definitions + enforcement rules (discovery)

1.6 Z2 Product cards — “Why this pick” microline (optional)
- Component: short “why” line per Top Pick / per rail winner
- Display: optional v1; recommended on Top Picks + rail cards (collapsed on mobile)
- Data: top reason codes (e.g., “top-rated”, “multi-language”, “best value”)
- Tokens: TOKEN_WHY_THIS_PICK_SHORT
- CMS-editable: yes (template)
- TBD: which reasons to expose; max length; localization

1.7 Z2 After Rails — Inline Trip Planner teaser (mini CTA)
- Component: 1-line teaser + CTA
- Display: always after Rails
- Tokens: TOKEN_TRIP_PLANNER_TEASER_1L + TOKEN_TRIP_PLANNER_CTA
- CMS-editable: yes
- TBD: modal vs scroll-to planner section behavior

1.8 Z3 Phase B Entry — “Show more options” + helper copy
- Component: CTA + short helper text (optional)
- Display: always
- Tokens: TOKEN_SEE_ALL_OPTIONS_IN_SEGMENT (or a dedicated “Show more options” token) + TOKEN_QUALITY_STANDARD_TOOLTIP (icon)
- CMS-editable: yes
- TBD: label naming consistency across templates

1.9 Z4 Trust Bridge — Live Snapshot title and microcopy
- Component: live metrics title + fallback label
- Display: show if city volume >= threshold, else fallback to 30-day
- Tokens: TOKEN_LIVE_SNAPSHOT_TITLE_CITY + TOKEN_LIVE_SNAPSHOT_FALLBACK
- CMS-editable: yes
- TBD: volume thresholds and which metrics shown (Phase D)

1.10 Z5 Safety net strip (optional)
- Component: “Still not sure?” strip with CTAs
- Display: at least once in the page (recommended: before FAQ or before comments)
- Tokens: TOKEN_STILL_NOT_SURE_STRIP + TOKEN_TRIP_PLANNER_CTA + TOKEN_ASK_EXPERT_CTA
- CMS-editable: yes
- TBD: exact placement

==
2) Placement Matrix — Things-to-do Page (/{lang}/{country}/{city}/things-to-do/)
==

Zone definitions (Things-to-do)
T1 Hero Zone
T2 Browse Curated Stack (Top Picks → Bundles → Rails → Private CTAs → Top attraction tours → Categories → Top activities feed)
T3 Phase B Entry
T4 Trust/SEO Zone (reviews, FAQ, content)

2.1 T1 Hero — Trust strip (optional)
- Component: Decision Platform Tagline
- Display: optional (recommended)
- Tokens: TOKEN_DECISION_TAGLINE_SHORT
- CMS-editable: yes

2.2 T2 Rails header — Why shortlist
- Same as City Page reference:
Use TOKEN_WHY_SHORTLIST_1L
- Display: always near Rails

2.3 T2 Rail header — Shortlisted X-of-Y (computed)
- Same as City Page (do not duplicate logic; reuse)
- Data: Y defined per rail context for things-to-do set
- Tokens: TOKEN_SHORTLISTED_X_OF_Y

2.4 T2 “Top activities feed” — “Curated feed” helper copy (TBD)
- Component: 1-line copy explaining this is curated, not full catalog
- Display: recommended above feed
- Tokens: (new) TOKEN_FEED_CURATED_EXPLAINER (add in CMS)
- TBD: final rule of feed sizing (band caps) and diversity enforcement (discovery)

2.5 T3 Phase B Entry
- Reference City Page Phase B spec (2.3.3)
- Tokens: TOKEN_PHASEB_TAB_* + TOKEN_QUALITY_STANDARD_TOOLTIP
- CMS-editable: yes

==
3) Placement Matrix — Attraction Page (/{lang}/{country}/{city}/{attraction}/)
==

Zone definitions (Attraction)
A1 Hero + snippet
A2 Live at POI strip
A3 Our 2 picks + intent mini-rails + ticket-first block
A4 “See more options for this attraction” (POI PACK entry)
A5 Entity content + visitor info + history + nearby
A6 Reviews/FAQ/comments + PDF hook

3.1 A3 “Our 2 picks” header — Shortlisted X-of-Y for this POI (computed)
- Component: “Shortlisted 2 of {Y} eligible ways to visit”
- Display: show when Y is known and >= 4; else optional
- Data: Y = eligible_count for this POI core after gates+dedup
- Tokens: TOKEN_SHORTLISTED_X_OF_Y (specialized template optional)
- TBD: definition of Y for POI page (no POI merge allowed)

3.2 A3 Ticket-first explanation (when ticket block shown)
- Component: small explainer
- Display: only when ticket-only block exists
- Tokens: TOKEN_TICKET_FIRST_EXPLAINER
- CMS-editable: yes

3.3 A3 Card badges + why-line
- Same as City Page: badges + tooltips + optional why line

- Tokens: TOKEN_BADGE_* + TOKEN_WHY_THIS_PICK_SHORT

3.4 A4 POI PACK entry
- Component: CTA label + standard tooltip
- Tokens: TOKEN_SEE_ALL_OPTIONS_IN_SEGMENT + TOKEN_QUALITY_STANDARD_TOOLTIP
- Reference: Phase B POI PACK rules in 2.3.3 (no duplication)

3.5 A6 PDF hook copy
- Component: hook copy tokens (new)
- Tokens: TOKEN_POI_PDF_HOOK_TITLE + TOKEN_POI_PDF_HOOK_TEASER (add in CMS)
- CMS-editable: yes
- TBD: localization + consent copy

==
4) Placement Matrix — Category Pages / Near-pages (City-level)
==

Scope
Includes city category pages and decision near-pages that surface curated choices by intent.

Requirement (v0)
Every category/near-page must include:
- decision narrative (why shortlist)
- shortlist proof (X-of-Y) when meaningful
- quality standard tooltip
- “see all options” CTA leading to Phase B overlay (no indexable variants)

Placements
- Near-page hero zone: TOKEN_DECISION_TAGLINE_SHORT + optional TOKEN_WHY_SHORTLIST_1L
- Above first curated list/rail: TOKEN_SHORTLISTED_X_OF_Y + TOKEN_QUALITY_STANDARD_TOOLTIP
- End of page: TOKEN_STILL_NOT_SURE_STRIP

TBD
Exact near-page template behavior must be reviewed and aligned with Model C (decision-first, not informational-only).

==
5) Placement Matrix — Product Page
==

Purpose
Product Page must reinforce “why recommended” and trust signals at purchase decision moment.

Required zones (v0)
P1 Near primary CTA (top fold)
- Badges + tooltips (Guaranteed Access / Best Value / Premium / Verified)
- Optional “Why this pick” microline
Tokens: TOKEN_BADGE_* + TOKEN_WHY_THIS_PICK_SHORT

P2 Near key specs (logistics and safety)
- Instant confirmation label (from data) + tooltip token (new)
- Cancellation framing (from data) + tooltip token (new)
Tokens: (new) TOKEN_INSTANT_CONFIRMATION_TOOLTIP,
TOKEN_CANCELLATION_TOOLTIP

P3 Reviews section
- Verified Review badge + tooltip
Tokens: TOKEN_BADGE_VERIFIED_REVIEW +
TOKEN_BADGE_VERIFIED_REVIEW_TOOLTIP

P4 “How we curate” micro-strip (optional)
- One-line decision platform tagline
Tokens: TOKEN_DECISION_TAGLINE_SHORT or dedicated
TOKEN_HOW_WE_CURATE_1L

TBD
Exact placements and UI styling to be finalized in discovery.

==
6) Placement Matrix — Phase B Overlay (B2C + Travel Agent variant)
==

B2C Phase B overlay
Required messaging:
- Tabs: Magic Finder / Browse Catalog (tokens)
- Quality standard tooltip
- On-your-request framing (when section shown)
- Expand region CTA (when triggered)
Tokens:
TOKEN_PHASEB_TAB_MAGIC_FINDER
TOKEN_PHASEB_TAB_BROWSE_CATALOG
TOKEN_QUALITY_STANDARD_TOOLTIP
TOKEN_ON_YOUR_REQUEST_FRAMING
TOKEN_PHASEB_EXPAND_REGION

Travel Agent (B2B) variant (authenticated)
Additional required messaging:
- Discount strip
- Below-standard warning (on-your-request)
Tokens:
TOKEN_AGENT_DISCOUNT_STRIP

TOKEN_AGENT_BELOW_STANDARD_WARNING

Hard rule:
All Phase B labels must be editable via CMS tokens.

==
7) Discovery Checklist (must be completed)
==

Discovery must finalize:
1) Placement Matrix (exact zones and mobile rules)
2) X-of-Y computation definitions and logging (eligible_count source and caching)
3) Badge proof policies + enforcement (Guaranteed Access, Curated by founders, Verified)
4) Copy token library coverage (add missing tokens for product page tooltips and PDF hooks)
5) Governance workflow for edits (roles + approval)
6) A/B testing and safe rollout plan (optional)

Brand Expertise + Reputation Strip for Country and City Pages

Status: ADD as a new, indexable body section in the page template.
Purpose: strengthen brand understanding for SEO and LLMs, position Rosotravel as a destination-specific decision-first travel expert, and combine three trust layers in one place:
1) what Rosotravel says about itself in the destination,
2) how travelers rate Rosotravel in that destination,
3) what travelers most value in Rosotravel there.

IMPORTANT
This is NOT a DEE snippet.
This is NOT hero copy.
This is NOT a dynamic runtime block.
This is a stable, indexable, body-content section generated during the normal content workflow and editable/lockable in CMS.

This section must work as:
- brand grounding for SEO,
- entity grounding for LLMs,
- and a trust/reputation explanation block for users.

CRITICAL SEO / SCHEMA NOTE
Aggregate rating structured data may improve machine readability and may help eligibility for enhanced search presentation, but visible stars in Google are never guaranteed.
Do not promise star display in Google.
Implement structured data only when:
- the rating is visible on the page,
- the review count is sufficient,
- the value is factual and updated,

- and the structured data matches the visible value.

C) Shared quality and compliance rules for Country and City

1. This content is indexable and part of the main body.
2. It must be unique enough to avoid a boilerplate footprint across destinations.
3. It must reinforce the same global identity:
Rosotravel = global experience expert + decision-first brand + curated operator model.
4. It must never contradict the global About Us.
5. It must follow the same truthfulness rule:
no fake local presence, no fake legal claims, no false ownership claims.
6. It should reference:
- expertise,
- curation,
- destination understanding,
- and decision support,
more than generic company marketing fluff.
7. It must NOT be generated as a tiny snippet.
It is a real explanatory body block.
8. All visible rating values must come from factual review aggregation only.
9. The page must never display a score if the sample is too small or misleading.
10. Top praise themes must come from clustered verified review data.
11. The explanatory text may be AI-generated, but:
- ratings,
- counts,
- and review-theme bullets
must be derived from data and only phrased by AI if needed.
12. The visible reputation strip and the schema values must always match.
13. Do not mix:
- global Rosotravel rating
with
- city-specific or country-specific rating
unless clearly separated.
14. This block must not overclaim:
- “best in the city”
- “highest rated”
- or similar leadership phrasing
unless separately proven and approved.

D) Shared generation guidance (AI draft + factual insertions)

The AI generation process for both Country and City sections must work in two layers:

Layer 1 — factual aggregation layer
System prepares:
- rating value
- rating count
- top review themes
- destination context fields
- product mix and operating model hints

Layer 2 — AI drafting layer
AI drafts:
- title
- intro
- explanatory paragraphs
- integration of rating/reputation into narrative

Important:
AI must not invent rating values, review volumes, or praise themes.
It may only compose them into human-readable text.

Add this note to both Country and City generation instructions:

"This section must combine three layers:
1) Rosotravel's own explanation of its role in the destination,
2) verified traveler reputation in that destination,
3) and the most valued service traits derived from local review evidence.
Together, this forms a complete brand understanding block for SEO, LLMs, and user trust."

E) Short workflow note (optional to paste elsewhere)

Country Page Brand Expertise + Reputation block:
- generated during destination-level "Run SEO Auto"
- explanatory text = AI drafted
- rating/count/themes = factual aggregation
- editable and lockable in CMS

City Page Brand Expertise + Reputation block:
- generated during city-level "Run AI Batch Processing"
- explanatory text = AI drafted
- rating/count/themes = factual aggregation
- editable and lockable in CMS

02. Landing pages detailed description and flow

HOMEPAGE TEMPLATE (Brand Hub + Decision Platform Entry)

Status: MUST VERIFY IN DISCOVERY (routing, metrics sources, Trustpilot integration, and URL patterns).

URL

/{lang}/

Intent

Top-level brand hub, discovery entrypoint, trust builder, segmentation entry, and conversion gateway to:

- cities (decision-first),

- curated categories / near pages (decision pages),

- attractions (POI hubs),

- and top-performing experiences,

while clearly positioning Rosotravel as a decision platform (curated picks, not an OTA catalog).

Primary Schema Types

- WebPage

- ItemList (top destinations, trending experiences)

- Organization

Global UI note

Homepage uses the global Header and Footer (already implemented). This template defines homepage content blocks only.

All CTAs must route to canonical, curated surfaces (no parameter-indexed duplicates).

INDEXING / SEO RULES (Homepage)

- Canonical: /{lang}/

- hreflang: standard language alternates

- No comments/Q&A block on homepage

- No filter parameters allowed to create indexable variants

2.0.1 Sections (Front-End, Desktop Order — FINAL)

1) Hero Block — Mission + emotional promise + Search

2) Decision Narrative Strip — “Why you see shortlists” + standards + “Who we are” link (MUST)

3) Rosotravel Live — proof block (MUST)

4) Top Destinations — curated geographic entry

5) Explore by Interest — category tiles (curated)

6) Trending Experiences — curated winners (global/local)

7) Social Proof — “Why travelers choose Rosotravel” + Trustpilot badge (if enabled)

8) Featured Guides / Experts — human face layer

9) Blog Highlights — editorial E-E-A-T

10) Email Capture Hook — “save picks / travel circle”

Footer

Mobile

Same order stacked. Hero search stays prominent.

2.0.2 HERO BLOCK (Mission + emotional promise + Search)

Ownership

Manual brand lines + prewritten snippet variants (AI-assisted drafting allowed offline) + DEE selection.

DEE does not generate text in real time.

A) Hero media

- Large hero photo or short looping video.

- CMS: homepage_hero_media_id

B) H1 (manual; brand line)

Recommended H1:

“Outstanding memories — chosen by people who live the experience.”

C) Subheadline (manual; emotional decision promise; no technical language)

Recommended:

“Feel the place, not the planning. We pick the experiences that turn a day into a memory — with clear booking and reliable delivery.”

D) Hero snippet (≤260 chars; prewritten variants aligned to Model C audiences)

Purpose:

Make the value clear instantly: curated decisions + realistic pacing + trust.

CMS fields (prewritten variants):

- homepage_snippet_default (fallback)

- homepage_snippet_first_time_visitor (Model C default baseline)

- homepage_snippet_family_planner

- homepage_snippet_comfort_easy_pace_traveler

- homepage_snippet_solo_social_traveler

- homepage_snippet_interest_deep_dive_traveler

- homepage_snippet_active_adventure_traveler

DEE selection rule (MVP)

1) If audience profile inferred with sufficient confidence → show that snippet variant.

2) Else → show homepage_snippet_default.

Recommended snippet copy (examples; final copy is CMS-editable)

- default:

"Rosotravel helps you choose with confidence: fewer options, better days. We focus on experiences that feel real, run smoothly, and leave you satisfied — from booking to the final stop."

- first_time_visitor:

"New to a destination? Start with our safest classics: clear logistics, strong reviews, and zero overwhelm. We shortlist the best ways to see the city — and keep your day realistic."

- family_planner:

"Family days should feel easy. We prioritize low-friction plans, clear meeting points, and options that work with kids — so you spend more time experiencing, less time figuring it out."

- comfort_easy_pace_traveler:

"Prefer an easier pace? We highlight options with simpler logistics, less strain, and reliable entry — so the day feels comfortable, not rushed."

- solo_social_traveler:

"Traveling solo? Find small-group experiences that feel welcoming and social — with clear confirmations and a smooth plan."

- interest_deep_dive_traveler:

"Want depth, not just photos? Choose experiences built around context, stories, and meaning — the kind that makes the destination 'click'."

- active_adventure_traveler:

“Chasing signature activities? We pick high-impact experiences with realistic pacing and clean booking — so the thrill doesn’t turn into chaos.”

E) Primary CTA + Search

- Primary CTA: “Start planning”

- Global search bar:

placeholder: “Where do you want to go?”

- Autocomplete uses internal index (cities, attractions/POIs, curated categories, products).

- Search routing must be deterministic (see Routing Table v1).

F) Quick shortcuts (curated)

- 4–6 tiles configured in CMS: homepage_quick_categories[]

Examples: “Top Cities”, “Top Attractions”, “Family planning”, “Less walking”, “Day trips”, “Skip-the-line”.

Hard rule:

These shortcuts route only to canonical curated pages (no filter parameter URLs).

2.0.3 DECISION NARRATIVE STRIP (MUST) — “Why you see shortlists” + Standards

Purpose

Prevent “small catalog” perception and communicate decision-platform logic in plain language.

Placement

Directly under hero.

UI structure (compact)

- Line 1: “Why you see shortlists” (CMS-editable token)

- Line 2: “Our standard” tooltip (computed thresholds; no hardcoded numbers in copy)

- Micro CTA: “How we choose” (opens a lightweight modal)

- Optional link: “Who we are” → /about/ (label CMS-editable)

Modal: “How we choose” (content is CMS-editable, versioned)

Must cover (short):

- We prioritize reliability and satisfaction (Fulfilment)

- We prefer clear logistics and realistic pacing

- We apply quality floors and avoid near-duplicates

- We separate “our picks” vs “on your request” options (where applicable)

Hard rule

Do not mention AI curation. Use “Rosotravel experts” / “Rosotravel standards”.

All text must be CMS-editable (SSOT tokens).

2.0.4 ROSOTRAVEL LIVE (MUST) — Proof block (rolling 7 days; fallback to 30)

Purpose

Show proof, not promises. Express human delivery and operational reality.

Placement

Immediately after Decision Narrative Strip.

Elements (display)

- “This week with Rosotravel”

- Tours run: X

- Travelers hosted: Y

- Average rating: Z

- Top cities: top 3 active cities (links)

- Guide avatars: last 3 guides who led tours (name + city)

- “What guests loved this week” (2–3 bullets; summary)

Fallback rule

If low volume in last 7 days (below thresholds):

- Switch to 30-day window and display: “Last month with Rosotravel”

Thresholds (configurable; defaults):

- tours_run_7d < 10 OR travelers_hosted_7d < 80 → use 30 days

DATA SOURCES & CALCULATION RULES (technical; required)

A) Data sources (system tables / events; names verified in discovery)

- Departures (tour instances):

source: operations_departures (or equivalent) with statuses

- Bookings:

source: orders + order_items + pax_count (and destination mapping)

- Reviews:

source: verified_reviews linked to booking_id (and city_id / product_id)

- Guides:

source: guide_assignments linked to departure_id

B) Definitions (hard)

- tours_run:

count of departures with status = COMPLETED within the window

exclude: CANCELLED, NO_SHOW, FAILED, TEST

- travelers_hosted:

sum(pax_count) across COMPLETED departures within the window

- average_rating:

weighted average of verified reviews associated to COMPLETED departures in the window

(final choice in discovery: by review_created_at vs by departure_date; must be consistent)

- top_cities:

top 3 cities by tours_run within the window

- guide_avatars:

last 3 unique guides with COMPLETED departures within the window (deduplicate by guide_id)

C) Computation frequency + caching (performance)

- Compute nightly aggregates (primary) + near-real-time incremental updates (optional).

- Cache layer:

homepage_live_metrics_{lang}_{region} TTL 15 minutes (recommended)

- If data unavailable:

show safe fallback copy and hide metrics (do not show zeros as proof).

D) Trust & fraud safety (hard)

- Only verified operational events count (no manual admin injection without logging).

- Only verified reviews count (linked to booking_id).

- Any manual override requires Admin-only action and audit log.

2.0.5 TOP DESTINATIONS (curated geo entry)

Purpose

Lead users into the geographic funnel:

Country → City → Things-to-do → Attraction → Product/Bundle.

Display

Grid of 6–12 destination cards (city-first by default).

Each card:

- image

- city name (+ country)

- snippet ≤120 chars (prewritten; editable)

- CTA: “Explore” → City Page

SELECTION & RANKING (must be concrete; no random)

Eligibility (hard)

- Only cities with published City Page (and minimum curated inventory) can be listed.

- Exclude cities flagged as “low-inventory sparse” unless manually pinned.

Ranking score (default formula; tunable in Product Ops config)

destination_score = 0.45 * bookings_28d_norm

+ 0.20 * conversion_rate_28d_norm

+ 0.15 * ctr_28d_norm

+ 0.10 * rating_norm

+ 0.10 * recency_activity_norm

Where:

- bookings_28d_norm: normalized bookings volume for city in last 28 days

- conversion_rate_28d_norm: city page → product page → checkout conversion

- ctr_28d_norm: SERP CTR (if available) or internal click-through

- rating_norm: aggregate rating of city's curated Set products (verified sources only)

- recency_activity_norm: recent operational activity indicator (departures / reviews)

Locale adjustment (optional)

If user locale is known:

- boost destinations within the same region (e.g., EU user sees EU cities higher)

- do not hide global destinations; only reorder.

Manual controls

- homepage_destinations_pins[]: pinned destinations always appear in top positions

- homepage_destinations_exclusions[]: remove destinations even if score is high

Caching

- Recompute daily (nightly job) and cache results.

2.0.6 EXPLORE BY INTEREST (Category tiles; curated)

Purpose

Help users who think in experiences, not geography.

Display

10–16 tiles, each linking to a curated discovery surface:

- either a governed global category hub (if exists), OR

- a city-level category chooser / near page suggestion flow.

Hard rule

No parameterized filter URLs. Tiles must route to canonical curated pages only.

CMS

homepage_interest_categories[] (taxonomy references + tile copy + icons)

DEE behavior

Reorder only (optional). Do not change destination of links in runtime.

2.0.7 TRENDING EXPERIENCES (Global / Local) — curated winners

Purpose

Show a small set of high-performing experiences without becoming an OTA catalog.

Eligibility (hard)

- Only curated inventory allowed:

products must be Set members in at least one destination scope OR be globally curated "approved winners".

- Exclude non-INSTANT items from main trending list (they may appear later in “more options” flows, not homepage trending).

- Exclude products below quality floors (as configured in Product Ops).

Selection pool

- Use published, curated products only (no RAW/pool-only products).

Ranking (default formula; tunable)

trending_score = 0.40 * bookings_28d_norm

+ 0.25 * revenue_28d_norm

+ 0.15 * conversion_rate_28d_norm

+ 0.10 * rating_component

+ 0.10 * review_volume_component

rating_component:

- normalized rating (post-quality-gate; no below-floor products)

review_volume_component:

- normalized log(review_count)

Output

- 10–16 cards (carousel), localized where possible:

if user region known → prefer nearby destinations first, else global mix.

Originals visibility (brand policy)

- Rosotravel Originals must be clearly labeled and can be pinned higher if within tolerance of top score.

- Partner products remain eligible when they are top performers and meet quality standards.

Caching

- Recompute daily + optional 6-hour refresh.

- Cache TTL 30 minutes.

2.0.8 SOCIAL PROOF (Why Travelers Choose Rosotravel) + Trustpilot

Purpose

Reinforce trust quickly with proof-led statements.

A) Review cluster bullets (3–5)

Source:

- Precomputed nightly summary from verified reviews (no invention).

Rules:

- bullets must remain generic and truthful (no numeric claims unless computed).

CMS:

homepage_social_proof_bullets[] (editable)

B) Trustpilot badge (optional; recommended if adopted)

- Placement: inside this Social Proof section (visible without scrolling too far).

- CMS: homepage_trustpilot_embed_html (embed snippet)

- Rule: embed must be editable and can be disabled by country/locale if needed.

2.0.9 FEATURED GUIDES / EXPERTS (human face layer)

Purpose

Support E-E-A-T and “human trust” signals.

Display

Carousel of expert cards:

photo, name, specialty (city), short tagline, link to expert profile, optional intro video.

Eligibility

- Only published expert profiles with minimum completeness (photo + bio + city coverage).

- Optional: rotate set monthly; allow manual pins.

CMS

homepage_featured_guides[] referencing /cms/experts/

2.0.10 BLOG HIGHLIGHTS

Purpose

Inspiration + SEO + authority.

Display

4–6 article cards:

image, title, teaser, reading time.

Rules:

- newest first

- allow pinned evergreen posts

CMS

homepage_blog_highlights[]

2.0.11 EMAIL CAPTURE HOOK (homepage)

Purpose

Lead capture without discount-first brand damage.

Offer (recommended positioning; CMS editable)

“Get curated picks and planning shortcuts — join the Rosotravel circle.”

Fields

- email + consent checkbox (GDPR)

- optional: destination interest picker (light)

Frequency

- once per new session until submitted (configurable)

CMS

/cms/hooks/homepage

2.0.12 ORGANIZATION SCHEMA (Trust)

Organization schema must include:

- Organization name
- official website
- contact point
- sameAs social profiles (if available)
- brand signals aligned with trust layer

2.0.13 CMS FIELDS OVERVIEW (Homepage)

Core

- homepage_hero_media_id (Human)
- homepage_h1 (Human)
- homepage_subheadline (Human)
- homepage_snippet_* (prewritten variants; AI-assisted drafting allowed; Human editable/lockable)

Decision narrative

- homepage_decision_strip_line1 (Human)
- homepage_decision_strip_line2 (Human)
- homepage_decision_strip_tooltip (Human)
- homepage_decision_modal_content (Human)

- homepage_about_link_label (Human) + homepage_about_link_url (System; /about/)

Live

- live_metrics_enabled (Human)

- live_metrics_window_thresholds (System config; editable in Product Ops config)

Destinations

- homepage_top_destinations[] (System + Human weighting)

- homepage_destination_pins[] (Human)

- homepage_destination_exclusions[] (Human)

Interests

- homepage_interest_categories[] (Human)

Trending

- homepage_trending_products[] (System; curated outputs only)

- homepage_trending_pins[] (Human; optional)

Social proof

- homepage_social_proof_bullets[] (AI precomputed; Human editable)

- homepage_trustpilot_embed_html (Human; optional)

Experts

- homepage_featured_guides[] (Human)

Blog

- homepage_blog_highlights[] (Editorial)

Hook

- homepage_hook_config (System + Human copy)

2.0.14 DISCOVERY VALIDATION (Developers must verify)

Devs must confirm in discovery:

- exact event/table sources for Live metrics and their final definitions
- product URL pattern and canonical logic
- whether global category hubs exist or only city-level categories
- Trustpilot embed strategy (locale/country conditions)
- caching strategy and performance budgets for homepage blocks

End of Homepage template.

"""

Recommended top-level marketing copy summary (for CMS drafting):

- Tone: premium, calm, confident, human.
- Avoid: “AI-curated”, “algorithm”, “shortlist” as technical terms in user-facing copy.
- Emphasize: Experiencing (real moments) + Fulfilment (a day that feels complete), plus Reliability (smooth delivery).

Country Page

URL format:
`/{lang}/{country}/`

Intent:
National-level SEO page, discovery hub, bridges user to city pages or things-to-do categories.

Schema:

- TouristDestination
- ItemList (top cities, attractions)
- FAQPage
- Organization

2.6.1 Sections

Hero + H1

Ownership: AI + manual
Elements:

- Hero image Nano-banan pro generated
- H1: “Discover {{Country_Name}}”
- Snippet 180–260 chars
    - Default + segment variants (DEE chooses version)
    - CMS:
        - Other according to DEE section etc.

DEE behaviour (MVP):

- If user shows family or couples intent, reorder cities/attractions accordingly.

Country Overview

Ownership: AI + manual edits

- 2–3 paragraphs describing culture, travel themes, seasonality.
- Stable, non-dynamic block (SEO).

Highlights (NEW, structured inspiration)

- 5–7 bullets
- AI-generated
- Stable block

Facts & Curiosities (NEW, factual)

- 5–7 verified bullets
- No AI hallucinations
- Tied to factual data sources

Top Cities

Ownership: Factual + AI snippet + DEE sorting
Elements:

- Grid of major cities in the country
- Each card:
  - Image
  - City name
  - Snippet ≤120 chars
  - “X tours available” (factual)
  - Link to city page

- DEE may reorder cities based on segment (families → kid-friendly cities first). snippet_default + segment variants
- used only for semantic personalization
- not indexed

CMS: `/cms/country/snippets/`

Top Regions

Ownership: Factual
Simple grid of regions.
Each card → region page.

Top Attractions in {{Country_Name}}

Ownership: factual + AI snippet
Elements:

- 6–12 cards
- Image, name, 120-char snippet
- Link to attraction page
- DEE can reorder to highlight segment-friendly attractions.

Things to Do in {{Country_Name}}

Ownership: taxonomy
Grid of categories: Food, Museums, Day Trips, Adventure, Family-Friendly, etc.

Other Sightseeing Options in {{Country_Name}}

Ownership: factual
List of all cities:

`City Name (X attractions)`
Links to city pages.

Blog & Guides

Ownership: editorial
List of articles tagged with this country.

Country-Level Reviews

Ownership: aggregated

- “Based on X reviews across Y tours in {{Country_Name}}”
- Sample review cards (photo, name, stars, date, product link).

Place it LOW on the page, after the main country discovery / destination sections and before FAQ / Comments / final bottom-of-page trust modules.

Recommended placement rule:

- insert after the last core discovery section of the Country Page,

- but before Country FAQ and before Comments / Ask our expert.

If the Country Page already has a blog/guides block, this section should appear either:

- after blog/guides and before FAQ,

or

- before blog/guides if blog is highly editorial and this section is intended to act as the main brand grounding block.

Preferred default:

main country discovery content → proof/trust → “About Rosotravel in {{Country}}” → FAQ → comments/footer.

NEW SECTION TITLE:

2.X COUNTRY PAGE — “About Rosotravel in {{Country_Name}}”

Alternative UI labels allowed:

- “Why travelers choose Rosotravel in {{Country_Name}}”

- “Rosotravel in {{Country_Name}}”

Recommended canonical label:

“About Rosotravel in {{Country_Name}}”

Purpose

This section explains Rosotravel’s role in the country as:

- a global experience expert,

- a decision-first travel brand,

- and a company that combines proprietary experiences, guide expertise, and selected partner inventory under one quality logic.

This section exists for three reasons:

1) SEO / LLM brand understanding:

it helps search engines and LLMs understand what Rosotravel is, how it works in this country, and what it specializes in.

2) Decision-platform narrative:

it clarifies that Rosotravel is not a generic catalog and does not simply list everything.

3) Trust:

it gives the traveler confidence that Rosotravel has a real operating logic in this country, even if it does not have a physical office there.

Generation timing

This section is generated during:

- destination-level “Run SEO Auto” for Country Pages.

This means:

- it belongs to the country-level automation flow,

- it is created together with other country-level metadata / snippets / FAQ drafts,

- it is versioned like other AI-generated destination content,

- and can be manually edited / locked in CMS.

Ownership

- AI-generated first draft

- Human editable

- Lockable in CMS

- If locked, it cannot be overwritten by regeneration unless explicitly re-enabled

A1) Country Reputation Strip (required sub-block)

The Country Page section must begin with a compact trust strip.

Place it:

- at the top of “About Rosotravel in {{Country_Name}}”,

- before the explanatory paragraphs,

- visibly inside the same section.

Required visible elements:

1. Country rating line

2. Country review-count line

3. “What travelers value most” bullets (2–4)

4. Optional link to review section or reviews page

Recommended visible wording examples:

- “Rosotravel is rated 4.8/5 by travelers across tours in {{Country_Name}}.”

- “Based on 1,248 verified reviews linked to experiences in {{Country_Name}}.”

- “Travelers most often value: knowledgeable guides, smooth organization, and clear booking.”

Country-level rating rules

The country-level visible rating must be based on:

- verified reviews only,

- products mapped to destinations inside that country,

- review aggregation jobs, not free-form AI estimation.

Recommended initial logic:

- country_brand_rating_value = weighted average rating across verified reviews linked to products mapped to {{Country_Name}}

- country_brand_rating_count = number of verified reviews included in the aggregate

- refresh frequency = nightly (or equivalent batch refresh)

- source of truth = review aggregation + product ↔ country mapping graph

Recommended visibility threshold

Show the numeric country rating strip only if:

- country_brand_rating_count >= 50 verified reviews

Else:

- hide the numeric score,

- optionally keep only “What travelers value most” if enough review-theme evidence exists,

- or hide the strip completely if evidence is too weak.

“What travelers value most” (Country Page)

This must come from clustered verified review themes, not generic copy.

Recommended logic:

- derive top positive themes from the country-level review corpus

- select themes that are:

- frequent,

- clearly positive,

- representative of Rosotravel’s traveler experience in that country,

- and destination-relevant

Examples of acceptable themes:

- knowledgeable local guides

- smooth organization

- easy booking and confirmations

- good pace for sightseeing

- strong value for money

- family-friendly planning

- memorable local stories

Rules:

- show 2–4 short bullets

- keep wording factual and calm

- do not overstate

- do not show themes if review volume is too low

Country-level schema requirement

Add structured data tied to the visible rating strip.

Recommended implementation direction:

- expose the visible Rosotravel-in-{{Country}} block in body content,

- attach machine-readable aggregate rating data to the Rosotravel destination-expertise entity referenced on that page,

- ensure the rating matches the visible number and review count.

Minimum schema fields to expose in country context:

- ratingValue

- reviewCount

- bestRating = 5

- worstRating = 1

Internal developer note:

“Aggregate based on verified reviews of products mapped to this country page.”

Important:

Do not mark up a rating that is not visible to users.

Do not show one number in the body and a different number in structured data.

Recommended CMS/system fields for Country Page

- country_brand_expertise_enabled
- country_brand_expertise_title
- country_brand_expertise_body
- country_brand_expertise_short_intro
- country_brand_expertise_lock
- country_brand_expertise_version
- country_brand_expertise_generated_at
- country_brand_rating_enabled
- country_brand_rating_value
- country_brand_rating_count
- country_brand_top_review_themes[]
- country_brand_rating_updated_at
- country_brand_rating_schema_enabled

A2) Country explanatory body (required)

After the Country Reputation Strip, the country-level explanatory block must contain 4 short sub-parts inside one continuous section:

1. Rosotravel's role in {{Country_Name}}

Explain:

- that Rosotravel helps travelers choose experiences across this country,

- that Rosotravel works as a decision-first expert, not a catalog,

- and that the country page is designed to guide travelers toward the right destinations, cities, and experience types.

2. How Rosotravel operates in {{Country_Name}}

Explain the operating model in a country-safe way:

- Rosotravel may run its own experiences in selected destinations,

- work with local guides / experts,

- and complement its own portfolio with carefully selected partner products.

Important:

Do not imply there is a branch, office, or legal entity in the country unless factually true.

3. How Rosotravel selects and presents experiences in {{Country_Name}}

Explain:

- that Rosotravel does not aim to show every possible option,

- that it organizes experiences into clearer decision paths,

- that it compares similar options under one quality logic,

- and that it helps travelers decide faster and more confidently.

4. Why this matters for travelers in {{Country_Name}}

Explain:

- clearer choices,
- better fit to traveler needs,
- less duplicate noise,
- stronger trust,
- and a more complete planning experience.

Length and content rules

Recommended total length:

- 500–900 words for Country Page version

Recommended paragraph count:

- 4–6 paragraphs

Optional:

- 3–5 short bullet points only if readability genuinely improves

Tone

- expert
- warm
- trustworthy

- calm

- useful

- not jargon-heavy

- not “AI platform” language

What this section MUST communicate

- Rosotravel is a global experience expert

- Rosotravel helps travelers choose, not just browse

- Rosotravel combines proprietary know-how with curated partner selection where relevant

- Rosotravel is destination-aware and city-aware inside the country

- Rosotravel’s country pages are built to support decision-making, not catalog sprawl

- Rosotravel is well-rated by travelers in this country

- Travelers consistently value specific strengths in Rosotravel here

What this section MUST NOT claim

Do NOT claim unless factually true:

- office / branch / headquarters / local company in that country

- legal presence in that country

- exact staff structure in that country

- that every product is operated directly by Rosotravel

- that Rosotravel is the biggest operator in that country

- any unverifiable market leadership claim

Allowed wording examples

- “Rosotravel helps travelers choose experiences across {{Country_Name}} through a mix of destination expertise, curated recommendations, and selected tour inventory.”

- “In {{Country_Name}}, Rosotravel combines its own experience standards with carefully selected local tours and attractions.”

- “Rather than showing an endless list, Rosotravel organizes {{Country_Name}} travel options into clearer decision paths.”

Forbidden wording examples

- “Rosotravel is based in {{Country_Name}}” (unless true)

- “Rosotravel has a local office in {{Country_Name}}” (unless true)

- “We personally run every tour in {{Country_Name}}” (unless true)

- “We are the #1 operator in {{Country_Name}}” (unless proven and approved)

Notes for generation quality

The country-level text must NOT be a boilerplate with only the country name swapped.

To keep it useful for SEO and LLM understanding, generation must vary across countries using:

- destination type profile,

- key cities within the country,

- major experience strengths,

- operating model hints,

- country-level portfolio mix,

- and country-relevant traveler framing.

Minimum variation sources for generation

Use at least:

- country name

- major destination structure in that country

- key travel strengths in that country

- country-specific product mix

- Rosotravel presence logic in that country (own tours / partner-heavy / mixed)

- high-level traveler intents common for that country page

- country rating value and top review themes

FAQ (Country Level)

Ownership: AI + manual

- Exactly 7 FAQs
- Topics:
    - visa
    - currency
    - safety
    - tipping
    - transport
    - best time to visit
    - cultural norms

Schema: FAQPage.

Comments & Ask Our Expert

Ownership: user + expert
Global thread referencing general country questions.
Moderated via CMS.

Hooks

Country Travel Starter Pack (PDF)

- Email capture
- Sends PDF with intro to the country, checklist, “top things to do”, etc.
- Configured: `/cms/content-assets/pdf-manager`.

Optional Exit Intent

- -15% one-time code.

Country Page CMS Fields

- `country_snippet_*` (AI + manual)
- `country_overview_html` (AI + manual)
- `country_top_cities[]` (System)
- `country_top_regions[]` (System)
- `country_top_attractions[]` (System)
- `country_categories[]` (taxonomy)
- `country_other_cities[]` (System)
- `country_blog_articles[]` (Editorial)

- `country_faq_items[]` (AI + manual)
- Hooks configuration

Region Page / Island other Entities

URL format:
`/{lang}/{country}/{region}/`

Intent:
Mid-level destination page connecting country → region → city.

Schema:

- TouristDestination
- ItemList (sights, cities)
- FAQPage

2.7.1 Sections

Hero + H1

Ownership: AI + manual

- H1: “Explore {{Region_Name}}”
- Snippet ≤260 chars
- Hero image Nano-banan pro generated

Dynamic Snippet:

- Small personalization possible via snippet variant.
  CMS: `/cms/region/snippets/`

Region Overview

Ownership: AI + manual
2–3 paragraphs describing landscapes, themes, cultural aspects.

Highlights (NEW, structured inspiration)

- 5–7 bullets
- AI-generated
- Stable block

Facts & Curiosities (NEW, factual)

- 5–7 verified bullets
- No AI hallucinations
- Tied to factual data sources

Top Sights in {{Region_Name}}

Ownership: factual + AI snippet

- Grid/gallery of attractions with counts.
- Each card links to attraction page.

DEE:

- Reorder to favor of 4.8.9 Stage F — Profile model

“Go beyond {{Region_Name}}”

Ownership: factual + taxonomy

- List of thematic or nearby clusters:

- Top day trips from region
- Cities near region
- Thematic lists: “Wine trails”, “Castles”, etc.

Things to Do in Region

Ownership: taxonomy
Grid of categories similar to country/city categories.

Top Cities in Region

Ownership: factual
Card layout: image + city name + snippet.

DEE ordering allowed.

Region-Level Reviews

Ownership: aggregated
Exactly same block style as country-level reviews.

FAQ (Region Level)

Ownership: AI + manual
7 items.
Schema: FAQPage.

Comments & Ask Our Expert

Same shared module as city/country.

2.7.2 Hooks

Region Circuit Planner (AI)

Placement: lower half of page.

Form fields:

- Cities user wants to visit
- Days available
- Interests
- Email address

AI output:

- Suggested route
- Day-by-day outline
- Key attractions

Sent by email + shown onscreen.

CMS configuration: `/cms/ai-tools/planner-settings-region`.

City Page

2.3 CITY PAGE TEMPLATE

URL format

/{lang}/{country}/{city}/

Example: /en/italy/rome/

Intent

City Page is a decision-first trip-planning hub (Model C). It is not a catalog-first page.

Main goals:

- Convert “first-time” city intent into a confident purchase decision through curated winners and plans.

- Keep the page compact and decisive in the upper half (main stake).

- Provide controlled exploration only after explicit user intent (Phase B).

- Maintain strong SEO authority and internal linking in the lower half (content zone).

- Adapt which precomputed blocks are shown (and their ordering) based on DEE profile, without creating indexable segment URLs.

Primary schema types

- TouristDestination (primary)

- ItemList (Bundles, Top Picks, Rails, Top attraction tours, Top Tickets)

- FAQPage (city-level)

2.3.1 Sections (Front-End, Desktop Order — FINAL)

Decision-first stack (main stake)

1) Hero block

2) Optional: “Top picks for you” micro-carousel (4–8)

3) Bundles (1/2/3-day itineraries) — 3 cards

4) Top Picks (Hero picks) — 6 products

5) Private Upgrade (CTA-only) — max 2 CTAs

6) Category Rails (Decision Rails) — max 4 rails × 2 products

6a) Inline teaser (mini CTA): City Trip Planner (after Rails) — compact

7) Top attraction tours in {{City}} — 3–6 product cards

8) Mini “Explore by category” chips (compact; max 8) — quick internal links

9) Top Tickets in {{City}} (ticket-only / city pass / city card) — 4–6 items

10) Phase B entry: “Show more options” (opens overlay; no indexable URLs)

Trust bridge (immediately after Phase B entry)

11) “This week in {{City}} with Rosotravel” (Live Snapshot)

SEO + content zone (below trust bridge)

12) City-level reviews (aggregated; compact)

13) Local tips

14) Best neighborhoods

15) Things to do categories (full block)

16) Blog / travel guides

17) FAQ (city-level; exactly 9)

18) City trip planner hook (full block; lead capture)

19) Comments & Ask our expert

Mobile

Same order, sections stacked. Phase B opens as full-screen overlay.

Mini CTA blocks must remain ultra-compact.

2.3.2 Sections — Detailed Specification (Decision-first stack)

A) Hero block

Ownership: AI + Human (lockable).

Elements:

1) Hero image

- City-level hero representing atmosphere and skyline or iconic view.

- Source: CMS media for the city; optimized via Cloudinary.

2) H1

Format (editable and lockable):

“Visit {{City_Name}} – Your starting point for [type of travel]”

AI-generated from City_Name, Country_Name and dominant city type; human can edit/lock.

3) Dynamic Snippet (DEE variant; NOT SEO body)

- 180–260 characters.

- High-density description of why the city is worth visiting (culture, highlights, vibe).

- Selected by DEE from a prewritten variant set (no realtime generation).

- Stored in CMS under /cms/city/snippets/.

City snippet variants (prewritten; stored in CMS):

- city_snippet_default (fallback when no audience can be inferred)

- city_snippet_first_time_visitor (baseline Model C default profile)

- city_snippet_family_planner

- city_snippet_comfort_easy_pace_traveler

- city_snippet_solo_social_traveler

- city_snippet_interest_deep_dive_traveler

- city_snippet_active_adventure_traveler

DEE snippet selection rule (MVP):

1) If an audience profile is inferred with sufficient confidence → show that audience snippet.

2) Else → show city_snippet_default (fallback).

4) Basic quick facts (optional strip)

- Country/region, best months (range), typical trip length (“Best for 2–4 days”), etc.

- Stored as factual city metadata.

5) Social share bar

- Share to WhatsApp, Email, Copy link (platform list configurable).

- Configured via /cms/social-sharing.

B) Optional micro-carousel: “Top picks for you in {{City}}”

Ownership: DEE + selection outputs (precomputed).

Purpose: a short personalized “spark” before bundles.

Rules:

- 4–8 cards (carousel desktop; 2×2 grid mobile).

- Cards can include: products, bundles, things-to-do entry, top attraction tour.

- Must be selected from precomputed allowed outputs only (DEE does not invent items).

- If not needed, this block may be removed without affecting Model C.

C) Bundles (1/2/3-day itineraries) — 3 cards (HARD)

Ownership: Bundle builder + DEE substitution (precomputed bundle variants).

Placement: after micro-carousel (or directly after hero if micro-carousel is disabled).

Hard rules:

- Always show exactly 3 cards: 1 day / 2 days / 3 days.

- Bundles are duration-based; no explosion into separate bundles for every segment combination.

- If DEE profile exists and profile-aware bundles exist, show those; else show default.

- Bundle feasibility enforced (Time Geometry + Difficulty + timed caps). Parameters are configured in Product Ops.

Card content:

- Title + subtitle

- Hero thumbnail

- 3 bullet highlights

- Trip DNA micro-strip (optional)

- CTA: “Preview itinerary”

On click:

- Opens a modal preview (no URL change).

- Preview includes: “Why these picks”, Anchor/Complement/Flex summary, and CTA:

“Check availability & adjust times”

- CTA navigates to:

/{lang}/{country}/{city}/itineraries/{segment-}{duration}-day/

D) Top Picks (Hero picks) — 6 products (HARD)

Ownership: Model C Set output + DEE ordering within allowed outputs.

Placement: after Bundles.

Purpose: safest default winners across archetypes.

Hard rules:

- Exactly 6 cards.

- Must be Set-eligible (INSTANT only).

- Enforce anti-duplicates:

- R1 no product twice on the page

- R2 avoid repeating the same attraction product across Top Picks

Card fields:

- Image, title, short snippet (≤120 chars), rating, review count, runtime price, CTA “View”.

DEE behavior:

- DEE may reorder within Set outputs based on profile.

- DEE must not introduce products outside Set outputs.

E) Private Upgrade (CTA-only) — max 2 CTAs (HARD)

Placement: after Top Picks.

Purpose: private is an upgrade path, not a catalog.

Hard rules:

- No list of products.

- Max 2 CTAs:

1) “Upgrade to Private”

2) “See all private tours”

- “See all private tours” opens Phase B overlay filtered to private (no indexable URLs).

F) Category Rails (Decision Rails) — max 4 rails × 2 products (HARD)

Definition: Rail (Decision Category)

A Rail is a Rosotravel decision category (not a provider category).

A rail groups one or more archetype cores to produce a compact shortlist and help users decide without catalog overload.

Examples (Iconic/Culture city type):

- City Highlights / Orientation

- Top Attraction Tour (guided experience)

- Museum & Culture (Ticket-first)

- Food & Tastings

Examples (Water/Cruise city type):

- Water Day (Primary)

- Water Day (Secondary)

- City Highlights / Orientation

- Food & Tastings or Night Views

Definition: Ticket-first rail

A “Ticket-first” rail is a decision rail where ticket/pass options are preferred as the primary surface.

Guided tours may exist as premium/upgrade winners but must not contaminate ticket-only sections.

Definition: Rail outputs

Each rail may show up to two winners:

- Best Value

- Premium

plus a CTA: “Choose by intent”.

Full value tier model (internal)

Products are categorized internally into:

- best_value_for_money

- balanced_core

- premium

- niche

Why the UI shows only two picks:

City Page must remain decisive. The system still uses a full tier model internally, but the City Page surfaces at most two choices per rail.

Rail display mapping (two visible slots)

- “Best Value” slot prefers best_value_for_money; may fall back to balanced_core.

- “Premium” slot prefers premium; may fall back to balanced_core.

- niche is not forced into the two slots; niche stays primarily for Phase B and long-tail discovery unless it is the only eligible option.

Hard rules:

- Max 4 rails on City Page.

- Two products only if materially different (“2-of-many must be real”, R6). If near-duplicate → show 1 pick + CTA.

- Rails must not repeat products already used in Top Picks (R3). If needed, select the next eligible product within the rail/archetype.

- Rails use Set-eligible inventory only (INSTANT).

F1) Inline teaser (mini CTA) — City Trip Planner (after Rails)

Purpose: provide an immediate “safety net” for users who are not ready to decide.

UI (compact, one-line + button):

- Text: “Not sure what fits your trip? Get a 1–3 day plan by email.”

- Button: “Plan my trip”

Behavior:

- Opens the Trip Planner modal (or scrolls to the full trip planner section below).

Hard rule:

- This teaser must be lightweight and must not expand the page height significantly.

G) Top attraction tours in {{City}} — 3–6 product cards (HARD)

Purpose:

Show curated tours/experiences tied to the city's top attractions, without duplicating Top Picks or Rails.

This is valuable because different audiences may want different attraction-led experiences.

Hard rules:

- Show 3–6 product cards (range prevents forced duplication).

- Must not repeat products already shown above (R1).

- Prefer one representative product per attraction (avoid multiple products for the same attraction).

- Must be Set-eligible (INSTANT).

Top attractions definition (TBD model; discovery-required)

We acknowledge that "top attractions" is not trivial:

- High review volume does not always mean "ideal for tours".

- Some attractions are primarily ticket-demand, not tours.

v1 approach:

- Use attraction feed aggregates (review_count + average_rating) as a base rank.

- Apply "tour suitability" filter:

- attraction qualifies only if it has Set-eligible products in relevant rails (Top Attraction Tour or Museum & Culture ticket-first).

- Choose top N attractions and pick 1 best representative product per attraction using Model C selection outputs.

The final "top attraction model" must be finalized during discovery and documented as a governed rule set.

H) Mini "Explore by category" chips (compact; max 8)

Purpose:

Provide quick internal links for browse-first users without converting the page into a catalog.

This also supports SEO internal linking and crawl discoverability.

Rules:

- Max 8 chips; no descriptions.

- Chips link to:

/{lang}/{country}/{city}/{category-slug}/

- DEE may reorder chips based on profile.

- This mini block is separate from the full categories section below.

I) Top Tickets in {{City}} — 4–6 items (HARD)

Purpose:

Provide a clean ticket surface (ticket-only / city pass / city card) without mixing guided tours.

Hard rules:

- Accept only products classified as: TICKET or CITY_PASS_CITY_CARD.

- Guided tours must never appear here.

- Must not repeat any product already shown above (R1).

- Must be Set-eligible (INSTANT).

J) Phase B entry — “Show more options” (HARD)

Placement: end of decision-first stack.

Behavior:

- Opens Phase B overlay (desktop) / full-screen overlay (mobile).

- Does not create indexable URLs; canonical remains the City Page URL.

2.3.3 Phase B — “More Options” UX (City Hub + Travel Agent) — FINAL v3

B2C Phase B (public City Page)

- As previously defined: Magic Finder (Topic/POI) + Browse Curated Catalog.

- Standard floors: rating ≥ 4.5

- On-your-request floors: Tours ≥ 4.2, Tickets ≥ 3.9

- POI PACK must be one-screen compact.

- No batch triggers; no indexable URLs.

B2B Phase B (Travel Agent logged-in variant)

Travel Agent Phase B is a role-gated variant of Phase B shown only when user is authenticated as Travel Agent (see Chapter VII Partner Programs).

Requirements:

- Discounts displayed prominently:

- Rosotravel Originals: X%

- Partner options: Y%

- Premium-first order:

1) 3 curated packages

2) Originals (group then private; highlighted)

3) Curated partner options (categorized)

- Expanded catalog limits by band:

0–299: up to 150

300–599: up to 250

600+: up to 400

- On-your-request access with warning; floors:

Tours ≥ 4.2, Tickets ≥ 3.9

Hard rules:

- No SEO/AI batch triggers from B2B views.
- No keyword retrieval or near-page creation from B2B browsing.
- No indexable URLs created.

2.3.4 Dynamic Experience Engine (DEE) — City Page Behavior (clarified)

DEE does NOT generate text in real time.

DEE selects from prewritten variants and precomputed sets.

DEE can:

- choose which precomputed snippet variant to show
- choose which precomputed Set variant to render (within allowed Set outputs)
- reorder within allowed outputs

DEE must NOT:

- create indexable segment URLs
- introduce products outside Set outputs

Snippet variants (prewritten; stored in CMS):

- city_snippet_default
- city_snippet_first_time_visitor
- city_snippet_family_planner
- city_snippet_comfort_easy_pace_traveler
- city_snippet_solo_social_traveler
- city_snippet_interest_deep_dive_traveler
- city_snippet_active_adventure_traveler

2.3.5 Trust bridge + SEO/content zone (detailed)

K) “This week in {{City}} with Rosotravel” — Live Snapshot (trust bridge)

Placement: immediately after Phase B entry.

Ownership: factual metrics + AI summarization (summary text can be precomputed nightly).

Elements:

- Tours run (rolling 7 days)
- Travelers hosted (rolling 7 days)
- Average rating (rolling 7 days)
- Top guides this week (own tours; optional)
- “Guests loved this week” (2–3 bullets; summarization)

Optional:

- “Live from last tours” photo strip (4–8 photos from verified recent tour media)

Fallback:

- If low volume: use last 30 days and label “Last month in {{City}}”.

L) City-level reviews (aggregated; compact)

Elements:

- City average rating

- “Based on X verified reviews”

- 3–6 sample review cards (recency + helpfulness weighted)

M) Local tips

3–5 short tips (160–260 chars) about transport, safety, seasonality, etiquette, money.

N) Best neighborhoods

3–8 neighborhood tiles.

Rule: neighborhood entities must exist in DB; AI cannot invent new neighborhoods.

O) Things to do categories (full block)

A larger category grid/chips block (beyond the mini chips above).

DEE may reorder categories; links go to city category pages.

P) Blog / travel guides

List/carousel of city-tagged articles.

Supports pinning evergreen hero guide.

NEW SECTION TITLE:

2.X CITY PAGE — “About Rosotravel in {{City_Name}}”

Alternative UI labels allowed:

- “Why travelers choose Rosotravel in {{City_Name}}”

- “How Rosotravel works in {{City_Name}}”

Recommended canonical label:

“About Rosotravel in {{City_Name}}”

Purpose

This section explains:

- what Rosotravel is in the context of the city,

- how Rosotravel builds and presents its city offering,

- why the page shows a curated shortlist instead of a giant catalog,

- and why that helps the traveler make a better decision.

This section is especially important for:

- LLM brand grounding at city level,

- trust in destination-specific expertise,

- and making the decision-platform logic explicit on the city page.

Generation timing

This section is generated during:

- city-level “Run AI Batch Processing”.

This means:

- it belongs to the city-level batch output,

- it is created together with the main city page content and city-level SEO assets,

- it is part of city content versioning,

- and it can be manually edited / locked in CMS.

Ownership

- AI-generated first draft

- Human editable

- Lockable in CMS

- If locked, it cannot be overwritten by regeneration unless explicitly re-enabled

---

B1) City Reputation Strip (required sub-block)

The City Page section must begin with a compact trust strip.

Place it:

- at the top of “About Rosotravel in {{City_Name}}”,

- before the explanatory paragraphs,

- visibly inside the same section.

Required visible elements:

1. City rating line

2. City review-count line

3. “What travelers value most in {{City_Name}}” bullets (2–4)

4. Optional link to city-level review section / reviews page

Recommended visible wording examples:

- “Rosotravel is rated 4.8/5 by travelers across tours in {{City_Name}}.”

- “Based on 436 verified reviews linked to experiences in {{City_Name}}.”

- “Travelers in {{City_Name}} most often value: guide storytelling, smooth organization, and clear meeting instructions.”

City-level rating rules

The city-level visible rating must be based on:

- verified reviews only,

- products mapped to {{City_Name}},

- review aggregation jobs, not free-form AI estimation.

Recommended initial logic:

- city_brand_rating_value = weighted average rating across verified reviews linked to products mapped to {{City_Name}}

- city_brand_rating_count = number of verified reviews included in the city aggregate

- refresh frequency = nightly

- source of truth = review aggregation + product ↔ city mapping

Recommended visibility threshold

Show the numeric city rating strip only if:

- city_brand_rating_count >= 20 verified reviews

Else:

- hide the numeric score,

- optionally keep top themes if enough review signals exist,

- otherwise hide the strip.

“What travelers value most” (City Page)

This must be derived from clustered verified review themes at city scope.

Examples:

- smooth meeting points

- local storytelling

- good pacing for first-time visitors

- family-friendly flow

- easy ticket handling

- great city overview

- helpful guides

Rules:

- 2–4 bullets maximum

- must feel destination-specific

- must come from real review evidence, not generic brand copy

City-level schema requirement

Add structured data tied to the visible city rating strip.

Recommended implementation direction:

- expose the visible Rosotravel-in-{{City}} block in body content,

- attach machine-readable aggregate rating data referencing the same visible value,

- keep schema synchronized with the visible block.

Minimum schema fields to expose in city context:

- ratingValue

- reviewCount

- bestRating = 5

- worstRating = 1

Internal developer note:

“Aggregate based on verified reviews of products mapped to this city page.”

Important:

Do not implement schema-only stars without the visible page strip.

Do not show a city rating if the evidence is too weak.

Recommended CMS/system fields for City Page

- city_brand_expertise_enabled
- city_brand_expertise_title
- city_brand_expertise_body
- city_brand_expertise_short_intro
- city_brand_expertise_lock
- city_brand_expertise_version
- city_brand_expertise_generated_at
- city_brand_rating_enabled
- city_brand_rating_value
- city_brand_rating_count
- city_brand_top_review_themes[]
- city_brand_rating_updated_at
- city_brand_rating_schema_enabled

B2) City explanatory body (required)

After the City Reputation Strip, the city-level explanatory block must contain 4 short sub-parts inside one continuous section:

1. Rosotravel’s role in {{City_Name}}

Explain:

- that Rosotravel helps travelers choose the right way to experience this city,

- that this city page is a decision-first city hub,

- and that Rosotravel specializes in organizing the city into clearer choices (top picks, attractions, bundles, rails, etc.).

2. How Rosotravel builds its {{City_Name}} offering

Explain:

- that Rosotravel may offer proprietary experiences in the city,

- work with local guide expertise,

- and include selected partner inventory where it improves the overall city offer.

Important:

This must be phrased as an operating model, not as a legal/local office claim.

3. Why Rosotravel does not show an endless list in {{City_Name}}

Explain:

- that the city page is designed to reduce duplicate noise,

- that similar products are compared and organized under one logic,

- that the goal is to help travelers choose faster and with more confidence,

- and that more options still exist through deeper browse flows where needed.

4. What travelers gain from using Rosotravel in {{City_Name}}

Explain:

- clearer planning,

- more relevant options,

- stronger trust signals,

- curated bundles and routes,

- and better alignment with traveler type and trip style.

Length and content rules

Recommended total length:

- 350–700 words for City Page version

Recommended paragraph count:

- 3–5 paragraphs

Optional:

- 3–4 bullets if needed, but default should remain paragraph-based

Tone

- city-specific

- expert but natural

- local-feeling without false local claims

- premium and useful

- not generic catalog copy

- not over-technical

What this section MUST communicate

- Rosotravel understands the city beyond “list of products”

- Rosotravel compares and organizes options under one decision logic

- Rosotravel may combine proprietary tours, guide expertise, and selected supplier inventory

- Rosotravel’s role is to help travelers choose the right city experiences, not browse duplicates

- Rosotravel is a trusted city-level expert even without claiming a physical branch there

- Rosotravel is well-rated by travelers in this city

- Travelers consistently value specific strengths of Rosotravel in this city

What this section MUST NOT claim

Do NOT claim unless true:

- local office / branch / physical point in the city

- exact city staff structure

- that Rosotravel directly operates every listed product

- that every guide shown is guaranteed on every departure

- any unverifiable city leadership claim

Allowed wording examples

- “Rosotravel helps travelers choose the right experiences in {{City_Name}} by organizing the city into clearer decision paths rather than endless duplicate listings.”

- “In {{City_Name}}, Rosotravel combines its own experience standards, guide expertise, and selected partner products to create a more useful city-level shortlist.”

- “This page is designed to help travelers compare the right kinds of experiences in {{City_Name}}, not browse every possible variation.”

Forbidden wording examples

- “Rosotravel is based in {{City_Name}}” (unless true)

- “Our {{City_Name}} office” (unless true)

- “We personally operate all tours in {{City_Name}}” (unless true)

- “We are the biggest tour company in {{City_Name}}” (unless proven and approved)

Notes for generation quality

The city-level block must be significantly more specific than the country version.

It must adapt to:

- city type (iconic / water / adrenaline / food-led / nature / sparse)
- dominant attraction logic
- typical bundle structure
- strength of own products vs partner-heavy logic
- traveler patterns in that city
- key themes that matter for choosing correctly in that city

Minimum variation sources for generation

Use at least:

- city name
- city type
- key attractions / major experience categories
- product mix and portfolio shape
- own vs partner mix
- strongest traveler intents relevant to the city
- proof/trust context if available
- city rating value and top review themes

Q) FAQ (city-level; exactly 9)

AI + manual reviewed.

FAQPage schema embedded.

R) City trip planner hook (full block)

Purpose: lead capture and return path for users not ready to buy.

Placement: below content blocks (as full form), plus inline teaser after Rails.

Form fields:

- Arrival date (optional)

- Number of days (1–7)

- Interests (multi-select)

- Budget level (low/medium/high)

- Email (required)

- Optional: traveling with children + ages

Rule:

AI can recommend only existing tours, attractions, and bundles (no invented products).

Store requests in user_trip_plans.

S) Comments & Ask our expert

City-level comment thread + “Ask our city expert” form.

Moderation via CMS..

THINGS-TO-DO PAGE TEMPLATE (City-level)

URL
/{lang}/{country}/{city}/things-to-do/
Example: /en/austria/vienna/things-to-do/

Intent
Broad discovery page for “things to do in {{City_Name}}” queries.
This page is browse-first but curated (Model C). It helps users discover categories, top experiences, attraction-led options, and curated shortlists — without becoming an OTA-style catalog.

Primary goals:
- Provide curated discovery across experience types (wider coverage than City Page).
- Keep decision outputs consistent with Model C (winners per archetype, no duplicates).
- Offer bundles as an optional “plan-first” shortcut.
- Provide controlled expansion only after explicit user intent (Phase B; no indexable URLs).

- Maintain strong SEO authority via structured content, internal linking, and FAQs.

Primary schema types
- TouristDestination (primary)
- ItemList (top picks, bundles, rails, top attraction tours, top tickets, top activities feed)
- FAQPage

2.4.1 Sections (Front-End, Desktop Order — FINAL)

Browse-first curated stack (main stake)
1) Hero + H1 + Dynamic snippet (DEE-selected, prewritten)
2) Quick jump links (anchors)
3) Top Picks (Hero picks) — 6 products (browse-first ordering)
4) Bundles (1/2/3-day itineraries) — 3 cards
5) Category Rails (Decision Rails) — 6 rails × 2 products
6) Private Upgrade (CTA-only) — max 2 CTAs
7) Top attraction tours in {{City}} — 3–8 product cards (broader than City Page)
8) Explore {{City}} your way (Category grid / chips)
9) Top activities feed — band-sized curated product list
10) Beyond the tourist spots (optional; explorers-oriented)
11) Nearby cities (optional; geo relations)
12) Phase B entry: “Show more options” (opens overlay; no indexable URLs)

Trust + SEO content zone
13) Reviews (things-to-do aggregated)
14) Insider videos / creator content (optional)
15) Local tips (optional)
16) Best neighborhoods (optional)
17) FAQ (things-to-do in {{City}}) — exactly 7
18) Comments & Ask our expert (optional)

Mobile
Same order, stacked. Phase B opens as full-screen overlay.
Rails and attraction tours should use compact layouts to avoid excessive scrolling.

2.4.2 Hero + H1 (Ownership: AI + Human; DEE selects variants)

Hero image
- Source: CMS media library
- Field: ttd_hero_media_id

H1
- Base text: “Things to do in {{City_Name}}”
- Editable in CMS: ttd_h1

Dynamic snippet (summary; 180–260 characters)
Important clarification:
- The baseline experience profile for Model C is first_time_visitor.

- “ttd_snippet_default” is a technical fallback used only when the system cannot reliably assign an audience profile.

Snippet variants (prewritten; stored in CMS):
- ttd_snippet_default (fallback when no audience can be inferred)
- ttd_snippet_first_time_visitor (baseline Model C default profile)
- ttd_snippet_family_planner
- ttd_snippet_comfort_easy_pace_traveler
- ttd_snippet_solo_social_traveler
- ttd_snippet_interest_deep_dive_traveler
- ttd_snippet_active_adventure_traveler

DEE selection rule (MVP):
1) If an audience profile is inferred with sufficient confidence → show that audience snippet.
2) Else → show ttd_snippet_default (fallback).

2.4.3 Quick jump links (anchors)

Purpose:
Help browse-first users jump to relevant sections without scanning the full page.

Anchors (recommended):
- Top picks
- Bundles
- Rails (Explore by style)
- Top attraction tours
- Categories
- Top activities
- Reviews
- FAQ

CMS field:
ttd_jump_links[] (anchor label + anchor id).
Generated automatically once; editor can reorder or hide.

2.4.4 Top Picks (Hero picks) — 6 products (browse-first)

Ownership:
Model C selection engine output (Things-to-do Set) + DEE ordering within allowed outputs.

Purpose:
Give users an immediate shortlist of high-confidence options.

Hard rules:
- Exactly 6 product cards.
- Must be Set-eligible (INSTANT only).
- No duplicates on page (anti-duplicate rules).

- Compared to City Page, ordering favors “browse diversity” (more variety in experience types).

Card fields:
- Image, title, snippet (≤120 chars), rating, review count, runtime price, CTA.

DEE behavior:
DEE may reorder Top Picks based on audience/intents, but must not introduce products outside the Things-to-do Set outputs.

2.4.5 City Bundles (“Ready-made 1–3 day plans”) — 3 cards

Ownership:
Bundle engine + DEE substitution (precomputed).

Location:
Directly under Top Picks.

Hard rules:
- Always show exactly three cards: 1-day, 2-day, 3-day.
- If profile-aware bundles exist, show them; else show default.
- Bundle feasibility enforced by bundle engine (Time Geometry + Difficulty caps).

Card content:
- Title (bundle.name), subtitle (bundle.subtitle)
- Mini hero image
- Trip DNA micro-strip
- 2–3 highlights
- Badge:
- AI bundle: “Curated with AI & local expertise”
- Human bundle: “Expert itinerary by {{curated_by_name}}”
- CTA: “View itinerary”

On click:
Open bundle preview modal (no URL change), then CTA to landing page:
- /{lang}/{country}/{city}/itineraries/{segment-}{duration}-day/

2.4.6 Category Rails (Decision Rails) — 6 rails × 2 products

Definition: Rail (Decision Category)
A Rail is a Rosotravel decision category (not a provider category). It groups one or more archetype cores to surface a compact shortlist.

Purpose on Things-to-do:
Rails provide broad discovery coverage while staying curated.

Hard rules:
- Exactly 6 rails on Things-to-do page.

- Each rail shows up to 2 winners (Best Value + Premium) only if materially different ("2-of-many must be real").
- If near-duplicate, show 1 pick + CTA.
- Rails must not repeat products already used in Top Picks; select next eligible within the rail.
- Rails use Set-eligible inventory only (INSTANT).

Rail examples (Iconic/Culture city type):
- City Highlights / Orientation
- Top Attraction Tour (guided experience)
- Museum & Culture (Ticket-first)
- Food & Tastings
- Day Trips / Nearby Must-See
- Night Experience / Views

Ticket-first rail clarification:
A "Ticket-first" rail prefers tickets/city cards as primary options. Guided tours may appear as premium/upgrade winners but must not contaminate ticket-only sections.

Value tiers (internal; full model)
Products are categorized internally into:
- best_value_for_money
- balanced_core
- premium
- niche
Things-to-do page still surfaces max 2 options per rail to remain curated.

2.4.7 Private Upgrade (CTA-only) — max 2 CTAs

Purpose:
Expose private as an upgrade path without turning Things-to-do into a private catalog.

Hard rules:
- No list of products.
- Max 2 CTAs:
1) "Upgrade to Private"
2) "See all private tours"
- "See all private tours" opens Phase B overlay filtered to private (no indexable URLs).

2.4.8 Top attraction tours in {{City_Name}} — band-sized attraction-led picks (TBD discovery

Purpose:
Attraction-led discovery for browse users. This section is broader than City Page (more cards allowed).

Use city band to determine the cap:

- Band A (0–299): show 3–4 items
- Band B (300–599): show 4–6 items

- Band C (600+): show 6 items (fixed cap)

Hard rules
- Must not repeat products already shown above (R1).
- Prefer one representative product per attraction (avoid multiple products for the same attraction).
- Must be Set-eligible (INSTANT).
- Selection must be stable (no random shuffle).
- If the city cannot produce enough eligible attraction-led products, show fewer items (do not fill with weaker products).

Top attractions definition (TBD; discovery required):
Use attraction graph aggregates as a base rank, but apply “tour suitability” filters (only attractions with eligible Set products in relevant rails).

Discovery requirement
Finalize the “Top Attractions ranking model” and “tour suitability” logic (some top attractions are not ideal tour anchors).
After approval, move the finalized selection model into Chapter 4.8 (portfolio choosing) as part of Set construction rules.

2.4.9 Explore {{City_Name}} your way (Category grid)

Ownership:
taxonomy-driven.

Elements:
Grid of category tiles (examples):
- Food & drink
- Museums & galleries
- Day trips
- Family-friendly
- Nightlife
- Hidden corners
- Boat cruises
- Walking tours

Each tile includes:
- Label + icon/image
- Short text (≤80 chars; AI prewritten)
- Link to:
/{lang}/{country}/{city}/{category-slug}/

DEE:
May reorder tiles based on audience profile (e.g., family shows family/day trips earlier).

2.4.10 Top activities feed — band-sized curated product list

Purpose
Provide a deeper browse list than City Page, still curated (not RAW catalog), tuned for "things to do" intent.

Status
The exact sizing and selection rules for this block are not fully defined in the current scope. This must be finalized during the 3-week discovery with engineering.
Until then, use the v1 proposal below.

V1 proposal (deterministic sizing by city band)
Use city band (based on RAW_count after validated attribution) to determine the cap:

- Band A (0–299): show up to 12 items
- Band B (300–599): show up to 16 items
- Band C (600+): show up to 20 items

Hard rules
- Items must come from Things-to-do curated outputs (Set-derived pool), not RAW.
- Must deduplicate against items already shown above on the page (Top Picks, Rails, Top attraction tours, Top Tickets).
- Must respect classification boundaries (do not mix ticket-only into guided-only surfaces incorrectly).
- Must be Set-eligible (INSTANT) for main stake placements on this page.

DEE behavior
DEE may reorder items within the allowed pool based on audience profile, but must not introduce products outside the allowed pool.

Discovery requirement (must be documented and then moved to Chapter 4.8 after approval)
Engineering must finalize:
- how the candidate pool is constructed (which Set outputs feed this list),
- how diversity across archetypes is enforced (to avoid 20 near-duplicates),
- exact band caps and any city-type-specific caps,
- how language coverage influences ranking in this list,
- how to prevent "randomness" (stable ordering + explainability).

2.4.11 Beyond the tourist spots (optional)

Ownership:
AI composition using curated entities only.

Elements:
3–7 tiles for less obvious ideas:
- lesser-known neighborhoods
- themed clusters ("Alternative art spaces", "Hidden courtyards")
Rules:
Entities must exist in curated database (no invented POIs).
DEE:

If audience = interest_deep_dive_traveler or explorers-like behavior, show higher or expand item count.

2.4.12 Nearby cities (optional)

Ownership:
geo relations.

Elements:
4–8 nearby city cards:
- Image, city name, travel time/distance, link to city page.

CMS:
city_nearby_cities[] computed from geo graph.

2.4.13 Phase B entry — “Show more options” (no duplication)

Phase B behavior for Things-to-do uses the same core rules as City Page Phase B.

Reference:
See 2.3.3 Phase B — “More Options” UX (City Hub + Travel Agent) — FINAL v3.

Notes for Things-to-do:
- Entry is the same (“Show more options” opens overlay; no indexable URLs).
- Mode 1: Magic Finder
- Mode 2: Browse Curated Catalog
- POI PACK rules apply when user selects POI.
- B2B Travel Agent variant applies only when authenticated as Travel Agent (Chapter VII).

2.4.14 Reviews (things-to-do aggregated)

Elements:
- “Based on X reviews across Y activities in {{City_Name}}”
- Average rating + distribution
- Review cards (same unified format)
DEE:
Only ordering may change slightly (e.g., family reviews earlier for family audience).

2.4.15 Insider videos / creator content (optional)

Ownership:
manual editorial.

Elements:
Carousel of short videos from social or own hosting:
- Thumbnail, title, creator name, city tag, “Watch” action.

CMS:

/cms/creators/videos/ filtered by city.

2.4.16 FAQ (things-to-do in {{City_Name}})

Ownership:
AI + manual review.

Rules:
- Exactly 7 FAQs.
- Focus: how to choose activities, what to book in advance, day structure expectations, safety, cancellations.
- Embed FAQPage schema.

CMS:
ttd_faq_items[] (question + answer + lock flag).

2.4.17 Comments & Ask our expert (optional)

Ownership:
UGC + expert moderation.
Simple comment thread + “Ask our expert” form.
Moderation via /cms/qa/.

2.4.18 Optional Hook — Experience finder quiz (recommended only if not duplicating Magic Finder)

Placement:
Below hero or as a floating bar.

Behavior:
6–8 questions about preferences (audience + intent signals).
At the end:
- Show recommended experiences (from curated outputs)
- Option: “Send my plan to email” (lead capture)

Important:
If Magic Finder is already implemented and performs the same role, this quiz should be treated as optional and can be postponed to avoid duplication.

Backend:
Config: /cms/quiz/
Lead storage: experience_quiz_results table.
Quiz answers reinforce DEE profile for the session.

Attraction page

URL

/{lang}/{country}/{city}/{attraction}/
Example: /en/austria/vienna/schonbrunn-palace/

Intent
Attraction Page is a hybrid:
- an entity hub and SEO target for "{{Attraction_Name}} + {{City}}" queries, AND
- a decision-first "2-of-many" shortlist for this specific attraction (Model C).

Primary goals:
- Answer "Is it worth it and how should I visit?" without OTA catalog overload.
- Surface the best 1–2 ways to experience this attraction (Best Value + Premium/Bucket-list if eligible).
- Provide a clean ticket-first surface when tickets are the primary purchase path.
- Provide controlled expansion only after explicit user intent ("See more options for this attraction").
- Maintain strong trust signals (live snapshot, verified reviews, factual visitor info).

Primary schema types
- TouristAttraction (primary)
- Place / GeoCoordinates / ImageObject
- ItemList (Our Picks, ticket block, related experiences)
- FAQPage
(Plus Product/Offer markup for selected picks, as defined in structured data section; see global schema rules.)

2.2.1 Sections (Front-End, Desktop Order — FINAL)

Decision-first stack (top)
1) Hero + H1 + Dynamic snippet (DEE-selected, prewritten)
2) Mini "Rosotravel Live at this attraction" strip (Phase D concept; trust)
3) Our 2 picks (2-of-many) — 1–2 cards (Best Value + Premium/Bucket-list)
4) Intent mini-rails (optional; only if eligible)
5) Private upgrade CTA (1 CTA; no list)
6) Tickets-only block (optional; if ticket-first applies)
7) "See more options for this attraction" (opens POI PACK overlay; simplified Phase B)

Entity hub content (SEO + practical)
8) Overview (contextual; no hard facts)
9) Top things to see (3–7)
10) Visitor information (factual-only)
11) History / significance (factual sources only)
12) Nearby attractions (4–8)

Commerce extensions (below)
13) Tours/activities related to this attraction (curated list)
14) Bundles including this attraction (optional; 1–3)
15) Reviews (attraction-level aggregated)
16) FAQ (attraction-level; exactly 7)

17) Comments & Ask our expert

Hook
18) Download visitor guide (PDF) — placed inside/under Visitor Information

Mobile
Same order, stacked. The decision-first stack must remain compact.
POI PACK overlay opens full-screen.

2.2.2 Hero + H1 + Dynamic snippet

Ownership: AI + factual, Human lockable.
DEE selects variants; DEE does not generate text in real time.

Hero image
- CMS: attraction_hero_media_id

H1
- Example: "{{Attraction_Name}} in {{City_Name}}"
- Editable: attraction_h1

Dynamic snippet (180–260 characters; prewritten)
Rule:
Factual statements (opening hours, ticket prices, address) must NOT be placed in the snippet.
Hard facts belong only in Visitor Information.

Snippet variants (prewritten; stored in CMS):
- attraction_snippet_default (fallback when no audience can be inferred)
- attraction_snippet_first_time_visitor (baseline Model C default profile)
- attraction_snippet_family_planner
- attraction_snippet_comfort_easy_pace_traveler
- attraction_snippet_solo_social_traveler
- attraction_snippet_interest_deep_dive_traveler
- attraction_snippet_active_adventure_traveler

DEE selection rule (MVP):
1) If an audience profile is inferred with sufficient confidence → show that audience snippet.
2) Else → show attraction_snippet_default.

2.2.3 Mini "Rosotravel Live at this attraction" strip (trust)

Purpose:
Provide a compact trust signal that Rosotravel actively operates or curates experiences around this attraction.

Elements (rolling 7 days):
- Tours run including this attraction

- Average rating from these tours
- “Guests loved:” 1–2 themes (summary; may be precomputed nightly)

Ownership:
aggregated data + AI summarization (summary text), no SEO impact (UX only).

Fallback:
If volume is low, show last 30 days and label “Last month”.

2.2.4 Our 2 picks (Model C — 2-of-many)

Purpose:
Provide the highest-confidence shortlist for this attraction.

Hard rules:
- Show up to 2 picks:
- Pick A: Best Value
- Pick B: Premium or Bucket-list (only if eligible and materially different)
- If only one option meets quality and eligibility → show 1 pick (do not fill with weaker products).
- Picks must come from POI-specific candidate sets (product_attraction_relations) and must be Set-eligible for main stake (INSTANT).

Card fields:
- Image, title, rating, review count, duration, runtime price, short snippet, CTA “Check availability”.

Anti-duplicate note:
These picks must not be near-duplicates. If near-duplicate, show only one pick + CTA “See more options”.

2.2.5 Intent mini-rails (optional; only if eligible)

Purpose:
Allow quick refinement for common intent needs without expanding into a catalog.

Behavior:
- Show intent chips and/or mini-rails only if eligible options exist for that intent.
- Max 3 intent chips shown by default (to stay compact).
Recommended intents:
- Family
- Less crowded
- Hidden gems (if applicable)
- Private upgrade (chip leads to private CTA)

Hard rules:
- No empty rails.
- No forced “2 picks” if only one materially different option exists.

2.2.6 Private upgrade (CTA-only)

Purpose:
Expose private as an upgrade path without listing many options.

Hard rules:
- Show 1 CTA only: “Upgrade to Private”
- Do not show a list of private tours here.
- CTA can open the POI PACK overlay filtered to private options.

2.2.7 Tickets-only block (optional; ticket-first)

Purpose:
When the attraction is primarily purchased as a ticket/city card, provide a clean ticket-first surface.

When to show:
- Show this block if the attraction has eligible ticket/city pass/city card products.
- Do not show if the attraction has no meaningful ticket-first options.

Hard rules:
- Ticket block may include only products classified as:
- TICKET
- CITY_PASS_CITY_CARD
- Guided tours (GUIDE) must never appear in tickets-only block.
- Must be INSTANT for main stake.

Display:
- Up to 2 ticket items (compact)
- If more exist, use “See more options for this attraction” for expansion.

2.2.8 “See more options for this attraction” (simplified Phase B)

Purpose:
Provide controlled expansion after explicit user intent.

Behavior:
- Button opens POI PACK overlay for this attraction.
- POI PACK rules, quality floors, and on-your-request separation follow the Phase B specification defined in City Page:
Reference: 2.3.3 Phase B — “More Options” UX (City Hub + Travel Agent).

Attraction page simplification:
- Attraction page uses POI mode only (no Topic mode by default).
- The overlay is designed to fit the POI PACK compact structure (one-screen constraint).

2.2.9 Overview (contextual)

Ownership: AI + human edits/locks.
1–2 paragraphs describing what the visitor experiences at this attraction.
Tone: “local guide speaking to a friend”.
Rules:
- No hard facts that must match APIs (hours, ticket prices, exact rules).
- Context only.

CMS:
attraction_overview_html

2.2.10 Top things to see (highlights)

Ownership: AI + factual sub-entities.
Elements:
3–7 list items:
- title (sub-place / highlight)
- 1–2 sentence description (≤220 chars)

Rules:
- Sub-entities must come from internal POI taxonomy and/or official sources.
- AI may describe; cannot invent elements that do not exist.

CMS:
attraction_highlights[] (title, description)

2.2.11 Visitor information (factual-only)

Ownership: factual only; AI may rephrase labels but not values.

Elements:
- Opening hours (normal + seasonal if available)
- Address
- Official website
- Ticket information (if any)
- Accessibility information (lift, steps, stroller, wheelchair)
- Dress code and rules (if applicable)

Sources:
- Google Places API
- Wikidata
- Official attraction website
- Manual override if conflicting/missing data

CMS structure:
attraction_visitor_info:
- opening_hours_raw
- address

- official_website
- ticket_info
- accessibility
- rules

2.2.12 History / significance (factual sources only)

Ownership: AI based on factual sources; human lockable.
Rules:
- 1–2 paragraphs with key dates and significance.
- Factual data must be sourced from:
- Wikidata
- Wikipedia structured summaries
- Official attraction website
AI cannot invent years or names; validation enforces that.

CMS:
attraction_history_html (lockable)

2.2.13 Nearby attractions

Ownership: geo graph + optional DEE reordering.
Elements:
4–8 attraction cards within defined radius:
- image, name, approx distance/time, link to attraction page

DEE:
May reorder based on audience (family-friendly first, etc.) but cannot invent new attractions.

CMS:
attraction_nearby_ids[] from geo engine.

2.2.14 Tours/activities related to this attraction (curated list)

Ownership: business logic + Model C ranking within POI context.

Source:
product_attraction_relations (tours/products where this attraction appears in itinerary mapping).

Display:
- Curated list of up to 6 cards (compact)
- Must avoid duplicating “Our 2 picks” products (no repeats on the same page)
- INSTANT only for main stake display

DEE:
May reorder within allowed set based on audience/intents.

Optional CTA:
“See more options for this attraction” (opens POI PACK overlay).

2.2.15 Bundles including this attraction (optional)

Ownership: bundles engine + DEE.
Placement: below related tours.

Elements:
- 1–3 bundle cards where contains_attraction_id = this attraction

Rules:
- If audience profile exists, show matching bundles; else show default.
- Modal preview + CTA to bundle landing page.

2.2.16 Reviews (attraction-level aggregated)

Ownership: aggregated from tours that include this attraction.

Elements:
- “Based on X verified reviews from tours and activities at {{Attraction_Name}}.”
- Average rating + distribution
- Review cards:
name/country, date, product name, rating, snippet, “See tour” link

DEE:
No text changes; only ordering may adjust slightly.

2.2.17 FAQ (attraction-level)

Ownership: AI + manual review.
Rules:
- Exactly 7 FAQs (global FAQ rules).
Topics:
visiting rules, best time, ticket policy, dress code, photography rules, lines/waiting time, accessibility.
Schema:
FAQPage

CMS:
attraction_faq_items[] (question, answer, lock flag)

2.2.18 Comments & Ask our expert

Ownership: user + expert moderation.
Comment thread for attraction-specific questions.
“Ask our expert” form moderated via /cms/qa/.

2.2.19 Hook — Download attraction visitor guide (PDF)

Placement:
Inside or directly under “Visitor information”.

Behavior:
Teaser:
“Want a printable {{Attraction_Name}} checklist and visiting tips? Get your free visitor guide.”
Email field + consent.

On submit:
- Save email + attraction id in lead_capture
- Send email with link to PDF (served via CDN)

CMS:
PDF manager /cms/content-assets/pdf-manager:
- pdf_id, title, attraction_id, language, file_url
Hook config in /cms/hooks/guide-download

DEE:
Hook is available for all audiences; no behavior changes in MVP.

End of Attraction Page template.

Product Page

URL format:
`/{lang}/{country}/{city}/{tour-slug}/`
Example: `/en/italy/rome/colosseum-guided-tour/`

Primary purpose:

- Convert motivated visitors into bookings.
- Surface human face (guides, recent guests).
- Show “Rosotravel Live — Last Tour Snapshot” as a core USP.
- Support Dynamic Experience Engine (segment-based snippets and highlights).
- Provide complete factual clarity on options, inclusions, itinerary and meeting point.

2.1.1 Sections (Front-End)

Order of sections on desktop - if it comes to UX schedule this is only proposition, we are open to see other possibilities of UX if it comes to design / placement on the website

1. Hero block
2. Booking panel (right column, sticky)
3. Rosotravel Live — Last Tour Snapshot
4. Options / Variants block (this section show options after choosing dat) example

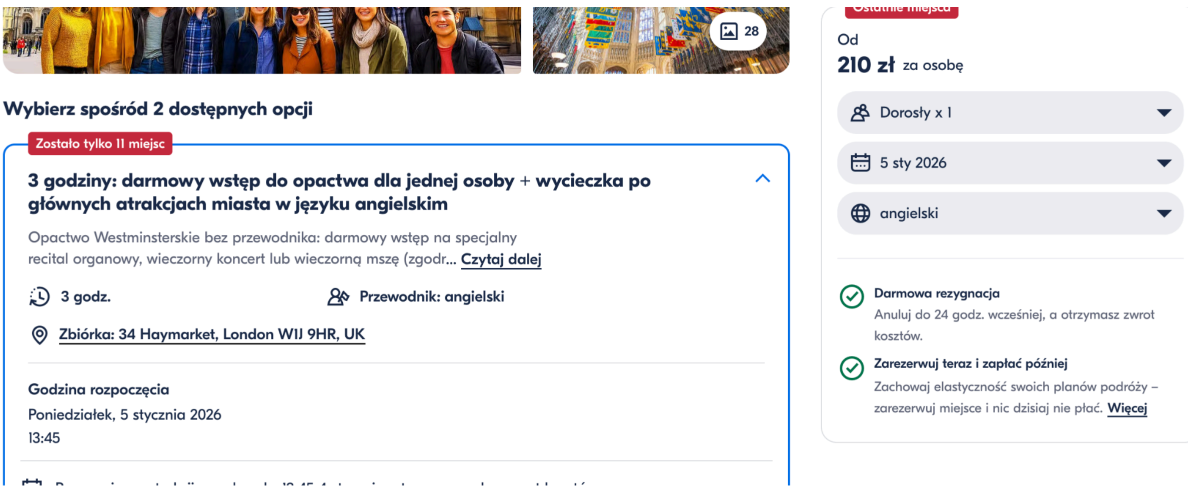

5. Long Descripton
6. Highlights
7. “Why travelers love this tour” (synthetic form reviews)
8. Included / Not included
9. Itinerary (full modal not changed by AI, non indexable)
10. Itinerary overviews (indexable)
11. Expert tip / Local insight
12. Meet your guides
13. Trip DNA (extended strip)
14. FAQ (product-level)
15. Reviews
16. Q&A (tour-specific)
17. Comments / “Ask our expert”

18. Related tours / “You may also like”

19. Hooks (lifetime code, wishlist, send to self, exit-intent)

Global live chat (tawk.to widget) appears in the bottom-right corner across all pages. It is configured once and is not part of this template.

Hero block

Ownership: factual + AI + manual.

Elements:

- Hero image carousel
    - Source: OTA feed or internal media library.
    - Managed in CMS; optimized via Cloudinary.
- Overlay avatar of last 5-star reviewer
    - Source: last review with 5-star rating for this product.
    - Data: reviewer name (or first name + initial), country flag, optional avatar photo; if photo missing, auto-generated avatar.
    - Placement: small circle avatar in bottom-left of hero image with small caption “Last 5-star review by [Name, Country]”.
- H1
    - AI-generated based on tokens: Tour_Name, Attraction_Name, City_Name.
    - Human can edit and lock in CMS.
- Dynamic snippet (in our product it is called short description)
    - Visible directly under H1.
    - Variants: snippet_first_time_visitor (default), snippet_family_traveler, snippet_couples_traveler, snippet_comfort_easy_pace_traveler, snippet_solo_social_traveler, snippet_interest_deep_dive_traveler, snippet_active_adventure_traveler
    - Selection: DEE intent mapping
    - Character limit: 180–260 chars
    - CMS: `/cms/content/product/snippets/`
    - Segment-based variant chosen by Dynamic Experience Engine (see 2.1.3).

    - Content: clear description of what the tour is, for whom, and why it is attractive.
- Rating summary
    - Aggregate rating and total review count (from reviews feed).
    - Short AI summary of themes (example: “Travelers love the small groups, smooth entry and passionate guides”).
    - Displayed as text under stars and review count.
- Trust badge
    - External rating logo (for example: Trustpilot) if configured.
    - Data comes from external trust system and is configured in CMS.
- Social share icons
    - Facebook, WhatsApp, Email, Copy link.
    - Share URLs and UTM templates configured in `/cms/social-sharing`.

Booking panel (right column, sticky)

Ownership: factual + UI.

Elements:

- “From” price (per adult or per main traveler type)
    - Source: product feed (current best price for default option).
- Note “Price may vary by date”
    - Static text configured in CMS.
- Date selector (calendar)
    - Standard availability calendar; uses same API as checkout.
- Travelers selector
    - Adults, children, other traveler types depending on product.

- Primary CTA
    - Label: “Check availability”.
    - Behaviour: scrolls to or opens availability options and sends request to availability API.
- Secondary CTA
    - Label: “Book now”.
    - Behaviour: direct step to booking if date and option already chosen.
- Quick facts
    - Duration
    - Language(s)
    - Free cancellation information
    - Instant confirmation badge (if available in feed)
- Scarcity / social proof banners (optional)
    - “Only X spots left for this date”
    - “Y people are viewing this tour now”
    - Values can come from real-time data or from configured range logic (for example: random range within [3–12]) with clear CMS setting for simulation.

Rosotravel Live — Last Tour Snapshot

- ➢ From ux perspective the short customer reviews should be shown in or near the option / booking box to strength the social proof

Ownership: Factual + AI summarisation
Fields:

- last_tour_date
- guide_name + guide_experience
- group_composition

- mood_score
- guests_loved[] (AI review clustering)
- photos_from_date (gallery filter)

CMS:
`/cms/reviews/last-tour-snapshot/`

Placement: directly under hero block, above Options / Variants section.

Purpose: show what actually happened on the most recent departure of this tour.

Display example:

Rosotravel Live — Last Tour Snapshot
• Last tour date: 12 June 2026, 10:00
• Guide: Anna K. (Blue Badge Guide, 8 years experience)
• Group: 12 travelers (6 couples, 2 solo)
• Mood: 4.9 / 5 (based on post-tour rating)
• Guests loved: “small group size”, “guide’s stories”, “smooth entry”
• [See photos and videos from this date] - it is very important that also movies can be added and are visible

Data logic:

- Last tour date
  - Derived from the most recent review’s travel_date for this product.
  - If no travel_date available, infer month from review date and display month-level (“Last tour in November 2026”).
- Guide
  - For own products: mapped via guide assignment table (tour_date → guide_id).
  - For OTA products without guide info: show “Rosotravel guide team”.
- Group
  - Group size and composition derived from booking data if available (adults, children).

- If not available: show only “Group size: X travelers”.
- Mood
    - Average rating for all reviews that share the same travel_date (or same week) as last tour.
- Guests loved
    - AI summarisation of top 2–3 themes from those reviews.
- See photos from this date
    - Link opens gallery filtered by review date equal to last_tour_date.
    - If no photos exist for that date, hide the link.

Fallback:

- If no recent tours in last 6 months:
    - Show text “Last tour: more than 6 months ago” and optionally “Low recent activity”.
    - Or fall back to last 30 days summary for MVP if you prefer not to show long inactivity.

Options / Variants block

Ownership: factual + AI-limited.

Location: under Rosotravel Live block, before long description sections.

Elements per option:

- Option name
    - Length ≤ 80 characters (global rule).
    - AI can propose based on duration, main feature and ticket type.
    - Human can edit and lock in CMS.
- Option description
    - Length ≤ 255 characters (global rule).

  - AI-generated from factual data; must clearly explain what is included and how it differs from other options.
  - Human can edit and lock.
- Duration
  - Source: product feed per option.
- Start time(s)
  - Feed-based list of start times or “flexible start time” text if applicable.
- Language(s)
  - Source: product feed.
- Ticket type
  - For example: “Skip-the-line entry”, “Private tour”, “Small group”.
- Option-specific inclusions and exclusions
  - Bullet lists; factual only.
- “View option details” button
  - Opens modal with option details: inclusions, exclusions, mini itinerary (if available) and practical notes.

Long Description

This is the main narrative description of the product page used by Google and AI systems to understand what the tour is, who it is for, and what the experience feels like.

Placement
Directly below “Options / Variants block”, before “Highlights”.

Ownership

- AI may generate the baseline, but the moment a human edits the block, the source becomes Manual.
- AI eligibility is controlled separately via lock and allow flags.

Content Rules

- 2–4 paragraphs (no bullet lists).

- Must explain:
    - what the experience is (in plain language)
    - the rhythm / pace (without timings)
    - who it is best for (optional final paragraph)
- Must NOT include: prices, exact times, opening hours, meeting point details, cancellation policy (those are factual blocks elsewhere).
- Must NOT invent facts. If a claim requires a data source, it belongs in factual sections, not here.
- Tone: “guide talking to a friend”, not OTA sales copy.

Highlights

Ownership: AI + manual.

Elements:

- 5 to 8 bullet points.
- Each bullet ≤ 85 characters.
- Each starts with a verb.
- Focus on core benefits (skip-the-line, expert guide, small group) rather than generic adjectives.
- Used both on product page and in product listing blocks on city and category pages.

“Why travelers love this tour”

Ownership: AI summarisation from reviews.

Elements:

- 3 to 4 short bullets.
- Derived from review clusters (guide quality, organization, value for money, atmosphere).
- No direct quotes; paraphrased summary.
- Example bullet: “Guests praise the guide’s storytelling and in-depth knowledge”.

Included / Not included

Ownership: factual (API / internal).

Elements:

- Two-column layout:
  - “What is included” list
  - “What is not included” list
- No AI rewriting.
- Only formatting of text for readability (capitalization, bullet formatting).

Itinerary

Elements:

- Full itinerary
  - Opens a modal with full itinerary text.
  - Modal contains all steps and details.
- For OTA products
  - Itinerary ingested from OTA feed as raw text, it is not indexable and non changeable
- For Rosotravel-owned products
  - AI can generate structured itinerary from internal data (meeting point, stops, durations), but exact times must not be invented.
  - Where times are not stable, use relative phrasing (“Morning visit to…”, “Afternoon walk through…”).
- Phase 1 rule
  - One global itinerary per product.
  - Phase 2: future not in MVP Per-option itineraries may be introduced in later phase; MVP uses common itinerary for all variants.

Itinerary Overview (indexable and created by AI for both OTA”s product and own product)

- Directly under itinerary, answer for a question, what you see, recommended for who, sightseeing rhythm, not include step by step and minutes
- sample

This tour is ideal for travelers who want a well-paced introduction to London's royal and political landmarks in a single day. The itinerary combines guided outdoor sightseeing with priority access to Westminster Abbey, balancing historical context, iconic views, and efficient routing. It is especially recommended for first-time visitors and those who value storytelling over rushed schedules.

Expert tip / Local insight

Ownership: AI + manual.

Elements:

- Short paragraph 250–350 characters.
- "Pro tip from our guide" style.
- Example: best time to arrive to avoid crowds, where to stand for the best view, how to use free time.
- AI proposes text; human can rewrite and lock.

Meet your guides

Ownership: manual for own tours, generic block for OTA feeds.

For Rosotravel-owned products:

- Carousel of guide profiles:
  - Guide photo
  - Guide name
  - Languages
  - Short bio (≤200 characters)
- Data sourced from `/cms/guides/` with relation table `product_id → guide_id`.

For OTA / Viator imports with no guide profiles:

- Generic guide block:
  - Static image representing guide team.

    - Short text about selection, training and quality standards for partners.

Trip DNA (extended strip)

Ownership: AI suggestion + human control.

Elements:

- Visual strip of 3–5 dimensions:
    - Walking level (low, medium, high)
    - Crowd level (calm, mixed, busy)
    - Vibe (local, touristy, mixed)
    - Best for (families, couples, solo, history lovers etc.)
- Rendered as icons with one-word labels and short text.
- Values stored in CMS and can be adjusted by marketing.

FAQ (product-level)

Ownership: AI + manual.

Elements:

- Exactly 10 FAQ entries (question + answer).
- Answers aligned with global FAQ rules and product factual data (inclusions, meeting point, cancellation policy, accessibility).
- Block placed after main description sections and before reviews.
- Embedded into FAQPage schema for structured data.

Reviews

Ownership: factual + unified UI.

Elements:

- Average rating and total review count.
- Rating breakdown chart (1–5 stars).

- Filters:
  - Sort by: Most recent, Highest rating, Lowest rating.
  - Filter by language.
  - Filter by rating bucket (5 stars, 4 stars etc.).
- Each review card includes:
  - Avatar (from user profile or auto-generated)
  - Name and country
  - Date of tour (travel_date)
  - Star rating
  - Short text snippet with “Read more” expander
  - User photos gallery (thumbnails; clicking opens lightbox)

Review sources:

- Own reviews and external reviews (for example Viator) normalized into unified format.
- Presentation is consistent regardless of source.

Q&A (tour-specific)

Ownership: user-generated questions + expert answers.

Elements:

- “Ask a question about this tour” form:
  - Question text
  - Email or login required for notification
- Display list:
  - Question text
  - Answer by Rosotravel expert or support team

    - Date answered

- No user-to-user answers in Phase 1 (only admin / expert can reply).
- All questions and answers moderated through `/cms/qa/`.

Comments / “Ask our expert” (general)

Ownership: user-generated + expert.

Elements:

- General comment thread for this page (separate from structured Q&A).
- Simple “Was this information helpful?” thumbs up / thumbs down widget.
- Expert responses can be tagged as “Expert reply” and visually highlighted.
- Moderation through CMS; spam protection required.

Related tours / “You may also like”

Ownership: AI + business logic.

Elements:

- Carousel or grid of 4–8 related tours.
- Selection logic: similarity of attraction, category, city, Trip DNA and price level.
- DEE can re-order suggestions per segment (family, couples etc.).

2.1.2 CMS Fields And Ownership

Ownership legend:

- AI: AI generates initial value; human can edit and lock.
- Human: created and maintained manually.
- System: computed or imported, not edited directly.

Core identification:

- `product_id` (System)
- `slug` (System)
- `language` (System)
- `source_type` (System: internal, external)

SEO:

- `h1` (AI, Human)
- `meta_title` (AI, Human, character rules from Chapter 1)
- `meta_description` (AI, Human)
- `canonical_url` (System)

Short description - Dynamic snippets (for DEE):

- `snippet_first_time_visitor, (it is default audience)`
- `family_traveler,`
- `couple_traveler,`
- `comfort_easy_pace_traveler,`
- `solo_social_traveler,`
- `interest_deep_dive_traveler,`
- `active_adventure_traveler`
-

Long Description (Indexable main body)

- `product_long_description_html` (AI baseline → Manual on edit)
- `product_long_description_locked` (System flag)
- `product_long_description_ai_allowed` (System flag)

Highlights:

- `highlights[]` (AI, Human; 5–8 items)

Why travelers love this tour:

- `love_bullets[]` (AI summarisation, Human)

Guide / human face:

- `primary_guide_id` (Human, for own tours)
- `guide_team_enabled` (Human)
- `guide_team_title` (Human)
- `guide_team_description` (Human)

Trip DNA:

- `trip_dna_walking_level` (AI proposal, Human final)
- `trip_dna_crowd_level` (AI proposal, Human final)
- `trip_dna_vibe` (AI proposal, Human final)
- `trip_dna_best_for[]` (AI proposal, Human final)

Rosotravel Live — Last Tour Snapshot:

- `last_tour_date` (System)
- `last_tour_time` (System)
- `last_tour_guide_name` (System or derived from guide assignment)
- `last_tour_group_size` (System)
- `last_tour_group_composition` (System; text such as “6 couples, 2 solo”)
- `last_tour_mood_score` (System or AI from reviews)
- `last_tour_guests_loved[]` (AI, from clustering)

- `last_tour_photos[]` (System, list of asset IDs or URLs)

Options / variants:

Per option:

- `option_id` (System)
- `option_name` (AI, Human)
- `option_description` (AI, Human)
- `option_duration` (System)
- `option_languages[]` (System)
- `option_ticket_type` (System or Human label)
- `option_includes[]` (System)
- `option_excludes[]` (System)

Itinerary:

- `itinerary_html` (AI for own or System for otas)
- `itinerary_overview_html` (AI, generated from full itinerary)
- `itinerary_collapsed_by_default` (System flag)

Included / Not included:

- `included_items[]` (System or Human)
- `not_included_items[]` (System or Human)

Expert tip / Local insight:

- `expert_tip` (AI, Human)
- `local_insight` (AI, Human)

Reviews:

- `reviews_source` (System)
- `reviews_average` (System)
- `reviews_count` (System)

FAQ:

- `faq_items[]` (AI, Human; always 10 items)

Q&A:

- `qa_enabled` (Human)

Comments:

- `comments_enabled` (Human)

Related tours:

- `related_products_first_time_visitor[] (AI / System)`
-
- `related_products_family_traveler[]`
-
- `related_products_couple_traveler[]`
-
- `related_products_comfort_easy_pace_traveler[]`
-
- `related_products_solo_social_traveler[]`
-
- `related_products_interest_deep_dive_traveler[]`
-
- `related_products_active_adventure_traveler[]`

Hooks configuration per product (optional overrides; global defaults exist in hook engine):

- `product_hooks_disable_discount`

- `product_hooks_disable_exit_intent`

2.1.3 Dynamic Experience Engine (DEE) For Product Page

Segments supported: according to **4.8.9 Stage F — Profile model**

Inputs:

- Entry keyword intent.
- Selected filters (family-friendly, romantic, budget).
- Referrer (blog article, city page, category page).
- User location and language where available.

Dynamic elements:

1. Dynamic snippet
    - Uses snippet fields described above; front-end selects the matching snippet based on inferred segment.
    - If no segment, uses `snippet_default`.
2. Dynamic highlights add-on
    - Segment-specific highlights fields:

      `Segment-specific highlights fields:`

        - `dynamic_highlight_first_time_visitor_1, dynamic_highlight_first_time_visitor_2`
        -
        - `dynamic_highlight_family_traveler_1, dynamic_highlight_family_traveler_2`
        -
        - `dynamic_highlight_couple_traveler_1, dynamic_highlight_couple_traveler_2`
        -

- dynamic_highlight_comfort_easy_pace_traveler_1, dynamic_highlight_comfort_easy_pace_traveler_2
-
- dynamic_highlight_solo_social_traveler_1, dynamic_highlight_solo_social_traveler_2
-
- dynamic_highlight_interest_deep_dive_traveler_1, dynamic_highlight_interest_deep_dive_traveler_2
-
- dynamic_highlight_active_adventure_traveler_1, dynamic_highlight_active_adventure_traveler_2

- At render time:
  - Start with base `highlights[]`.
  - If segment-specific dynamic highlight items are non-empty, append up to 2 of them at the end of the list.
  - If empty, show base highlights only.

3. Dynamic “Why this experience” add-on
   - Segment-specific paragraph fields:
     - wwhy_addon_first_time_visitor
     -
     - why_addon_family_traveler
     -
     - why_addon_couple_traveler
     -
     - why_addon_comfort_easy_pace_traveler
     -
     - why_addon_solo_social_traveler
     -
     - why_addon_interest_deep_dive_traveler
     -
     - why_addon_active_adventure_traveler
   - At render time:
     - Show base “Why travelers love this tour” bullets.

            - If matching `why_addon_*` is present for the current segment, show it as one extra sentence under the bullets.

4. Dynamic related tours ordering
    - DEE chooses which related list to display based on segment.
    - If `related_products_{segment}` is non-empty, use it.
    - Otherwise fallback to `related_products_general`.

Non-dynamic content:

- Title, long description, itinerary, FAQ, inclusions and meeting point remain constant for SEO and factual stability.

2.1.4 Hooks On Product Page

Hook 1: Lifetime -15% discount code

- Placement: under booking panel or in right sidebar.
- Visibility: non-logged-in users only; hidden once user has claimed code.
- Content: short message explaining that user can lock a one-time lifetime discount code for their first booking.
- Behaviour:
    - Captures email.
    - Stores lead with campaign tag “lifetime_code_product_page”.
    - Sends email with code and explains where to apply it.
- Frequency control:
    - At most one impression per user per 7 days (cookie or user account based).
- Configuration:
    - Default text and code behaviour defined in `/cms/hooks/discount-code`.

Hook 2: Add to Wishlist → Email My Trip Plan

- Placement:
    - Heart icon near tour title in hero block.
    - Additional “Add to wishlist” link near booking panel.
- Behaviour:
    - Click on heart icon adds product to wishlist in the current session.
    - After user adds one or more products, system can show a modal:
        - “Want to save your shortlist? Enter your email and we will send you a private link.”
    - User can also click a “Share wishlist” link which generates a public link; shareable via WhatsApp, Email and copy link.
- Backend and CMS:
    - Global settings in `/cms/wishlist/settings` for:
        - Minimum number of items before asking for email.
        - Default email template.
        - Whether shareable links are enabled.
    - Social share copy and UTM patterns in `/cms/social-sharing`.

Hook 3: Send this tour to myself

- Placement: below the intro or in right sidebar under booking panel.
- Behaviour:
    - Simple email capture form.
    - Sends email containing:
        - Tour title and link.
        - Snippet.

        - Highlights.
        - Optional note about discount code if applicable.
- Backend and CMS:
    - Configured in `/cms/hooks/send_to_self` (template, subject, from-name).

Hook 4: Exit-intent lifetime code

- Placement: modal overlay.
- Trigger:
    - Mouse leaving window at top, or browser back action on mobile (where detectable).
- Behaviour:
    - Show once per 7 days:
        - “Before you go: lock your -15% Rosotravel lifetime code for your first booking. Enter your email to receive it.”
    - On success: same code flow as Hook 1.
- Configuration:
    - Enabled or disabled globally in `/cms/hooks/exit-intent`.
    - Per-product override time if needed.

2.1.5 Schema Types

Primary schema types rendered for product pages:

- TouristTrip (main entity)
- Offer (for each option / variant)
- Review and aggregateRating
- FAQPage

- Organization (Rosotravel and provider where relevant)

Schema content is generated from tokens and CMS fields, not free text AI, and respects the global Schema Mapping rules defined elsewhere in the specification.

➢ **HowTo Module for Product (UI + Schema)**

Purpose

Provide a short, step-by-step **"How to book and join this tour"** guide that:

- reduces booking friction,
- supports AEO/featured snippets,
- enables agents/bots to understand the booking flow **without inventing facts**.

Placement in Product Page Template

Place it **between**:

- Booking panel / options selection, and
- FAQ/Reviews sections
  (after "Included/Not included" and before "FAQ".)

Eligibility (when the module is allowed)

The **HowTo module is allowed only if** all conditions are true:

1. Page type = **Product** and page is **bookable** (CTA exists).
2. At least **Meeting_Point OR Redemption_Instructions OR Check-in_Instructions** exists as factual input.
3. The steps can be written **without referencing hidden itinerary details** (itinerary modal is non-indexable and must not be summarized as "fixed schedule").

If these conditions are not met → do not show the module and do not emit HowTo schema.

UI spec (exact structure)

Module title (fixed): "How to book and join this tour"
Content consists of:

- **3–6 steps** (hard limits)
- Each step has:
  - `step_title` (max **60 characters**)
  - `step_text` (max **220 characters**, ideally 120–200)
  - optional `step_duration_hint` (optional, display-only, e.g., "~2 min")
- Optional "What to bring" checklist:
  - 0–8 items, each max **60 characters**
  - must be based on factual tokens (e.g., voucher required, ID required, dress code), not invented.

Ownership & editing rules

- Default generation: **AI draft** (Marketing can edit).
- Each step and checklist item is **editable + lockable** in CMS.
- If any step is locked, AI cannot overwrite it unless unlocked.

Content constraints (to prevent unsafe/misleading HowTo)

Steps must **not**:

- promise specific time slots,
- restate full itinerary content,
- claim "skip-the-line" unless it's an included item token,
- claim accessibility/age constraints unless those tokens exist,
- mention external provider constraints unless allowed by policy.

Recommended step set (canonical baseline for products)

This is the default structure; steps may be omitted only if the factual input is missing.

1. **Choose a date and option**

    - Text must reference the booking UI: calendar → timeslot/variant.

2. **Review key details before booking**
    - Included / not included, duration, language, cancellation summary (must match existing modules).
3. **Confirm meeting point**
    - Uses `{{Meeting_Point}}` and (optional) `{{Place_Address}}`.
4. **Arrive and check in**
    - Uses voucher/ticket requirement token if present.
5. **During the tour** *(optional)*
    - One line guidance (e.g., “Follow your guide; keep your ticket/voucher handy.”)
6. **If plans change** *(optional)*
    - Uses cancellation policy text (must match cancellation block).

Data fields (for dev implementation)

Store as `howto_steps[]` object on Product record:

- `howto_enabled` (bool)
- `howto_steps[]`: `{title, text, duration_hint, locked}`
- `bring_list[]`: `{text, locked}`
- `howto_last_generated_at`
- `howto_last_edited_at`
- `Howto_version`

Product HowTo schema output (JSON-LD rules)

Emit **HowTo** as a node referenced by the WebPage or Product entity, containing:

- @type: HowTo

- name: “How it works: {{Entity_Name}}”
- inLanguage
- step: array of HowToStep with name and text (and optional url anchors)

Important: If you include a step link (url), it must point to anchors that exist in the DOM

NEAR-PAGE TEMPLATE (City-Level Category + Programmatic Decision Pages)

URL patterns

A) Category pages

/{lang}/{country}/{city}/{category}/

e.g. /en/italy/rome/day-trips/

B) Programmatic “keyword near pages”

/{lang}/{country}/{city}/{keyword-slug}/

e.g. /en/italy/rome/colosseum-underground-tours/

Intent

Transactional or mixed-intent SEO pages capturing long-tail keywords such as:

- best colosseum tours

- rome family activities

- vienna day trips

- london walking tours

These pages must convert into tours/tickets/day trips while explaining Rosotravel’s decision platform approach.

Primary schema

- ItemList

- FAQPage

(Additional schema types may apply depending on page type; see global schema rules.)

Indexing rules

Near pages are indexable SEO targets.

DEE variants must NOT create separate indexable URLs (no segment URLs).

Phase B overlay must not create indexable URLs (canonical remains the near page URL).

2.5.0 Decision Platform Messaging (mandatory, SSOT reference)

Every near page must include the global “Decision Platform Messaging & Trust Layer” components:

- Why shortlist (1–2 lines)

- Shortlisted X-of-Y proof (computed)

- Quality standard tooltip (computed; no hardcoded thresholds)

- “See all options” CTA leading to Phase B overlay

- Badges with tooltips (only when proof policy satisfied)

Reference:

GLOBAL UI / TRUST COMPONENTS — Decision Platform Messaging & Trust Layer (SSOT).

Do not hardcode copy. All text/labels/tooltips must be CMS-editable tokens.

2.5.1 Page Types (supported)

Type 1 — Category Near Page

A city category hub (e.g., Day trips, Walking tours, Food tours).

Type 2 — Keyword Near Page

A programmatic landing page driven by keyword_map retrieval/ownership (e.g., “colosseum-underground-tours”).

Both types use the same template but may differ in:

- H1 decision angle pattern,

- comparison criteria,

- bundle relevance,

- map relevance.

2.5.2 Sections (Front-End, Desktop Order — FINAL v0)

Decision-first stack

1) Hero + H1 (decision-first) + snippet (DEE-selected, prewritten)

2) Decision messaging strip (mandatory):

- Why shortlist

- Shortlisted X-of-Y (computed)

- Quality standard tooltip

- “See all options” CTA

3) Expert byline / signature (mandatory for decision authenticity)

4) Overview (2 short paragraphs; decision framing)

5) “How to choose” decision guide (criteria bullets) (mandatory)

6) Comparison section (Top 3–5 products) (recommended; toggle)

7) Top choices list (band-sized; curated) + “See all options” CTA

8) Bundles related to this category (optional)

9) Map / SRP summary (optional)

10) Related categories (taxonomy)

11) Reviews (near-page scoped) (optional v1; recommended later)

12) FAQ (near-page; exactly 7)

13) Comments & Ask our expert

14) Hook (email capture) (optional)

Mobile

Same order; comparison becomes swipeable cards; decision strip must remain visible early.

2.5.3 Hero + H1 (decision-first)

Ownership: AI + manual (lockable).

DEE selects snippet variants; DEE does not generate text in real time.

Hero image

- nearpage_hero_media_id (generated / curated; optimized)

H1 (decision-first; derived from keyword_map)

H1 must be driven by keyword_map primary keyword + city name, but should be phrased as a decision.

H1 patterns (examples; final allowlist must be defined in discovery):

A) “Best / shortlist intent”

- “Best {{Keyword}} in {{City}} — our short list”

- “Our top picks: {{Keyword}} in {{City}}”

B) “Family / comfort / constraints intent”

- “{{Keyword}} in {{City}} for families — what’s worth booking”

- “{{Keyword}} in {{City}} with low walking — easiest options”

C) “Decision / comparison intent”

- “{{Keyword}} in {{City}} — which option is right for you?”

- “Is {{Attraction}} worth it with kids? Better alternatives in {{City}}” (when applicable)

Important:

The system must never invent facts in H1. It must remain a framing headline, not a factual claim.

CMS field:

nearpage_h1 (generated via keyword_map; editable; lockable)

Dynamic snippet (180–260 chars; prewritten)

Snippet is UX-only, not meta description.

Snippets must be prewritten variants stored in CMS.

Snippet variants (aligned to Model C audiences; prewritten):

- nearpage_snippet_default (fallback when no audience can be inferred)

- nearpage_snippet_first_time_visitor (baseline)

- nearpage_snippet_family_planner

- nearpage_snippet_couples_planner

- nearpage_snippet_comfort_easy_pace_traveler

- nearpage_snippet_solo_social_traveler

- nearpage_snippet_interest_deep_dive_traveler

- nearpage_snippet_active_adventure_traveler

DEE selection rule (MVP):

1) If an audience profile is inferred with sufficient confidence → show that audience snippet.

2) Else → show nearpage_snippet_default.

2.5.4 Decision Messaging Strip (mandatory)

Purpose

Make the decision-platform model explicit and prevent “low inventory” perception.

Must include:

A) Why shortlist (CMS token)

B) Shortlisted X-of-Y proof line (computed)

C) Quality standard tooltip (computed thresholds; no hardcoding)

D) CTA: “See all options in this segment” (opens Phase B overlay)

Computation requirement (TBD discovery):

- Define how Y is computed:

eligible_count_after_gates_dedup for this page context (keyword cluster / category mapping).

- Log this computation for Explainability (“why shown X-of-Y”).

All copy must be CMS-editable tokens (SSOT).

2.5.5 Expert byline / signature (mandatory)

Purpose

Increase authenticity and trust: “this is curated by real people with a role”.

Display (example format):

“Reviewed by {{Expert_Name}} — {{Expert_Role}} at Rosotravel”

Where Expert_Role is one of:

- First-time Visitor Expert
- Family Traveler Expert
- Couple Traveler Expert
- Comfort & Easy Pace Expert
- Solo Social Expert
- Deep Dive / Culture Expert
- Active Adventure Expert

(or a controlled role set)

Rules:

- Expert must be a real profile entity in CMS (no invented names).
- Expert assignment can be:
- manual pin per near page (recommended for top markets),
- or auto-assigned by audience intent mapping (discovery-defined).
- The byline text must be CMS-editable tokens; names pulled from expert profile.

2.5.6 Overview (2 short paragraphs; decision framing)

Ownership: AI + manual lock.

CMS: nearpage_overview_html

Rules:

- 2 concise paragraphs:

- relevance for the city,

- value proposition,

- decision-making hints (what matters when choosing).

- No hard facts that must match APIs (hours, prices, etc.).

2.5.7 “How to choose” decision guide (mandatory)

Purpose

Convert “search intent” into a structured decision.

Elements:

- 3–6 bullets (stable; AI-generated but reviewable) covering criteria such as:

- duration (short vs half-day vs full-day)

- group size / private upgrade

- ticket-first vs guided (when relevant)

- walking level / accessibility

- best time of day

- language coverage (if relevant)

Rules:

- Must not include invented factual claims.

- Must be consistent with product data and Model C concepts.

CMS:

nearpage_decision_guide[] (title + short text; lockable)

2.5.8 Comparison Section (Top 3–5 products) (recommended; toggle)

Ownership: AI + listing engine

CMS: nearpage_comparison_block_enabled (toggle)

Purpose

Provide a fast “compare” view (the core conversion block for decision queries).

Content:

Table or comparison cards for top 3–5 products.

Recommended columns:

- Duration
- Group size / private
- Ticket-first vs guided
- “Best for” labels
- Cancellation summary
- “Why recommended” microline (optional)

DEE behavior:

Products remain the same; ordering may change per audience profile.

Do not introduce new products outside the candidate set for this near page.

Hard rules:

- Must not conflict with Top Picks already shown above.

- Must be explainable (reason codes available in Explainability).

2.5.9 Top Choices list (curated) + “See all options” CTA

Ownership: listing engine + DEE ordering within allowed outputs.

CMS: nearpage_top_list_limit (band-sized defaults; tunable)

Sizing (v1 proposal; finalize in discovery):

Use city band (RAW_count) to set caps:

- Band A (0–299): up to 10

- Band B (300–599): up to 15

- Band C (600+): up to 20

Rules:

- Products come from keyword cluster mapping (keyword_map → eligible products / archetypes).

- DEE may reorder within the allowed pool; never expand the pool.

- Must avoid duplicates with comparison block and earlier picks.

- Must be stable ordering (no randomization).

- Must include CTA at the end:

“See all options in this segment” → opens Phase B overlay filtered to this keyword/category context.

Note:

Exact pool construction and diversity enforcement (avoid 20 near-duplicates) must be finalized in discovery and then moved into Chapter 4.8 selection logic.

2.5.10 Bundles related to this category (optional)

Ownership: bundles engine + DEE.

Display only if category/near-page aligns with bundles (e.g., day trips, walking tours, family activities).

Up to 2–3 bundle cards:

- Title, duration, subtitle, Trip DNA strip, curated-by badge

- CTA → modal preview → bundle landing page

DEE:

If a profile-aware bundle exists, show it; else show default.

2.5.11 Map / SRP Summary (optional)

Ownership: factual + UI widget

CMS: nearpage_map_enabled (toggle)

Elements:

- Map widget with pins representing tours/attractions (based on product coordinates)

- Text: “See all {{keyword}} options on the map”

DEE:

No personalization in MVP.

Discovery requirement:

Define coordinate sources and fallbacks; ensure map does not mislead (only pin what is truly relevant).

2.5.12 Related categories

Ownership: taxonomy

CMS: nearpage_related_categories[]

Elements:

Tag-style links to related near pages or categories.

DEE:

Reordering only.

2.5.13 Reviews (optional v1; recommended later)

Ownership: aggregated from products in this near page context.

If implemented:

- show short “Based on X verified reviews”

- 3–6 cards

DEE:

Order only.

2.5.14 FAQ (near page)

Ownership: AI + manual

CMS: nearpage_faq_items[]

Rules:

- Exactly 7 FAQs (Global FAQ rules)

- Based on SEMrush + GSC cluster

- No duplication with product or city FAQs

Schema: FAQPage

2.5.15 Comments & Ask our expert

Standard block:

- comment thread

- “Ask our expert” form

- moderation in /cms/qa/

2.5.16 Hooks (optional) — Long-tail Intent Capture

Hook: Long-tail Intent Capture

Placement: under H1 or after overview

Behavior:

Email capture exchange for value:

“Get €10 travel credit for top {{keyword}} experiences — receive your code by email.”

CMS:

/cms/hooks/longtail-intent (copy, enable/disable, frequency)

DEE:

Hook always shown for long-tail pages; no personalization in MVP.

2.5.17 Phase B reference (no duplication)

Near pages must use the same Phase B overlay system as City Page.

Reference:

2.3.3 Phase B — “More Options” UX (City Hub + Travel Agent)

Important:

If user is authenticated as Travel Agent, Phase B overlay uses the Travel Agent variant (Chapter VII).

End of Near Page template (v0). Finalize in discovery.

Blog Page

URL:
`/{lang}/stories/{topic}/`

Intent:
Inspiration + content marketing.
Supports internal linking, brand authority, and SEO.

Primary schema:

- Article
- Person
- FAQPage

2.9.1 Sections

Header Block

Ownership: manual
Elements:

- Title (`blog_title`)
- Short description (`blog_subtitle`)
- Author name + profile link (internal)
- Publication date
- Last updated date (only if different from published)
- Estimated reading time
- Social share buttons (CMS-controlled text for UTM templates)

Hero Image

Ownership: manual
CMS: `blog_hero_media_id`

Table of Contents

Ownership: auto-generated

- Based on H2/H3 headings in article body
- Shown right sidebar or top-floating TOC

Article Body

Ownership: manual (editorial)
AI allowed only for grammar or style suggestions.

Elements:

- Paragraphs, images, headings
- Inline product suggestions:
    - Cards or text blocks linking to tours
- Inline attraction links

Rules:

- Must avoid keyword cannibalization
- Must follow structured internal linking rules

“How to Visit” Snippets (optional)

Ownership: AI + manual
Elements:

- Blocks for major attractions discussed in article
  - Attraction name
  - 2–3 sentence summary
  - CTA: “View tours” (deep link to near-page or attraction tour list). Product suggestions inside the article must use DEE ordering rules.

Author Bio

Ownership: manual
Fields:

- Photo
- Short biography (≤300 chars)
- Links to expert page

Related Articles

Ownership: editorial + AI recommendations
CMS: `blog_related_articles[]`

FAQ (article-specific)

Ownership: AI + manual
Rules:

- Exactly 5 FAQs
- Highly contextual; not to conflict with product or city FAQs

---

Comments & Ask our expert

Standard component
Moderated via `/cms/qa/`

2.9.2 Hooks

Hook: “Send this guide + map to my email”

Placement: sticky bar after 25–35% scroll

Behaviour:

- Email capture
- Sends:
  - Article link
  - Optional PDF map (if uploaded)
  - Recommended tours list

CMS: `/cms/hooks/blog-email-share`

Rosotravel Expert / Guide / Bloger Page

Hero + H1

Ownership: manual + AI snippet
Elements:

- Expert photo (`expert_profile_media_id`)
- H1 = “{{Expert_Name}}, {{Role}} in {{City_Name}}”
- Dynamic Snippet (AI 180–260 chars) describing expertise and travel style
  - intent consistent with City, Country, Region, Attraction, Category, Product.
- **Author name (clickable to profile)**
- **Published date**
- **Last updated date** (only if different from published)
- Social share buttons (CMS-controlled text for UTM templates)

Bio

Ownership: manual
CMS: `expert_bio_html`
2–3 paragraphs about the expert’s:

- background
- certification
- favorite places
- guiding philosophy

Tours Authored / Guided

Ownership: factual
CMS: relational data `expert_product_relations[]`

Elements:

- Product cards for all tours where the expert is:
  - lead guide
  - content author
  - curator

DEE:

- Ordering may reflect user segment (e.g., families see family-friendly tours first)

Articles Authored

Ownership: editorial
CMS: `expert_article_ids[]`

Displayed as article cards.

Short Q&A Block

Ownership: manual
CMS: `expert_qa_items[]`
3–5 hand-picked Q&A, e.g.:

- “What’s your favorite hidden spot?”
- “When is the best time to visit the Old Town?”

FAQ (expert)

Ownership: AI + manual
Rules:

- 5–7 FAQs
- Topics: guiding style, languages, accessibility, groups, certifications

Schema: FAQPage

Comments & Ask our expert

Standard block
User → CMS moderation → published answer

2.10.2 Hooks

Hook: “Subscribe to {{Expert_Name}}’s insider tips”

Placement: after Q&A block or sidebar

Behaviour:

- Email capture to receive monthly insights from this expert
- Lead stored in `expert_subscribers`
- Controlled via `/cms/hooks/expert-tips`

Package / Expert Bundle

2.8 PACKAGE / BUNDLE PAGE (CITY ITINERARIES)

2.8.0 System Overview

- A “bundle” is a curated 1-, 2- or 3-day itinerary in a given city.
- A bundle is only a combination of existing products (tours, tickets, activities).
    - Price = sum of prices of the selected variants of those products (no special bundle pricing).
- Every city should have at least three bundles for the general segment:
    - 1-day itinerary
    - 2-day itinerary
    - 3-day itinerary
- Bundles are created:
    - Automatically by AI for all cities (default, auto_generated = true).
    - Manually by human experts for priority cities (auto_generated = false, overrides AI versions).
- Bundle visibility:
    - City page: dedicated “Bundles / Itineraries” section near the top.
    - Attraction page: “Itineraries including {{Attraction_Name}}” section (if that attraction appears in a bundle).

    - Bundle has its own landing URL and behaves like a “multi-product” product page.

2.8.1 URL & Page Intent

- URL pattern (canonical):
    - /{lang}/{country}/{city}/{package-slug}/
    - For segmented itineraries: /{lang}/{country}/{city}/itineraries/{segment}-{duration}-day/ (e.g. /en/italy/rome/itineraries/family-2-days/).
- Intent: themed bundles / itineraries (“2 days in Rome”, “Family weekend in Vienna”).

2.8.2 Sections (Front-End Structure)

Each bundle landing page follows this structure.

1. Hero + H1
 Ownership: AI + manual

- H1 (name): e.g. “3 Days in Rome: Classic Highlights Package”.

Curator Block

- avatar (if human) or AI icon
- “Curated with AI & local expertise”
- or “Expert itinerary by {{name}}”
- Author name (clickable to profile)
- Published date
- Last updated date (only if different from published)
- Dynamic Snippet: 1–2 sentences summarising who the package is for and what you see.
- Hero image: main city/attraction collage or key product image.

2. Overview
 Ownership: AI + manual

- Short paragraph answering:
    - Who is this itinerary for (families, first-timers, couples, etc.)
    - What is the overall rhythm (easy / intense, central / mixed).

- Integrates Trip DNA tags visually (walking level, crowds, vibe, “best for”).

3. Included tours (ItemList block)
   Ownership: factual

- List of all tours/tickets in the bundle (data from products):
    - Product name
    - Thumbnail
    - Duration
    - Category (attraction, walking tour, food, day trip, etc.)
    - Base “from” price (from product feed)
- Clicking product name opens product details in new tab (optional).

4. Suggested day-by-day itinerary
   Ownership: AI based on included tours, editable in CMS

- Day 1 / Day 2 / Day 3 breakdown.
- For each day:
    - Day title (optional): e.g. “Classics in the historic centre”.
    - Short day description.
    - Ordered list of products assigned to that day.
    - Optional suggested_time text per product (informational only, does not affect availability).
- No hard schedule enforcement; text is descriptive, not binding.

5. Booking details panel (availability selection)
   Ownership: system + factual

- For each product in the bundle, the landing page embeds the existing availability selector component:
    - Select date (calendar)
    - Select timeslot / option (variants)
- Behaviour is equivalent to Tiqets “booking_details” step:
    - User chooses dates and times *before* going to cart.

- Availability is always fetched from the underlying product APIs.

- If a product is unavailable on the chosen date:

    - Show "Sold out on this day – choose another date/time".

    - User can change the date or optionally de-select that product (future phase).

6. Highlights of the package
   Ownership: AI + manual

- 3–7 short bullets summarising main benefits, e.g.:

    - "All must-see highlights in 2 easy days"

    - "Skip-the-line entries where possible"

    - "Balanced walking time and breaks"

7. FAQ (package)
   Ownership: AI + manual

- 7 FAQs focused on:

    - How flexible the plan is

    - Whether activities can be moved across days

    - Ticket validity windows

    - Meeting points, accessibility, refund policy for multi-day plans.

- Structured for FAQPage schema, respecting global FAQ rules.

8. Comments & Ask Our Expert
   Ownership: users + moderators

- Standard cross-site "Comments & Ask our expert" module:

    - Users can leave comments about the bundle.

    - Users can ask pre-booking questions ("Is this suitable for small kids?").

    - Admin or designated experts answer; answers are displayed on page.

9. Schema
   Ownership: system

- Primary: ItemList (for included tours) and TouristTrip (for the itinerary).
- Secondary: FAQPage.

2.8.3 Booking and Availability Flow (Bundle vs Cart)

- First customer clicks on the product card on the Product page or Attraction page and a pop-up with shortage is shown. Then the customer can click the button to show the availability and is redirected to the landing page of the bundle. Customer can also click X and browse normal tours
- The bundle product page should be unique ulr which contain the general description of the bundle (remember that the screenshot and link shows general idea but the front end and description should include all hooks, cta, communication, human face which we use) - [tiqets link example]()
- 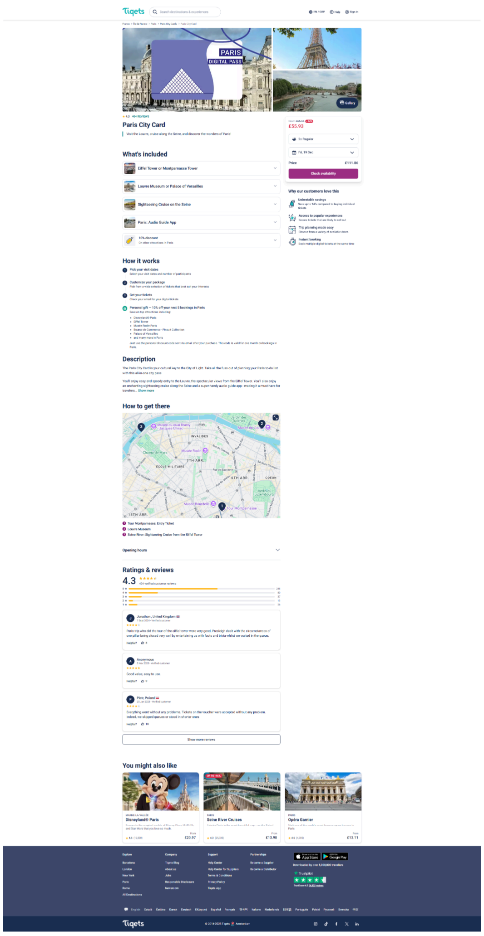

- Customer choose general availability and is redirected to availability / option page
- For each product in the bundle:
  - The user selects date and timeslot using the standard product availability UI (same component as on normal product pages), all required fields should be

possible to choose, date, hour, language, options, depends on the product feed

- [Tiqets example link]()
- 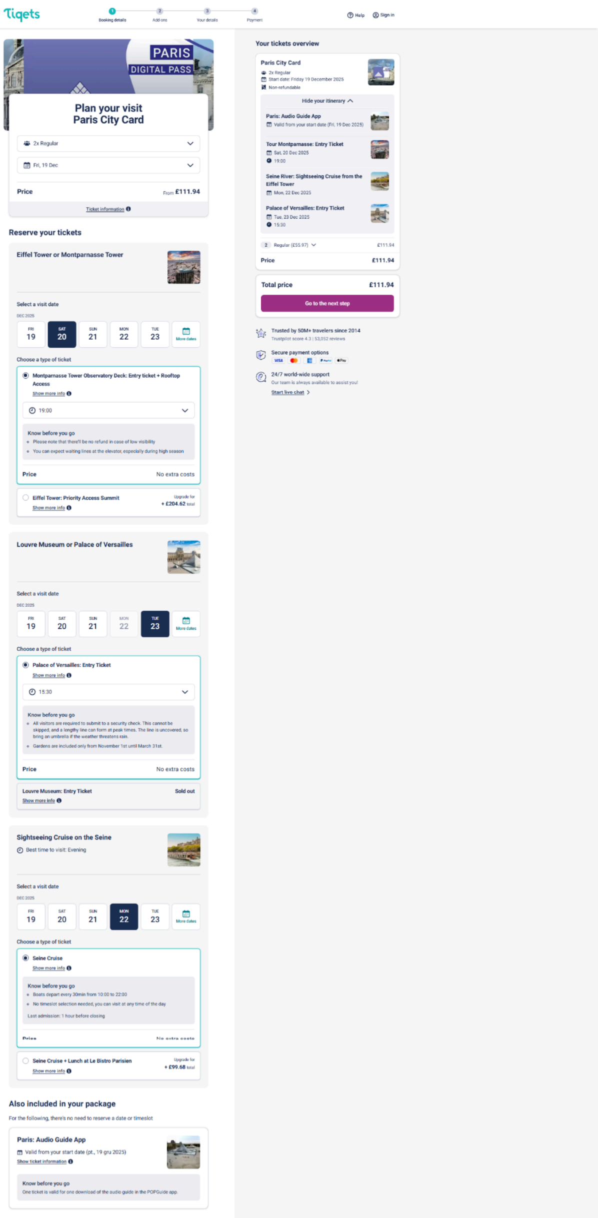

- After the user has selected dates/times for all required products, clicking the main CTA (“Go to next step”) does the following:
  - Adds each selected product to the cart as a normal line item (with its chosen variant/date/time).
  - Marks all those line items with a shared bundle_session_id and bundle_id for analytics and UX (e.g. grouping in cart).
    Redirects the user into the existing checkout flow.

- Price calculation:
  - Total bundle price displayed on the right of the landing page = sum of all selected variants.

    - No special discount logic in MVP (later we may add promo codes or bundle-specific discounts using existing promotion engine).

2.8.4 Bundle Entity — CMS Fields and Responsibilities

Bundle is a separate entity type (bundle_itinerary).

Identification / meta

- id (system)
- city_id (human – select existing city)
- language (system/human – e.g. en, de, pl)
- slug (system or human – unique URL, e.g. vienna-3-days-classic-itinerary)
- status (human) – draft | published | archived
- auto_generated (system) – true for AI bundles, false for expert bundles

Type / target

- duration_days (system for auto; human for manual) – 1 | 2 | 3
- target_segment (system default = general; human can override)
- bundle_type (system) – fixed_itinerary (MVP)

Branding and main copy

- name (AI for auto bundles, human may override)
- subtitle (AI or human)
- short_description (AI or human) – 2–3 sentences for cards and popup
- long_description (AI or human) – used in Overview and page body
- highlights[] (AI or human) – 3–6 items (title + optional description)

Curated-by / attribution

- curated_by_type (system/human) – ai_assisted | human_expert

- curated_by_name
    - AI bundles: default “Rosotravel Travel Intelligence”
    - Manual bundles: e.g. “Julia, local guide”
- curated_by_bio_short (AI for AI bundles, human for expert bundles)
- curation_reasoning_short (AI or human) – 1–2 sentences “Why these places”

Trip DNA and intent tags

- trip_dna.walking_level (AI or system; editable) – low | medium | high
- trip_dna.crowd_level (AI or system; editable) – calm | mixed | busy
- trip_dna.vibe (AI or system; editable) – local | touristy | mixed
- trip_dna.best_for[] (AI or system; editable) – families, first_timers, couples, general
- intent_tags[] (AI or system; editable) – e.g. ["family_friendly","history","food"]

Days and products

- days[]:
    - day_number (1, 2, 3)
    - title (optional)
    - description (optional)
    - products[]:
        - product_id (system for auto; chosen by human for manual)
        - order (display order for that day)
        - suggested_time (optional, free text; informational only)
        - note (optional extra comment to user)

SEO

- seo_title (AI or human)

- seo_description (AI or human)
- canonical_url (optional override)

System fields

- created_at, updated_at, generated_at (system)
- created_by_user_id (system for manual bundles)

Responsibilities

- Human editor:
    - chooses city, duration, segment
    - selects products per day (for manual bundles)
    - approves/edits name, descriptions, highlights, Trip DNA, “Why these places”
    - sets curated_by_type = human_expert and curated_by_name
    - publishes bundle
- AI generator:
    - drafts name, subtitle, descriptions, highlights, curation_reasoning_short, Trip DNA, suggested_time and notes
- Backend/system:
    - selects products for auto bundles
    - sets auto_generated = true, curated_by_type = ai_assisted
    - ensures manual bundles (auto_generated = false) are never overwritten

2.8.5 Front-End on City and Attraction Pages (Cards + Popup)

City page

- Placement: dedicated “Curated Itineraries” section near the top, above normal tours.
- It is presented as a longer product card but normal tours are visible below at the same screen
- By default show 3 cards in a row (1, 2, 3-day bundles for the relevant segment).

- Each card displays:

    - name
    - subtitle
    - mini hero image
    - 2–3 highlights
    - Trip DNA micro-strip (icons for walking/crowds/best-for)
    - “Curated with AI & local expertise” or “Expert itinerary by {{name}}” badge

- Clicking a card opens a modal popup, not a new page.

Popup content

- Title = name
- Subtitle and curated-by badge
- Hero thumbnail
- Short_description
- 3 highlights
- Trip DNA strip
- “Why these places” paragraph (curation_reasoning_short, plus optional global exclusion_reason)
- Main CTA button: “Check availability and adjust times” → navigate to bundle landing URL
- Small note under CTA: “You’ll choose dates and times for each experience on the next step.”
- X icon to close and return to city page.

Attraction page

- Optional smaller section: “Itineraries including {{Attraction_Name}}”.
- Same card and popup pattern; only display bundles that contain at least one product at that attraction.
- Packages here are the separate section below the normal attraction products, customer should mainly see the single products

2.8.6 Backend Auto-Generation Logic (General Segment, MVP)

Inputs per city

- Each product has:

    - category: attraction | activity | walking_tour | food | day_trip

    - rating_average

    - reviews_count
    - popularity_score (or similar)
    - optional tags (family_friendly, romantic, etc.)

Scoring

- score = rating_average * w1 + log(reviews_count + 1) * w2 + popularity_score * w3
- Construct sorted lists:
    - top_attractions, top_activities, top_walking, top_food, top_day_trip

Composition rules

1-day bundle

- A1 = first(top_attractions)
- B1 = first(top_activities not equal A1). If none, use second attraction.
- days[1].products = [A1, B1]

2-day bundle

- Day 1: A2 = next attraction; B2 = next activity
- Day 2: C2 = top walking tour; D2 = top food experience (with fallbacks)
- days[1].products = [A2, B2]
- days[2].products = [C2, D2]

3-day bundle

- Day 1: A3 = top unused attraction; B3 = top unused activity
- Day 2: C3 = top unused walking tour; D3 = top unused food experience
- Day 3: E3 = top day trip (or similar full-day product)
- days[1].products = [A3, B3]

- days[2].products = [C3, D3]
- days[3].products = [E3]

AI content generation

- After selecting products per day, backend calls the AI Drafting Layer with entity_type = "bundle" and passes:
    - city metadata
    - selected products per day
    - preliminary Trip DNA metrics
    - target_segment = general (MVP)
- AI returns: name, subtitle, short_description, long_description, highlights[], curation_reasoning_short, Trip DNA.
- Backend writes fields into the bundle record and marks:
    - auto_generated = true
    - curated_by_type = ai_assisted
    - curated_by_name = "Rosotravel Travel Intelligence"
- If a manual bundle exists for (city_id, duration_days, target_segment) with auto_generated = false, skip generation for that slot (never overwrite manual).

2.8.7 Brand Copy for “Curated by” and “Why These Places”

AI bundles

- Card badge: “Curated with AI & local expertise”.
- Landing intro component:
    - “This itinerary was created by Rosotravel Travel Intelligence – our AI system that looks at ratings, distances and traveller reviews to build a balanced plan using trusted tours and tickets.”
- Default base for “Why these places”:

- "From many options in {{City_Name}} we selected attractions that combine must-see highlights, easy logistics and consistently high traveller ratings."

Human expert bundles

- Card badge: "Expert itinerary by {{name}}".
- Landing intro component:
    - "This itinerary was hand-crafted by {{name}}, a local expert in {{City_Name}} who combined real traveller reviews, timing and logistics so your day simply works."
- Default base for "Why these places":
    - "I chose these places to show you both the classic {{City_Name}} and its local side – without rushing or wasting time in queues."

2.8.8 Intent / Segment Logic and URLs

- Each bundle has target_segment (**4.8.9 Stage F — Profile model**).
- Each bundle gets its own URL including city, segment and duration if needed, e.g.:
    - /en/italy/rome/3-days-classic-itinerary/
    - /en/austria/vienna/itineraries/family-2-days/
- A simple rules engine infers the user segment from:
    - search query (keywords like "with kids", "romantic")
    - active filters (family-friendly, budget)
    - language/region and browsing behaviour (if allowed).
- On city pages:
    - If bundles exist for inferred segment → show those 3 first.
    - Otherwise → show general bundles.
- All bundle URLs are indexable; personalization affects only which ones are surfaced and in which order.

2.8.9 Hooks (Lead Capture)

Hook: "Download your {{City_Name}} {{duration_days}}-day checklist (PDF)"

- Placement: in the bundle landing page, near top or in sticky sidebar.
- Behaviour:
    - User enters email to get a PDF checklist aligned with the bundle (places, tips, packing list).
    - Email + bundle_id + city_id stored in lead_capture table.
    - PDF file is managed via CMS "Content Assets / PDF Manager" and attached to the bundle by pdf_id.
    - Email with download link is sent via transactional email provider.

This hook integrates with the global hook engine and PDF manager defined in Chapter 2 (hooks and assets).

➢ **Package / Expert Bundle Page: HowTo Module (UI + Schema)**

Purpose

Packages are purchasable and include multiple products. The HowTo explains how to select dates/options for each included tour and complete checkout, without implying a fixed schedule.

Placement in Package Template

Render inside or directly below the Booking Details panel (where availability selection happens), because it is a booking-instructions module.

Eligibility (when allowed)

Allowed only if:

1. Page type = Package/Bundle and bundle checkout exists.
2. Included tours list exists (ItemList renders).
3. The page supports selecting date/option per included tour (or clearly states the current phase behavior).

UI spec (exact structure)

Module title (fixed): "How to book this package"
Steps: 3–7 steps (hard limits)

- `step_title` max 60 chars
- `step_text` max 240 chars (120–220 recommended)

- optional “What to prepare” checklist: 0–8 items

Recommended step set (canonical baseline for packages)

1. Review what’s included

○ references the ItemList module (included tours).

2. Select dates and options for each tour

○ references per-product calendar + timeslot/variant selection UI.

3. Resolve sold-out items

○ “If a tour is sold out on your chosen day, change date/time.”

○ If future phase allows de-selecting items, mention it only when implemented.

4. Check the total and proceed to checkout

○ Must match displayed price aggregation rules.

5. Receive confirmations per tour

○ Clarify that confirmations are issued for each included product booking.

6. Changes and cancellations *(optional)*

○ Must align with package cancellation rules (or per-product policies if that is how you implement it).

Ownership & constraints

- AI draft + manual edit (curator/expert can tune).
- Must not present “Day 1/Day 2 schedule” as binding; it is descriptive only.

Data fields (for dev implementation)

Store as `package_howto_steps[]`:

- howto_enabled
- howto_steps[]

- prepare_list[]
- same timestamps/version fields as product

EXPLORE ATTRACTION MAP / AI FEED CONTEXT

drdd
2.11 EXPLORE MAP / GLOBAL DISCOVERY GRAPH / AI FEED CONTEXT

2.11.0 Purpose and strategic role
The Explore Map is not just a visual map feature.
It is Rosotravel’s global discovery and routing surface built on top of the destination and attraction graph.

Its role is threefold:
1. User discovery layer
It helps users who do not yet have a closed destination or attraction decision move from inspiration to the correct decision-first page.
2. Routing layer
It routes users from a map interaction into the correct page type:
country, city, attraction, things-to-do, near-page, or product.
3. AI / machine-readable graph layer
It exposes Rosotravel’s destination, city, attraction, and relationship graph in a structured form for:
- AI sitemap,
- external LLMs,
- internal AI drafting and retrieval systems,
- future agentic and recommendation use cases.

This means the Explore Map is not a separate “mini-product”.
It is an upstream discovery and graph interface for the rest of the Rosotravel ecosystem.

2.11.1 URL and intent
Front-end map URL:
/{lang}/explore/

Machine-readable AI dataset URL:
/{lang}/explore/ai-dataset/

Intent of the front-end Explore page:
- inspire users who are still open-ended,
- help users navigate from “I do not know where to start” to the correct city / attraction / category page,
- expose Rosotravel’s global destination and attraction knowledge through an interactive map,

- support top-of-funnel and mid-funnel traffic without turning Rosotravel into a generic map catalog.

Intent of the AI dataset feed:
- expose a machine-readable destination and attraction graph,
- support AI sitemap and AI-summary outputs,
- support external LLM understanding of Rosotravel's destination coverage and expertise,
- support internal prompt context, retrieval, and structured routing.

2.11.2 Strategic business value
The Explore Map has business value only if it is treated as a discovery router and graph surface, not as a decorative world map.

The Explore Map creates value in five ways:

A. Top-of-funnel destination discovery
It captures users who are still choosing a country, city, or attraction and are not yet ready to land directly on a single city page.

B. Better routing into decision-first pages
It moves users from map exploration into the correct high-conversion destination pages:
- country pages,
- city pages,
- attraction pages,
- things-to-do pages,
- selected near-pages,
instead of leaving them inside a map-only experience.

C. Stronger brand understanding
It shows that Rosotravel understands destinations globally, not only through listings but through structured destination and attraction relationships.

D. LLM and machine-readability value
It acts as the visible and machine-readable layer of Rosotravel's global destination and attraction graph.

E. Internal reuse of graph logic
The same graph that powers the Explore Map can also power:
- nearby logic,
- attraction relationships,
- internal linking,
- AI retrieval,
- future destination recommendation systems.

2.11.3 Page type and schema
Front-end page type:
Explore / Collection / Discovery UI

Primary schema:
CollectionPage

Secondary schema:
ItemList
Organization

Important:
The front-end Explore page is not a transactional product page and not a map-only widget. Its schema should reflect discovery and collection logic, not a catalog of all possible bookable items.

2.11.4 Core product principle
The Explore Map must not become an OTA-style “everything on a map” interface.

Therefore:
- the global map should prioritize countries, cities, attraction clusters, and top POIs,
- product/tour pins must be tightly controlled and must not overload the world or country view,
- the map must route toward decision-first pages instead of trying to solve the whole booking journey inside the map.

Hard rule:
The Explore Map is a discovery-and-routing interface first, not a full search results page on a map.

2.11.5 Explore page — front-end structure
The Explore page should contain the following sections in this order:

1. Hero / H1 / Intro
2. Global Interactive Map
3. Map support controls and curated overlays
4. Top Places Worldwide
5. Top Cities Worldwide
6. Recommended Destinations
7. Blog / Inspiration Block
8. Explore PDF Hook
9. Optional quiz: Find your perfect destination
10. Optional final CTA / support layer

Comments & Ask Our Expert:
not recommended as a default block on the Explore page in MVP.
Reason:
the Explore page is too broad and global for a standard comments thread to create high-value interaction.
If retained, it should be Phase 2 only and heavily moderated.

2.11.6 Hero / H1 / Intro

Ownership:
AI-drafted with manual override

CMS fields:
- explore_h1
- explore_snippet
- explore_intro_enabled

Rules:
H1 example:
“Explore the world with Rosotravel”

Snippet example:
“Discover destinations, attractions and curated travel ideas on an interactive map designed to help you move from inspiration to the right next step.”

Important:
The hero copy must describe the page as:
- a global discovery entry point,
- not a marketplace catalog,
- and not an “AI feature”.

The intro must explain the Explore page in plain language:
- browse countries, cities, and top attractions,
- zoom into places that interest you,
- and continue into the correct destination page when ready.

2.11.7 Global Interactive Map
Ownership:
factual + graph-driven + system controls

Primary source inputs:
- country coordinates / centroids
- city coordinates
- attraction coordinates
- destination hierarchy graph
- attraction ↔ city / country relationships
- selected prominence / popularity signals
- optional product location signals (only where allowed by zoom rules)

2.11.8 Zoom-layer logic (mandatory)
The Explore Map must not show the same entity density at every zoom level.
It must use zoom-layer logic.

Level 1 — World view
Primary entities:
- countries
- top global cities

- landmark / attraction clusters
Do not show product pins by default.

Goal:
help users choose where to go, not overwhelm them with detail.

Level 2 — Country / regional zoom
Primary entities:
- cities
- major attractions
- major attraction clusters
Optional:
- theme overlays (museums, landmarks, day trips, food)
Do not show full product pin density.

Goal:
help users choose which city / region / landmark area to open next.

Level 3 — City zoom
Primary entities:
- attractions
- selected category overlays
- selected route / cluster overlays
Optional:
- tightly controlled product pins if city density and performance allow it

Goal:
move the user into the right city / attraction / things-to-do decision page.

Level 4 — Deep local zoom
Primary entities:
- detailed attractions
- nearby attraction relationships
- optionally selected product pins
- route context for local discovery
This level is optional in MVP and may be simplified.

Hard rule:
Product / tour pins must never dominate the global or country view.

2.11.9 Entity types shown on map
Supported entity types:
- countries
- cities
- attractions / POIs
- selected attraction clusters
- optional category/theme overlays
- optional product pins (only in controlled zoom contexts)

Recommended MVP:
- countries
- cities
- attractions

Optional in Phase 1.5:
- theme overlays
- local product pins
- nearby attraction clusters

Not recommended in MVP:
- dense global product pins
- world-scale “all tours on the map”

2.11.10 Filters and controls
Required filters:
- Country
- City
- Attraction type
- Category / theme

Recommended category examples:
- museums
- landmarks
- day trips
- food & wine
- family-friendly
- boat cruises
- walking tours

Controls should remain light and decision-oriented.
This must not feel like a search engine replacement.

Recommended additional controls:
- reset map
- zoom to current country / region
- show top places only
- show cities only / attractions only

2.11.11 Hover / click behavior
Hover or tap on an entity should show a compact info card.

Mini card fields:
- entity name
- entity type
- destination context
- optional rating / popularity signal

- optional “top place” or “editorial” label
- primary CTA

Click routing rules:
- country pin → country page
- city pin → city page
- attraction pin → attraction page
- attraction cluster → relevant city page or themed discovery page
- product pin (if enabled) → product page
- category/theme overlay → relevant things-to-do page or selected near-page

Hard rule:
Clicks must route to decision-first pages, not trap the user inside the map.

2.11.12 Curated overlays (recommended)
The Explore Map should not rely only on pins.
It should support curated overlays that make discovery more useful.

Recommended overlay types:
- Top Places
- Top Cities
- Museums
- Landmarks
- Day Trips
- Family-Friendly
- Food & Wine
- Hidden Gems (later phase)
- Bucket List Icons

Purpose:
These overlays give users a more meaningful way to explore than raw pin density.

2.11.13 Top Places Worldwide
Ownership:
system logic + AI snippet

CMS fields:
- explore_top_places_enabled
- explore_top_places_limit
- explore_top_places_title

Recommended limit:
12–20 cards

Selection logic:
Top places should be selected from the global attraction graph using:
- attraction prominence,
- aggregated review strength,

- popularity score,
- destination importance,
- internal editorial curation weight if applicable.

Each card should show:
- thumbnail
- attraction name
- short snippet
- city / country
- CTA to attraction page

Important:
This section should be treated as a curated shortcut derived from the same graph as the map, not as a separate ranking product.

2.11.14 Top Cities Worldwide
Ownership:
system logic + AI snippet

CMS fields:
- explore_top_cities_enabled
- explore_top_cities_limit
- explore_top_cities_title

Selection logic:
Cities should be selected using:
- number of experiences
- city-level review reputation
- city prominence
- internal strategic importance
- editorial weight if needed

Each card should show:
- city image
- city name
- short snippet
- CTA to city page

Important:
This section acts as a non-map shortcut for users who prefer cards over pin exploration.

2.11.15 Recommended Destinations
Ownership:
system logic + light DEE / geo-based logic

Purpose:
help users who are open-ended but likely closer to a region or travel style.

Inputs may include:
- approximate region / country
- selected language
- soft market signal
- seasonality
- global popularity
- manually defined fallback set

Recommended visible limit:
4–8 cards

Important:
This block must remain light-touch and inspirational.
It must not feel like Rosotravel already decided everything for the user.

2.11.16 Blog / Inspiration Block
Ownership:
editorial

CMS fields:
- explore_blog_articles[]

Rules:
- 4–6 articles
- tagged global / top picks / editor’s choice
- this block supports inspiration and depth, but should not overshadow the map

2.11.17 PDF Hook — World Travel Starter Pack
Hook name:
explore_world_travel_starter_pack

Purpose:
capture upper-funnel users who are still planning globally.

Placement:
mid-page or near bottom, after map and curated discovery sections

Elements:
- headline
- subline
- email field
- consent checkbox

On submit:
- save lead with source = explore_pdf
- send PDF via asset manager / email

CMS:

- pdf-manager
- explore hook config
- email template mapping

2.11.18 Quiz — Find your perfect destination
Placement:
below map or lower on page, never before the map

Purpose:
help users who are open-ended and want a structured path

Behavior:
- 5–7 lightweight questions
- travel style
- budget
- trip length
- season
- interests

Results:
- recommended cities
- selected themed destination pages
- optional email me this plan

Integration:
Answers may seed session intent signals for later pages, but the Explore page itself should not become quiz-first.

Hard rule:
No forced quiz at page entry.

2.11.19 Comments / Ask Our Expert
Current recommendation:
Do not enable by default in MVP.

Reason:
The Explore page is too global and too early in the journey for a standard comments thread to create clear value.

If this module is kept for future phases:
- it should be optional,
- strongly moderated,
- and reframed more as “Need help choosing where to start?” than as open-ended comments.

2.11.20 AI dataset feed — strategic role
The machine-readable feed at:
/{lang}/explore/ai-dataset/

must not be treated as a simple export.
It should be the structured global destination and attraction graph of Rosotravel.

This feed should support:
- AI sitemap
- AI-summary feed
- external LLM understanding
- internal AI prompt context
- future agent / recommender systems

2.11.21 AI dataset feed — minimum entity schema
The AI dataset feed should expose at minimum:

For each entity:
- stable entity_id
- entity_type (country / city / attraction / optional product)
- canonical_name
- language
- canonical_url
- coordinates
- parent_id / hierarchy references
- related_entities[]
- destination context
- prominence score
- review / trust indicators where allowed
- updated_at

For cities:
- country reference
- city type
- top attractions references
- city page URL

For attractions:
- linked city_id
- country_id
- attraction category/type
- attraction page URL
- related nearby attractions

For countries:
- top cities references
- country page URL

Optional later:
- bundle references
- category page references
- nearby clusters

- selected product references

2.11.22 AI dataset feed — output principles
The feed must be:
- machine-readable
- lightweight enough for consumption
- stable in IDs
- regularly refreshed
- consistent with front-end routing
- aligned with the entity graph used internally

Hard rule:
The AI dataset feed must not expose broken or non-canonical routing.

2.11.23 Performance and pin governance
Because map performance can degrade quickly, the Explore page must define:

CMS/system controls:
- entity type toggles
- max pins per viewport
- max pins by zoom level
- default zoom
- cluster behavior
- lazy load rules
- fallback card mode if map performance degrades

Hard rule:
The system must prefer clustering and curated prominence over raw density.

2.11.24 Discovery and routing KPI
The Explore Map must be measured as a discovery router, not as a vanity interaction tool.

Primary KPIs:
- CTR from map to country/city/attraction pages
- assisted conversions from Explore entry
- continuation rate into decision-first pages
- quiz completion rate
- PDF hook conversion
- map-to-destination routing quality
- graph coverage quality for AI feed

Secondary KPIs:
- pin click rate
- zoom interactions
- overlay usage

Not sufficient on their own:
- time on map

- map scrolls
- hover interactions

2.11.25 What this page is not
The Explore page is not:
- a replacement for destination pages
- a full OTA map catalog
- a giant product results page
- a pure marketing gimmick
- only a machine-readable feed wrapper

It is:
- a discovery layer
- a graph surface
- and a routing hub.

2.11.26 Recommended MVP scope
For MVP, prioritize:
- countries
- cities
- attractions
- world/country/city zoom logic
- strong routing
- Top Places
- Top Cities
- PDF hook
- optional quiz
- AI dataset feed with stable entity graph

Defer or tightly restrict:
- dense product pins
- comments module
- advanced personalization
- overly complex overlays
- transactional logic inside the map

2.11.27 Business recommendation
The Explore Map should be positioned internally not as “global map feature”, but as:
Rosotravel’s global discovery graph and routing interface.

This framing is strategically stronger because it gives the feature:
- business value,
- SEO/LLM value,
- clearer KPI ownership,
- and a better role inside the platform.

End of 2.11

Offer Comparison Tables

Purpose
Clarify differences between variants (private vs group; morning vs evening) for UX and crawlability.

Data/Schema

- **HTML tables + optional JSON mirror; fields: price, duration, group size, pickup, refund.**

Backend

- **Build comparison datasets from variant definitions; update incrementally.**
- **Localize labels.**

Front-End

- **Fast, responsive tables; accessible (headers, captions).**
- **Optional sticky compare.**

Validation

- **No missing required variant fields; responsive rendering.**

KPIs

- **Increased conversion on pages with comparisons; reduced bounce.**

Acceptance (DoD)

- **Tables present where variants exist; data matches source**

“Data at a Glance”

Purpose
Compact factual block for users and crawlers.

Data/Schema

- City, Price, Duration, Languages, Meeting Point; mapped to `TouristTrip` / `Offer`.

Backend

- Populate from feed; fallbacks; localization per locale.

Front-End

- Prominent, fast rendering above the fold; icons optional, text guaranteed.

Validation

- No empty fields without an explicit placeholder.

KPIs

- Improved time-to-first-interaction; better CWV due to minimal block.

Acceptance (DoD)

- Present on all product pages

04. Long-Form Quality & TOC Rules (Template-Specific)

a. Global hard rules (apply to all templates)

1. **One H1 only** (visible).
2. **No empty Title/Meta/H1** (publishing gate handled in IV/4.15).
3. If a module is rendered (FAQ, ItemList, HowTo), schema must match it (no hidden schema).

b.TOC (Table of Contents) exact templates + behavior. TOC is allowed ONLY on Article-based templates:

- **Blog Page**
- **Travel Guide** pages (if implemented inside Blog template family)

TOC is **NOT used** on:

- Homepage, Country/Region/City destination pages, Things-to-do listings, Attraction pages, Product pages, Near-pages, Expert pages, Explore map.

c.TOC trigger thresholds (when to render automatically)

Render TOC automatically when **either** condition is met:

- `article_body_length >= 4,500 characters` (including spaces), **OR**
- number of H2 headings `>= 5`

If neither condition is met → no TOC.

TOC source and heading rules

- TOC is generated from **H2 headings only** (H3 optional as collapsible nested items).
- Do not include these sections in TOC (even if they have headings):

    - “FAQ”
    - “Comments / Ask our expert”
    - “Related tours”
- Heading anchors must be deterministic:
    - `anchor_id = slugify(heading_text) + '-' + stable_counter`
    - must be stable across minor edits to prevent broken deep links.

d.TOC UI behavior (dev-ready)

05. Desktop: TOC appears **after the intro** and before the first H2 section.

06. Mobile: TOC appears as a **collapsible “On this page”** block.

07. Clicking an item scrolls to the section with **offset** for sticky header.

08. The currently visible section is highlighted (optional; nice-to-have)

e.Only for Blog Page (Article template)

- **TOC rules apply** (see above).
- Required blocks:
    - Title, author box (clickable), Published date, Updated date (if different), Reading time
    - Intro paragraph (target **600–1,200 characters**)
    - Key Takeaways (3–6 bullets) if article length ≥ 4,500 chars
    - Minimum structure:
        - at least **3 H2 sections** if length ≥ 4,500 chars
- Article length guidance:
    - standard blog: **4,500–12,000 characters**
    - deep guide: **10,000–25,000 characters**
- Freshness UX:

- show “Updated” only if dateModified differs from datePublished.
- show “Reviewed” only if you actually editorially reviewed it (token exists).

09. H1, H2, H3 tags

1) HOMEPAGE

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	Outstanding memories — chosen by people who live the experience.	CMS (manual brand line)	No	None
H2	Why you see shortlists	Static i18n label	No	None
H2	Rosotravel Live	Static i18n label	No	None
H2	Top Destinations	Static i18n label	No	None
H2	Explore by Interest	Static i18n label	No	None
H2	Trending Experiences	Static i18n label	No	None
H2	Why travelers choose Rosotravel	Static i18n label	No	None
H2	Meet our experts	Static i18n label	No	None
H2	Travel guides	Static i18n label	No	None

---

2) COUNTRY PAGE

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	Discover {{Country_Name}}	Tokens + CMS override (country_h1 optional)	Yes (baseline), then Manual if edited	Medium (page-owned keyword may shape phrasing)
H2	Country Overview	Static i18n label	No	None
H2	Highlights	Static i18n label	No	None
H2	Facts & Curiosities	Static i18n label	No	None
H2	Top Cities	Static i18n label	No	Low (ordering may reflect intent; label not)
H2	Top Regions	Static i18n label	No	None
H2	Top Attractions in {{Country_Name}}	Tokens + Static i18n label	No	Low (token only)
H2	Things to Do in {{Country_Name}}	Tokens + Static i18n label	No	Low (token only)
H2	Other Sightseeing Options in {{Country_Name}}	Tokens + Static i18n label	No	Low (token only)
H2	Blog & Guides	Static i18n label	No	None
H2	Country-Level Reviews	Static i18n label	No	None
H2	FAQ	Static i18n label	No	Medium (FAQ topics can be keyword-shaped; label stays static)
H2	Comments & Ask Our Expert	Static i18n label	No	None
H2	Hooks	Static i18n label	No	None

3) REGION PAGE

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	Explore {{Region_Name}}	Tokens + CMS override (region_h1 optional)	Yes (baseline), then Manual if edited	Medium
H2	Region Overview	Static i18n label	No	None
H2	Highlights	Static i18n label	No	None
H2	Facts & Curiosities	Static i18n label	No	None
H2	Top Sights in {{Region_Name}}	Tokens + Static i18n label	No	Low
H2	Go beyond {{Region_Name}}	Tokens + Static i18n label	No	Low
H2	Things to Do in Region	Static i18n label	No	Low
H2	Top Cities in Region	Static i18n label	No	Low
H2	Region-Level Reviews	Static i18n label	No	None
H2	FAQ	Static i18n label	No	Medium
H2	Comments & Ask Our Expert	Static i18n label	No	None
H2	Hooks	Static i18n label	No	None

---

4) CITY PAGE

Headin g Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)

H1	Visit {{City_Name}} — curated tours, tickets & day plans	Tokens + CMS override (city_h1)	Yes (baseline), then Manual if edited	Medium (city-owned keyword can shape tagline)
H2	Top picks for you in {{City_Name}}	Tokens + Static i18n label	No	Low (DEE affects selection/order; label fixed)
H2	{{City_Name}} in 1–3 days	Tokens + Static i18n label	No	Low
H2	Private upgrades	Static i18n label	No	Low
H2	Choose your style	Static i18n label	No	Low
H2	Top attraction tours in {{City_Name}}	Tokens + Static i18n label	No	Low
H2	Top tours and activities	Static i18n label	No	Low
H2	Best neighborhoods in {{City_Name}}	Tokens + Static i18n label	No	None
H2	Local tips for {{City_Name}}	Tokens + Static i18n label	No	None
H2	Explore {{City_Name}} by category	Tokens + Static i18n label	No	Low
H2	This week in {{City_Name}} with Rosotravel	Tokens + Static i18n label	No	None
H2	Travel guides for {{City_Name}}	Tokens + Static i18n label	No	None
H2	What travelers say about {{City_Name}} with Rosotravel	Tokens + Static i18n label	No	None

H2	FAQ about {{City_Name}}	Tokens + Static i18n label	No	Medium (FAQ topics keyword-shaped; label fixed)

---

5) THINGS-TO-DO PAGE (City-level)

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	Things to do in {{City_Name}}	Tokens + CMS override (ttd_h1)	Yes (baseline), then Manual if edited	High (this page is keyword-owned)
H2	Top picks for you in {{City_Name}}	Tokens + Static i18n label	No	Low
H2	Ready-made {{City_Name}} plans	Tokens + Static i18n label	No	Low
H2	Top attraction tours in {{City_Name}}	Tokens + Static i18n label	No	Low
H2	Insider videos / creator content	Static i18n label	No	None
H2	Explore {{City_Name}} your way	Tokens + Static i18n label	No	Low
H2	Top activities	Static i18n label	No	Low
H2	Beyond the tourist spots	Static i18n label	No	Low
H2	Nearby cities	Static i18n label	No	None
H2	Reviews	Static i18n label	No	None

H2	FAQ	Static i18n label	No	Medium
H2	Comments & Ask our expert	Static i18n label	No	None
H2	Hooks	Static i18n label	No	None

---

6) ATTRACTION PAGE (POI)

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	{{Attraction_Name}} in {{City_Name}}	Tokens + CMS override (attraction_h1)	Yes (baseline), then Manual if edited	High (attraction name queries)
H2	Our 2 picks	Static i18n label	No	Low
H2	Intent options	Static i18n label	No	Low
H2	Private upgrade	Static i18n label	No	Low
H2	Tickets-only	Static i18n label	No	Low
H2	Visitor information	Static i18n label	No	None (factual section)
H2	Overview	Static i18n label	No	Low (content may reflect intent; label fixed)
H2	History / significance	Static i18n label	No	None

H2	Top things to see	Static i18n label	No	Low
H2	Nearby attractions	Static i18n label	No	Low
H2	Tours related to this attraction	Static i18n label	No	Low
H2	Bundles including this attraction	Static i18n label	No	Low

---

7) PRODUCT PAGE

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	{{Tour_Name}} (optionally with {{Attraction_Name}}, {{City_Name}} tokens)	Tokens + CMS override (h1)	Yes (baseline), then Manual if edited	**Low–Medium (SEO may influence, but clarity of offer is primary)**
H2	Rosotravel Live — Last Tour Snapshot	Static i18n label	No	None
H2	Options / Variants	Static i18n label	No	None
H2	Highlights	Static i18n label	No	Low
H2	Why travelers love this tour	Static i18n label	No	None
H2	Included / Not included	Static i18n label	No	None
H2	Itinerary Overview	Static i18n label	No	Low
H2	Expert tip / Local insight	Static i18n label	No	None
H2	Meet your guides	Static i18n label	No	None
H2	Trip DNA	Static i18n label	No	None

H2	FAQ	Static i18n label	No	Medium (questions can reflect intent)
H2	Reviews	Static i18n label	No	None
H2	Q&A	Static i18n label	No	None
H2	Comments / Ask our expert	Static i18n label	No	None
H2	Related tours / You may also like	Static i18n label	No	None
H2	Hooks	Static i18n label	No	None

Note (intentional): the booking panel “What’s included” etc. are section headings; we do not treat UI labels inside widgets as separate SEO headings unless they are rendered as actual H2/H3 elements.

---

8) NEAR-PAGE (Category / Keyword-driven city page)

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	{{Primary_Keyword}} in {{City_Name}} (or keyword + city decision framing)	Tokens + CMS override (nearpage_h1)	Yes (baseline), then Manual if edited	**High (this page owns the keyword)**
H2	Overview	Static i18n label	No	Medium (overview content keyword-driven; label fixed)
H2	Why these picks	Static i18n label	No	Medium (generated from selection proof; label fixed)

H2	Trade-offs / What we skipped	Static i18n label	No	Medium (generated from selection proof; label fixed)
H2	Comparison (Top 3–5 picks)	Static i18n label	No	Low
H2	Top X choices	Static i18n label	No	Low
H2	Bundles related to this category	Static i18n label	No	Low
H2	Map / SRP summary	Static i18n label	No	None
H2	Related categories	Static i18n label	No	Low
H2	FAQ	Static i18n label	No	High (FAQ topics keyword-driven)

---

9) BLOG PAGE (Article / Stories)

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	{{blog_title}}	CMS (manual by author)	No	High (blog targets keywords, but title is editorial manual)
H2	Table of Contents	Static i18n label (auto module label)	No	None
H2	FAQ	Static i18n label	No	Medium
H2	Comments & Ask our expert	Static i18n label	No	None
H2	Related Articles	Static i18n label	No	None
H2	Hooks	Static i18n label	No	None

Important: article body headings (H2/H3 inside the article) are **author-defined** and are not part of a global template heading map (they live inside rich text).

---

10) EXPERT PAGE

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	{{Expert_Name}}, {{Role}} in {{City_Name}}	Tokens + CMS	No (prefer manual)	Low
H2	Bio	Static i18n label	No	None
H2	Tours Authored / Guided	Static i18n label	No	None
H2	Articles Authored	Static i18n label	No	None
H2	Short Q&A Block	Static i18n label	No	None
H2	FAQ	Static i18n label	No	Low
H2	Comments & Ask our expert	Static i18n label	No	None
H2	Hooks	Static i18n label	No	None

---

11) PACKAGE / BUNDLE PAGE (City itineraries)

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	{{bundle.name}}	CMS (AI baseline allowed for auto bundles), then Manual if edited	Yes (baseline), then Manual if edited	Medium (theme keywords may influence)

H2	Overview	Static i18n label	No	Low
H2	Why these picks	Static i18n label	No	Low (generated from bundle selection proof)
H2	Trade-offs / What we skipped	Static i18n label	No	Low (generated from bundle selection proof)
H2	Included tours	Static i18n label	No	None
H2	Suggested day-by-day itinerary	Static i18n label	No	Low
H2	Booking details	Static i18n label	No	None
H2	How to book this package	Static i18n label	No	None
H2	Highlights of the package	Static i18n label	No	Low

---

12) EXPLORE ATTRACTION MAP / AI FEED CONTEXT

12A) Explore Attraction Map (front-end page)

Heading Level	Heading Label (UI)	Heading Source (Label)	AI Label Allowed	Keyword Dependency (Label)
H1	World Attraction Map	CMS (explore_h1)	Yes (baseline), then Manual if edited	Low
H2	World Attraction Map	Static i18n label	No	None

H2	Top places worldwide	Static i18n label	No	None
H2	Top cities worldwide	Static i18n label	No	None
H2	Recommended destinations	Static i18n label	No	Low (personalization affects list/order; label fixed)
H2	Blog / Inspiration	Static i18n label	No	None
H2	Email capture hook — World Travel Starter Pack	Static i18n label	No	None
H2	Comments & Ask Our Expert	Static i18n label	No	None

12B) Explore AI Dataset (machine-readable; not a normal heading page)

No HTML heading map required (JSON endpoint). If you still render a human-facing debug page, headings would be static labels only.

Notes:

10. This is **not SEO text**, but **LLM fuel**
11. Enables:
    a. AI Sitemap
    b. External LLM consumption
    c. Internal AI grounding
12. Zero risk of cannibalization

05. Relations & Graph Signals (SEO + Agents) UI-Bound Rules

The platform exposes a destination-and-products graph for search engines and transaction agents.
Hard rule: any relation emitted in JSON-LD (or AI feeds) must be supported by visible UI modules (lists, breadcrumbs, “Nearby”, “Included tours”, etc.). No hidden relations.

A) Country Destination Page — required relations

- **BreadcrumbList**: Home → Country
- **hasPart** (or ItemList to Cities/Regions): lists the child destinations rendered on page
- **Organization/provider** (Rosotravel identity)
  UX must show: breadcrumbs + visible “Top cities/regions” list.
  Do not emit child lists if the page does not render them.

B) Region Destination Page — required relations

- **BreadcrumbList**: Home → Country → Region
- **isPartOf**: Region → Country
- **hasPart** / ItemList: Region → Cities (only if rendered)
  UX must show: breadcrumbs + visible city list.

C) City Destination Page — required relations

- **BreadcrumbList**: Home → Country → Region (if exists) → City
- **isPartOf**: City → Region/Country
- **hasPart** / ItemList: City → (Top Attractions, Top Tours) *only if rendered*
- **nearby**: optional, only if “Nearby” UI module exists
  UX must show: breadcrumbs + the lists that schema claims exist.

D) City “Things to do” / Category Listing Page — required relations

- **BreadcrumbList**: Home → … → City → Things to do (or Category)
- **isPartOf**: Listing → City

- **ItemList**: list of products rendered (must match UI ordering + count constraints)
  UX must show: the exact product list (or paginated subset rules documented).

E) Attraction (POI) Page — required relations

- **BreadcrumbList**: Home → … → City → Attraction
- **containedInPlace** (or isPartOf): Attraction → City
- **GeoCoordinates** (when available)
- **sameAs**: optional but recommended (Wikidata/Wikipedia/official site when known)
- **tourBookingPage / relatedLink**: optional, only if page renders "Book tours for this attraction"
  UX must show: key facts (address/coordinates if used), and any "Book" links that schema exposes.

F) Tour / Product Page — required relations

- **BreadcrumbList**: Home → … → City (and/or Attraction) → Tour
- **offers**: Offer must match the visible price/currency/availability CTA
- **provider**: Organization identity
- **potentialAction** (ReserveAction/BuyAction): only if a direct booking action exists
- **isPartOf / relatedLink**: optional (tour belongs to city/attraction) if UI shows those parents
  UX must show: price, currency, availability state, cancellation summary, booking CTA.

G) Package / Bundle Page — required relations

- **BreadcrumbList**: Home → … → City → Package (or Hub → Package)
- **ItemList**: included tours must match UI list
- **offers**: packages are purchasable → must include Offer (price/currency/availability if applicable)
- **isPartOf / relatedLink**: package ↔ city/attraction cluster (only if rendered)
  UX must show: included tours + bundle price/CTA.

H) Blog / Guides / Itinerary Articles — required relations

- **author (Person)** + visible attribution (name + profile link)
- **datePublished / dateModified** visible (Published / Updated)
- **about / primaryTopic / mentions**: only if the article visibly targets a destination/POI
- **BreadcrumbList** (if blog uses breadcrumbs)
  UX must show: author box + dates + (optional) TOC if long-form (see content quality section).

I) Expert Page — required relations

- **Person** schema with identity fields
- **worksFor/provider** (Organization)
- **sameAs** optional
  UX must show: identity + bio.

J) Global Experiences Hub — required relations

- **CollectionPage** + **ItemList** of categories/countries that are visible entry points
- **hasPart** optional (hub → category pages) if the hub renders them
  UX must show: the category/country lists schema claims.

2.4 Marketing USP Rules

01. DYNAMIC EXPERIENCE ENGINE (DEE) – UX principles & adaptation rules

0) Purpose (what DEE is)

The Dynamic Experience Engine (DEE) is a runtime presentation layer that adapts how Rosotravel content is shown to users so pages feel personally relevant and decision-first, while remaining index-safe and governance-safe.

DEE exists to deliver three outcomes:
A) Decision clarity: different traveler types see different emphases and “next best steps” (without turning pages into a catalog).
B) Trust: users understand why they see a shortlist (“why these picks” framing, proof blocks,

and consistent logic).
C) Consistency: the same rules drive City / Things-to-do / Attraction / Near pages / Bundles, so Rosotravel feels like one expert system.

DEE is NOT:

- a search engine replacement,
- SEO logic (keyword ownership, indexation decisions),
- a black-box recommender,
- a runtime AI writer.

DEE IS:

- controlled, explainable, rule-based,
- variant-based (selects pre-generated variants, does not generate SEO text at runtime),
- constrained to “safe levers” (ordering + emphasis + variant selection).

---

1) Non-negotiables (DEE invariants)

DEE MUST NOT change:

- URL structure, canonical tags, hreflang, indexation state
- keyword ownership locks and page-type ownership rules
- H1 and other SEO-critical locked fields
- schema hierarchy and JSON-LD content that is part of the indexed contract
- pricing, availability logic, checkout logic
- the set membership / winners / portfolio selection rules (Model C selection is upstream)

DEE MAY change ONLY:

- which pre-generated snippet variant is shown (audience-based)

- ordering of lists (within the already eligible pool)
- visibility/emphasis of selected UX blocks (hooks, proof components, “more options” entry)
- which bundles are featured (choose among precomputed bundles by profile)
- which reviews/photos/human proof are surfaced first (reorder and summarize from precomputed clusters)
- microcopy add-ons that are explicitly marked “DEE-safe” (non-indexed UX copy)

---

2) Relationship to Model C (where DEE gets “what to show”)

DEE does not decide which products are eligible. Model C does that upstream (Phase A + Phase C: pool, set, archetypes, winners, dedup, slot budgets).

DEE consumes Model C outputs per template:

- City Page: bundles (1/2/3), Top Picks (6), Private Upgrade CTAs, Rails (4×2), Top Tickets (4–6), Phase B entrypoint
- Things-to-do Page: Top Picks (6), bundles (3), rails (6×2), Top Tickets (4–6), Phase B entrypoint
- Attraction Page: Our 2 picks, mini intent rails, private upgrade CTA, ticket-only block (if applicable)
- Near Pages: keyword-owned list, comparison block candidates, related categories, hooks
- Bundles: 1/2/3-day bundles and their internal “plan blocks”

DEE then selects how to present these elements for the current dee_state (audience/intents/constraints).

---

3) DEE state model (what DEE computes)

DEE computes a session-level state:

dee_state = {
 audience_primary: <one of Model C audiences>,
 intents: <0..2 of Model C intents>,

```
constraints: <0..n of Model C constraints>,
confidence: {
audience: 0..1,
intent: 0..1
},
flags: {
planner_mode: boolean,
decision_difficulty: boolean
},
sources: [ {type, value, weight, timestamp} ],
updated_at: timestamp
}
```

Terminology clarification (avoid confusion):

- Model C baseline audience is “first_time_visitor”.
- “snippet_default” fields remain for backward compatibility, but semantically mean “snippet_first_time_visitor”.

---

4) Model C taxonomy used by DEE (reference)

DEE uses Model C tags and does not redefine them here (full taxonomy and selection logic are defined in 4.8 Portfolio Choosing / Product Classification).

Audiences (primary, exactly 1 active):

- first_time_visitor (default),
- family_traveler,
- couples_traveler,
- comfort_easy_pace_traveler,
- solo_social_traveler,
- interest_deep_dive_traveler,
- active_adventure_traveler

Intents (0..2 active):

- best_value
- premium_comfort

- hidden_gems
- less_crowded
- bucket_list_icons
- deep_dive_learning
- activity_intensity
- family_compatibility
- private_upgrade

Constraints (0..n active):

- private_only
- small_group_preferred
- low_walking_required
- stairs_sensitive
- stroller_required
- wheelchair_access_needed
- short_time
- half_day_only
- hotel_pickup_preferred
- free_cancellation_preferred
- strict_budget_cap
- weather_sensitive

---

5) Signal groups and precedence (how DEE detects “who is who”)

DEE uses 5 signal groups in strict precedence order:

G1) Explicit user actions (hard override; highest confidence)

- quiz result (if implemented)
- explicit “travel style” chip selection (if implemented)
- explicit filters that map to constraints (e.g., wheelchair, private-only)

G2) Landing context (hard seed; very strong)

- Near-page theme (see Section 7) ← strongest non-explicit seed
- Bundle landing page entry
- Travel Agent / Partner context (authenticated)
- Explicit “mode” entry (Magic Finder / Browse curated catalog)

G3) Query & routing signals (strong)

- landing keyword slug
- internal search query
- organic query (if available)
- page type + route (weak prior only)

G4) On-site behavior (medium; refines intent/constraints, may set flags)

- clicks on rails / tickets-only / private upgrade CTA
- bundle previews/landing behavior (L1/L2/L3)
- compare behavior (decision difficulty)
- save actions / wishlist (if implemented)

G5) Locale & language priors (soft tie-breakers; capped influence)

- GeoIP-derived country/region (no raw IP stored)
- UI language selected

- device type (weak modifier)

Hard principles:

- Explicit actions override everything.
- Near-page theme seeding overrides generic “city page defaults”.
- Locale/language priors never override explicit signals.

---

6) Query & keyword mapping (audience + intents + constraints)

DEE parses (in order): near-page theme → query/slug → internal search → referrer tags → on-site behavior.

6.1 Audience pattern mapping (examples; extendable)

Patterns	audience_primary
“with kids”, “family”, “child”, “children”, “stroller”, “teen”, “kid-friendly”, “for families”	family_traveler
“couple”, “for two”, “romantic”, “honeymoon”, “anniversary”, “date night”, “sunset cruise”, “romantic dinner”, “romantic tour”	couple_traveler
“easy pace”, “accessible”, “wheelchair”, “stair-free”, “low walking”, “slow pace”, “relaxed”, “senior”, “comfortable pace”	comfort_easy_pace_traveler
“solo”, “alone”, “meet people”, “social tour”, “group vibes”, “join a group”, “solo traveler”	solo_social_traveler

“deep dive”, “history”, “architecture”, “art lovers”, “culture”, “museum lovers”, “WW2”, “Jewish history”, “expert-led”, “learn more”	interest_deep_dive_traveler

6.2 Intent pattern mapping (0..2 intents; additive)

Patterns	intents
“best value”, “cheap”, “affordable”, “worth it”, “budget”, “good price”, “save money”, “low cost”	best_value
“premium”, “comfortable”, “premium comfort”, “small group comfort”, “better experience”, “higher-end”, “more comfortable”	premium_comfort (only if explicit)
“VIP”, “private”, “exclusive”, “just for us”, “private guide”, “private experience”	private_upgrade (or private_only if explicit and restrictive)
“hidden gems”, “local”, “underrated”, “off the beaten path”, “secret spots”, “non-touristy”	hidden_gems
“less crowded”, “avoid crowds”, “quiet”, “peaceful”, “not busy”, “crowd-free”	less_crowded
“skip the line”, “reserved”, “timed entry”, “priority access”, “fast track”	bucket_list_icons (ticket-first / iconic-access bias; fine-grain TBD)

“learn”, “guided”, “expert guide”, “stories”, “explained”, “historian”, “educational”	deep_dive_learning
“intense”, “extreme”, “active”, “physical”, “adrenaline”, “challenging”	activity_intensity

6.3 Constraint mapping (set only when explicit)

Patterns	**constraints**
“wheelchair”, “wheelchair accessible”, “accessible entrance”	wheelchair_access_needed
“stroller”, “pushchair”, “pram”	stroller_required
“stairs”, “no stairs”, “many steps”, “avoid stairs”, “stair-free”	stairs_sensitive
“hotel pickup”, “pickup”, “pick-up included”, “from hotel”	hotel_pickup_preferred
“short time”, “few hours”, “limited time”, “quick visit”, “only half day”	short_time (and/or half_day_only where applicable; TBD)
“half day”, “morning only”, “afternoon only”	half_day_only

“free cancellation”, “cancel for free”, “refundable”	free_cancellation_preferred

Resolution rules:

- audience_primary: choose strongest match else first_time_visitor.
- intents: maximum 2; never infer hidden_gems unless explicit.
- constraints: only if explicit query/filter/selection.

---

7) Near-page theme seeding (HARD RULE) — decision platform consistency

Near pages are not generic category pages. They must carry a theme/profile so the landing itself communicates “this is for you”.

Each near page MUST store:

- nearpage_theme_id (enum; e.g., family_activities, skip_the_line, day_trips, walking_tours, hidden_gems, etc.)
- target_audience_default (optional but strongly recommended)
- intent_bias[] (0..2)
- constraint_bias[] (0..n; soft unless explicit)
- proof_style (enum): decision_shortlist | comparison_heavy | ticket_first | bundle_first
- dee_seed_strength (default: high)

Runtime behavior:
Landing on a near page seeds dee_state strongly:

- audience_primary = target_audience_default (if present) else first_time_visitor
- intents add = intent_bias (up to 2)
- constraints add = constraint_bias (soft)

- confidence increases (user self-selected the theme via query)
  This seed persists across the session unless user explicitly changes filters/quiz.

Example:
Near page: “what to do with family in london”

- theme = family_activities
- audience_default = family_traveler
- intent_bias = [family_compatibility]
- proof_style = decision_shortlist

---

8) Referrer signals (practical + campaign-controlled)

Referrer signals set a weak prior unless overridden by near-page theme or explicit actions.

Referrer → bias

- From Travel Guides (editorial) → audience: first_time_visitor (baseline), intent: deep_dive_learning (weak)
- From paid ads:
  - if utm_audience/utm_intent present → treat as explicit (strong)
  - else bias intent toward best_value OR bucket_list_icons depending on campaign taxonomy
- From CRM/email:
  - set planner_mode soft bias (bundles + save-plan hook), but do not change audience unless other signals confirm

UTM rules (treated as explicit):

- utm_audience=<audience_tag> sets audience_primary (high confidence)
- utm_intent=<intent_tag> adds intent (up to 2)
- utm_constraints=<...> sets constraints (if allowed by governance)

9) On-site behavior signals (what changes, where, and how)

Behavior does not auto-filter inventory by default.
Behavior changes emphasis, ordering tie-breaks, hooks, and prompts.

9.1 Filters (explicit)

If user applies a filter:

- map filter → intent/constraint explicitly
- show an “active context chip” that explains what drives ordering
- user can clear it (“Reset preferences”)

9.2 Click behavior (rails/picks/tickets-only)

Examples:

- click “Top Tickets” module → intent confidence boost: bucket_list_icons (+bounded)
- click “Private Upgrade” CTA → intent confidence boost: private_upgrade (+bounded)
- repeated engagement with “Less Crowded” rail → boost less_crowded (+bounded)

Definition: “boost intent” means:

- increase tie-break weight in ordering profile (bounded)
- enable an optional intent add-on line (if confidence threshold met)
- prioritize matching proof blocks (decision explanation, logistics clarity)
  It does NOT:
- remove products
- rewrite SEO content
- change URLs

9.3 Bundle interaction levels (L1/L2/L3)

L	Name	Trigger	Effect
L	light planner signals	user opens ≥1 bundle preview modal OR clicks “View itinerary” once	increase bundle prominence (see 10.3); show “Save this plan” mini-CTA; do not reorder Top Picks drastically (only highlight bundles + related CTAs)
L	confirmed planner signals	user opens ≥2 bundle previews OR clicks bundle landing page once	increase bundle prominence further; show a 1-line confirmation prompt (non-blocking; see 10.3); optionally surface “Compare plans” micro-CTA (if implemented)
L	active planner mode	user clicks bundle landing page ≥2 times OR starts selecting date/time inside bundle	set flags.planner_mode = true; make “continue planning” the primary next step on City/TTD returns; keep jump links anchored to bundles and highlight “plans-first” navigation

9.4 High exploration / compare behavior (“decision difficulty”)

If user:

- opens 6+ product cards OR
- toggles between 2–4 products repeatedly OR
- spends >X seconds on comparison block
 Then set flags.decision_difficulty = true

Effect:

- surface “Why these picks” and “Trade-offs” blocks earlier
- surface “Still not sure?” safety net CTA earlier
- reduce aggressive personalization (keep ordering stable to avoid confusion)

---

10) Variant-based personalization (DEE does not write at runtime)

DEE does not generate content on the fly (MVP).
Variants are created during drafting and stored in CMS fields. Runtime selects only.

10.1 Snippet variants (primary lever)

Every page stores snippet variants aligned to Model C audiences:

- snippet_default (= first_time_visitor)
- snippet_family_planner
- snippet_comfort_easy_pace_traveler
- snippet_solo_social_traveler
- snippet_interest_deep_dive_traveler
- snippet_active_adventure_traveler

Rules:

- 180–260 characters
- 1 paragraph
- no CTA, no keyword stuffing
- UX-only (not meta tags, not indexed)

10.2 Review theme emphasis (precomputed, not rewritten)

Reviews are clustered offline into themes (embeddings + rules), e.g.:

- guide quality, pacing, organization, value, crowds, accessibility, atmosphere, group dynamics

DEE:

- never changes reviews
- reorders and surfaces theme summaries based on audience/intents

Output:

- 3–4 summary bullets “What guests loved” (non-quoted), per audience

10.3 Ordering profiles (deterministic)

Ordering profiles define how lists are sorted, not which items exist.
DEE chooses ordering_profile_id based on dee_state.

Required profiles (initial):

- order_first_time_visitor (baseline)
- order_family_planner
- order_comfort_easy_pace
- order_solo_social
- order_deep_dive
- order_active_adventure
  Tie-break intents can add bounded weights.

Important:

- No ML ranking in MVP
- Scores used are precomputed upstream or derived from safe fields (rating, reviews_count, language_coverage, availability_state, etc.)

10.4 Bundle selection variants

Bundles exist as:

- general (first_time_visitor baseline)
- optional audience/intents (profile-aware)

DEE selects:

- if profile bundles exist → show those 3 (1/2/3-day)
- else show baseline 3

10.5 Hooks and prompts (predefined; DEE selects)

DEE can change:

- which hook appears (save plan / discount / safety net / checklist)
- where it appears (within allowed template placements)
  DEE cannot invent new hook copy at runtime unless it is explicitly designated "DEE-safe microcopy".

---

11) Runtime application (how pages adapt)

At render time:

Step 1) Resolve dee_state

- apply G1..G5 signals
- ensure 1 audience, up to 2 intents, constraints as applicable
- set flags (planner_mode, decision_difficulty)

Step 2) Select levers

- snippet_variant_id (audience-based)
- optional intent add-on line (if intent confidence threshold met)
- ordering_profile_id (audience + intent tie-break)
- bundle_prominence level (0..3) based on behavior L1/L2/L3

- hook priority order (template-defined)
- review theme weights (audience-based)
- proof_style emphasis (near-page proof_style or decision_difficulty)

Step 3) Render (client/server/edge)

- No regeneration occurs at runtime.
- Only presentation changes occur.

---

12) What “boost visibility”, “show a prompt”, “intent strengthened” mean (UI concrete)

12.1 Boost bundles visibility (bundle_prominence 0..3)

Allowed actions (template-defined only):

- subtle highlight of bundle header
- move bundle jump link higher
- show “Save this plan” inline teaser
- on return from bundle landing: show a “Continue planning” chip or auto-scroll to bundles anchor (one-tap)

12.2 1-line confirmation prompt (Level 2)

Definition:

- a non-blocking microcopy line under the bundles header (City/TTD) or small sticky bar (mobile)
- includes max 2 lightweight actions

Example structure:
Line: “Planning 2–3 days? See a ready plan matched to your pace.”
Actions: [View plans] [Not now]
No modal, no interruption.

12.3 Intent strengthened (e.g., Top Tickets clicked)

Definition:

- bounded intent confidence increases
- enables:
  - tie-break weight adjustments
  - optional intent add-on line
  - proof blocks aligned to that intent (ticket certainty, entry clarity)

It does NOT:

- remove items
- change SEO content
- change indexable URLs

---

13) Locale & language priors (soft, configurable; discovery-validated)

These priors are hypotheses and must be validated via analytics in discovery.
They are used only as tie-breakers with capped influence.

13.1 Country priors (initial 20+ set; to validate)

Country	Soft preference profile (examples)
United States (US)	convenience + bundles + premium comfort; private upgrades accepted
Canada (CA)	similar to US; clarity + comfort
United Kingdom (UK)	culture + walking; guided highlights acceptance

Ireland (IE)	culture + walking; value clarity
Germany (DE)	strong value_for_money + structure + inclusions clarity
Austria (AT)	culture/museums + clarity; comfort/value balance
Switzerland (CH)	premium comfort + punctuality; smaller groups
Netherlands (NL)	independent travelers; logistics clarity; smaller groups
Belgium (BE)	culture + museums; value clarity
France (FR)	culture + premium narrative; quality framing matters
Spain (ES)	social/group orientation; iconic highlights
Italy (IT)	iconic highlights + ticket certainty; crowd management
Portugal (PT)	value + relaxed pacing; scenic experiences
Poland (PL)	group/value; price clarity; simplified choices
Czechia (CZ)	value + group; practical details

Denmark (DK)	small groups; quality signals; calm pacing
Norway (NO)	small groups + outdoor; quality-first
Sweden (SE)	small groups; clarity; sustainability signals
Finland (FI)	calm pacing; clarity; comfort
Turkey (TR)	family/group; confirmation clarity + trust blocks
Israel (IL)	family planning + clarity; reliability proof
UAE (AE)	private upgrades + convenience; premium comfort
India (IN)	value + group; inclusions clarity; confirmation clarity
Japan (JP)	punctuality + details; private/small groups; reliability and steps
South Korea (KR)	convenience + trendy highlights; private/small groups often preferred

China (CN)	iconic + photo moments; certainty of entry; step-by-step clarity
Taiwan (TW)	quality + calm pacing; smaller groups
Brazil (BR)	highlights + social; value clarity
Mexico (MX)	highlights + value; simple decision cues

13.2 Language priors (soft)

Language	Soft preference profile
Deutsch	structure, inclusions/exclusions, value justification
Polski / Čeština / हिंदी	value clarity + group options + lower cognitive load
日本語 / 한국어	reliability, punctuality, private/small groups, detailed “how it works”
简体中文 / 繁體中文	iconic highlights, photo moments, certainty of entry, clear steps
ישראל / العربية	family logistics + private upgrades, reliability proof
Español / Italiano / Português	social highlights + clear availability framing

Русский / Türkçe	logistics and confirmations clarity; trust blocks earlier

13.3 How priors are applied (bounded)

Locale/language priors may nudge:

- ordering tie-break weights (max small effect)
- bundle prominence initial nudge (max small effect)
- proof_style preference (decision_shortlist vs discovery_first)
- hook prominence preference (save plan / safety net)

They must never override explicit query/near-page/filter signals.

---

14) IP/Geo signals (GDPR-safe)

DEE may use GeoIP to derive:

- country_code
- inferred locale

Rules:

- never store raw IP
- store only derived country_code for analytics and routing
- user language selection overrides inferred language

---

15) Interaction with near-pages, FAQs, and dynamic near-page creation

15.1 Near pages (runtime)

DEE uses near-page theme seeding (Section 7). Near-page theme is the strongest non-explicit signal.

15.2 Dynamic near-pages creation (offline; not runtime DEE)

Every 5 weeks (post-publish observation window):

- system analyzes GSC queries + performance gaps
- generates new decision near-pages aligned to intents/audiences

DEE does not “create” pages at runtime; it only ensures that once a near page exists, the session stays consistent with its theme.

15.3 FAQ adaptation (variant-based, not runtime-generated)

FAQ content is authored/generated offline and stored in CMS.

DEE may:

- choose which FAQ set to surface first (theme-based)
- reorder FAQ entries by audience relevance (if allowed)

DEE must not generate new indexed FAQ content at runtime.

---

16) Interaction with bundles (planner behavior)

DEE supports planner-mode without rewriting the site.
Bundles are not a separate funnel; they are a “plan-first” decision surface.
DEE behaviors are defined via L1/L2/L3 (Section 9.3) and concrete UI levers (Section 12).

---

17) Analytics & feedback loop (MVP-safe)

DEE must log:

- dee_state (audience/intents/constraints + confidence)
- sources used (nearpage theme, query patterns, filters, locale priors)
- variant IDs shown (snippet_variant_id, ordering_profile_id, hook_ids)
- user actions (bundle L1/L2/L3, filters, CTA clicks)
- outcomes (CTR to product, add_to_cart, booking, bounce)

Governance:

- No auto-learning ranking in Phase 1.
- Improvements happen via controlled tuning tables (Product Ops) and rules updates.
- Any changes to mapping rules must be versioned and auditable.

---

18) Roadmap (high level)

Phase 1 (MVP):

- rule-based dee_state
- audience snippet variants
- ordering profiles
- bundle prominence L1/L2/L3
- hook priority selection
- near-page theme seeding

Phase 2:

- richer explicit quiz/travel-style chips
- persistent profiles (opt-in)
- richer review theme summaries per audience
- expanded near-page theme library (more granular but governed)

Phase 3:

- predictive intent improvements (still governed)
- seasonal intent shifts
- deeper personalization across journeys (without breaking index safety)

UX PRINCIPLES FOR AN ADAPTIVE INTERFACE

DEE RENDERING & UX PRINCIPLES — Adaptive Interface Implementation (Model C) — SSOT v1.0

(How pages render adaptively, without breaking SEO stability or trust)

Scope

This section defines:

- UX principles for adaptive behavior (what is allowed vs forbidden),

- a concrete rendering plan (server/client/edge behavior),

- the allowed “DEE levers” per template and how they apply in-session,

- traceability (why something changed) without user-facing “personalized for you”.

Important: Nothing below replaces Model C selection rules in 4.8.

DEE consumes Model C outputs and adapts presentation only.

UX PRINCIPLES FOR AN ADAPTIVE INTERFACE (Rosotravel DEE)

X.01 Core UX Philosophy

Rule

Rosotravel does not change what the product is. Rosotravel changes how it is presented to match traveler intent.

What this means in practice (dev-usable)

- Selection is upstream (Model C): which products exist in Top Picks / Rails / Bundles is decided before DEE.

- DEE may only adjust:

(a) which variant microcopy is shown (snippet, reassurance line, review-theme summary),

(b) ordering inside already-approved lists (no list membership change),

(c) emphasis (badges, small helper text, which CTA is visually primary),

(d) which “decision proof” block is surfaced first (where multiple proof blocks exist).

Interface must feel:

- calm, not reactive
- helpful, not intrusive
- adaptive, not “personalized in an obvious way”

User perception goal:

“This page highlights what matters to me — without asking me to configure anything.”

X.02 Adaptation Must Be Subtle and Non-Disruptive

Hard UX constraints

Adaptive behavior must never:

- reload the page
- visibly “switch modes” for the user (no abrupt reshuffle after they started reading)
- show banners like “We customized this for you”
- break content hierarchy or layout stability

Adaptation happens through:

- emphasis (what is visually primary)
- ordering (stable membership, deterministic order)
- selection of pre-generated variants
- summarization (precomputed review clusters; audience-weighted)

Not through:

- hiding core information

- removing sections

- changing URLs

- changing canonical content

Implementation rule (rendering)

- DEE changes may apply immediately only when the user performed an explicit action (filter, Magic Finder choice, near-page theme landing).

- If DEE wants to apply behavior-based improvements (non-explicit), it applies them on the next navigation within the session (or via non-disruptive highlight only), not by reordering content the user is currently reading.

X.03 One Page, One Stable Structure

Every landing page type has:

- a fixed section order

- a predictable layout

- consistent interaction patterns

DEE can only:

- reorder items within a section (if user has not “committed” to reading state; see X.05)

- swap copy variants within a defined slot

- change which blocks are highlighted first (within the same fixed structure)

DEE must never:

- add/remove core sections

- change H1/H2 structure

- create layout jumps between users

Why

This guarantees:

- usability

- trust

- SEO stability

- cognitive comfort

Concrete rule for dev

Each template must define a “DEE Slot Map”:

- FIXED sections: cannot move/remove

- DEE slots: only these can change (variant/ordering/emphasis)

If a slot is not listed, DEE is not allowed to touch it.

X.04 Adaptation Layers (UX Hierarchy)

Adaptation is applied in layers, from least to most visible:

Layer 1 — Micro-copy adaptation (lowest risk)

- snippets (audience variants)

- short reassurance lines (“what to expect”)

- review theme summaries (“what guests loved”)

Layer 2 — Ordering adaptation (medium risk)

- reorder cards inside a list

- reorder rails within the rail set (only if rail set itself is fixed in count and purpose)

Layer 3 — Emphasis adaptation (medium risk)

- highlight which CTA is primary

- add small helper labels

- prioritize which trust component appears first (when multiple exist)

Layer 4 — Recommendation adaptation (highest risk)

- reorder related items

- reorder “people also booked”

- reorder suggested next steps

MVP rule

Only one or two layers should be active at the same time.

Recommended MVP combo:

- Layer 1 + Layer 3

Optionally Layer 2 only when:

- user has not scrolled past the section OR

- user triggered an explicit segment/intent action.

X.05 No Surprises Principle

Users should never feel:

- manipulated

- “tested on”

- confused why something is shown

Therefore:

- The same product always has the same core description (SEO body stays stable).

- Adaptation highlights different angles, not different facts.

- Conflicting messages across sessions are not allowed.

Determinism rules (dev)

- DEE ordering must be deterministic given the same dee_state.

- Once a list is rendered and the user scrolls into it, its order is “locked” for that page view (no reshuffle).

- If intent cannot be inferred confidently → fall back to neutral presentation (X.06).

Adaptive Recommendations (keep, but constrain)

“Travelers like you also booked…” may exist, but in MVP:

- list membership is fixed per page type (precomputed candidates)

- DEE only reorders, never injects new unknown items

- do not show this block above primary decisions (Top Picks / Our 2 picks)

Adaptive Writing Tone (keep, but clarify)

Tone does not change at runtime.

Tone differences exist only in pre-generated microcopy variants (snippet, short helper lines).

SEO body and long descriptions remain stable and governed by locks.

X.06 Default Is Always Neutral

Every adaptive element must have:

- a default variant used when confidence < threshold

Default characteristics:

- balanced tone

- general traveler focus

- no segment-specific language

Rule

If in doubt, show default.

Concrete thresholds (initial; tuneable)

- audience_confidence >= 0.65 to use non-default audience snippet

- intent_confidence >= 0.75 to show an “intent add-on line”

- constraints only apply if explicit (filters / query / theme)

X.07 Human Face Is an Anchor, Not Decoration

Human elements (guides, last tour snapshot, expert quotes) are:

- trust anchors

- reality checks

- grounding elements for adaptive content

Rules

- Human elements stay consistent across segments (they exist regardless of audience).

- Adaptation can change which human insight is highlighted first (e.g., which review theme), not invent people.

- No fake personalization of people (“your guide”) unless factual.

- Human presence must feel factual, earned, calm — not promotional.

Rendering implication

Human blocks must be:

- server-renderable from factual data (no runtime AI generation)

- consistent in layout and placement across segments (slot may change emphasis, not existence)

X.08 Adaptation Must Be Explainable Internally

Every adaptive outcome must be:

- traceable to signals

- explainable to a human reviewer

- reversible

Implication

No black box behaviors. Deterministic rules in MVP.

Dev requirement

Each page render generates an internal trace object:

```
dee_decision_trace = {

dee_state: {audience, intents, constraints, confidence},

seed_sources: [...],

levers_applied: [

{slot_id, action: "variant_selected", value, reason_codes[]},

{slot_id, action: "ordering_profile", value, reason_codes[]},

{slot_id, action: "emphasis", value, reason_codes[]}

],

locked_lists: [...],

timestamp

}
```

This trace is visible only in debug mode (admin/ops), never to users by default.

X.09 Adaptation Should Reduce Cognitive Load

Goal: less thinking, not more options.

Adaptive UX should:

- reduce scrolling
- reduce comparison effort
- surface reassurance early
- answer objections sooner

If adaptation adds complexity → it is wrong.

Concrete targets

- City/TTD “more options” must remain controlled (Phase B).
- POI PACK must fit one screen (no vertical scroll) by design.
- When “decision difficulty” behavior detected → show explanation/trust components earlier, not bigger lists.

X.10 Progressive Disclosure Over Aggressive Personalization

Rosotravel should reveal relevance progressively:

1) neutral overview
2) relevant highlights
3) tailored recommendations
4) optional planning tools

Never start with:

- quizzes
- popups

- forced segmentation

Rule

DEE earns attention gradually; the UI never forces it.

X.11 Respect Reading vs Booking Modes

Users arrive in different mental states:

- browsing
- comparing
- planning
- booking

Adaptive UX must:

- not force booking signals too early
- not block reading with conversion elements
- allow “save / come back later” paths

Bundle behavior (planner mode) must be implemented as emphasis + prompts, not as a page “switch”.

Wishlist, email-to-self, and bundles support this principle.

X.12 Adaptation Is Session-Based in MVP

In MVP:

- adaptation applies per session
- no strong cross-session assumptions
- no hard user profiles

Consistency rule

Within one session, behavior must stay consistent.

Between sessions, neutral reset is acceptable.

Storage rule

- store only dee_state tokens (no PII)
- expiry: end of session (or short TTL if using cookies)

X.13 Success Metric for Adaptive UX

Adaptive UX is successful if:

- users scroll less but read more
- fewer comparisons are needed
- more users save/wishlist before booking
- fewer repeated questions in Q&A
- booking confidence increases

Not measured by:

- novelty
- complexity
- number of adaptive rules

X.14 Anti-Patterns to Avoid

Do not implement:

- visible “personalized for you” labels
- sudden layout changes
- segment-specific jargon

- hard gating based on assumed intent

- conflicting recommendations on refresh

Heuristic:

If adaptation feels clever → probably wrong.

If it feels obvious → probably too aggressive.

If it feels natural → correct.

RENDERING PLAN — How DEE is applied without layout jumps (dev plan)

R.01 “DEE Slot Map” (mandatory per template)

Each template must declare:

A) Fixed sections (cannot be removed/moved)

B) Allowed DEE slots (variant/ordering/emphasis)

Example slot types:

- SLOT_SNIPPET (variant swap)

- SLOT_INTENT_ADDON_LINE (optional; only if confidence)

- SLOT_LIST_ORDER (ordering profile, stable membership)

- SLOT_BLOCK_EMPHASIS (visual priority: primary CTA, highlight)

- SLOT_TRUST_COMPONENT_ORDER (which trust micro-block appears first)

- SLOT_REVIEW_THEME (which 3–4 bullets to show first)

- SLOT_HOOK_PRIORITY (which hook appears in the allowed position)

- SLOT_BUNDLE_PROMINENCE (planner levels L1/L2/L3; emphasis only)

If a change is not in Slot Map → DEE cannot do it.

R.02 Rendering strategy (Server-first for stability, client for session updates)

Goal: no visible “switch”, no reload, index-safe.

We use a 2-step approach:

1) Initial render (SSR / server-side render)

- server computes dee_state_seed from strongest signals available at request time:

- near-page theme (hard)

- explicit query/slug mapping (strong)

- UTM tags (explicit)

- selected UI language (explicit)

- geo country code (soft)

- server selects initial variants + ordering profiles

- server renders the page with the chosen variant content in DEE slots

2) In-session refinement (client-side, non-disruptive)

- client tracks behavior signals (bundle interactions, clicks)

- client updates dee_state (confidence + flags), stores in session

- client applies ONLY non-disruptive levers:

- highlight bundles / show save-plan teaser

- show 1-line prompt (planner L2) inside an already reserved slot area

- update hook priority for the next section below fold

- never reorder a list the user is currently viewing (list-lock rule)

Navigation rule

On navigation to another page in the same session:

- server reads dee_state token (cookie/session header) and renders the new page already aligned.

This avoids visible “switching” on the new page.

R.03 List-lock rule (prevents “I was reading and it moved”)

A list becomes “locked” for the page view when:

- it enters viewport for the first time OR

- user interacts with any card inside it

Once locked:

- DEE cannot reorder that list during this page view.

- DEE may still change emphasis (e.g., add a small label), but not reorder.

Implementation

- front-end emits events: list_viewed(list_id), list_interacted(list_id)

- client marks locked_lists[list_id] = true

- server does not need to know locks; locks are per-page-view.

R.04 “Boost” and “Prompt” (exactly what it means in rendering)

When we say “boost bundles visibility” (planner behavior), it means:

Allowed visual actions (no structural change):

- add a subtle highlight style to the Bundles section header

- show a reserved one-line “save plan” teaser directly under Bundles header

- elevate bundle jump link in the jump links row (same row; no new row)

- on return from bundle landing: show a small “Continue planning” chip near the top (same location reserved for that chip)

When we say “show a prompt”, it means:

- a non-modal, single-line microcopy rendered INSIDE a reserved slot area (so layout is stable)

- with 1–2 actions (View plans / Not now)

- never as an interrupting popup

R.05 Ordering application (deterministic, membership-stable)

DEE ordering must work like this:

- the backend provides a stable candidate list (membership is fixed)

- DEE selects ordering_profile_id

- ordering sorts the candidates deterministically

Critical constraint

Ordering must never turn a curated shortlist into a catalog.

If a list is “Top Picks = 6”, ordering only reorders these 6.

It does not expand beyond the 6.

R.06 Variants storage (pre-generated, CMS-backed)

DEE variants exist as stored fields:

- snippet variants per audience

- review theme summary variants (derived from clusters)

- optional intent add-on lines (short; UX-only; pre-generated or rule-composed from approved templates)

Runtime rule

DEE selects variants; it does not write new variants at runtime.

SEO safety rule

DEE-controlled microcopy blocks that are not meant for SEO (snippets) should be wrapped with a “non-snippet” marker (implementation detail) so they don’t influence search snippets, while remaining visible to users.

R.07 Near-page theming (required for decision consistency)

Each near page must store:

- nearpage_theme_id

- default audience (optional, recommended)

- intent_bias (0..2)

- proof_style

DEE must treat near-page theme as a session seed stronger than generic city defaults.

Outcome

If user lands on “family in London” near page:

- their session should remain family-consistent across City/TTD/Product surfaces,

without visibly announcing personalization.

R.08 Country/language priors (apply as tie-breakers only)

DEE may apply locale and language priors to:

- tie-break ordering within a list

- bundle prominence nudge (small)

- which reassurance line to show (clarity vs comfort vs logistics)

Hard cap

Priors cannot override:

- near-page theme

- explicit filters

- explicit UTMs

- explicit internal search queries

R.09 Debugging & internal explainability (rendering hooks)

Add a debug mode (admin-only):

- query param: ?dee_debug=1 (or internal toggle)

- shows dee_state, confidence, reason codes, and applied levers

- shows which lists are locked and why

- never visible publicly

R.10 Analytics events required (to tune safely)

DEE must log:

- dee_state_seed (source)

- levers applied per slot (variant chosen, ordering profile chosen)

- list locks

- planner L1/L2/L3 triggers

- outcomes: CTR, add_to_cart, booking, bounce, time-on-page

This is required to tune priors and mappings during discovery.

R.11 What we consciously reject (so it is not implemented by accident)

We explicitly reject:

- runtime rewriting of SEO body text (“AI rewrites at runtime”) in MVP

- any “personalized for you” banners

- page reloads or layout mode switching

- reshuffling lists while user is reading them

If these are desired later, they require a new governance pass and a new spec version.

02. Human-Face and Live Experience Layer

1) Purpose

The Human-Face and Recent Experience Layer is the trust and authenticity layer of Rosotravel.

Its purpose is to make experiences feel:

- real,

- recent,

- human,

- and operationally alive,

before the traveler books.

This layer exists to support the three core Rosotravel promises:

A) Expert Decision Platform

Users are not looking at a random catalog. They are looking at a curated choice made by Rosotravel.

B) Human Face

There are real people, real guides, and real standards behind the experience.

C) Recent Experience Proof

The experience is not “timeless catalog copy”; it is active, recent, and validated by fresh traveler signals.

This layer must strengthen:

- trust,

- conversion confidence,

- premium brand perception,

- and differentiation vs anonymous OTAs.

2) Core principle

Rosotravel must never make a stronger factual claim than the data can support.

Therefore:

- if exact departure-level linkage exists, we may communicate exact recent evidence carefully,

- if only review-date or media-date evidence exists, we must communicate it as recent guest/review/media proof,

- if product-level recent proof is weak, we must fall back to attraction-level or city-level recent proof.

This is a truthfulness rule, not a UX preference.

3) What this layer is NOT

This layer is NOT:

- fake “live” marketing,

- invented human storytelling,

- exact per-tour reconstruction unless exact data exists,

- decorative guide avatars without logic,

- or “social proof noise”.

It must be:

- factual,

- recent,

- calm,

- premium,

- and clearly connected to the product / attraction / city.

4) Strategic role in the Rosotravel brand

This layer is one of the three major differentiators of Rosotravel:

1. Expert Decision Platform

Rosotravel curates the portfolio and explains why a traveler should choose these options.

2. Dynamic Experience Engine (DEE)

Rosotravel adapts presentation and emphasis to the traveler’s audience, intent, and constraints.

3. Human-Face and Recent Experience Layer

Rosotravel proves that the experience is real, recent, and backed by people.

Together, these three layers make Rosotravel feel:

- curated, not crowded,

- human, not anonymous,

- current, not stale.

5) Human-Face and Recent Experience Layer — structure

This layer is composed of four connected components:

A) Guide / Host Anchor

B) Latest Tour Snapshot

C) Trip DNA

D) Recent Activity / Transparency blocks across City and Attraction pages

All four components may coexist on the same product page, but they have different data rules and different UX roles.

A) Guide / Host Anchor

---

6) Purpose of the Guide / Host Anchor

The Guide / Host Anchor answers:

“Who stands behind this experience?”

Its purpose is to:

- humanize the product,

- create premium trust,

- reduce marketplace anonymity,

- and connect the traveler with Rosotravel standards.

This anchor is not required to prove who led the latest exact departure.

It is a stable human reference attached to the product / city / guide team.

7) Allowed Guide / Host Anchor modes

The system must support three modes:

Mode A — Exact Guide Known

Use only when the guide identity is genuinely tied to the product or departure logic.

Displayed elements:

- guide photo
- guide name
- short role / specialization
- optional short quote

Example:

“Hosted by Anna K., local historian”

Mode B — Guide Team Anchor (recommended default for many products)

Use when a product belongs to a Rosotravel team or city guide team, but exact departure guide is not guaranteed.

Displayed elements:

- guide team photo or representative guide photo
- team label
- short trust statement

Examples:

- “Hosted by our London guide team”
- “Led by Rosotravel guides in Rome”

Mode C — Curated by Rosotravel (partner / non-guide-specific)

Use when no reliable guide identity can be attached to the product.

Displayed elements:

- Rosotravel curator / standards label

- short statement explaining selection logic

Examples:

- “Curated by Rosotravel”

- “Selected against our quality standards”

8) UX rules for Guide / Host Anchor

Rules:

- this block must feel factual, not decorative,

- it must not imply the exact person will lead the booked tour unless that is operationally true,

- it must remain calm and premium,

- it must not use manipulative language like “your guide” unless already confirmed at booking level.

Allowed fields:

- guide / team photo

- guide / team name

- specialty or city expertise

- optional guide quote

- optional CTA to Guide / Experts page

Forbidden claims unless exact data exists:

- “This guide led the latest tour”

- “Your guide”

- “Meet the guide from this exact day”

9) Recommended front-end position

Product page:

- directly under hero or integrated into the hero trust strip

City page:

- optional smaller version inside “This week in {{City}}”

Attraction page:

- optional if strongly relevant

Near pages:

- optional, only when connected to a real expert page or city guide team

---

B) Latest Tour Snapshot

---

10) Purpose of Latest Tour Snapshot

The Latest Tour Snapshot is the key “recent proof” component.

It answers:

“What have travelers been experiencing recently around this product?”

Its purpose is to create the feeling that:

- this experience is active,

- real guests are sharing real signals,

- and Rosotravel surfaces recent, human evidence before booking.

11) Important conceptual decision

The Latest Tour Snapshot is not necessarily an exact “single departure snapshot”.

For MVP and Phase 1, it is a date-clustered recent proof component.

This means:

- signals are grouped by date,

- the system selects the freshest valid cluster,

- the traveler sees a coherent recent snapshot,

- but the UI must not falsely claim that every photo/comment/guide element comes from one exact departure unless that linkage is known.

12) Snapshot truth levels (internal logic)

The system must support three source confidence tiers internally:

Tier A — Verified Experience-Date Snapshot

Use when exact booking/departure linkage exists.

Typical source:

- Rosotravel internal reviews tied to a booking / departure

Possible confidence:

highest

Tier B — Review-Date / Media-Date Snapshot

Use when the strongest available evidence is:

- review submission date,

- review publish date,

- media upload date,

for a given product.

Typical source:

- Viator/external reviews

- external media attached to reviews

Possible confidence:

medium-high

Tier C — Recent Window Snapshot

Use when exact day-clustering is weak or insufficient.

Typical source:

- last 7 days / last 30 days / current month summary

Possible confidence:

fallback

Important:

The UI may still use the same component label “Latest Tour Snapshot”, but the wording inside must match the truth level.

13) Snapshot cluster logic

For each product, the system builds date-based clusters:

snapshot_cluster(product_id, snapshot_date)

A cluster may include:

- reviews associated to that date,

- media associated to that date,

- derived review themes,

- optional guide/team anchor,

- optional freshness indicators.

Date selection rules:

- for internal reviews with booking linkage: use experience_date
- for external reviews: use review publish/submission date
- for external media: use media upload date
- if no strong date-cluster exists: build a recent-window aggregate

Minimum valid cluster requirement:

A cluster should contain at least one of:

- 1+ review
- 1+ media item

Preferred:

- review + media together

Guide/team anchor alone is not enough to create a snapshot cluster.

14) Snapshot selection logic

At render time:

1. Prefer Tier A cluster if available and sufficiently recent.
2. Else use Tier B cluster.
3. Else use Tier C recent-window summary.
4. If no valid product-level snapshot exists:
   - fall back to attraction-level recent proof (if available),
   - else city-level recent proof,
   - else hide the component.

15) Why this model is correct

This model gives the user:

- recency,

- visual freshness,

- emotional proof,

- and confidence,

without making unprovable exact-tour claims.

It creates a strong “this is happening now” feeling while remaining factually safe.

16) Allowed UI labels for the component

Primary recommended label:

- Latest Tour Snapshot

Alternative acceptable labels:

- Latest Traveler Snapshot

- Recent Tour Snapshot

- Recent Guest Moments

Recommendation:

Keep one main brand label across product pages for consistency:

Latest Tour Snapshot

17) Wording rules inside the snapshot

The inner wording must depend on data confidence.

Allowed wording when exact experience-date data exists:

- “Latest Tour Snapshot — 12 March 2026”

- “Guest review highlights from 12 March”

- “Guest media from 12 March”

Allowed wording when only review/media date exists:

- “Latest traveler snapshot — 12 March 2026”

- “Latest review highlights from 12 March”

- “Guest media added on 12 March”

- “What guests shared on 12 March”

Allowed wording when only recent window exists:

- “Recent guest moments”

- “Latest reviews this month”

- “From recent tours”

- “Recent traveler feedback”

Forbidden wording unless exact departure linkage exists:

- “Last tour took place on 12 March”

- “These photos are from the tour on 12 March”

- “Guide Anna led the tour on 12 March”

- “This was the last departure”

18) Snapshot content structure (recommended MVP)

The component should display:

- component title

- date or time-window label

- 2–3 “Guests loved” bullets

- 1–3 media thumbnails (photo / video)

- optional link to recent reviews

- optional side guide/team anchor (separate logic)

Recommended product page layout:

Left side:

- Latest Tour Snapshot title

- date label / freshness label

- “Guests loved” bullets

- recent media thumbnails

- CTA: “See recent reviews”

Right side:

- Guide / Host Anchor

- photo/avatar

- role/team label

- short trust line

- optional CTA: “Meet our guides”

This creates a strong premium block without claiming too much.

19) What “Guests loved” means

“Guests loved” is not a quote block.

It is a summary of clustered review themes from recent reviews in the selected cluster.

Examples:

- “smooth entry”

- “guide’s stories”

- “manageable pace”

- “small-group feel”

- “great family flow”

Rules:

- max 2–3 bullets in the snapshot block

- calm, factual wording

- no exaggerated marketing language

- derived only from real review data

20) Photos and video for snapshot

Primary source (recommended):

- customer-uploaded media attached to reviews

Secondary source:

- review-linked external media (if available and permitted)

Tertiary source:

- admin-uploaded recent media tied to the product and date cluster

Guide profile images must NOT be treated as “snapshot media”.

Guide images belong to the Guide / Host Anchor component.

21) Date display logic

If exact internal experience date exists:

- exact date may be shown in the snapshot label

If only review-date/media-date exists:

- date may still be shown,
- but wording must reference review/media/freshness, not departure

If data is weak:

- use month or recent-window label instead of exact day

Examples:

Strong:

- “Latest Tour Snapshot — 12 March 2026”

Medium:

- “Latest review highlights — 12 March 2026”

Fallback:

- “Recent guest moments — March 2026”
- “Latest reviews this month”

22) What we reject in MVP

We explicitly reject:

- exact public departure reconstruction from weak data
- exact “group composition” claims like “6 couples, 2 solo”
- “mood score” as a public standalone metric unless methodology is clear and stable

- guide identity claims tied to a specific day without operational proof

23) Traveler mood / satisfaction treatment

Instead of public “mood score”, use:

- “Recent satisfaction”

- “Highly rated in recent reviews”

- or simply let “Guests loved” do the work

If a score is shown, it must be clearly tied to review methodology and recent window.

Otherwise, do not show it.

24) Group size / group composition treatment

If exact group size per verified departure exists:

- may be used later

For MVP:

- do not show exact public numbers unless verified

- prefer derived signals:

  - “small-group feel”

  - “popular with couples”

  - “works well for families”

These should come from review themes or curated product metadata, not invented assumptions.

C) Trip DNA (Experience Profile)

25) Purpose of Trip DNA

Trip DNA helps users understand the nature of an experience quickly.

It supports the decision platform by reducing reading effort and comparison burden.

Trip DNA answers:

- What kind of experience is this?
- How demanding is it?
- Who is it best for?
- What is the emotional / practical profile of the tour?

26) Recommended Trip DNA dimensions

Recommended standard dimensions:

- Walking level
- Crowd level
- Local vs iconic balance
- Pace / emotional intensity
- Best for

Example:

- Walking: moderate
- Crowds: busy
- Local vibe: mixed
- Pace: calm
- Best for: first-timers, history lovers