


27) Rules for Trip DNA

- fixed order across the site
- compact visual system
- same dimensions everywhere
- values derived from product metadata, itinerary, reviews, and controlled AI proposals
- human-editable in CMS
- never generated live at runtime

Trip DNA is not part of the snapshot.

It is a stable product-profile component.



D) Real Transparency Across Pages



28) Product page

Product page should include the strongest trust stack:

1. Guide / Host Anchor
2. Latest Tour Snapshot
3. Trip DNA
4. Review section with recent emphasis
5. “Why these picks” / trade-offs decision layer (separate system)

29) City page

City pages should include:

- “This week in {{City}}” or fallback “Recently in {{City}}”

Possible metrics:

- tours run
- travelers hosted
- satisfaction signal
- active guide names
- recent media strip
- recent guest themes

Purpose:

to make Rosotravel feel present and operationally alive in the city.

30) Attraction page

Attraction pages should include:

- “Recent traveler moments at {{Attraction}}”

Possible elements:

- recent themes
- optional media strip
- review freshness emphasis
- wording such as:

“From recent tours that included {{Attraction}}”

This is safer and stronger than pretending a single exact last visit exists.

31) Other surfaces

Optional extensions:

- email / CRM:
 - “Recently in Rome”
 - “Latest guest moments this month”
- expert pages:
 - guide photos and city guide team presence
- Phase B / Magic Finder:
 - not primary place for snapshot; use only lightly if useful

Data model and source-of-truth rules

32) Review entity requirements

Recommended fields:

- review_id
- linked_product_id
- linked_city_id
- linked_attraction_ids[]
- source_type (internal / viator / external)
- submitted_at
- published_at
- experience_date (nullable but strongly recommended for internal)
- rating
- review_text

- media_ids[]
- verification_status

33) Media entity requirements

Recommended fields:

- media_id
- source_type (customer_review / admin_snapshot / guide_profile / partner / stock)
- linked_product_id (nullable)
- linked_city_id (nullable)
- linked_attraction_id (nullable)
- submitted_at
- experience_date (nullable)
- moderation_status
- media_kind (photo / video)

34) Guide entity requirements

Recommended fields:

- guide_id
- guide_name
- guide_photo
- guide_city_ids[]
- guide_specialties[]
- guide_quote
- profile_url
- is_active

35) Snapshot entity (new required object)

Recommended object:

recent_experience_snapshot

Fields:

- snapshot_id
- entity_type (product / attraction / city)
- entity_id
- source_tier (verified_experience_date / review_date / recent_window)
- snapshot_date
- date_precision (exact_day / month / window)
- review_ids[]
- media_ids[]
- guide_anchor_id (nullable)
- summary_topics[]
- freshness_score
- trust_confidence_score
- proof_label_variant_id

36) Snapshot generation rules

Generation job should:

- group reviews and media by date
- create snapshot candidates
- select freshest valid candidate
- store fallback recent-window snapshot if exact day is weak
- generate topic summaries from real review content only

- never invent guide/date/media linkage

Wording and compliance policy

37) Truthfulness matrix

Rule:

The wording must match the strongest verified data level available.

A) Exact experience date known

Allowed:

- "Latest Tour Snapshot — 12 March 2026"
- "Guest review highlights from 12 March"
- "Guest media from 12 March"

B) Review/media date only

Allowed:

- "Latest traveler snapshot — 12 March 2026"
- "Latest review highlights from 12 March"
- "Guest media added on 12 March"

C) Only recent aggregate

Allowed:

- "Recent guest moments"

- "Latest reviews this month"
- "From recent tours"

Forbidden unless exact departure proof exists:

- "Last departure"
- "Tour took place on..."
- "Guide X led this tour on..."
- "Photos from the last tour"

38) Editorial tone

This layer must sound:

- factual
- calm
- premium
- human
- current

It must not sound:

- sensational
- gimmicky
- overly "social media"
- fake-live
- over-precise without proof

Rollout plan

39) MVP

Implement:

- Guide / Host Anchor (Guide / Team / Curated-by modes)
- Latest Tour Snapshot using date-clustered review/media logic
- Trip DNA
- City-level “This week / Recently in {{City}}”
- Attraction-level recent traveler moments
- customer review media as main snapshot media source
- guide profile images from guide profiles
- internal snapshot object generation

40) Phase 1.5

Add:

- admin-uploaded snapshot media tied to product/date
- stronger city/attraction-level aggregation
- improved proof labels by source confidence
- better guide anchor linking

41) Later phase

Add:

- exact guide-to-departure linkage where operationally possible
- exact booking/departure-driven proof for more products
- richer recent activity dashboards
- optional advanced CRM/email recent-proof blocks

2.5 Internal linking rules by page type (manual/UX rules)

01. Header & Footer

a. Landing pages description:

➤ [link](#)

b. Header (only sales elements)

➤ Destinations

- /destinations/
- /destinations/{country}/
- /destinations/{country}/{region}/
- /destinations/{country}/{region}/{city}/
- /destinations/{country}/{region}/{city}/things-to-do/
- /destinations/{country}/{region}/{city}/attractions/{attraction}/
- /destinations/{country}/{region}/{city}/tours/{tour-slug}/
- /destinations/{country}/{region}/{city}/packages/{package-slug}/

➤ Experiences

- /experiences/
- /experiences/{category}/
- /experiences/{category}/{country}/
- /experiences/{category}/{country}/{city}/

➤ Travel Guides

- /travel-guides/
- /travel-guides/{continent}/
- /travel-guides/{country}/
- /travel-guides/{city}/
- /travel-guides/{topic}/

➤ Experts

- /experts/
- /experts/{country}/
- /experts/{country}/{city}/
- /experts/{country}/{city}/{expert-slug}/

➤ Global Attractions Map

- /explore/global-map/

➤ Sign In

- /sign-in/

CTA:

- Start planning / Search Trips (search bar)
 - /start-planning/

- c. Footer (brand + E-E-A-T + support))

About Rosotravel

- Who We Are -> /about/who-we-are/
- Mission & Values -> /about/mission-values/
- How We Work -> /about/how-we-work/
- Sustainability -> /about/sustainability/
- Careers -> /careers/
- Press & Media -> /press/

Explore

- Destinations
- Experiences
- Travel Guides
- Experts
- Creators
- Explore Map

- /destinations/
- /experiences/
- /travel-guides/
- /experts/
- /creators/
- /explore/global-map/

For Businesses

- Travel Agents / DMC
- Affiliate Program
- Content Creator Program

- /b2b/travel-agents/
- /b2b/affiliate/
- /b2b/creators/

Support

- Contact Us
- Help Center
- Manage Booking
- Refund Policy
- Safety Standards

- /contact/
- /help-center/
- /manage-booking/
- /refund-policy/
- /safety-standards/

Legal & Technical

- Terms & Conditions
- Privacy Policy
- Cookies
- Sitemap
- AI Sitemap
- API Feed
- /legal/terms/
- /legal/privacy/
- /legal/cookies/
- /sitemap/
- /ai-sitemap/
- /api-feed/

Internal Linking Rules by Page Type

This is where navigation meets SEO and semantic layer.

1. Country page → Region, City, Things-to-do, Guides

Each **Country** page should automatically include:

- “Top regions in {{Country_Name}}”

- from DB: regions with highest number of products or manually flagged `is_featured`.
- “Top cities in {{Country_Name}}”
 - from DB: cities under that country, ranked by product_count or revenue.
- “Featured experiences”
 - from semantic layer: top categories (e.g., Day trips, Family-friendly) for that country.
- “Travel guides for {{Country_Name}}”
 - from travel_guides table: country-level guides.

2. Region page → City, Attraction, Guides

Each **Region** page links to:

- Cities in this region.
- Highlights (attractions aggregated from cities).
- Region-level travel guides.

3. City page → Things-to-do, Attractions, Products, Guides

Each **City** page automatically includes:

- CTA to `/destinations/{country}/{region}/{city}/things-to-do/`
- List of **top attractions** in this city.
- List of **featured tours** (based on bookings/revenue).
- “Travel guide to {{City_Name}}” if exists.
- “Nearby cities” from poi_vectors / geo.

4. Things-to-do in City

The **Things to do in {{City_Name}}** page links to:

- Experience categories (Day trips, Food tours, etc.).
- Top products per category.
- “City guide” and “Best-of” near-pages.
- Link back to main city page.

5. Attraction page

Each **Attraction (POI)** page includes:

- Products (tours) that include this attraction.
- “Nearby attractions” (from poi_vectors & geo).
- Things-to-do in the same city.
- Travel guide sections or blog posts about this attraction.
- “Map view” linking to Explore Map with this POI focused.

6. Product page

Each **Product** page includes:

- Link back to its attraction (if any).
- Link back to city page and things-to-do page.
- Up to 3–5 “Similar tours” (semantic + category-based).
- Link to **package/bundle** if this product is part of a package.
- Link to relevant near-pages (“Best Florence walking tours”).
- Link to city travel guide.

7. Near-page

Each **Near-page** (“Best museums in Rome with kids”) includes:

- Internal links to relevant products.
- Link to city page and things-to-do page.

- Optional link to travel guide if relevant.

8. Travel guide / Blog articles

Every article must link to:

- City / Country pages.
- Attractions it mentions (via schema `about` / `subjectOf`).
- 3–6 strongly relevant products.
- Relevant near-pages (for deeper exploration).

ROUTING TABLE v1 (WHOLE SITE) — Model C — MUST VERIFY IN DISCOVERY

Status: This routing contract must be validated by developers during the 2–4 week discovery phase.

Goal: deterministic routing to curated surfaces only; **no parameter-indexed duplicates**.

0) URL TYPE GLOSSARY (patterns; confirm exact routes in discovery)

URL Type	Pattern	Indexing (default)	Discovery Notes
Homepage	<code>/lang/</code>	INDEX	confirm lang routing + hreflang strategy
Country	<code>/lang/{country}/</code>	INDEX	confirm country pages exist
Region (optional)	<code>/lang/{country}/{region}/</code>	INDEX (if used)	decide if region layer is live

City (City Page)	/lang/country/city/	INDEX	City Page is decision-first
Things-to-do (City)	/lang/country/city/things-to-do/	INDEX	must exist per city
City Category Page	/lang/country/city/category/	INDEX (governed)	avoid param filters; confirm category taxonomy
Attraction (POI)	/lang/country/city/attraction/	INDEX	ensure POI slug governance + dedup
Near Page (keyword slug)	/lang/country/city/keyword-slug/	INDEX (governed)	must follow keyword_map ownership rules
Product Page	/lang/tours/country/city/product-slug/	INDEX	TBD: confirm actual product path
Bundle Page (baseline)	/lang/country/city/itineraries/duration-day/	INDEX	canonical bundle URL
Bundle Page (segment variant)	/lang/country/city/itineraries/segment-duration-day/	NON-INDEX	self-canonical to baseline; no parameter variants

Guides hub	/ {lang} /travel-guides/	INDEX	
City guide	/ {lang} /travel-guides/ {country} / {city} /	INDEX	align with SEO governance
Theme guide	/ {lang} /travel-guides/themes/ {theme} /	INDEX (capped)	avoid thin pages
Experts hub	/ {lang} /experts/	INDEX	E-E-A-T surface
Local experts (optional)	/ {lang} /experts/ {country} / {city} /	INDEX (if complete)	strict caps; avoid thin pages
Global map	/ {lang} /explore/global-map/ (or /explore/world-map/)	INDEX	naming TBD; must be canonical
Partners hub	/ {lang} /b2b/	INDEX (if public)	TBD
Travel agents	/ {lang} /b2b/travel-agents/	INDEX	
Affiliate	/ {lang} /b2b/affiliate/	INDEX	
Creators	/ {lang} /b2b/creators/	INDEX	
Contact	/ {lang} /contact/	INDEX	

Help center	<code>/{{lang}}/help-center/</code>	INDEX	
Manage booking	<code>/{{lang}}/manage-booking/</code>	INDEX	
Refund policy	<code>/{{lang}}/refund-policy/</code>	INDEX	
Safety standards	<code>/{{lang}}/safety-standards/</code>	INDEX	
About hub	<code>/{{lang}}/about/</code>	INDEX	
About subpages	<code>/{{lang}}/about/who-we-are/ /mission-values/ /how-we-work/ /sustainability/ ...</code>	INDEX	confirm full list
Legal	<code>/{{lang}}/legal/terms/ /privacy/ /cookies/</code>	INDEX	
Sitemap	<code>/{{lang}}/sitemap/</code>	INDEX	
AI sitemap	<code>/{{lang}}/ai-sitemap/</code>	INDEX	
API feed/docs	<code>/{{lang}}/api-feed/</code>	INDEX (or restricted)	TBD: exact path + access policy

INDEXING SAFETY DEFAULTS (global)

Rule

Default

Parameter-based URLs **NON-INDEX** (no ?filters=, no UTM canonicals)

Segment variants **NON-INDEX** and self-canonical to baseline URLs

Phase B overlays do not create new indexable URLs

Canonicals always point to clean template URL

1) GLOBAL HEADER ROUTING (Site-wide)

Source (Header)	User intent / input	Resolver logic (deterministic)	Target URL type	Indexing rule	Notes (Discovery)
Logo click	“Go home”	direct	Homepage	INDEX	-
Destinations (menu)	browse geo	user selects country	Country	INDEX	confirm country pages exist
Destinations (menu)	browse geo	user selects city	City Page	INDEX	City Page is decision-first
Destinations (menu)	browse geo	user selects “Things to do in {City}”	Things-to-do	INDEX	must exist per city
Destinations (menu)	browse geo	user selects “Top	Guides or Map or	INDEX	choose 1 canonical

		attractions” (global)	Attractions hub		surface; avoid duplicates
Experience s (menu)	browse by activity	user selects category	Global category hub OR City category chooser	INDEX only if real SEO page	confirm whether global category hubs exist
Experience s (menu)	quick link: Best for first-timers	route to curated surface	Near Page OR Guide OR Category page	INDEX only if governed	must NOT route to parameter filters
Experience s (menu)	quick link: With kids	route to curated surface	Near Page / Category / Guide	INDEX only if governed	prefer near-page decision surface
Experience s (menu)	quick link: Skip-the-line	route to curated surface	Near Page / Category	INDEX (governed)	avoid param filters
Travel Guides (menu)	inspiration	user selects city guide	City guide	INDEX	align with SEO governance
Travel Guides (menu)	inspiration	user selects theme	Theme guide	INDEX (capped)	prevent thin pages
Experts (menu)	trust/human face	“Meet our experts”	Experts hub	INDEX	E-E-A-T

Experts (menu)	trust/human face	"Local experts in {City}"	Local experts page	INDEX if complete	apply caps; avoid thin pages
Global Map (menu)	visual exploration	open map	Global map	INDEX	also acts as discovery surface
Start Planning (search)	typed query + optional dates/style	run intent/entity resolver	Section 2	N/A	must be deterministic

2) START PLANNING / SEARCH ROUTING (Deterministic resolver)

Resolver pipeline (must be implemented and verified in discovery)

Step	Name	What happens	Output
S1	Normalize query	language detection is UI-level only; remove punctuation; canonicalize synonyms	normalized_query
S2	Entity detection	detect entity type; highest confidence wins	entity_type + confidence + resolved entity
S3	Route decision priority	apply deterministic priority rules	target_url_type + target_url

Entity types (search)

Entity type	Meaning
-------------	---------

CITY city entity

COUNTRY / REGION country/region entity

ATTRACTION / POI POI entity

CATEGORY curated taxonomy category

PRODUCT internal/partner product

BUNDLE intent "2 days in Rome", "3-day itinerary Vienna"

GUIDE intent "Rome travel guide", "best time to visit Vienna"

Route decision priority (recommended)

Priority	Condition	Route
1	Exact PRODUCT match	Product Page
2	Exact ATTRACTION match (+ city)	Attraction Page
3	BUNDLE intent (city + duration)	Bundle baseline page (if exists) else Things-to-do
4	Exact CITY match	City Page

5	CATEGORY intent (with city)	City Category Page OR Near Page if keyword_map exists
6	CATEGORY intent (global, no city)	Global category hub (if exists) else prompt city selection
7	GUIDE intent	relevant Guide page
8	Fallback	Search results page (NON-INDEX) with curated suggestions

Search routing table

Search result entity	Example query	Target URL	Indexing rule	Notes
PRODUCT	“Schonbrunn skip the line tour”	Product Page	INDEX	product canonical only
ATTRACTION/POI	“Westminster Abbey”	Attraction Page	INDEX	requires city disambiguation
BUNDLE intent	“2 days in Vienna”	Bundle baseline 2-day	INDEX	segment variant NON-INDEX
CITY	“Vienna”	City Page	INDEX	decision-first
CATEGORY + CITY	“Rome day trips”	/rome/day-trips/ OR near-page slug	INDEX if governed	never parameter filters

CATEGORY (global)	"food tours"	/experiences/food-tours/ (if exists)	INDEX only if real page	else require city selection
GUIDE intent	"things to do in Rome with kids"	guide OR near-page decision page	INDEX	choose canonical surface (avoid duplicates)
Fallback	unknown	Search results (internal)	NON-INDEX	must not create duplicate SEO

3) PAGE TEMPLATE CTAs ROUTING (Key in-page journeys)

CITY PAGE CTAs

Source (City Page)	CTA	Target	Indexing rule	Notes
City hero / navigation	"Things to do in {City}"	Things-to-do	INDEX	primary browse hub
Bundles cards	"View itinerary" / modal CTA	Bundle baseline URL	INDEX	segment shown ≠ new indexable URL
Rails (decision categories)	"Choose by intent"	Near Page OR Category decision page	INDEX if governed	must not create segment/param URLs
Top Attractions block (nav-only)	"View attraction"	Attraction Page	INDEX	if using tours instead, see next row

Top attraction tours block (product cards)	“Check availability”	Product Page	INDEX	dedup vs Top Picks
Phase B trigger	“Show more options”	Overlay (same URL)	NON-INDE X	must not create new URL
Phase B Magic Finder	Topic/POI selection	Results within overlay	NON-INDE X	POI Pack is one-screen
Phase B Browse Catalog	tab switch	same overlay	NON-INDE X	curated band list
City trip planner hook	“Plan my trip”	on-page submission	NON-INDE X	email capture rules apply

THINGS-TO-DO PAGE CTAs

Source	CTA	Target	Indexing rule	Notes
Categories grid	category tile	City Category Page / Near Page	INDEX if governed	avoid param filters
Top attraction tours block	“Check availability”	Product Page	INDEX	band caps apply
Bundles cards	“View itinerary”	Bundle baseline	INDEX	-
Phase B trigger (if present)	“Show more options”	Overlay	NON-INDEX	same Phase B rules as City Page

ATTRACTION PAGE CTAs

Source	CTA	Target	Indexing rule	Notes
Our 2 picks	"Check availability"	Product Page	INDEX	2-of-many decision surface
Tickets-only block	"Check availability"	Product Page	INDEX	ticket-only classification enforced
Bundles including attraction (optional)	"View itinerary"	Bundle baseline	INDEX	optional block
"See more options" (optional)	open POI Pack / Magic Finder	overlay	NON-INDEX	must not explode list

NEAR PAGE / CITY CATEGORY PAGE CTAs

Source	CTA	Target	Indexing rule	Notes
Top list	"Check availability"	Product Page	INDEX	list is curated; ordering can vary
Comparison block	"See details"	Product Page	INDEX	do not create filter URLs
Related categories	tag link	Near page / Category	INDEX	avoid loops/cannibalization

PRODUCT PAGE CTAs

Source	CTA	Target	Indexing rule	Notes
"Check availability"	open availability selector	same page	INDEX	no new URL
"Book now"	checkout	Checkout session URL (private)	NON-INDE X	transactional
"Explore more in {City}"	link	City / Things-to-do	INDEX	internal linking

BUNDLE PAGE CTAs

Source	CTA	Target	Indexing rule	Notes
"Continue to checkout"	add items to cart	Checkout	NON-INDE X	bundle_session_id stored
"Download checklist (PDF)"	email capture	same page action	NON-INDE X	hook engine
"How to book this package"	expand module	same page	N/A	instruction module

EXPERTS / GUIDES CTAs

Source	CTA	Target	Indexing rule	Notes
--------	-----	--------	---------------	-------

Expert profile	“See tours in {City}”	City / Things-to-do	INDEX	trust → conversion
Guides	contextual internal links	City / Category / Product	INDEX	no paid/param links

4) FOOTER ROUTING (Any page → footer)

Footer columns and targets

Footer Column	Canonical targets (examples; confirm in discovery)
A) About Rosotravel	/about/; /about/who-we-are/; /about/mission-values/; /about/how-we-work/; /about/sustainability/; /about/careers/; /about/press-media/
B) Explore	/destinations/ (if exists) OR Destinations hub; /experiences/ (if exists) OR Experiences hub; /travel-guides/; /experts/; /creators/ (if public); /explore/global-map/
C) For Businesses (Partners)	/b2b/travel-agents/; /b2b/affiliate/; /b2b/creators/
D) Support	/contact/; /help-center/; /manage-booking/; /refund-policy/; /safety-standards/
E) Legal & Technical	/legal/terms/; /legal/privacy/; /legal/cookies/; /sitemap/; /ai-sitemap/; /api-feed/

Hard rule: Footer links must always route to canonical pages (no parameter URLs).

5) BREADCRUMB RULES (Canonical hierarchy)

Canonical breadcrumb hierarchy

Home → Country → Region (optional) → City → (Things-to-do / Category / Attraction) → Product/Bundle

Breadcrumb rules

Rule	Requirement
Canonical only	Breadcrumb must always use canonical URLs (no segment variants, no parameter URLs).
Segment bundle handling	If user is on a segment bundle URL, breadcrumb must still point to baseline bundle URL.
Near page linking	Near pages must link back to City and to the relevant Category hub only (avoid cannibal loops).

6) DISCOVERY VALIDATION CHECKLIST (for developers)

Area	Must verify in discovery
URL patterns	exact URL patterns for products and bundles
Global categories	whether global category hubs exist (/experiences/{category}/) or only city-level categories
Canonical / noindex rules	segment variants; Phase B overlays; search results pages

Search resolver	entity types, disambiguation behavior, confidence thresholds, tie-break rules
Experts governance	caps and thin-page prevention logic
Breadcrumb rendering	canonical correctness, correct hierarchy, baseline bundle linking
“No duplicate SEO”	ensure no parameter URLs, no duplicate canonical surfaces for the same intent

2.6 General description of blocks on landing pages (reviews, ask expert, live chat, etc.)

01. General description of blocks on landing pages

a. Reviews module

All review blocks (product, city, attraction, region, country) use the same card layout:

- Avatar
- Name + country
- Date of tour (not date of posting)
- Star rating
- 180–240 char summary
- Thumbnail gallery (if photos)
- “Read more” expand

Data sources: Viator/OTA feeds + internal Rosotravel reviews.

b. Global comments and “Ask our expert”

All landing pages except homepage have:

- Comments thread (user comments)
- “Ask our expert” form (submits a private question)

Product pages additionally have structured Q&A.

Phase 1: answers only from admins/experts, not other users.

c. Global live chat

Live chat (tawk.to) is always shown in the bottom-right corner as a floating widget. It is not part of template content; it is a global layout element.

- Already implemented on the website

02. Global AI Content Framework Rules

This chapter defines **global, cross-template rules** that all AI-generated content must follow. The purpose is to maintain **quality, consistency, non-duplication, factual accuracy, and SEO safety** across 300k+ programmatic pages in 22 languages.

Tone Rules (Global Standard)

Tone Requirements

- Human, expert, conversational.
- “Local guide speaking to a traveler”, not OTA style.
- Neutral, factual, descriptive; avoids hype.
- No clichés or AI phrasing patterns.

Prohibited Phrases (Global Blocklist)

- “hidden gem”, “iconic”, “breathtaking”, “nestled in the heart of...”
- “vibrant”, “picturesque”, “perfect for everyone”
- “unforgettable”, “journey like no other”
- Filler connectors: “In conclusion”, “Ultimately”, “Overall”

Implementation

- Stored in GLOBAL_TONE_RULES and PROHIBITED_PHRASES.
 - Checked by post-generation validator before merge.
-

Character-Based Length Rules (Multilingual Safe)

Why character count, not word count:

- Languages like German/Polish use long compounds → distort word limits.
- Translations can bloat or shrink by $\pm 20\text{--}35\%$.
- SERP truncation is character-based, not word-based.
- AI handles token-length easier when expressed in **character windows**.

Global Character Targets

Element	Character Target	Notes
Meta Title	60–75 chars	Hard limit for SERP
Meta Description	140–160 chars	No CTA spam
Snippet / Summary	180–260 chars	Used in city/attraction/product
Intro Paragraph	350–550 chars	1 paragraph
Long Description	1,600–2,400 chars	2–4 paragraphs. Backend: AI injects context from tags (local guide, family-friendly, hotel pickup).
Option Name (Variant)	≤ 80 chars	Hard limit (OTA-compatible)
Option Description	≤ 255 chars	Mandatory factual clarity
Highlight bullet	≤ 85 chars	Starts with verb
FAQ Answer	220–350 chars	Concise, factual

Sentence & Paragraph Structure Rules

Sentence Rules

- 12–20 words preferred (approx. 65–110 chars).
- Active voice $\geq 80\%$.
- Max 1 comma per sentence.

Paragraph Rules

- 1 concept per paragraph.
- No paragraph longer than **550–600 characters**.
- No sentence longer than **180 characters**.

Implementation

- Validator: sentence_length_check, active_voice_check.
-

Keyword Rules (Ownership & Safety)

Primary Keyword Requirements

- Must appear in:
 - H1
 - Meta Title
 - First 100 chars of description
- Density target: **0.8%–1.8%**
- Never repeated unnaturally.

Secondary Keywords

- 1–3 per page.
- Must stay within embedding similarity 0.75–0.92 of primary.
- Auto-block if they collide with another page's ownership.

Anti-Cannibalization Logic

- AI is forbidden to use any keyword locked by another URL.
 - Validator checks all generated content against keyword_map.
-

03. Global Meta Title & Meta Description Rules

Meta Title (60–75 chars)

Must include:

- Primary keyword
- City name (if applicable)
- Page type context (“Tickets”, “Guided Tour”, “Things to Do”)
- All landing pages needs meta title & description
- Other languages has metas translated automatically from english

Prohibited:

- Years (e.g., “2025 version”)
- Adjective spam
- Clickbait (“Must-see!”)

Formula Examples:

- **Product:** *{{Tour_Name}} — {{Attraction_Name}} | {{City_Name}} Guided Tour*
- **City:** *Things to Do in {{City_Name}} — Top Attractions & Tours*
- **Attraction:** *{{Attraction_Name}} — Visitor Guide & Tips*

Meta Description (140–160 chars)

Must include:

- Key value (skip-the-line, expert guide, entry included)
- City and attraction/context
- Soft emotional tone

Avoid:

- CTA spam
- Keyword stuffing
- Lists of features

Example Pattern:

“Discover {{Attraction_Name}} with a local guide. Learn the stories behind {{City_Name}}’s most famous site, enjoy smooth entry, and explore at your pace.”

Snippet / Summary Rules (Global)

Snippet Requirements (180–260 chars)

- 1 paragraph
- High information density
- No adjectives chains
- No “marketing tone”
- Summarizes: what, where, why it matters

Used inside:

- product listings
 - city “top picks”
 - attraction summaries
 - internal linking blocks
-

04. Global Rules for Product Options / Variants

Every option has **two required fields**:

Option Name Rules

- ≤ 80 characters (hard limit)
- Must begin with:
 - duration OR
 - distinguishing feature OR
 - ticket type
- No fluff, no adjectives
- Mandatory disambiguation

Examples:

- ✓ “4-hour: Skip-the-line Westminster Abbey & City Highlights (EN)”
 - ✓ “2-hour: Guided Tour + Abbey Tickets”
 - ✗ “Amazing Westminster Abbey Tour with Best Guide Ever”
-

Option Description Rules

- ≤ 255 characters (hard limit)
- Must explain what this option includes or excludes
- Must highlight differences from other variants
- Must match actual product data (no hallucinations)
- No CTA

Example (good):

“Explore Westminster with a certified guide and skip-the-line entry to the Abbey. Includes Big Ben, Buckingham Palace (exterior) and group headset.”

Variant Factual Consistency Rules

Option description must strictly reflect:

- inclusions

- exclusions
- duration
- ticket type
- skip-the-line or not
- guide language
- meeting point differences
- special access (organ recital, evensong, etc.)

Any mismatch → **block AI output and require human confirmation.**

05. Global hook engines rules

a. Hook Trigger Rules (Global)

- Never show more than **1 hook within 20 seconds**
- Never show a hook **twice** if user dismissed it
- Don't show hooks to logged-in users
- Don't show discount hooks **on checkout pages**
- Hook logic is placed in: `HOOK_ENGINE.js` (centralized)
- "Frequency: 1× weekly" means: per user/device, the same hook should not be shown more than once within 7 days. Implementation: cookie/localStorage or user_id-based suppression.

b. Hook Frequency Control

- Discount hook → max 1× per 7 days
- Exit intent hook → max 1× per 7 days
- PDF/Checklist hooks → unlimited but show after scroll 35%

c. Backend Requirements for All Non-Content Modules

Feature	Backend Module	Notes
Wishlist	<code>/cms/wishlist/settings</code>	Enable/disable; default email template

PDF Generator	/cms/content-assets/ /	Upload PDF; Assign pdf_id, page_type, language
Trip Planner	/cms/ai-tools/planner-settings	AI prompt + form fields
Hook Engine	/cms/hooks/	Toggle hooks per page type, Copy text, labels, Frequency logic (per-user suppression)
Quiz Builder	/cms/quiz/	Define questions, options, logic, Mapping to categories/products
Exit Intent Offer	/cms/hooks/exit-intent	Enable/disable + Text, discount code
Sharing (Social Links)	/cms/social-sharing	Default share text + UTM templates
Expert Q&A	/cms/qa/	Approve answers; moderate user questions, assign to expert

d. PDF + Checklist + Guide Handling (Technical)

1. Marketing team uploads PDFs in **CMS** → **Content Assets** → **PDF Manager**
2. Each PDF receives:
 - pdf_id
 - title
 - page_type (city, attraction, package)

- language
3. Hook references the PDF via pdf_id
 4. When user submits email → system:
 - Saves to lead_capture table
 - Sends email with PDF link via transactional email provider (SendGrid/Mailgun)
 - Logs event in analytics

No heavy files stored in DB; always served via Cloudinary CDN.

e. Workflow of hooks

All hooks store leads in a single table lead_capture with fields like: email, page_type, page_id, hook_type, created_at, metadata (JSON).

PDF download hooks::

1. Marketing uploads PDF in CMS → PDF Manager.
2. PDF gets: pdf_id, title, page_type, language, file_url (Cloudinary/CDN).
3. Hook on page references pdf_id.
4. User submits email → system:
 - Saves lead in lead_capture
 - Sends transactional email containing PDF link
 - Logs event in analytics.

No PDF binary data stored in DB, only URL.

2.7 Booking flow for AI Agent (headless) + human booking

2.8 URL Architecture (tree, language layer, formats)

01. URL Architecture

a. URL hierarchy shortage

URL hierarchy

- Country
- City
- Attraction
- Tour
- Category page (per city)
- Things to do (per city)
- Packages (per city)
- Travel Guides (global + per city)
- Experts (global → country → city → profile)

Information hierarchy (deep)

- Continent
- Country
- Region
- City
- POIs
- Products
- Bundles
- Guides
- Experts
- Stories

Breadcrumbs

User sees:

Home > Europe > Italy > Lazio > Rome > Colosseum > Tour


Schema

Google sees relationships fully.

URL

Browser shows a **short SEO-friendly path**.

b. URL Tree Global Architecture






```
|
├─ /legal/terms/
├─ /legal/privacy/
├─ /legal/cookies/
|
├─ /sitemap/
├─ /ai-sitemap/
└─ /api-feed/
```

c. URL Semantic Concept

01. Philosophy

- Rostravel uses a compressed, semantic URL system:
- URLs must stay short, human-readable, and machine-friendly.
- Full geographical hierarchy (Continent → Country → Region → City) exists in breadcrumbs and schema, not in URLs.
- No technical IDs (133, tc1094) are used in paths.
- Maximum 3 segments after language code.

02. Language Layer

`/ {lang} /`

Example: `/en/`, `/de/`, `/pl/`

d. URL formats for specific pages

01. Destination Pages

Page Type	URL Format	Example
Country	<code>/ {lang} / {country} /</code>	<code>/en/italy/</code>
City	<code>/ {lang} / {country} / {city} /</code>	<code>/en/italy/rome/</code>

Things to Do	<code>/ {lang} / {country} / {city} / things-to-do /</code>	<code>/ en / italy / rome / things-to-do /</code>
Category (per city)	<code>/ {lang} / {country} / {city} / {category} /</code>	<code>/ en / italy / rome / day-trips /</code>

02. Attraction Pages (POIs)

`/ {lang} / {country} / {city} / {attraction} /`

Example:

`/ en / italy / rome / colosseum /`

03. Tours / Product Pages

`/ {lang} / {country} / {city} / {tour-slug} /`

Example:

`/ en / italy / rome / colosseum-guided-tour /`

04. Packages / Bundles

`/ {lang} / {country} / {city} / {package-slug} /`

Example:

`/ en / italy / rome / 3-day-family-city-break /`

05. Travel Guides

Type	URL
Hub	<code>/ {lang} / travel-guides /</code>
City Guide	<code>/ {lang} / travel-guides / {country} / {city} /</code>

06. Blog / Stories

`/ {lang} / stories / {topic} /`

Example: `/ en / stories / best-old-towns-europe /`

07. Experts

Page	URL
Experts Hub	<code>/ {lang} /experts/</code>
Country Experts	<code>/ {lang} /experts/ {country} /</code>
City Experts	<code>/ {lang} /experts/ {country} / {city} /</code>
Expert Profile	<code>/ {lang} /experts/ {country} / {city} / {name} /</code>

08. Explore / Global Attraction Map (UI)

`/ {lang} /explore/`

09. Brand & Business Pages

Page	URL
About	<code>/ {lang} /about/</code>
Contact	<code>/ {lang} /contact/</code>
For Business	<code>/ {lang} /for-business/</code>

10. Technical Feeds (not in menu)

`/sitemap.xml`

`/ai-sitemap.xml`

/ai-summary.json
/api/tours/feed

11. AI Schema Architecture

AI / LLM feed: each level (country, region, city, attraction, package) should have a context-aware JSON containing the fields `geo`, `hasPart`, `isPartOf`, and `nearby`, enabling LLMs to understand the semantic structure.

➤ Structured data:

- `TouristDestination` for regions,
- `TouristAttraction` for attractions,
- `hasPart` / `nearby` relations between layers.
- others

➤ Linking:

- AI can automatically generate “nearby links” based on coordinates (geohash / bounding box).
- All interconnected with structured data for `TouristTrip`, `Offer`, `Review`, `Organization`, `FAQPage`, and the Rosotravel Knowledge Graph.

3.2 Semantic URLs

Purpose

Human-readable and keyword-rich structure that encodes country → city → tour type → topic.

Data/Schema

- Examples:
 - `/uk/london/private-guided-tour-british-museum`
 - `/italy/rome/vatican-museum-skip-the-line`
 - `/spain/barcelona/gaudi-sagrada-familia-tour`

Backend

- Slug generator with locale awareness; collision handling; redirects on rename.
- Rules per page type (product, city, attraction, long-tail).
- Store canonical path; update sitemaps on change.

Front-End

- Enforce canonical link; breadcrumb rendering.
- Hreflang alternates for localized paths.

APIs/Integration

- URL registry `POST /urls/reserve`, `PUT /urls/rename`.
- Redirect service (301) with audit log.

Validation

- One canonical per page; no duplicate paths; path length limits respected.

KPIs

- Crawl errors (4xx/5xx) near zero; CTR uplift on semantic paths.

Acceptance (DoD)

- All templates generate correct semantic URLs with working hreflang.

2.9 Images, photo & Reviews rules:

Purpose

Rosotravel must build a destination and product media layer that:

- improves visual quality and brand consistency,

- supports premium presentation and stronger conversion,
- provides machine-readable image context (ALT text, ImageObject),
- and remains compliant with source licensing, indexing rules, and partner certification requirements.

This section replaces the previous simplified image rules and introduces a governed media model.

2.9.0 Core principle

Rosotravel must distinguish between:

- A) Rosotravel-owned or Rosotravel-cleared media assets and
- B) Viator-sourced or otherwise license-restricted assets.

This distinction is mandatory because indexing, ALT text policy, derivative processing, and hero video generation may differ depending on source rights.

Hard rule:

No media asset may be treated as indexable or SEO-eligible unless Rosotravel has the necessary rights to do so.

2.9.1 Media source classes

Class A — Rosotravel-owned / Rosotravel-cleared assets

These assets may include:

- Rosotravel original media,
- guide / expert profile images,
- admin-uploaded destination or product media,
- customer-uploaded review media where Rosotravel terms grant reuse rights,
- other assets with confirmed commercial and SEO-safe usage rights.

Class B — Viator-sourced assets

These assets may include:

- supplier photos returned in Viator product content endpoints,
- traveler photos retrieved through /reviews/product,
- review-linked external media,
- other assets governed by Viator API content restrictions.

Class C — AI-derived assets

AI-derived assets are media processed or generated using:

- Nano Banana 2 / Nano Banana Pro (image editing / generation),
 - Google Veo (image-to-video / video generation),
- from a source asset.

AI-derived assets inherit the rights risk of the source asset unless Rosotravel has explicit rights that allow derivative commercial and indexable usage.

2.9.2 Legal and certification rule for Viator-derived media

Viator's official partner documentation states that:

- review content is protected proprietary information and must not be indexable by search engines,
- traveler photos linked to /reviews/product are treated as protected review content and must also remain non-indexed,
- certification requires Protected Viator Unique Content to be non-indexed.

:contentReference[oaicite:3]{index=3}

Therefore, until Rosotravel obtains explicit written permission from Viator that allows indexable derivative media usage:

- Viator-sourced media must remain non-indexable,
- review-linked traveler photos must remain non-indexable,
- and AI-derived media generated from Viator-sourced assets must also be treated as non-indexable by default.

This is the default compliance-safe rule.

2.9.3 Rosotravel-owned / cleared media policy

Rosotravel-owned or Rosotravel-cleared assets may be:

- processed by Nano Banana 2 / Nano Banana Pro,
- converted into derivative hero visuals,
- transformed into short hero videos with Google Veo,
- assigned AI-generated ALT text,
- indexed,
- and exposed in ImageObject / VideoObject schema where applicable.

This is the preferred long-term media model for destination and product pages.

2.9.4 Product hero visual strategy (target state)

Target visual strategy for product pages:

A) Hero asset

The hero asset for a product should be a short video generated from the first two approved hero-source images.

Target duration:

- approximately 8–9 seconds

Audio:

- background music / native audio may be added where permitted by tooling and brand rules.

Primary tools:

- Nano Banana for image preparation,
- Google Veo for hero video generation.

B) Supporting gallery assets

The next three approved product images should be processed through Nano Banana to produce:

- brand-consistent visual quality,
- improved composition / clarity,

- stronger premium feel,
- generated ALT text,
- and indexable Rosotravel-owned derivative assets only when source rights allow it.

C) Rendering order

Recommended product media order:

1. Hero video (derived from approved first two source images)
2. Supporting image 1
3. Supporting image 2
4. Supporting image 3
5. Additional gallery assets (if available and permitted)

Important:

This target state applies fully only to Rosotravel-owned / cleared media.

It must not be automatically applied to Viator-sourced images until rights are confirmed.

2.9.5 Interim implementation rule (before legal clearance)

Until licensing / certification implications are explicitly cleared:

For Country / Region / Other editorial surfaces

- Use Rosotravel-owned / AI-generated media only
- Nano Banana-generated or edited assets allowed
- Indexable
- ALT text mandatory
- Eligible for schema

For Blog / Guide posts

- Images from Rosotravel Media Library only (approved_processed assets)
- Indexable
- ALT text mandatory
- External hotlinks and unverified uploads forbidden

For Cities / Attractions / Products using Viator-sourced images

- Source images may be displayed operationally,
- but must remain non-indexable if they fall under Viator-protected or license-restricted content,
- no SEO ALT text should be applied to restricted images,
- images should be treated as presentational / accessibility-safe only when required.

For AI-derived assets from Viator sources

- Do not assume they become indexable just because they were transformed,
- do not expose them as indexable SEO assets by default,
- keep them non-indexable until written approval / contractual clarification is obtained.

2.9.6 Final implementation rule (only after legal/licensing approval)

If Rosotravel obtains explicit written approval from Viator or otherwise contractually valid rights allowing derivative indexable use of Viator-sourced product images, then the following model may be activated:

Cities, Attractions, Products

- first two approved source images are transformed into a short hero video (8–9 seconds) using Google Veo,
- next three images are processed through Nano Banana,
- all five resulting Rosotravel-processed assets receive ALT text,
- the hero video becomes the default product hero asset,
- all processed assets become indexable,
- schema may expose ImageObject / VideoObject accordingly.

Important:

This rule must remain conditional until rights are validated in discovery and approved legally.

2.9.7 ALT text rules

ALT text policy must follow media rights and page purpose.

A) Rosotravel-owned / cleared assets

- ALT text mandatory
- ALT text may be AI-generated, then optionally editor-reviewed
- ALT text should describe the visible scene and travel relevance
- ALT text must not be keyword-stuffed

B) Restricted Viator-sourced assets

- No SEO ALT text if the asset is non-indexable and license-restricted
- If needed for accessibility, ALT may be empty or purely functional depending on implementation
- Presentational handling (aria-hidden / role="presentation") may be used where appropriate

C) AI-derived assets

- ALT text allowed only when the derivative asset is legally approved for indexable use
- Otherwise treat like restricted media

2.9.8 Indexing policy by media type

Rosotravel-owned / cleared product images:

- indexable
- ALT text required
- eligible for schema

Rosotravel-owned / cleared hero video:

- indexable
- eligible for VideoObject
- should include title, description, thumbnail, duration metadata

Viator review content and traveler photos:

- non-indexable
- source-restricted
- not eligible for SEO indexing exposure

Viator-sourced supplier photos:

- do not treat as indexable by default unless rights are explicitly cleared
- current safe implementation = non-indexable operational display only

2.9.9 Reviews & photos — trust and machine readability

Purpose

Reviews and photos support:

- trust,
- E-E-A-T / reputation signals,
- Human-Face / Latest Tour Snapshot,
- image and review comprehension by search systems and LLMs.

Visible UI requirements

Reviews:

- source tagging mandatory (Rosotravel / Viator / Tripadvisor / other)
- verified booking label where applicable
- aggregate rating may combine multiple sources if transparently handled

Photos / media:

- source classification mandatory
- freshness / attribution logic where applicable
- only approved assets can enter product galleries and snapshot modules

Backend requirements

- aggregate reviews from allowed sources
- generate AI summaries of “what travelers value most”
- cluster recent review themes
- store media metadata (source_type, linked entity, uploaded_at, moderation status)
- generate ALT texts only for rights-cleared assets
- convert approved indexable assets to modern formats (e.g., WebP where relevant)
- support hero video generation pipeline for approved products/assets

Front-end requirements

- product pages support hero video rendering as primary hero asset when available
- supporting gallery supports image ALT rendering for approved assets
- restricted images remain non-indexable and must not be presented as SEO assets
- optional social share remains secondary and rights-dependent

Validation and governance

Before launch of the final hero-video model, discovery must validate:

- licensing and derivative-use rights for Viator-sourced images
- API certification implications
- whether derivative AI assets are considered restricted content under partner terms
- acceptable indexing policy for processed media
- moderation workflow for customer-uploaded review media
- operational cost and latency for Nano Banana + Veo processing
- fallback behavior when video generation fails
- content QA rules for ALT text and hero-video quality

2.9.10 Recommended rollout sequence

Phase 1

- keep Viator-sourced images non-indexable
- use Rosotravel-owned / cleared assets as indexable media
- enable ALT text only for cleared assets
- prepare media source classification and metadata layer

Phase 1.5

- enable Nano Banana processing for Rosotravel-cleared product assets
- enable ALT text automation for processed assets
- enable schema exposure for processed images

Phase 2

- enable hero video generation from approved product hero assets with Veo
- expose VideoObject where appropriate
- activate only for products whose source rights are cleared

Phase 3 (conditional)

- expand processed/indexable product media coverage only if partner-license review confirms that derivative usage from supplier assets is allowed

Reviews & Photos

Purpose

Trust + E-E-A-T signals; improve image and review comprehension by search/LLMs.

Data/Schema

- **Reviews: source tagging (**reviewSource**), **aggregateRating**.**
- **Photos: **imageObject** with EXIF/author/date; ALT texts.**

Backend

- **Aggregate reviews from sources; AI summary + top 3 highlights.**
- **Store EXIF/geo when available; generate AI alt texts; convert to WebP; lazy-load flags.**

Front-End

- CSS-only star ratings; gallery with alt rendering; share to social (optional).

Validation

- Deduplicate reviews; ALT required for all images.

KPIs


- Engagement with media; star rendering time; schema coverage.

Acceptance (DoD)

- Reviews + images present with correct schema and performance.

III. Marketing & SEO Operations (CMS & governance)

3.1 CMS Panel new menu



The CMS MUST expose the following panels. If a panel does not exist, the system is considered **non-compliant** with SEO governance.

1) Keyword Governance Panel (Critical)

1) CMS Requirements Keywords

On each Product / Attraction / City / Category page, add:

1. Keyword Ownership Panel

Fields:

- **primary_keyword**
- **secondary_keywords[]**
- **locked_keywords[]**
- **intent classification**
- **source (GSC / SEMrush / AI)**

Buttons:

- **Unlock keyword (admin only)**
- **Re-run keyword allocation**

Purpose

Single source of truth for keyword ownership, routing, and conflict resolution.

Marketing MUST be able to:

- view all keywords fetched from SEMrush & GSC
- see intent, volume, embedding similarity, entity mapping
- see which URL owns each keyword
- see lock status (locked / unlocked)
- see ownership history and timestamps
- approve or reject new ownership assignments
- manually lock or unlock keywords (with audit log)
- mark keywords as:
 - deprecated
 - seasonal

- escalated for review

Marketing MUST NOT be able to:

- assign transactional keywords to non-product pages
 - bypass intent allowlists
 - delete locked transactional keywords
-

2 Page SEO State Panel (Per-URL Control)

Purpose

Full transparency of why a page is indexed, canonicalized, or blocked.

For every URL, Marketing MUST see:

- page_type
- intent_type
- primary keyword
- owned keywords
- owner_url
- index_state (index / noindex)
- canonical_target
- hreflang_cluster
- entity_flag
- seasonal_flag
- GSC metrics (impressions, clicks, CTR – last 28d)
- rewrite cooldown status
- near-page lifecycle status (draft / contained / eligible / graduated)

Marketing MAY edit (explicit buttons):

- request index promotion (if eligibility rules met)
- request index suppression (force noindex)
- request canonical override (ONLY to owner_url)
- trigger ownership re-evaluation
- mark page as flagship (packages only)
- approve near-page graduation

All actions require:

- justification note
- timestamp
- user ID

-

5 AI Rewrite & Optimization Panel (Guarded)

Purpose

Ensure AI optimization does not destabilize rankings.

Marketing MUST see:

- last rewrite date
- rewrite count this quarter
- cooldown timer
- impression threshold status
- sections eligible for rewrite

Marketing MAY:

- approve or reject proposed AI rewrites
- trigger manual rewrite (within allowed scope)
- lock sections permanently from AI

Marketing CANNOT:

- trigger rewrites below impression threshold
- unlock forbidden sections (H1, SEO body, canonical)

3.15 UI Main panel Add: All Destinations (Entities) (isn't part of marketing panel)

All Destinations is the **single operational view** for managing SEO Auto processing across **all destination entities** (Country, Region, Island, Other). It supports **bulk processing**, **issue-driven triage**, and **safe regeneration** without fragmenting logic across multiple tabs.

Pagination & Scale Rules

- **100 entities per page** (fixed).
 - **Server-side pagination, filtering, and sorting only.**
 - Default load: first page only.
 - All counters/statuses are read from a **materialized Destinations List View** (no live computation).
-

Bulk Processing Rules

- **Run SEO Auto** can be executed for **up to 50 selected entities at once**.
- Selection can be:
 - manual (checkboxes), or
 - filtered selection + **Select all matching (max 50)**.
- If a filter returns >50 entities:
 - system allows **batching automatically** (50 → next 50 → queue).
 - Important: system call SEO Database (e.g. Semruhs) one by one it is different process than processing a City/Product batch. Here each entity is processed

separately. For example you select Spain, German and Poland. The process is going one by one first Spain -> finished -> move to under entity

- **Run SEO Auto button:**
 - visible to all
 - **clickable only for SEO Manager / Admin**
 - requires **one confirmation modal per action**, even for multi-batch runs.

Confirmation modal (required):

- Action type: Initial run / Regeneration
 - Reason (mandatory, free text)
 - Scope summary (e.g. "50 countries, filter: No SEO Auto run")
 - Actor (auto-filled)
 - Timestamp (auto)
-

Regeneration Strategy (Critical)

Regeneration is **never global** and **never blind**.

How regeneration works

- Regeneration is allowed **only via filters**, never "Re-run all".
- Typical regeneration filters:
 - SEO Auto Status = Completed
 - AND Issues Severity $\geq$ High
 - OR Quality Score < threshold
 - OR Missing Required Blocks
 - OR Manual Review Required = true

- User selects affected entities → **Run SEO Auto (Regeneration)**.

Why this is safe

- You never regenerate entities that:
 - passed validation
 - were manually locked and unchanged
 - Regeneration always has:
 - explicit scope
 - explicit reason
 - audit trail
-

Left Sidebar — Filters (Primary Control)

Core Filters

- **Entity Type**
 - Country
 - Region
 - Island
 - Other
- **Parent Relation**
 - Parent Country
 - Parent Region
- **SEO Auto Status**
 - Not Run
 - Queued

- Processing
 - Completed
 - Failed
- **SEO Auto Version State**
 - Up to date
 - Outdated vs RAW
 - Regeneration required
- **Issues**
 - Has issues (any)
 - Critical only
 - High / Medium / Low
- **Quality Signals**
 - SEO score range
 - Missing required blocks
- **Manual Interaction**
 - Has manual edits
 - Has locked fields
- **Assets**
 - Hero image present
 - Hero image missing (indexing blocker)
- **Products Dependency**
 - Has published products
 - No published products (blocking)
- **Last Activity**

- Updated in last X days
- Not reviewed since X days

Saved Views

- Predefined (system):
 - Countries — No SEO Auto run
 - Countries — SEO Outdated
 - Missing Hero Image
 - Blocking Issues Only
 - Ready for Publication
 - User-defined:
 - Save current filter set
 - Shareable (role-based)
-

Main Table — Horizontal Overview (Core UI)

Each row = one destination entity

Columns (left → right)

1. Select

- Checkbox

2. Entity

- Name
- Entity type badge
- Parent path (e.g. Europe → Italy)

3. SEO Auto

- Status badge (Not run / Completed / Failed / Outdated)
- Last run timestamp

4. Coverage

- Blocks processed (icons or %)
- Tooltip with missing blocks

5. Quality

- SEO score
- Color-coded

6. Issues

- Count by severity
- Red/yellow dot if blocking

7. Manual

- Manual edits present (icon)
- Locked sections count

8. Assets

- Hero image status (OK / Missing)

9. Products

- Published products count

10. Actions

- View entity
- Open issues
- (Managers only) Run / Regenerate SEO Auto

- **Select all on page (100)**
 - **Select all matching filter (max 50)**
→ if >50, system shows “This will run in X batches”.
 - **Run SEO Auto** (role-gated)
 - **Clear selection**
 - Active filter summary (read-only)
-

Notifications & Visual Alerts (Mandatory)

Entity-level alerts

- New RAW update detected → entity row highlighted (red dot)
- SEO Auto outdated vs RAW → yellow warning
- Blocking issue → red badge

Global alerts

- Bell icon (global, always visible)
- Shows:
 - Failed runs
 - Entities awaiting review
 - New blocking dependencies (e.g. product added without destination readiness)

Acknowledgement

- Clicking an alert marks it as reviewed.
 - Reviewed alerts disappear from list but remain in audit log.
-

Entity Relationships (Visibility Rule)

- Each entity shows:
 - Parent entity (if exists)
 - Child count (e.g. Regions under Country)
 - Blocking rule surfaced in UI:
 - If a product cannot attach due to missing destination readiness → destination row flagged.
 - Navigation:
 - Clicking parent/child count opens filtered All Destinations view.
-

Why There Are No Separate Tabs (Countries / Regions)

- All destinations share:
 - the same SEO Auto lifecycle
 - the same blocking logic
 - the same regeneration rules
 - Separation is achieved via:
 - filters
 - saved views
 - This avoids:
 - duplicated logic
 - inconsistent batch handling
 - hidden failures across tabs
-

Summary

- 100 entities per page, always.
- Max 50 entities per SEO Auto run (auto-batched if needed).
- Regeneration only via filtered, intentional selection.
- Full audit trail for every run and regeneration.
- Clear visual signals for:
 - outdated SEO
 - blocking issues
 - manual locks
- One table, one logic, no hidden states.

Description of statutes -

01. Internal view when you click e.g. Country: Entity detailed view

(Country / Region / Island / Other – unified model)

This view is used exclusively for **top-level destination entities** (countries, regions, islands, other non-city entities).

Cities, attractions, products, and near-pages are handled in separate marketing pipelines and are **not edited here**.

1. Core Concept (Single Unified Model)

All destination entities share:

- the **same UI**
- the **same governance rules**
- the **same AI / manual control model**

Differences between Country / Region / Island are handled via:

- template rules
- enabled / disabled sections
- data sources

No separate UI variants are created.

2. Header (Sticky)

Displays at all times

- Entity name + entity type
- Breadcrumb:
 - Country → Country
 - Region → Country > Region
 - Island / Other → Country > Region > Entity
- Status indicators:
 - Active / Inactive (based on product linkage)
 - Run SEO Auto status:
 - NOT RUN
 - RUN
 - FAILED
 - OUTDATED
- Last SEO run timestamp
- AI version ID (destination-level)

Primary Actions (role-gated)

- **Run SEO Auto / Regenerate SEO Auto**
 - SEO Lead / Admin only
 - Confirmation modal required
 - Reason required (stored in audit log)
- Preview public page
- View full audit log

3. Tabs Overview

1. **SEO Content**
2. **Assets**
3. **Relations**
4. **Manual Overrides**
5. **History / Audit**

4. TAB: SEO Content

Purpose: manage all **indexable SEO-owned content** for the entity page.

4.1 Keywords (Read-only by default)

- Primary keyword (exactly 1)
- Secondary keywords (max 2)

Attributes per keyword:

- Source: AI / Manual
- Version ID
- Generated at

Rules:

- Editable only by SEO Manager
 - Manual edit does NOT trigger new SEO run
 - Keyword ownership rules still apply globally
-

4.2 Meta (Field-level control)

Fields:

- Meta title
- Meta description

Per field:

- Source badge: AI / Manual
 - Lock status
 - **Chain status** (see section 7)
 - Regenerate (AI) button (only if AI-enabled)
-

4.3 Snippet / Intro

- 180–260 characters
- AI-generated by Run SEO Auto
- Editable manually

Supports:

- Default snippet
 - Segment variants (DEE only, non-indexed)
-

4.4 Overview

- 2–3 paragraphs
- Stable SEO block
- Ownership: AI + manual

Used for:

- semantic authority
 - internal linking anchor block
-

4.5 Highlights (Structured Inspiration)

- 5–7 bullet points
 - AI-generated
 - Stable block (not dynamic)
 - Editable manually
-

4.6 Facts & Curiosities (Factual Block)

- 5–7 verified factual bullets
 - **AI is NOT allowed to invent facts**
 - Sourced from factual data modules
 - Manual correction allowed
-

4.7 FAQ

- Exactly 7 FAQs
- Informational only
- AI-generated by Run SEO Auto

Per Q/A:

- Edit manually
- Lock

- Remove

Regenerate FAQ (AI) affects ONLY FAQ block.

5. TAB: Assets (Index-safe Images)

Purpose: ensure **SEO-safe, licensable imagery**.

Rules

- Country / Region / Island:
 - Images generated via **Nano Banana Pro**
 - Indexable
 - Alt text generated automatically by Run SEO Auto
- Supplier / Viator images:
 - NEVER used at this level

UI

- Hero image:
 - Upload
 - Replace
 - Generate via AI
 - Compliance status:
 - OK
 - Missing
 - Non-compliant
-

6. TAB: Relations (Read-only)

Purpose: visibility only, no editing.

Shows:

- Parent entity
- Child entities (counts only)
- Inventory linkage summary:
 - “X published products linked indirectly via cities”

No actions allowed.

7. CHAIN / UNCHAIN MODEL (CRITICAL)

This is the missing piece that makes the entire system usable.

7.1 Block Ownership States

Each editable SEO block (meta, snippet, overview, highlights, FAQ) can be in **one of three states**:

1. **AI-Controlled (CHAINED)**

- Default state
- AI can regenerate
- No manual override

2. **Manual Override – LOCKED**

- Block edited manually
- AI is blocked for this block only

3. **Manual Override – UNCHAINED**

- Manual edit exists
- AI is re-enabled explicitly

7.2 UNCHAIN Action (Re-enable AI)

Available per block.

Rules:

- SEO Lead / Admin only
- Confirmation required
- Reason required:
 - Cosmetic edit
 - Temporary test
 - Outdated manual copy
 - Align with new SEO run

Effects:

- Manual content remains
- AI may overwrite on next regeneration
- Full audit trail preserved

7.3 Important Clarifications

- Manual edits **do NOT block the entire entity**
- Lock / Unchain is **per block**
- Unchain does NOT trigger regeneration automatically
- Next Run SEO Auto respects chain state per block

8. TAB: Manual Overrides

Purpose: transparency and recovery.

Table:

- Field name
- Edited by
- Edited at
- Lock state (Locked / Unchained)
- Reason

Actions:

- Unlock (Admin only)
- Unchain (Admin only)

9. TAB: History / Audit

Timeline includes:

- SEO runs (with reasons)
- Regenerations
- Manual edits
- Lock / Unchain actions
- Status transitions:
 - NOT RUN → RUN → OUTDATED → FAILED

Immutable log.

10. Country vs Region Differences (Template-only)

No UI difference.

Handled via:

- Different H1 templates
- Different enabled sections
- Different hooks (PDF vs Planner)
- Different relations depth

Same governance, same controls.

02. General Rule of Logic Content Ownership & AI Control

1. Core Principle

Every content field in the system is governed by **two independent dimensions**:

1. **Source** — who authored the current published value
2. **AI Permission** — whether AI is allowed to modify this field in the future

These two dimensions must **never be merged, inferred, or overloaded into one state**.

2. Source (Content Ownership)

Source defines authorship only.

It answers the question:

“Who created the content that is currently published?”

Each field has exactly **one Source**:

- **AI**
Content generated by AI (initial generation or AI regeneration)
- **Manual**
Content edited or written by a human in CMS

- **Factual**
Verified system data (API, structured sources, validated facts)
- **Aggregated**
Synthetic content derived from multiple inputs (e.g. review summaries)

Source is **descriptive, not behavioral**.
It does **not** control future AI behavior.

3. AI Permission (Behavioral Control)

AI Permission is the only mechanism that controls AI behavior.

It answers the question:

“Is AI allowed to regenerate this field?”

Each field has exactly **one AI Permission state**:

- **AI Enabled**
AI may regenerate this field in future runs
- **AI Disabled**
AI must never modify this field

AI Permission is **explicit, intentional, and user-controlled**.

4. Default States (System Rules)

4.1 AI-Generated Content

Source = AI

AI Permission = Enabled

AI-generated fields are **regenerable by default**.

4.2 Manual Edit (No Lock)

Source = Manual

AI Permission = Enabled

A manual edit **does not block AI** unless the user explicitly disables it.

This supports cosmetic or corrective edits without breaking automation.

4.3 Manual Edit + AI Disabled

Source = Manual

AI Permission = Disabled

AI is fully blocked for this field.

This is the **only state that prevents future AI changes**.

4.4 Explicit AI Lock (Rare Case)

Source = AI

AI Permission = Disabled

Used when AI content is considered final and must not be regenerated.

4.5 Factual Fields

Source = Factual

AI Permission = Disabled

AI is **never allowed** to modify factual values
(hours, prices, locations, policies, coordinates).

AI may only rephrase **labels**, never factual values.

5. Manual Edit Workflow (Mandatory UX Rule)

After saving a manual edit, the system must explicitly ask:

“Should AI be allowed to update this field in the future?”

Options:

- **Keep AI Enabled** (default, recommended)
- **Disable AI for this field**

This choice is always **explicit** and **reversible**.

6. AI Regeneration Rules

AI regeneration applies **only** to fields where:

AI Permission = Enabled

Fields with AI Disabled are:

- automatically skipped
 - never overridden
 - never partially updated
-

7. Version Mismatch & Alerts

If source data (RAW / factual) changes and:

- **AI Permission = Enabled**
→ AI may regenerate the field
- **AI Permission = Disabled**
→ Field remains unchanged
→ System creates a **version mismatch**
→ Alert is raised: *“Update skipped due to AI disabled”*

This behavior is deterministic and auditable.

8. Key Guarantees of This Model

- Manual edits never block AI by accident
 - AI behavior is always predictable
 - Locks are field-level, never global
 - Version mismatches are explainable
 - No “hidden states” or inferred logic
-

9. One-Sentence Mental Model (for teams)

Source defines who wrote the content.

AI Permission defines whether AI may change it again.

J.09 CMS Requirements for Navigation & Linking

1. Navigation Admin (Menus)

CMS must provide:

- Menu builder for Header & Footer.
- Ability to pin specific countries/experiences/guides as “featured” in mega menus.
- Language-aware labels but single logical menu per site.

2. Internal Links Section on Product / Attraction / City / Article Edit Pages

Fields:

- “Override related links” (checkbox)
 - If checked, marketer can manually choose related products/attractions/articles; AI suggestions become secondary.
- Manually curated links list:
 - Related tours
 - Related attractions

- Related articles / guides
- Highlighted packages
- “Exclude from automatic linking” (checkbox)
 - e.g. for temporary or low-priority pages.

3. Entity Parent Configuration

On each entity (product, attraction, city, region, country):

- parent_country_id
- parent_region_id
- parent_city_id
- parent_attraction_id (for multi-attraction product)

Used directly by navigation + breadcrumbs + schema.

J.10 Validation & Monitoring

To keep the navigation graph healthy, system must monitor:

1. Orphan pages

- Pages with no inbound links (excluding sitemap).
- Report for SEO to fix manually or via auto-linking rules.

2. Overlinked pages

- Node appears in too many “related” blocks → cap number per page or adjust similarity threshold.

3. Broken links

- Any internal link returning 404 or redirected chain.

4. Navigation depth

- Excessive click depth (e.g. more than 4 levels from home) flagged for UX/SEO review.

5. **A/B tests** (future)

- Variation of internal link blocks + measure impact on CTR and conversion.

J.11 How This Integrates With Existing Layers

The navigation layer consumes:

- From **Data & Versioning layer** (A–B):
 - Entities, versions, parents.
- From **Semantic & Keyword layer** (C):
 - `pages_vectors`, `poi_vectors`, keyword ownership.
- From **Content & Schema layer** (D–H):
 - URLs, entity types, schema relations.
- From **Localization layer** (I):
 - language-specific URLs and hreflang for cross-locale linking.

Output:

- Fully wired internal linking structure.
- Correct breadcrumbs and menus.
- Smart “related” and “nearby” blocks.
- Healthy link graph for SEO + LLMs.

Backend takes this, populates **ai_content_sections**, **faq_entries**, and **schema_snapshots**.

3.2 UI Main panel Add: Marketing Panel

3.2.1 Pipeline view & console - UI Option of Marketing Panel

01. Global AI & SEO Operations Console

0. Purpose of Pipeline View (Top-Level Summary)

Pipeline View is the operational control center for all AI-driven SEO and content processing.

It is designed to:

- Run AI Batch Processing for cities (products, attractions, near-pages)
- Run AI processing for single records (new or late-arriving products)
- Monitor RAW → AI → Published state changes
- Detect version mismatches, locks, and SEO risk
- Enable safe, large-scale AI regeneration without manual page-by-page review
- Surface only what requires action

Pipeline View is NOT:

- a CMS
- a content editor
- a destination manager (countries / regions)

Manual content editing never happens here.

1. Mental Model — Two Different Worlds (Critical)

A. PIPELINE VIEW (this screen)

Process-first, reactive

Used to:

- run AI batch processing
- run AI processing for single records
- handle updates, regenerations, retries
- assess SEO and index risk
- perform mass AI operations safely

B. CONTENT & STRUCTURE VIEW (separate screen)

Structure-first, editorial

Used to:

- browse content by hierarchy
- manually edit text
- lock / unlock blocks
- inspect detailed diffs
- resolve issues record-by-record

👉 These two views must never overlap in responsibility.

2. Core AI Batch Processing Concept (Missing Before — Now Explicit)

2.1 What “AI Batch Processing” Means

AI Batch Processing is the primary entry point into Pipeline View.

A batch represents:

- a city-level processing run

- applied to all eligible records in that city
- executed as independent jobs per city

Examples:

- *Vienna Products Batch*
- *Rome Attractions Batch*
- *Paris Near-Pages Cycle #2*

Rules:

- Max 10 cities selected at once
 - Each city creates its own isolated batch
 - No shared state between city batches
 - Failure in one batch never blocks another
-

2.2 What Happens in a City Batch

When a city batch is started:

1. System collects eligible records:
 - Products
 - Attractions (POI)
 - Near-pages (depending on scope)
2. For each record:
 - RAW version is evaluated
 - AI generation is executed where allowed
 - Manual locks are respected. In Pipeline View, a “lock” always means: AI Permission = Disabled for that content block. Source (AI vs Manual) is not modified here and is not inferred.

- A new AI version may be created

3. Batch results appear in Pipeline View:

- per-record status
- per-record risk
- per-record version state

Pipeline View shows results, not content.

3. Single Record Processing (Also Missing Before)

Pipeline View also handles single-record AI processing.

3.1 When Single Record Processing Is Used

Examples:

- A new product arrives after Vienna batch already completed
- A single POI was skipped earlier
- A manual unlock requires AI reprocessing of one record

3.2 How It Works

- The record appears in Pipeline View with status:
NOT PROCESSED
- User selects the record
- User clicks Run AI Processing
- The record is processed independently
- It does not re-open or modify past batches
- It logically “joins” the city set but remains its own job

👉 This prevents re-running entire cities just for one new product.

4. What a “Record” Means in Pipeline View

Pipeline View operates on Records, not destinations.

A Record can be:

- a Product
- an Attraction (POI)
- a Near-page

Records may belong to a Batch, e.g.:

- *Vienna Products Batch*
- *Rome Near-Pages Cycle #3*
- *Viator Feed Import 2026-01-15*

But records can also exist outside of any batch (single processing).

5. Pipeline View — Screen Layout (Top to Bottom)

5.1 Top Bar (Above the Table)

Purpose: global context + bulk actions

Elements:

- Selected records counter (0 / X)
- Active batch indicator (when filtered), e.g.
Batch: Vienna Products (47 records)
- Primary actions (disabled until selection):
 - Content Actions
 - Acknowledge
- Global search (record name / ID)

5.2 Left Sidebar — Filters (Always Visible)

Purpose: reduce noise, surface risk

Filter groups:

Record Type

- Product
- Attraction
- Near-page

Batch

- Batch name (e.g. Vienna Batch)
- Feed source
- Cycle ID

Status

- OK
- Attention
- Review Required
- Blocked
- Processing

Risk Level

- Low
- Medium
- High

Index Impact

- None
- Review
- Block

Content State

- Has manual locks
- Has version mismatch
- Has unpublished AI changes

Recency

- Updated last 24h / 7d / 30d

Filters are applied before selection.

6. Main Pipeline Table (Center)

This table answers:

“What exactly is going on, and is it dangerous?”

6.1 Table Columns (Operational Summary)

Column	Meaning
Record	Name + type (Product / Attraction / Near-page)
Batch	Batch name (if applicable)
RAW	Latest RAW version

AI	Latest AI version
Published	Published version
Status	Computed pipeline status
Risk	SEO risk level
Locks	Manual locks present (Yes / No)
Index	Indexing impact
Last Change	Timestamp
Actions	View / Inspect

6.2 Version Mismatch Visibility (Critical)

Mismatch is shown directly in the table:

RAW: v2

AI: v1

Published: v1 

- Status automatically becomes Review Required
 - No guesswork, no hidden state
-

7. Row-Level Actions

From a single row you can:

- View change summary (read-only)
- Jump to Content & Structure View
- Inspect audit log

 You cannot edit content here.

8. Selection Model (After Filtering)

8.1 Selection Rules

- Select individual records
 - Select all (filtered results)
 - Max 50 records per operation
 - Selection always happens after filters
-

9. Content Actions — Batch-Level AI Control

9.1 Purpose

Pipeline is the operational view for **processing RAW batches** and executing **mass AI actions** on selected records **without manual editing** and without changing keyword ownership.

9.2 Default layout: Batch list

On entering Programmatic SEO / Pipeline, user sees a list of RAW batches.

Each Batch row **MUST** include:

- Checkbox (multi-select)
- Batch name + batch_id

- Record type (Product / City / Near-page / etc.)
- Record count (total)
- Flags summary (OUTDATED / CONFLICT / MISSING / LOCKED counts)
 - **OUTDATED_RAW** (raw_version mismatch OR any of: content_hash/factual_hash/media_hash/relationship_hash changed since published_for_raw_version)
 - **OUTDATED_OFFER** (offer_version mismatch OR offer_hash changed since published_for_offer_version)
 - **OUTDATED_REALTIME** (realtime_offer_version mismatch OR realtime_offer_hash changed since published_for_realtime_offer_version)
 - **OUTDATED_REALTIME never requires AI regeneration**
- Button: “Open batch contents”
- Button: “Delete batch” (role-gated, with confirmation)
- Status pill: Draft / Processing / Completed / Failed

Global controls (top + bottom of list):

- Select all on page
- Select all matching results (across pages) — see 11.6
- Bulk Actions bar (disabled until at least 1 item selected)

10. Batch Contents View — Record List (inside one Batch)

10.1 Entry rule

Button “Open batch contents” is enabled only when **exactly one batch** is selected. (If 0 or >1 selected → disabled.)

10.2 Record list must mirror Batch list UX

-

The Record list view must be **visually and functionally identical** to the Batch list:

- Checkbox per record (multi-select)

- Quick “Open record details” (read-only preview)
- Same bulk action bar
- Same “Select all” mechanics

Each Record row MUST include:

- Checkbox
- Record name (e.g., Product title / City name)
- Canonical URL
- Page type
- AI status: OK / OUTDATED / NEEDS REVIEW
 - **Sync status (Domains)** – one field with 3 substatuses
 - RAW: In sync / Outdated
 - OFFER: In sync / Outdated (PRODUCT only)
 - REALTIME: In sync / Outdated (PRODUCT only)
 - **Blocking flags (count + top reason)** – “Eligible records” in overlay:
 - RELATIONSHIP_MISSING
 - OFFER_MISSING
 - FACTUAL_MISSING (CITY/ATTRACTION)
 - MEDIA_INVALID (warning)
 - AI_GENERATION_FAILED (warning)
- Locks summary (how many blocks locked)
- “Delete record from batch” (role-gated, confirmation)

Global controls (top + bottom):

- Select all on page
- Select all matching results (across pages)
- Bulk Actions bar

11. Bulk Actions Bar — Rules (applies to both views)

11.1 Activation rule

Bulk Actions bar is disabled until at least **one** item (batch or record) is selected.

11.2 Scope rule

Actions apply to the **current scope**:

- In Batch list: actions apply to selected batches (and implicitly to all records inside them, if action supports it)
- In Record list: actions apply to selected records only

11.3 “Select all” behavior (two-level)

The UI must support BOTH:

- A) Select all on current page
- B) Select all matching results across pages (across the active filters/search)

11.4 Required controls in Bulk Actions bar

- Content Actions (opens overlay)
 - Mark as Reviewed
 - “Mark as Reviewed” resolves **change signals only** (awareness/workflow). It must NOT clear:
 - OUTDATED_RAW / OUTDATED_OFFER / OUTDATED_REALTIME
 - Blocking_flags
 - index_state restrictions
 - Delete selected (role-gated)
 - Clear selection
-

12. Content Actions — Mass Regeneration Overlay (the key expansion)

12.1 Purpose

Content Actions overlay exists to **regenerate selected content blocks** across many records:

- for fixing OUTDATED blocks
- AND for proactive refresh after prompt improvements
- WITHOUT triggering keyword research
- WITHOUT changing keyword ownership

- WITHOUT manual editing
- WITHOUT auto-publishing

12.2 Open behavior

When user clicks “Content Actions”, open an overlay (not a new page) showing:

- Selected scope summary (records or batches)
- Filters summary (page type, language, batch)
- Block selection + caps + preview of impact

12.3 Header (mandatory)

Header must show:

- “Batch Content Control”
- Selection scope:
 - “Records selected: X” (or “Batches selected: X → Records affected: Y”)
- “Eligible records (after locks & validation):”
 - **Excluded reasons breakdown** (must be shown):
 - excluded_due_to_blocking_flags (count)
 - excluded_due_to_manual_lock (count per block)
 - excluded_due_to_missing_dependencies (count)
- Warning: “Actions apply to ALL selected eligible records.”

13. Content Block Library — What can be regenerated (explicit list)

This **MUST** be a fixed registry (developer enforced), not dynamic guesswork.
Blocks are selectable independently and each has its own cap.

13.1 Product Page — Regeneratable blocks

- Title (H1) → CMS: h1
- Meta Title → CMS: meta_title
- Meta Description → CMS: meta_description

- Snippet (Short description, DEE variants) → CMS: snippet_default + snippet_family/couples/...
- Highlights → CMS: highlights[]
- Long Description (Main body) → CMS: long_description_html (or product_overview_html — final naming must be one)
- “Why travelers love this tour” bullets → CMS: love_bullets[]
- Itinerary Overview (indexable) → CMS: itinerary_overview_html
- FAQ (product-level) → CMS: faq_items[]
- Trip DNA labels (ONLY if defined as AI-suggested fields) → CMS: trip_dna_* (optional; role-gated)

13.2 Other landing page types — blocks (must exist in registry)

City Page:

- City H1
- Meta Title / Meta Description
- Snippet variants (DEE)
- Overview (stable SEO block)
- Local tips
- FAQ

Things-to-do Page:

- H1
- Meta Title / Meta Description
- Snippet variants (DEE)
- Overview
- FAQ

Attraction Page:

- H1
- Meta Title / Meta Description
- Snippet variants (DEE)
- Overview
- History (AI allowed only if sources pass validation — optional)
- FAQ

Near-page:

- H1
- Meta Title / Meta Description
- Snippet variants (DEE)
- Overview
- Comparison intro copy
- FAQ

Country / Region:

- H1
- Meta Title / Meta Description
- Snippet variants (DEE)
- Overview
- Highlights bullets
- FAQ

Blog Page (editorial safety):

- Meta Title / Meta Description ONLY (optional; role-gated)
- NO body regeneration in Pipeline

Expert Page:

- Snippet
- Meta Title / Meta Description
- FAQ

Package / Bundle Page:

- H1 (bundle name)
- Meta Title / Meta Description
- Short description
- Long description
- Highlights
- FAQ

Explore Map / AI Feed Context:

- H1
- Snippet
- Meta Title / Meta Description (optional)

Hard rule: If a block is not in this registry → it cannot be regenerated via Pipeline.

13.3 Global dependency rules

Blocks eligible for AI regeneration require:

- RAW domain present and valid (published_for_raw_version exists OR entity is draft)
- No blocking flags for required dependencies of the block

Product-only:

- Content blocks (Title/Snippet/Long Description/FAQ/Highlights/Itinerary Overview) depend on **raw_version**
- Offer-related fields are **never regenerated by AI** (offer_version is handled by ingestion)
- Realtime is **never regenerated by AI** (realtime_offer_version is ingestion/caching)

14. Central Control Table — Block Selection + Caps (replace the old simplified table)

The overlay must include a table like below - link to table

 IT: AI & programmatic SEO - schemes a

Cap rules:

- Caps are per single “Regenerate” job request (per block).
 - If user selects multiple blocks, system enforces the lowest cap and auto-splits into multiple jobs.
 - Blocks in OFFER or REALTIME domains are visible as read-only status items (OUTDATED_OFFER / OUTDATED_REALTIME) but cannot be selected for AI regeneration.
-

15. Selection Mechanics inside the Overlay (must be explicit)

15.1 Block selection is independent from record selection

User selects records first (checkboxes).

Then inside overlay selects which blocks to regenerate for those records.

15.2 “Select all blocks” convenience

Overlay must provide:

- Select all visible blocks
- Clear block selection

15.3 Validation before enabling “Regenerate”

Button “Regenerate Selected Blocks” is enabled only if:

- ≥ 1 record selected AND
 - ≥ 1 block selected AND
 - $\text{selected_records_count} \leq \text{cap_for_each_selected_block}$ (or system auto-splits into jobs)
-

16. Pagination, visibility and “Select all across results” (must be unambiguous)

16.1 Records per page

Pipeline lists must support:

- 50 / 100 / 200 per page selector (default 100)

16.2 Select all on page

Selects only currently displayed rows.

16.3 Select all matching results (across pages)

After “Select all on page”, UI shows a banner:

“You selected 100 items on this page. Select all 20,413 items matching your filters?”

If user clicks “Select all matching”:

- selection becomes a query-based selection (filter snapshot stored)
- UI must display “All matching selected: 20,413”
- any change of filters clears the selection (with warning)
- filter snapshot + domain filters (OUTDATED_RAW/OFFER/REALTIME)
- $\text{selection_count_by_entity_type}$ (so caps can be computed safely)

16.4 Hard safety cap

Even if user selects 20,413 records:

- system MUST enforce per-block caps and auto-split into jobs (see 17)
- user sees job count before confirming

17. Execution Model — Job Queue (how 20k title regen actually works)

17.1 Why job queue is required

Large selection (e.g., 20k Titles) must not be executed as a single synchronous operation. Pipeline must create background jobs (server-side) with consistent chunk sizes.

17.2 Chunking rule

For each selected block:

- `chunk_size = min(default_cap_per_job, system_hard_max_chunk)`
- system creates N jobs:
 - `job_1`: records `1..chunk_size`
 - `job_2`: records `chunk_size+1..2*chunk_size`
 - etc.

Example:

Regenerate Title (H1) for 20,000 products

`cap = 20,000` → 1 job

Regenerate Long Description for 20,000 products

`cap = 2,000` → 10 jobs

17.3 Confirmation screen must show job plan

Before user confirms, show:

- blocks selected
 - total eligible records
 - locked/skipped records counts per block
 - total jobs to be created
 - estimated unit/tokens cost (if you track it)
 - “This will not publish automatically.”
-

18. Locks, skips, and deterministic outcomes (must be shown)

18.1 Lock rule

Locks are per-block:

- If Long Description is locked on a record → skip only that block for that record.
- Do not skip other blocks for the same record.

18.2 Skip reporting

After job completion, system must report:

- regenerated_count
- skipped_locked_count
- failed_count
- validation_failed_count (e.g., FAQ uniqueness)

18.3 Output versioning

Every regen produces:

- new AI version per block
- previous version retained for diff in Content & Structure View
- status: Draft (not published)

19. Allowed Actions (expanded and explicit)

19.1 Regenerate Selected Blocks (AI)

Creates new AI versions for selected blocks on selected records.

Must NOT:

- change keyword ownership
- re-run keyword research

- change canonical URL
- override locks
- auto-publish

19.2 Mark as Reviewed

Clears attention flags for selected records/blocks.
No content changes.

19.3 Unlock Selected Blocks (role-gated)

SEO Lead/Admin only.

Requires reason.

Unlock applies ONLY to chosen blocks, not the entire record.

19.4 Delete selected (role-gated)

Batch delete:

- deletes batch files OR removes records from a batch (depending on scope)
Requires confirmation.

20. Navigation shortcuts — “regenerate only Titles” (your key requirement)

Pipeline must support a quick filter mode:

- View mode dropdown:
 - All blocks
 - Titles only
 - Meta only
 - Snippets only
 - Long-form only (Long Description + FAQ)

When “Titles only” is selected:

- list shows the Title block status prominently

- Content Actions overlay opens with “Title (H1)” pre-selected
 - cap displayed clearly (e.g., 20,000)
-

21. Important invariant (keep your current rule, but make it visible)

Pipeline View:

- no typing
- no manual edits
- no creating locks

Manual changes happen only in Content & Structure View:

- edit text
- lock blocks
- compare diffs
- approve/publish

Pipeline executes actions; Content & Structure resolves content.

15. Where Version Mismatch Is Seen

Pipeline View

- high-level mismatch badge
- status = Review Required

Content & Structure View

RAW version: v2

AI version: v1

Published version: v1 ⚠

Reason: RAW updated, manual lock present

Pipeline = detection

Content & Structure = resolution

16. Final One-Sentence Summary

Pipeline View tells you what is happening and lets you act safely at scale; Content & Structure View is where humans actually edit and decide.

11.2 Pipeline caps per block (must be enforced)

Each block type must have its own batch cap to allow large-scale re-runs safely:

- Title / H1: high cap (e.g., up to 20,000 per action)
- Snippet variants: medium cap
- Long description: lower cap
- FAQ: lower cap
- Meta title/description: medium cap
Reason: cost + latency + risk differ by block size.

3.2.2 CONTENT & STRUCTURE VIEW - UI Option for marketing panel

UX PRINCIPLE

Marketing must never guess. If something is automated, the CMS must explain it. If something is risky, the CMS must slow it down.

01. Structure-First Content Management Console

0. Purpose of Content & Structure View

Content & Structure View is the **editorial and structural workspace** of the system.

It exists to answer one core question:

“Where is this content, what exactly does it contain, and how do I edit it safely?”

This view is used to:

- Browse content by **geographic and logical hierarchy**
- Inspect **individual records** (products, attractions, near-pages)
- Manually edit text fields
- Lock or unlock **specific content blocks**
- Review **detailed versions, diffs, and audit history**
- Resolve issues surfaced by Pipeline View

Content & Structure View is **NOT** used for:

- Batch processing
- AI job orchestration
- Mass regeneration across records

Those responsibilities belong **exclusively** to **Pipeline View**.

Important about inventory: Inventory governance (RAW/Pool/Set), eligibility thresholds, and algorithm tuning are NOT exposed in Content & Structure View.

Editors may see a read-only “Inventory status badge” (e.g., Pool/Set membership and INSTANT/MANUAL) for awareness, but all inventory actions and tuning must be performed exclusively in the Product Operations Panel.

1. Core Mental Model (Critical)

There are two clearly separated worlds in the system:

Pipeline View answers:

“What changed and what needs action at scale?”

Content & Structure View answers:

“Where is this content and what exactly does it contain?”

You enter Content & Structure View in two ways:

- 1. Navigation-first — browsing the content tree**
- 2. Problem-first — deep-link from Pipeline View to resolve a specific record**

This separation is intentional and non-negotiable.

2. Where Batches Are Handled (Clarified)

✓ Batches are always initiated, selected, and controlled in Pipeline View

Examples:

- Vienna Products Batch**
- Rome Attractions Batch**
- Near-pages Cycle #4**

✗ Content & Structure View:

- does not run batches**
- does not show batch controls**
- does not orchestrate AI**

Content & Structure works inside the result of a batch, record by record.

3. High-Level Layout of Content & Structure View

3.1 Global Layout

The screen is divided into three permanent zones:

Left: Hierarchical navigation tree

Center: Table or editor (context-dependent)

Top: Context bar (location + actions)

This layout never changes.

4. Left Sidebar — Structure Tree (Navigation)

This is the structure-first, Google Drive–like model you originally defined.

World



Each node displays:

- Status badge (OK / Attention)
- Record count (e.g. Products: 124)
- Alert indicator (if any child record has unresolved issues)

The sidebar is pure navigation:

- no actions
- no AI
- no editing

5. City-Level View — Main Working Table

When a user clicks:

Country → City → Products / Attractions / Near-pages

They enter a City Records Table.

This table answers:

“Which individual records in this city need attention?”

5.1 City Records Table (Example: Products)

Column	Meaning
Name	Record name
Source	Viator / Internal
Published	Yes / No
Content Status	OK / Needs Review
Manual Edits	Yes / No
Locks	Number of locked blocks
AI Version	vX
Published Version	vY
Alerts	 / 
Actions	Open

Rules:

- This table is read-only

- No batch actions here
- Clicking Open enters Record Detail View

5.1A Product Inventory Modes (All Products / Pool / Set)

The City → Products table must support 3 modes (same entity list; no duplicates):

- All Products (RAW): all provider products validated for this city (read-only payload)
- Pool: eligible candidates (dynamic; auto add/remove)
- Set: curated selections (stable; flagged-only on threshold drops)

Mode selection is a UI filter (not a new dataset).

Products are single entities; modes represent membership.

Add columns to City Records Table (Products):

- Pool Member: Yes/No
- Set Member: Yes/No
- confirmationType: INSTANT / MANUAL (from bookingConfirmationSettings.confirmationType)
- Classification: TICKET / PASS / TOUR (system classification; locked logic)
- Archetype Core ID (read-only; Phase C output)
- Flags: Non-instant / Threshold breach / Structural change

Row actions (role-gated, Product Ops):

- Add to Pool (manual override; logged)
- Remove from Pool (logged)
- Promote to Set (auto Pool membership; logged)
- Remove from Set (logged; requires reason)

- Pin / Unpin (Set only; logged)
- Block / Unblock (Pool/Set; logged)
- Open (Record Detail View)

Important stability rule:

- Threshold drops remove product from Pool automatically.
 - If a product is in Set, threshold drops only FLAG it (no auto-removal)
-

6. How Manual Work Happens (Single Record)

6.1 Entry Point

User clicks Open on a row
→ **enters Record Detail View**

This is the only place where manual editing occurs.

7. Record Detail View (Product / Attraction / Near-page)

Record Detail View is deep, explicit, and type-aware.

It is where:

- **all content lives**
 - **all locks exist**
 - **all diffs are reviewed**
-

7.1 Sticky Header (Always Visible)

Header shows:

- **Record name**
- **Record type (Product / Attraction / Near-page)**

- Breadcrumb (Country → City → Record)
- Version summary:

RAW v2 | AI v1 | Published v1 ⚠

- Status badge: OK / Mismatch
- Primary actions:
 - Preview public page
 - View audit log

No AI actions here.

8. Tabs in Record Detail View

Tabs are consistent across record types, but content differs.

Tabs:

1. Overview
 2. Content
 3. Factual
 4. SEO
 5. Hooks
 6. History / Audit
-

9. Overview Tab (Read-Only Status)

Purpose: situational awareness

Shows:

- **Record health summary**
- **Version alignment / mismatch**
- **Lock summary**
- **Last AI run**
- **Last manual edit**

No editing allowed.

10. Content Tab (Primary Editorial Workspace)

This is where manual editing happens.

10.1 Content Sections Table

Each record has a section-level control table.

Section	Source	Status	Locked	AI Allowed
Title (H1)	AI	OK	No	Yes
Meta Title	AI	OK	No	Yes
Meta Description	AI	OK	No	Yes
Snippet (default)	AI	Review	No	Yes
Snippet (family)	AI	OK	No	Yes
Full Description	AI	OK	No	Yes

Highlights	AI	OK	No	Yes
Itinerary Overview	AI	OK	No	Yes
Expert Tip	AI	OK	No	Yes
FAQ	AI	OK	No	Yes

Rules:

- Locks are per section
- Manual edits affect only that section
- No section blocks the entire record

10.2 Section Actions

For each section:

- Edit (manual)
- Lock / Unlock (role-based)
- View diff
- Regenerate (single-section AI)

- ✗ No “regenerate all” here
- ✗ No keyword changes here

11. Factual Tab (Read-Only by Default)

Purpose: protect facts

Contains:

- **Opening hours**
- **Address**
- **Coordinates**
- **Duration**
- **Meeting point**
- **Cancellation rules**
- **Accessibility**

Rules:

- **AI cannot modify factual values**
 - **Manual override available (role-gated)**
 - **Overrides are explicitly flagged and versioned**
-

12. SEO Tab

Purpose: SEO clarity without risk

Shows:

- **Primary keyword (read-only)**
- **Secondary keywords (read-only)**
- **Canonical URL**
- **Index / noindex state**
- **Schema preview**

Keywords are never editable here.

13. Hooks Tab (Editorial Control)

Purpose: manage conversion hooks, not SEO.

Editable texts and toggles for:

- Discount code hooks
- Wishlist / Send to self
- PDF downloads
- Exit intent copy

Hooks are:

- independent from AI batches
 - independent from keywords
 - safe to edit manually
-

14. History / Audit Tab

Purpose: full traceability

Timeline shows:

- AI generations
- Manual edits
- Locks / unlocks
- Version transitions
- Pipeline-triggered changes

Nothing is hidden.

15. How Issues Flow Between Views

Example: Broken titles after Vienna Batch

1. **Pipeline View detects issue**
2. **User regenerates titles in bulk (Pipeline View)**
3. **Some records still need polish**
4. **User clicks record → Content & Structure**
5. **Manual fix + lock in Content tab**

This is by design.

16. Processing a New Single Product (End-to-End)

Scenario:

- **Vienna batch already completed**
- **New product arrives later**

Flow:

1. **Product appears in Pipeline View**
2. **Status: NOT PROCESSED**
3. **User clicks Run AI Processing**
4. **AI creates first AI version**
5. **Product appears under:
City → Products**
6. **Manual work (if needed) happens here**

No re-running city batches.

17. Button Naming (Final, Consistent)

Pipeline View

- **Run AI Processing**
- **Regenerate Selected Content**
- **Mark as Reviewed**

Destination Management

- **Run SEO Auto**

Content & Structure

- **Edit**
- **Lock / Unlock**
- **Regenerate (single section)**

Buttons never overlap across views.

18. Why This Design Works

This structure:

- **prevents accidental SEO damage**
- **enforces clear responsibility boundaries**
- **scales to thousands of records**
- **allows AI to work at scale**
- **allows humans to work precisely**

No screen does more than one job.

19. Final One-Sentence Summary

Content & Structure View is where humans understand, edit, and lock individual records; Pipeline View is where AI is safely controlled at scale.

02. MANUAL OPTIMIZATION (CMS)

Inside CMS, user may:

Page-level actions

- Approve AI suggestions
- Regenerate meta manually
- Regenerate intro manually
- Regenerate highlights manually
- Add or modify FAQ manually
(locked fields override AI forever until unlocked)
- Unlock AI for specific field
- Assign trending keyword
- Approve/Reject near-page draft
- Approve canonical merge suggestion
- Force re-run of optimization

4.1.1 Dynamic Experience Engine Controls (CMS, MVP)

Editors must be able to manage dynamic variants without touching SEO-core content.

A) Variant Fields (editable + lockable per field)

For each entity/page (product, city, attraction, region, country, things-to-do, near-page, blog):

- Dynamic Snippet variants: snippet_default / snippet_family / snippet_couples / snippet_seniors / snippet_budget / snippet_solo / snippet_explorers
- Review Summary variants (where reviews exist): review_summary_default / review_summary_family / review_summary_couples / review_summary_seniors / review_summary_budget / review_summary_solo / review_summary_explorers
Each field supports: edit, regenerate (AI), lock/unlock, and “revert to AI”.

B) Segment Preview (required for QA)

CMS must include a preview switcher: according to **4.8.9 Stage F — Profile model**

Preview must show exactly what the user would see for:

- snippet variant
- review summary variant
- ordering profile result (top list ordering + related tours ordering + bundles ordering if present)

C) Ordering Profile Overrides (safe, non-destructive)

Editors can:

- pin items to the top for a specific segment (limited slots)
- blacklist items from a specific segment feed
- force fallback to default for a specific segment
All overrides must be stored separately from the base ranking model and must not change canonical URLs.

4.1.2 Human-Face / Live Layer Controls (CMS, MVP)

Editors must be able to manage guide identity, last-tour freshness, and photo eligibility without rewriting reviews.

A) Guide Assignment (Rosotravel-owned products only)

CMS must allow linking product → one or more guides from /cms/guides/ and setting:

- primary_guide_id
- guide_carousel_order
- guide_display_enabled (on/off)
If no guide assigned (or OTA import) → render the default “Rosotravel Guide Team” block.

B) Last Tour Snapshot Governance

CMS must include:

- toggle: show_last_tour_snapshot (on/off)
- fallback mode selector if snapshot is stale: (1) show month/year + disclaimer, (2) show “last 30 days summary”, (3) hide block
- moderation for last-tour photos: approve/hide specific guest images from appearing in the “See photos from this date” gallery
- manual “rebuild snapshot” action (forces recomputation from review/tour-date data)

C) Review Integrity Rules (editor-safe)

Editors can hide an individual review or image from UI (moderation), but cannot edit the original review text from OTA feeds. Any UI-level summary must be regenerated from the remaining visible reviews only.

Product detail view for SEO

(Content & Structure View / Single Record Workspace)

0. Purpose of Product Record Detail View

Product Record Detail View is the single source of truth for one product page.

It is used to:

- **inspect *everything* that appears on the product page**
- **manually edit AI-generated text**
- **lock / unlock individual content blocks**
- **review version alignment and diffs**
- **configure hooks and personalization inputs**
- **safely resolve issues surfaced by Pipeline View**

It is NOT used to:

- **run batches**
 - **re-run keyword research**
 - **regenerate many products at once**
 - **manage destinations (country / region / city)**
-

1. Sticky Header (Always Visible)

1.1 Identity & Context

- **Product Name (current H1)**
- **Product Type:**
 - **Internal (Rosotravel)**
 - **External (Viator / OTA)**
- **Breadcrumb:**
Country → City → Product

1.2 Version & Status Summary

Displayed clearly and permanently:

RAW version: v2

AI version: v1

Published version: v1 ⚠

Status: MISMATCH

Reason: RAW updated, manual lock present

1.3 Primary Actions

- **Preview public page**
- **View full audit log**

- ✗ **No AI actions**
- ✗ **No regeneration buttons**
- ✗ **No publishing actions**

This header is informational only.

2. Tabs Overview (Product-Specific)

Tabs are ordered by editorial workflow, not by front-end layout:

1. **Overview**

2. Content
 3. Factual
 4. SEO & Keywords
 5. Personalization (DEE)
 6. Hooks
 7. History / Audit
-

3. Overview Tab — Product Health Snapshot

Purpose: fast understanding, no editing.

Displays:

- **Publication state (Published / Draft / Inactive)**
- **Source (Internal / OTA)**
- **Last AI generation timestamp**
- **Last manual edit timestamp**
- **Lock summary:**
 - **“3 sections locked”**
- **Indexing state (Index / Noindex)**
- **Warnings (read-only):**
 - **Missing hero image**
 - **Manual locks blocking updates**
 - **Outdated AI vs RAW**

No actions allowed here.

4. Content Tab — Core Editorial Control (Most Important)

Fundamental Rule (must be stated explicitly)

Each content block has exactly one current Source: AI, Manual, or System. I may generate the baseline, but the moment a human edits the block, the Source becomes Manual.

Definitions

Indicates who produced the currently active version of a content block.

Possible values:

- **AI** – last version generated by AI
- **Manual** – last version edited by a human
- **System** – factual or operational data coming from feeds / APIs

Exactly one value is always present.

Lock

Controls whether a block is protected from AI overwrites.

- **Locked = Yes** → block cannot be regenerated by AI
- **Locked = No** → block may be regenerated (if AI Allowed)

Locks are per block, never global.

AI Allowed

Controls whether AI is permitted to generate or regenerate this block.

- **AI Allowed = Yes** → AI regeneration possible
- **AI Allowed = No** → AI regeneration forbidden

AI Allowed can be No even if the block is not locked (e.g. System blocks).

Final Ownership & Control Table

Final Ownership & Control Table							
Content Block	Initial Source	Editable by Human	Source After Manual Edit	Lockable	AI Allowed	Notes	
Title (H1)	AI	Yes	Manual	Yes	Yes	Manual edit switches source to Manual; AI may still regenerate if not locked	
Meta title	AI	Yes	Manual	Yes	Yes	SEO-sensitive, lock recommended after optimization	
Meta description	AI	Yes	Manual	Yes	Yes	Same rules as meta title	
Dynamic snippets (DEE variants)	AI	Yes	Manual	Yes	Yes	Each variant handled independently	
Full description	AI	Yes	Manual	Yes	Yes	Long-form SEO content	
Highlights	AI	Yes	Manual	Yes	Yes	Structured bullets	
Why travelers love this tour (review summary)	AI	Yes	Manual	Yes	Yes	AI summarization; no raw reviews	
Itinerary overview (SEO)	AI	Yes	Manual	Yes	Yes	Narrative only, not step-by-step	
Expert tip / Local insight	AI	Yes	Manual	Yes	Yes	Starts AI, often manually refined	
Guide team description	AI	Yes	Manual	Yes	Yes	AI baseline, brand-safe copy	
Trip DNA attributes	AI (proposal)	Yes	Manual	Yes	Yes	Human has final say	
FAQ	AI	Yes	Manual	Yes	Yes	Exactly defined count per page	
Included / Not included	System	Yes (override)	Manual	Yes	No	Values are factual; AI never rewrites	
Meeting point	System	Yes (override)	Manual	Yes	No	AI cannot change factual values	
Price	System	No	System	No	No	Feed-only	
Availability	System	No	System	No	No	Feed-only	
Cancellation policy	System	Yes (override)	Manual	Yes	No	Text labels may be rephrased, values not	
Duration	System	Yes (override)	Manual	Yes	No	Factual	
Start times	System	No	System	No	No	Factual	



➤ Tabele in excel  IT: AI & programmatic SEO - schemes

State Transition Rules (Critical)

1. AI → Manual

- **Happens immediately when a human edits a block**
- **Source switches from AI → Manual**
- **Previous AI version remains in history/audit log**

2. Manual + Not Locked

- **Block source = Manual**
- **AI Allowed = Yes**
- **AI regeneration can overwrite manual text (intentional)**

Use case:

Minor cosmetic edit, still want AI to iterate later.

3. Manual + Locked

- **Block source = Manual**
- **AI Allowed = No**
- **AI regeneration must skip this block**

Use case:

Final SEO copy, legal text, carefully crafted messaging.

4. System Blocks

- **Source = System**
- **AI Allowed = No**
- **Locking controls API overwrite, not AI**

AI may:

- **rephrase labels**
- **generate narrative summaries elsewhere**

AI may NOT:

- **change factual values**
-

How This Appears in UI (single block)

Example:

- **Title (H1)**
- **Source: Manual**
- **Locked: No**
- **AI Allowed: Yes**
- **Last edited by: Anna K. (SEO)**

CTA buttons:

- **Edit**
 - **Lock**
 - **Regenerate (AI) ← visible only if AI Allowed = Yes**
-

4.2 Editing a Section

Clicking Edit opens an inline editor with:

- **Current published text**
- **AI source indicator**
- **Last modified by / when**
- **Character limits (where applicable)**
- **Inline validation (SEO rules)**

Available actions per section:

- **Save (manual edit)**
 - **Lock / Unlock (role-gated)**
 - **View diff vs previous version**
 - **Regenerate (AI, single block)**
-

5. Factual Tab — Immutable Truth Layer

Purpose: protect correctness.

5.1 Factual Fields (Read-Only by Default)

- **Duration**
- **Languages**
- **Meeting point**
- **Cancellation policy**
- **Accessibility**
- **Ticket type**
- **Included / Not included**
- **Options / variants factuals**

5.2 Rules

- **AI may format labels, never values**
- **AI may not invent or change facts**
- **OTA facts override AI**
- **Manual override available:**
 - **Role-gated**
 - **Explicit “Manual factual override” flag**
 - **Appears in audit log**

Factual changes do not regenerate SEO text automatically.

6. SEO & Keywords Tab

Purpose: transparency without risk.

Displays (Read-Only):

- **Primary keyword (locked)**
- **Secondary keywords (locked)**
- **Keyword ownership (this product)**
- **Canonical URL**
- **Index / noindex state**
- **Schema preview:**
 - **TouristTrip**
 - **Offer**
 - **FAQPage**
 - **Review / AggregateRating**

Rules

- **Keywords are assigned upstream**
- **No keyword edits here**
- **No re-clustering**
- **No ownership changes**

7. Personalization (DEE) Tab

Purpose: control Dynamic Experience Engine inputs.

7.1 Snippet Variants

Editable blocks:

- **snippet_XX**

**DEE selects at render time.
SEO remains stable.**

7.2 Dynamic Highlights Add-Ons

Per segment:

- up to 2 optional highlights
- appended after base highlights

7.3 Segment-Specific “Why This Experience” Add-Ons

Optional 1–2 sentence blocks per segment.

All fields:

- AI suggested
 - Human editable
 - Lockable
-

8. Hooks Tab — Conversion Layer

Purpose: manage monetization hooks safely.

Editable content & toggles for:

- Lifetime discount hook
- Wishlist / Email My Trip
- Send this tour to myself
- Exit-intent offer
- Hook copy overrides

Hooks:

- do NOT affect SEO
 - do NOT affect AI versions
 - are safe to change anytime
-

9. History / Audit Tab — Full Traceability

Timeline includes:

- **RAW ingestion events**
- **AI generations (what block, when)**
- **Manual edits (who, what, why)**
- **Locks / unlocks**
- **Publishing events**
- **Pipeline-triggered changes**

This is immutable.

10. OTA vs Internal Product Differences (Important)

External (Viator / OTA)

- **Full itinerary: system-owned, non-indexable**
- **Images: non-indexable, no alt text**
- **Guide info: generic block only**
- **Reviews: read-only, summarized only**

Internal (Rosotravel)

- **Itinerary may be AI-generated**
- **Guide profiles editable**
- **Full personalization enabled**
- **Higher manual control**

Same UI, different permissions.

11. How This View Relates to Other Views

Action	Where
Run batch	Pipeline View
Regenerate many products	Pipeline View
Fix one product	Content & Structure
Edit text	Content & Structure
Lock sections	Content & Structure
Resolve mismatch	Content & Structure
Keyword changes	SEO Admin / Destination flows

12. Final One-Sentence Summary

Product Record Detail View is a precise, section-level editorial workspace that exposes the full product page while enforcing strict separation between AI automation, factual truth, and human control.

Attraction detail View (only differences)

TAB: Overview

(Differences only)

Purpose

High-level understanding of **what this attraction is, where it lives, and how it is monetized indirectly.**

Additional elements vs Product

Attraction Overview Summary (read-only):

- Attraction name
- City / Country
- Attraction type (museum, landmark, site, etc.)
- SEO enabled: Yes / No
- Total related tours (count)
- Total reviews (aggregated across tours)

No booking / price / options panel.

TAB: Content

(Most important differences)


Content Sections Table — Attraction

Content Block	Initial Source	Editable	Lockable	AI Allowed	Notes
H1 (Attraction in City)	AI	Yes	Yes	Yes	Pattern-based title
Dynamic snippets (DEE variants)	AI	Yes	Yes	Yes	Segment-based

Overview description	AI	Yes	Yes	Yes	Visitor experience
History / significance	AI (from factual)	Yes	Yes	Yes	Facts validated
Top things to see	AI	Yes	Yes	Yes	Internal POI elements
Visitor tips	AI	Yes	Yes	Yes	Non-factual advice
Nearby attractions	System + AI order	No	No	No	Graph-based
Tours related to attraction	System + DEE	No	No	No	Relation-driven
FAQ	AI	Yes	Yes	Yes	Exactly 7
Reviews summary	AI	Yes	Yes	Yes	No raw reviews

What is NOT here (vs Product)

-  Included / Not included
-  Itinerary overview
-  Options / variants
-  Booking CTAs
-  Guide profiles



TAB: Factual

(Much more important than in Product View)

Purpose

Single source of truth for **verified attraction facts**.

Factual blocks (read-only by default)

Block	Source	Manual override	AI Allowed
Opening hours	Google Places	Yes	No
Address	Google Places	Yes	No
Coordinates	Google Places / Wikidata	Yes	No
Official website	Google Places / Wikidata	Yes	No
Ticket information	Official site	Yes	No
Accessibility	Official / manual	Yes	No
Rules / dress code	Official / manual	Yes	No

Rules:

- AI **never changes values**
- Manual override possible
- Overrides require reason

- Overrides show warning badge

TAB: SEO

(Different intent, same mechanics)

Differences vs Product

- Primary keyword = attraction name + city
- No transactional keywords allowed
- Focus on:
 - entity clarity
 - factual consistency
 - internal linking strength

SEO Blocks

Block	Editable	Lockable	AI Allowed
Meta title	Yes	Yes	Yes
Meta description	Yes	Yes	Yes
FAQ schema	Yes	Yes	Yes
Internal links (preview)	No	No	No

TAB: History / Audit

(Identical mechanics, different signals)

Additional logged events:

- Factual data refresh (Google Places / Wikidata)
 - Factual conflict resolution
 - SEO enable / disable toggle
 - Relation graph updates (nearby attractions, tours)
-

4. What Makes Attraction View UNIQUE (Conceptual)

1. Factual authority > persuasion

Attraction pages **build trust**, not urgency.

2. Indirect monetization

Conversion happens via:

- related tours
- bundles
- nearby activities

3. Stronger coupling to data integrity

Wrong facts = SEO & UX damage.

4. No pipeline regeneration of facts

Facts flow **into** AI, never the other way.

5. Summary

Attraction Record Detail View is structurally identical to Product View, but content-wise optimized for informational authority rather than conversion. It replaces itinerary and booking logic with richer factual blocks, attraction-specific narratives, and indirect monetization via related tours.

The same AI/manual ownership, locking, and versioning rules apply, ensuring consistency across the entire system.

Near-Page Record Detail View

(Programmatic SEO Page – City-Level Category / Keyword Page)

1. Purpose and Role

A **Near-Page** is a programmatic SEO record designed to capture **long-tail, city-specific search intent** (e.g. *“Rome underground Colosseum tours”*, *“Vienna family activities”*).

A Near-Page is **not** a product page and **not** an attraction page.
Its primary role is to:

- capture keyword-driven demand at scale,
- connect users to relevant products or bundles,
- operate with minimal manual intervention,
- remain structurally consistent across hundreds or thousands of pages.

Near-Pages are **AI-generated by default** and are intentionally restrictive in terms of manual editing.

2. Entry Points

A Near-Page Record Detail View can be opened:

- from **Pipeline View** (to resolve batch-level issues or review generated output),
 - from **Content & Structure View** (via city → near-pages navigation),
 - from **Issue notifications** (indexing, lifecycle, or performance alerts).
-

3. Header (Sticky)

The header is fixed and visible on all tabs.

Elements:

- Near-Page title (H1 preview)
- Record type: *Near-Page*
- Breadcrumb:
Country → City → Near-Page
- Version summary:
RAW vX | AI vY | Published vZ
- Status badge:
OK / Review Required / Draft / Archived
- Primary actions:
 - Preview public page
 - View audit log

No content editing is available in the header.

4. Tabs Overview

Near-Page Record Detail View contains the following tabs:

1. Overview
2. Content
3. SEO
4. History / Audit

There is **no Factual tab** and **no Operational tab** for Near-Pages.

5. Tab: Overview

Purpose

To explain **what this Near-Page represents**, how it was generated, and whether it is ready for indexation.

Fields (Read-Only)

Keyword Context

- Primary keyword
- Keyword cluster ID
- Search intent: *Transactional / Mixed*
- Data source: *GSC / SEMrush*
- Generation type: *Immediate / GSC-based cycle*

Lifecycle Status

- Draft (noindex)
- Candidate
- Published (indexable)
- Archived

Parent Context

- City
- Parent taxonomy category (if applicable)
- Canonical target (if defined)

This tab contains **no editable fields**.

6. Tab: Content

Core Principle

Near-Page content is **AI-owned by default**.

Manual edits are:

- allowed only for **SEO Lead or Admin**,

- discouraged,
- always flagged as SEO risk,
- tracked explicitly in the audit log.

Content Sections Table

Content Block	Source	Editable	Lockable	AI Allowed	Notes
H1	AI	Limited	Yes	Yes	Keyword-driven
Dynamic snippets (DEE variants)	AI	No	No	Yes	UX-only
Overview (2 paragraphs)	AI	Limited	Yes	Yes	Intent explanation
Comparison block intro	AI	No	No	Yes	Structural
Product list	System	No	No	No	Auto-generated
Bundles section	System	No	No	No	Optional
Map summary text	AI	No	No	Yes	Short
FAQ (7 items)	AI	Limited	Yes	Yes	Global FAQ rules

“Limited Edit” Rule

A limited edit:

- requires justification,
- does not change keyword ownership,
- does not affect batch logic,
- triggers a warning:
 “Manual override on programmatic page – monitor performance.”

Near pages must support an “Expert byline / signature” to increase authenticity and trust (“Reviewed by ... at Rosotravel”).

This byline is a governed UI component and must be manageable at scale.

A) Expert Profile entity (CMS)

Near pages reference an Expert Profile stored in CMS.

Expert Profile fields (minimum):

- expert_profile_id (stable)
- display_name
- role_label (controlled role set: Family Planner Expert / Senior Comfort Expert / Deep Dive Expert / Active Adventure Expert / First-time Visitor Expert / Romantic Trip Expert according to **4.8.9 Stage F — Profile model**)
- avatar_image_id (optional)
- bio_short (optional)
- languages (optional)
- is_active (boolean)

B) Near Page fields (per page)

Add fields to Near Page record:

- expert_byline_enabled (boolean; default false)
- expert_profile_id (nullable; required if enabled)
- reviewed_by_user_id (system field; audit)

- reviewed_at (timestamp)
- byline_text_override (optional; CMS token override; default uses SSOT token + expert profile fields)

Display rule:

If expert_byline_enabled = true and expert_profile_id is set:

Show byline on the near page:

“Reviewed by {{Expert_Name}} — {{Expert_Role}} at Rosotravel”

(copy text is CMS-editable via SSOT tokens; name/role from profile)

C) UI placement in Near Page editor

Add a dedicated panel in Near Page editor:

“Expert Signature / Review”

- Toggle: Enable expert byline
- Dropdown: Select expert profile (searchable)
- Preview: shows final rendered byline text
- Optional: “Mark as reviewed” action (sets reviewed_at + reviewed_by_user_id)

D) Bulk Actions (Near Pages list)

Add bulk actions:

1) Auto-assigned to expert

Rules to be defined during discovery

2) “Assign expert byline”

- choose expert_profile_id
- set expert_byline_enabled = true
- set reviewed_at and reviewed_by_user_id

3) “Remove expert byline”

- set expert_byline_enabled = false
- clear expert_profile_id (optional; or keep for history; define in implementation)

Bulk action must show:

- count of selected pages
- confirmation modal with summary
- audit log entry

E) Governance & Audit

- Only SEO Lead and Admin (or assigned Content Governance role) can assign/enable expert bylines.

- All changes must be audit-logged:

actor, role, timestamp, affected page IDs count, previous/new expert_profile_id, enable flag changes.

F) Discovery note

The final role set and auto-assignment rules (if any) are finalized in discovery.

This section defines the required UI capability and bulk workflow.

7. Tab: SEO

Purpose

Ensure consistent, predictable indexation behavior at scale.

SEO Fields

Field	Editable	Lockable	AI Allowed	Notes
-------	----------	----------	------------	-------

Meta title	Limited	Yes	Yes	Template-based
Meta description	Limited	Yes	Yes	Template-based
Canonical URL	No	No	No	System-controlled
Index state	No	No	No	Lifecycle-driven
FAQ schema	Yes	Yes	Yes	AI-generated

Manual SEO edits are **strongly discouraged** for Near-Pages.

8. Tab: History / Audit

Purpose

Provide a full, immutable timeline of Near-Page lifecycle events.

Recorded Events

- Near-Page generated (cycle ID)
- Keyword promoted from candidate
- Index state changes (noindex → index)
- Performance threshold reached
- Archived due to low performance
- Manual override applied or removed
- AI regeneration triggered

No actions are available in this tab.

9. Explicit Exclusions (By Design)

Near-Pages do **not** contain:

- factual visitor information,
- booking or pricing blocks,
- reviews editing,
- guide or expert sections,
- itinerary content,
- included / excluded lists,
- manual merchandising.

Near-Pages must never evolve into full content pages.

10. Comparison Summary

Feature	Product	Attraction	Near-Page
Primary goal	Conversion	Authority	SEO capture
Manual editing	Common	Common	Rare
Factual data	Medium	High	None
Booking	Yes	No	No
Reviews	Direct	Aggregated	None
AI regeneration	Frequent	Moderate	High-scale

Index control Per record Per record Lifecycle-driven

11. Operational Rule (Critical)

If a Near-Page requires repeated manual editing,
the problem must be fixed at the **prompt, template, or system-rule level**,
not at the individual page level.


The UI enforces this principle by design.

12. Final Statement

The Near-Page Record Detail View is a controlled, AI-first editor for programmatic SEO pages, intentionally designed to minimize manual intervention and ensure scalable, predictable performance across large keyword sets.

01. Move Blog from current General Administration to Marketing Panel

02. AI Proc&Publi: Cities, PD, POI, NP



- This panel has 4 layers. First is the entity layer (country), second is a (city country) and last is a product)
- City level batch for example Rome
- List of each product with product code. Version,
- Button Start AI Processing (can process 1 or all products)
- “Country” Raw Feed - > button
- List of each product with product code. version, status locking,
- 03. Option: Keyword ownership panel
- clearly show what keywords are scheduled to which landing page
- Marketer can unlock and edit
- 04. Alerts:
- 01. This is linked to inbox it is just shown also in this panel

02. Semantic Diff Viewer

Shows difference between:

- RAW → transformed
- transformed → AI draft
- AI draft → published
- published → optimized

03. Entity Graph Viewer

Visual nodes:

- parent entities
- children entities
- nearby attractions
- related pages
- topic cluster membership

03. AI Draft Layer

F.12 CMS Integration

The CMS should expose validation results per page:

1. Validation summary widget
 - traffic-light badge: Green (ok), Yellow (needs_review), Red (blocked)
 - list of top 3 issues (by severity)
2. Detailed validation tab
 - grouped by category:
 - Structure
 - Content
 - SEO
 - Factual
 - Duplication
 - FAQ
 - Schema

Each issue:

- code, message, field, severity
- “View context” (opens the field with highlighted text)

3. Manual override (admin-only)

- ability to mark certain warnings as “acknowledged”

ability to override some blocking rules (e.g. schema optional fields) with justification text stored in audit log.

G.08 CMS UI Requirements for Publishing

G.08.1 Publishing panel per entity (Product, City, Attraction, etc.)

On product / city / attraction edit screens, add a “Publishing” box with:

- Status dropdown: **Draft** / **Scheduled** / **Published** / **Unpublished** / **Noindex**.
- Language tabs (EN, DE, PL, etc.) with separate status per locale.
- **Publish now** button (immediate).
- **Schedule publish at [datetime]** field.
- **Unpublish** button (with reason).
- **Force 301 redirect to [URL]** field for deprecating pages.

G.08.2 Preview and diff

- “Preview draft” link must render the page with draft content, not yet live.
- “Compare with live” link shows diff between current draft and published version (titles, H1s, meta, FAQ, highlights).

G.08.3 SEO summary widget in CMS

On every entity, display:

- Primary keyword and secondary keywords (from keyword ownership).
- FAQ count and “collision risk” flag.
- Internal linking hints (how many tours link to this city, etc.).
- Indexability status (index/noindex) and current canonical target.

G.08.4 Respect legacy CMS behavior

Given current Laravel CMS + WordPress frontend:

- For countries/cities already created manually, existing fields (title, meta, content) must be imported into the new model and preserved as v1.
- New AI-generated content can be stored as v2+ but must not overwrite human content without explicit marketer action (“Replace with AI version”).

Page SEO State Panel (per URL)

PAGE SEO STATE PANEL PER URL

UX: TABS (top navigation)

TAB 1: Overview

- URL
- Page type
- Language
- Source (AI / human)
- Primary keyword
- Owned keywords
- Intent
- Entity flag
- Seasonal flag

Badge:

INDEXED / NOINDEX / BLOCKED

TAB 2: Indexation & Canonical

- Index state (read-only + reason)
- Canonical target (read-only + rule explanation)
- ter (list languages + status)

Explanation box (mandatory):

“This page is noindex because volume < threshold and no entity override.”

UX MUST show **why**

TAB 3: GSC Validation

- Impressions 28d
- Clicks 28d
- CTR
- Queries (top 10)
- Distinct query class: YES / NO

Buttons:

- [Request promotion to index](#)
- [Request demotion to noindex](#)
- Buttons active only if conditions from G.07 are met

TAB 4: Near-Page Lifecycle

- Parent page
- Creation date
- Stage: draft / contained / eligible / graduated
- GSC proof attached (link)

Buttons:

- [Approve graduation](#)
- [Demote to contained](#)

TAB 5: AI Optimization

- Last rewrite date
- Rewrite count this quarter
- Cooldown timer
- Eligible sections (checkbox list)

Buttons:

- [Approve AI rewrite](#)
- [Reject rewrite](#)
- [Lock section permanently](#)

TAB 6: History & Audit

Timeline:

- Ownership changes
- Canonical changes
- Index state changes
- Rewrites
- Redirects

Każda zmiana:

- user
- timestamp
- reason

Near-Page Control Panel (High Risk)

Purpose

Prevent near-page explosion and cannibalization.

Marketing MUST be able to:

- see all near-pages by city/region
- see creation date, parent, intent, keywords
- see current lifecycle stage:
 - draft
 - contained (noindex)
 - eligible
 - graduated (indexed)
- see GSC proof attached to graduation
- approve or block graduation
- manually demote a near-page back to contained

Hard limits enforced by system (read-only):

- max 1 new near-page per city/week
- no transactional ownership ever

Marketing cannot override these limits.

UX: City → Near-pages list

Columns:

- URL
- Parent
- Intent
- Volume
- Age (days)
- Status
- Conflict flag

Header badge:

“Max 1 new near-page / city / week enforced”

Actions:

- Approve graduation
- Block graduation
- Demote
- No "Create manually" option

Canonical & Duplicate Inspector**Purpose**

Prevent silent SEO disasters.

Marketing MUST be able to:

- see canonical decision reasoning (rule-based explanation)

- see index/noindex reasoning
- see loser/winner resolution explanation
- simulate changes (dry-run) without publishing
- see warnings if:
 - canonical conflicts with ownership
 - index violates volume or intent rules
 - hreflang cluster is broken

Marketing CANNOT:

- set arbitrary canonicals
- canonicalize products upward
- override language rules
-

UX Flow

- Select URL
- System shows:
 - similarity matches
 - winner URL
 - reason for win
 - required resolution

Buttons:

- Approve canonicalization
- Approve redirect (301)
- Keep both (not allowed) → disabled

UX must **prevent wrong actions**, nie tylko ostrzegać.

11 CMS NAVIGATION STRUCTURE (FINAL)

CMS

- |— SEO Control Center
- |— Keyword Governance
- |— Pages
 - | — [Page SEO State]
- |— Near-Pages
- |— Canonical & Duplicates
- |— AI Optimization
- |— Seasonality
- |— Crawl Budget
- |— Redirects
- |— SEO Memory
- |— Alerts

3.2.3 Keyword Ownership - UI Option of marketing panel

01. Keyword Ownership change: UI Sub-option of Keyword Ownership

0. Purpose (Why This View Exists)

The **Static SEO / Keyword Ownership UI** is the **single strategic authority** for managing how the platform targets search demand.

It is the only place where keywords can be:

- **assigned** to canonical pages
- **reassigned** (ownership changes)
- **reserved** (blocked from use)
- **locked** (protected from changes)
- **audited** (full traceability)
- **validated** (conflict + intent rules)
- **versioned** (ownership rules + SEMrush strategy versions)

This UI exists to:

- prevent **keyword cannibalization** across Near-pages, Things-to-do pages, City pages, Attraction pages, and Blog
- ensure **AI regeneration is intentional**, scoped, and predictable
- keep **keyword editing out of all other panels** (Pipeline, Content & Structure, Product records)

This view is **not operational** and **not editorial**.

1. Scope (What This UI Controls)

1.1 Page Types Covered (Ownership + Governance)

This UI governs keyword ownership for:

- **Country pages**
- **Region pages**
- **City pages**
- **Attraction (POI) pages**
- **Things-to-do pages**
- **Near-pages** (programmatic keyword pages)
- **Blog articles** (visibility + reservation/locking status only; editing happens in Blog editor)

1.2 Product Pages (Important Clarification)

Product pages **do not “own” keywords** (they are not canonical owners).

However, keyword changes **can still affect some product-level SEO blocks** (via regeneration rules), because products must stay consistent with the platform’s intent strategy.

So:

- **Ownership:** no

- **Keyword-dependent regeneration:** yes, but only for a limited set of blocks (see Section 8)

1.3 Explicitly NOT in Scope

This UI does **not**:

- edit content text blocks
 - regenerate AI text
 - run operational batches
 - bypass Pipeline View
 - change product inventory, prices, availability, or factual feed fields
-

2. Non-Negotiable Rules

1. **Keywords cannot be edited anywhere else.**
Pipeline View, Content & Structure View, and Product/Attraction record editors must not expose keyword fields.
 2. **Any ownership change invalidates keyword-driven AI outputs.**
No exceptions. The system must mark affected blocks as **OUTDATED** until regenerated and approved.
 3. **Acknowledging a notification never clears an issue.**
Notifications are “seen/unseen.”
SEO issues (OUTDATED/CONFLICT/MISSING/BLOCKED) persist until fixed.
 4. **H2/H3 section headings are not keyword-driven.**
They are **static i18n labels** (global dictionary). Only H1/meta/FAQ/internal links are keyword-driven.
-

3. Core Concepts & Data Model

3.1 Keyword Entity (normalized row)

Each keyword row stores:

- normalized keyword string
- locale/language
- intent class (navigational, informational, commercial, transactional, mix of 2)
- demand signals (SEMrush volume, KD, trend snapshot)
- current ownership state + history

3.2 Ownership States (Strategic)

- **Locked:** cannot be reassigned (SEO Lead/Admin only may unlock with reason)
- **Unlocked:** editable by SEO Lead/Admin
- **Reserved:** blocked from assignment (future use or policy)
- **Legacy-locked:** migrated ownership from legacy system (protected)
- **Rejected:** banned/unusable keyword (policy or quality)

3.3 SEO States (Operational outcome of strategy)

- **OK:** ownership is valid and AI outputs are in sync
- **OUTDATED:** ownership changed and keyword-driven blocks must be regenerated
- **CONFLICT:** overlapping ownership detected (two pages competing)
- **MISSING:** required keyword missing (e.g., Near-page without primary keyword)
- **BLOCKED:** cannot be indexed until resolved (hard rule)

4. High-Level Layout (Screen Structure)

Static SEO / Keyword Ownership UI
Strategic control – high impact changes
[GLOBAL SEO HEALTH BAR: X Critical Issues] (sticky)

[LEFT FILTER SIDEBAR] [MAIN KEYWORD OWNERSHIP TABLE]

(Click row → opens RIGHT “Ownership Drawer”)

Global SEO Health Bar (sticky, always visible):

- shows counts of **CRITICAL (red)** and **WARNING (yellow)** issues
 - clicking opens an **Issues Drawer** listing unresolved items
 - this bar exists here **and** appears in Pipeline + Content & Structure (read-only there)
-

5. Left Sidebar Filters (Always Visible)

Filters apply before selection.

Filter Groups

- **Page Type:** Country / Region / City / Attraction / Things-to-do / Near-page / Blog
- **Hierarchy:** Country → Region → City
- **Batch / Cycle Context**
 - batch name (e.g., “Vienna City Batch”)
 - near-page cycle id
 - migration batch id
- **Ownership State:** Owned / Reserved / Candidate / Legacy-locked / Rejected
- **SEO State:** OK / OUTDATED / CONFLICT / MISSING / BLOCKED
- **AI Sync:** In sync / Regen required / Regen running / Regen failed
- **Index Impact:** None / Review required / Blocked
- **Search box**
 - keyword text
 - page URL

- entity name
- batch name

6. Main Keyword Ownership Table (Authoritative)

Table Question: “Who owns which keyword, and is it still valid + in sync?”

Column	Meaning
Keyword	Normalized keyword
Locale	Language/market (en/de/pl etc.)
Intent	navigational, informational, commercial, transactional, mix of 2
Volume	SEMrush snapshot value
KD	SEMrush snapshot value
Owner Page	Page name + type
Canonical URL	Canonical owner URL
Hierarchy	Country > Region > City context
Batch / Provenance	Which batch introduced or last changed it
Ownership State	Locked / Unlocked / Reserved / Legacy-locked / Rejected
SEO State	OK / OUTDATED / CONFLICT / MISSING / BLOCKED
AI Sync	In sync / Regen required / Running / Failed
Index Impact	None / Review / Block
Keyword-Driven Blocks	List of impacted blocks if changed (computed)
Coverage Checks	pass/warn score (computed validators)
Internal Link Target	City / Attraction / Near-page / Things-to-do (governance)
Last Change	timestamp + actor
Actions	Manage / Open owner / Resolve conflict

7. Changing Keyword Ownership (Controlled Flow) — Expanded (Key Section)

7.0 What “Ownership Change” Means

A keyword ownership change is a **manual strategic reassignment** of a keyword to a different canonical page (or to Reserved/Unassigned).

It always:

- invalidates keyword-driven AI outputs
- triggers persistent SEO issues until regenerated and approved
- may change index gating and internal linking governance

Ownership changes are **never automatic**.

7.1 Entry Points (How SEO Lead Starts)

Entry A — Page-first

- Table → filter to a page → click **Manage**

Entry B — Keyword-first

- Search bar → paste keyword → open keyword row → **Manage**

Entry C — Batch-first

- Filter by batch/cycle → view all keywords introduced/owned → use batch governance actions (Section 10)
-

7.2 Ownership Drawer (Right Side Panel)

Drawer Header

- Page name + page type
- Canonical URL
- Hierarchy context

- Provenance summary (introduced by batch, last change, who)

Section A — Current Ownership (Read-only snapshot)

- Primary keyword (required for Near-pages / Things-to-do)
- Secondary keywords (0–2 where allowed)
- Ownership state (Locked/Unlocked/Legacy-locked)
- SEO State + AI Sync
- Index Impact
- “Affected blocks if ownership changes” (computed list)

Section B — Edit Ownership (Manual Controls)

Visible only if:

- user role = SEO Lead/Admin
- ownership state is **Unlocked** (or explicitly unlocked with reason)

Fields:

- **Primary keyword** (required where policy requires)
- **Secondary keywords** (0–2; only for certain page types)
- **Reserve old keyword after change** (toggle, default ON)
- **Internal link target** (governance field; blog guidance)
- **Reason** (required on save)

Section C — Impact Preview (Read-only, computed)

Before save, show:

1. **Keyword-driven blocks that will become OUTDATED** (by page type)
2. **Pages potentially affected** (conflict candidates)

3. **Indexing impact** (Review / Block rules)
4. **Pipeline tasks that will be created** (regeneration queue markers)
5. **Coverage validators that will flip to FAIL** until fixed

This preview is computed using the **Keyword** → **Blocks Impact Map** (Section 8), not ad-hoc logic.

Section D — Conflict & Policy Validator (Mandatory)

Live validation runs while editing.

Hard blockers (cannot save):

- keyword already owned by a **Locked or Legacy-locked** owner of higher priority
- intent rule violated (e.g., transactional keyword assigned to Country page)
- keyword is in Rejected list
- Near-page missing primary keyword after edit

Soft warnings (save allowed but requires extra confirmation):

- similarity overlap above threshold with another owned keyword
- near-page keyword too close to things-to-do keyword
- blog keyword conflict (blog locked keywords are protected)

7.3 Save Action (Mandatory Confirmation Modal)

On Save, always show:

Title: Confirm Keyword Ownership Change

This action WILL:

- update keyword ownership immediately
- mark impacted blocks as **OUTDATED**
- create persistent issues + pipeline regeneration markers

- potentially block sitemap/index promotion until resolved

This action WILL NOT:

- regenerate content automatically
- change product ownership (products do not own keywords)
- override manual locks in content blocks
- publish anything automatically

Required fields

- Reason (text)
- Optional incident/ticket link

Buttons:

- Cancel
- Confirm change

7.4 After-Save System Behavior (Exact Rules)

Immediately after confirmation:

1. Ownership updates

- keyword_map updated
- old keyword becomes Reserved if toggle ON

2. SEO state transitions

- page SEO State → **OUTDATED**
- AI Sync → **Regeneration required**
- Index Impact → **Review required** (or Block for hard rules)

3. Persistent issue creation

- Issue type: KEYWORD_CHANGE_OUTDATED
- Issue is visible in:
 - Keyword ownership (this UI)
 - Pipeline View (marker)
 - Content & Structure (read-only banner)

4. Pipeline handoff

- Creates regeneration markers (not actions) for:
 - affected page(s)
 - affected keyword-driven blocks only (Section 8)

5. Index/Sitemap safeguards

- Page remains live with old content
- but **promotion/sitemap rules** may be blocked until AI Sync returns to OK

Issue persistence rule:

The OUTDATED issue remains until:

- regeneration completed AND
- an approval step marks it “SEO Sync OK”

7.5 “Where keywords matter” — Full Impact Across Elements (No Mixing)

7.5.1 Headings (labels only)

- **H2/H3 section labels** (e.g., “Highlights”, “Included / Not included”)
 - Source: Static i18n dictionary
 - Keyword dependency: **None**
 - AI regeneration: **Never**

- **H1**

- Keyword dependency:

- **Near-pages / Things-to-do:** High (must align with keyword ownership)
- **City/Attraction/Country/Region:** Medium (may shape phrasing)
- **Product:** Low–Medium (offer clarity first; keyword is a constraint, not the output)

7.5.2 Content blocks (what actually regenerates)

These are handled by **Section 8** (Keyword → Blocks Impact Map). The ownership UI must surface them in Impact Preview and issue tracking.

8. Keyword → Blocks Impact Map (Built Into Ownership UI)

This is the “Content Blocks Map” embedded inside the ownership system.

It controls: **what becomes OUTDATED** when ownership changes.

8.1 Global Rule

Only **keyword-driven blocks** are invalidated on keyword ownership change.

Keyword-driven blocks are limited to:

- Meta title
- Meta description
- H1 (page-dependent guardrails)
- Keyword-driven FAQ (where applicable)
- Internal link mapping targets + anchor rules
- Near-page overview / comparison intros (because these pages exist for that keyword)

Everything else must remain **offer/UX driven**, not keyword driven:

- highlights, options, inclusions, itinerary details, guide descriptions, review summaries, Trip DNA, etc.

8.2 Impact Table (Authoritative)

Page Type	Keyword Ownership Importance	Keyword-Driven Blocks (invalidate on change)	Blocks NOT driven by keywords
Near-page	High	H1, Meta title, Meta description, Overview, Comparison intro, FAQ, Internal link rules	Product cards, prices, inclusions, map pins
Things-to-do	High	H1, Meta title, Meta description, FAQ, Internal link rules	Listings, bundles ordering, reviews module
Attraction	High	H1 (tokened), Meta title, Meta description, FAQ, Internal link rules	Visitor information, history facts, nearby list
City	Medium	H1 tagline (guarded), Meta title, Meta description, FAQ, Internal link rules	neighborhoods labels, tips labels, listing UI
Region	Medium	H1 tagline (guarded), Meta title, Meta description, FAQ, Internal link rules	sights grid UI labels
Country	Medium	H1 tagline (guarded), Meta title, Meta description, FAQ, Internal link rules	top cities UI labels
Blog (Article)	High but editorial	No editing here (display only). Blog editor sets: title/meta/targets. Ownership UI only enforces “locked/reserved.”	Article body headings/content
Product (non-owning)	Low–Medium	Meta title, Meta description, optional “SEO H1 modifier” (guarded), Internal links, optional keyword-driven FAQ (if enabled)	Snippet (DEE short description), Long description, Highlights, Options text, Inclusions, Itinerary, Guides, Reviews

Guardrail for Product H1:

If product H1 is regenerated due to upstream keyword strategy, it must remain **offer-clarity token based** (Tour/Attraction/City), not keyword-stuffed.

9. Notifications, Visibility & Persistent Errors (Critical)

9.1 Two separate systems (must not be confused)

A) Notifications (awareness)

- can be acknowledged
- can be marked seen/unseen
- do not represent “fixed”

B) SEO Issues (state)

- OUTDATED / CONFLICT / MISSING / BLOCKED / REGEN_FAILED
- cannot be dismissed
- persist until resolved

9.2 Global SEO Health Bar (Sticky, Red Until Fixed)

Every strategic keyword problem generates a persistent issue displayed as:

- a **red banner** at top: “Critical SEO issues: X”
- clickable → opens **Issues Drawer**
- the Issues Drawer lists:
 - issue type
 - page(s) affected
 - keyword(s) involved
 - required fix action
 - status (pending / regen running / awaiting approval)
 - CTA links

This bar appears in:

- Keywords Ownership UI (editable)
- Pipeline View (actionable)
- Content & Structure View (read-only warning + deep link)

9.3 Visibility rules (where issues show)

When a keyword changes:

- Keywords Ownership UI: page row becomes **OUTDATED (red/yellow)** + “Regen required”
- Pipeline View: page(s) appear with “Regeneration Required (Keyword change)”
- Content & Structure: record shows a persistent warning banner:
 - “SEO OUTDATED — keyword changed — regeneration required”
 - no keyword editing possible there

9.4 Clearing rules (what closes the red banner)

An OUTDATED issue clears only when:

1. Pipeline regeneration finished for required blocks AND
2. an approval step marks content as SEO-safe AND
3. AI Sync returns to **OK**

Acknowledging notifications does nothing to issue state.

9.5 Escalation & “cannot forget” behavior

For critical issues (BLOCKED / CONFLICT / MISSING):

- Issue stays red
- Index Impact stays Blocked
- The UI must prevent “promotion” actions until fixed
- Optional: assign issue owner (SEO Lead) with SLA

10. Batch-Level Fixes (When “the whole batch went wrong”)

This belongs here (strategic layer), not in Pipeline.

10.1 Batch Keyword Audit Table (Batch-first Mode)

When filtered by batch, show a batch table:

Keyword	Owner	Page	Ownership	SEO	Provenance	Last	Actions
d	Page	Type	State	State		Change	
...	...	...	...	...	Batch X	...	Manage / Resolve

10.2 Batch Strategic Actions (SEO Lead/Admin only)

A) Clear batch assignments

- removes ownership assignments introduced by that batch
- converts them to Candidate/Unassigned
- does not touch Legacy-locked keywords
- You can select up to 50 batches at once or choose 1 or more

B) Re-run keyword assignment (batch-level)

This is not content regeneration. It is an ownership recompute that must be justified by strategy changes.

Outputs:

- **Keyword Assignment Draft** (preview-only)
- unit estimate + conflicts preview
- requires “Apply” confirmation

C) Apply draft

- applies ownership changes in bulk
- marks impacted pages OUTDATED
- creates pipeline regeneration markers

12. Blog Governance (Correct Model)

12.1 Where blog keywords are entered (NOT here)

Blog keywords are declared in the **Blog editor** at creation time:

- target_keyword (required)
- optional secondary targets
- meta title + meta description (manual by author)

12.2 What ownership UI does for blog

This UI:

- displays blog keyword targets as **Locked (Blog-locked)**
- prevents those keywords from being assigned to Near-pages / Things-to-do / other SEO pages
- allows SEO Lead/Admin to resolve conflicts only via governance actions (not by editing blog keywords here)

Blog keyword rows appear as:

- Ownership State: **Legacy-locked / Locked**
- Owner Page: blog URL
- Eligible for reassignment: **No**

13. Migration Support (Legacy + Blog)

For migrated content:

- import blog posts with declared target keyword(s)
- lock ownership to the blog URL
- generate conflicts report (must be resolved before Near-page expansion)

14. Permissions (Role Gating)

- Viewer: read-only
- SEO Specialist: can propose changes (draft), cannot apply
- SEO Lead/Admin: can edit, unlock, apply drafts, clear batch, approve strategy re-runs

Unlocking a Locked/Legacy-locked keyword:

- requires reason + optional ticket
- creates an audit event

15. Audit & Traceability (Immutable)

Every change logs:

- who / when
- old vs new ownership
- affected pages
- strategy_version_id (if batch re-run)
- regeneration status (pending/running/done/failed)
- approval state

Audit log is immutable.

16. What This UI Explicitly Does NOT Do

- regenerate text
- edit copy blocks

- unlock content blocks
 - publish pages
 - run operational content batches
 - bypass Pipeline View
-

Strategy Settings tab: UI Sub option of Keyword ownership

Layout:

1. SEMrush Governance (units, caps, endpoints summary)
2. Intent & Ownership Routing (page type → allowed intents, Mixed rules)
3. Scoring & Similarity (weights, thresholds, dedupe rules)
4. Scale Controls (caps per city/cycle, hub thresholds, demand allowlists)
5. Versioning & Audit (strategy_version_id history, diffs, approvals)

a. SEMrush Strategy Settings (Required for meaningful re-run in keyword ownership ui)

Keyword ownership re-runs are **not** an operational action. They change the “who owns what” map and can create:

- OUTDATED pages (regen required)
- conflicts (cannibalization risk)
- new near-page candidates (scale + index impact)
- increased SEMrush cost

Therefore the system must provide a **Strategy Settings** layer that:

- defines the inputs that drive keyword discovery + assignment
- is versioned and auditable
- is role-gated
- is the only legitimate reason to run keyword re-assignment at scale

Hard rule: **No strategy change** → **no re-run**, unless an explicit override reason is provided and logged.

a. Roles and edit permissions

01. Marketing Editor (NO strategy access)

Marketing can:

- edit content blocks (within content UIs)
- lock/unlock content blocks (where allowed)
- review notifications

Marketing cannot:

- change intent rules
- change SEMrush usage rules
- run ownership re-assignment drafts

02. SEO Admin (bounded strategy access in UI)

SEO Admin can edit **bounded, safe strategy parameters** in a dedicated UI tab:

- intent routing matrix (page type eligibility)
- Mixed-intent handling rules (allowed pairs)
- scoring weights (volume vs KD vs relevance etc.)
- similarity thresholds (duplicate/variant/new topic)
- caps and throttles (near-page caps, topic caps)
- unit caps and abort rules (within bounds)

03. Developer / Platform Admin (integration-only controls)

Developer controls anything that affects:

- SEMrush endpoint allowlist
- forbidden endpoints
- per-endpoint display hard caps enforcement logic

- any connector behavior (API keys, retry logic, rate limits)

Rule: If a setting changes which SEMrush endpoints are called, it is **Dev-only**.

b. Strategy Inputs Table

Full table in excel file - [IT: AI & programmatic SEO - schemes](#)

Additional clarification:

SEMrush budgeting (units, not keyword count)

The system **MUST** budget SEMrush usage in units, not “number of keywords”.

Known unit costs (allowed endpoints):

- Batch Keyword Overview / Keyword Overview: 10 units / keyword line
- Broad Match (phrase_fullsearch): 20 units / keyword line
- Related Keywords: 40 units / keyword line
- Phrase Questions: 40 units / keyword line

Hard rule: Any historical mode or historical endpoint is treated as ~5× cost and is forbidden.

Where changed:

- Unit costs: **Dev-only** (read-only in UI)
- Run/month caps: **SEO Admin** (bounded) + **Dev enforcement**
- Endpoint allowlist + forbidden list: **Dev-only**

Endpoint allowlist + display limits enforcement (Dev-only)

Developer **MUST** enforce:

- allowed endpoints (only these 5)
- forbidden endpoints (all historical endpoints)
- per-call display limits (hard caps)
- truncation + warning logs

UI behavior:

- Strategy Settings shows these as **read-only** with a “Managed by Dev” label.

Trend handling without history (12-month trend only)

System **MUST NOT** fetch historical datasets.

Trend is computed using the 12-month trend values returned in Keyword Overview responses and stored as snapshots.

Where changed:

- Enable/disable snapshot storage: **SEO Admin**
- Any “history endpoint” usage: **not allowed** (Dev-enforced ban)

Unit caps

CONFIG must be editable (SEO Admin within bounds):

- `max_units_per_run` = 200,000
- `max_units_per_month` = 4,000,000
- `max_monthly_keyword_budget` = 50,000
- `max_new_topics_per_entity` = 20
- `top_demand_cities_only` = true
- `top_demand_cities_allowlist_max_size` = 300

Abort behavior:

If `run_units_estimate` > `max_units_per_run` → abort before calling SEMrush.

Where changed:

- SEO Admin edits caps in Strategy Settings (bounded)
- Dev enforces monthly/global caps regardless of UI

Intent & assignment strategy — where intent parameters are changed

Where this is configured

Inside Strategy Settings → “Intent & Ownership Routing”

SEO Admin can edit:

- Page type → allowed intents matrix
- Mixed intent allowed pairs
- Product 1:1 fit gating (and threshold)
- Bundle eligibility patterns (bounded patterns)

Developer must enforce:

- the matrix constraints are not bypassed by any pipeline
- “blocked assignment” is deterministic and logged

Mixed intent (SEMrush “mix of 2”) handling (hard UI rule)

The system must treat “mix of 2” as **two intents**, not a new intent type.

The UI must force SEO Admin to declare:

- which mixed pairs are eligible per page type (allowlist)
- which mixed pairs are always blocked (denylist)

c. Strategy versioning

Every re-run stores:

- strategy_version_id
- changed parameters
- author + timestamp
- estimated units and actual units
- resulting ownership draft hash

Hard rule:

If no strategy changes since last run, re-run requires explicit override reason.

-

SEO Governance Principle - CMS UX

KEYWORD GOVERNANCE PANEL (CORE PANEL)

UX Layout

Split view: LEFT = keywords, RIGHT = owning page

LEFT – Keyword Table

Columns:

- Keyword
- Intent (SEMrush)
- Volume
- Similarity band
- Entity match (icon)
- Seasonal (icon)
- Status: free / locked / conflict

Filtry:

- intent
- volume range
- entity_flag
- seasonal_flag
- locked / unlocked
- conflict only

RIGHT – Ownership Detail

- Owning URL

- Page type
- Ownership timestamp
- Lock status
- History (timeline)

Actions (buttons)

- Lock ownership
- Request reassignment
- Escalate conflict
- Mark seasonal
- Deprecate keyword

UX rule:

Not available for editing – only user actions + reason field.

3.3 UI main panel: Product Ops (not part of marketing Panel)

Purpose

Product Operations Panel is a dedicated backend UI area for inventory governance and curated selection under Model C. It is separate from Marketing Panel.

Product Operations Panel is used to:

- Manage product inventory membership (All Products / Pool / Set)
- Apply manual overrides safely (logged)
- Monitor eligibility and Set stability flags
- Configure safe algorithm parameters (numeric tuning only)
- Resolve inventory alerts (rating drops, band changes, INSTANT violations)

Product Operations Panel is NOT used to:

- Edit SEO keywords ownership
- Approve near-pages
- Run global AI batch pipelines (Pipeline View owns this)

- Publish content

Access

Only Admin and Product Operations Lead have write access. SEO Lead has read access to Set membership and flags.

Product Ops panel MUST provide the following sub-modules:

1) Inventory Explorer

- All Products (RAW)
- Pool (Eligible)
- Set (Curated)

2) Product Record Detail (Product Ops view)

- Includes Inventory Tab (actions + membership + flags)
- Includes Version History (RAW vs AI)
- Includes Restore RAW (role-gated)

3) Algorithm Configuration (Model C Tuning)

- Numeric parameter tables (bounded)
- Config versioning
- Impact preview
- Approval workflow

4) Inventory Alerts

- Rating/reviews threshold drops
- INSTANT violations (non-instant)
- Structural change flags (classification risk)
- Band changes and coverage risk

- Manual override audit alerts

Product Governance & Selection Panel

This panel is responsible for inventory control and selection logic across three operational layers:

1. All Products (RAW view)
2. Pool (eligible candidates)
3. Set (curated collections)

This module is strictly operational.

Marketing and SEO teams do not control inventory eligibility or folder transitions.

3.X.1 Product Identity Model (Single Source of Truth)

A product exists only once in the database.

There are no duplicated copies across RAW, Pool, and Set.

Folder states represent membership flags, not duplicated records.

Structure:

- RAW layer = provider payload (always refreshed from API)
- AI layer = transformed version used for publishing
- Membership flags = pool_member / set_member / blocked / pinned

Provider refresh updates RAW payload only.

AI layer is versioned separately.

3.X.2 All Products (RAW View)

Purpose:

Operational mirror of provider inventory.

Capabilities:

- View provider payload
- View current AI-transformed version
- Compare RAW vs AI
- View version history
- View folder membership status

Actions allowed:

- Manually add to Pool (override eligibility flag)
- Flag for review
- Restore RAW version (Admin / CMO only)
- Trigger single-product AI run (if already in Set and batch completed)

Restrictions:

- No direct editing of provider fields
- No silent threshold bypass without logging

Inventory Explorer (All Products / Pool / Set)

Core rule:

Products exist as a single entity. Inventory “layers” are memberships/flags (no duplication).

Modes:

A) All Products (RAW)

B) Pool (Eligible Candidates)

C) Set (Curated Selections)

Common table columns (all modes)

- productCode (provider stable id)
- provider (viator/tickets/rosotravel/other)
- status (active/inactive)
- bookingConfirmationSettings.confirmationType (INSTANT/MANUAL)
- rating (combinedAverageRating)
- reviews.totalReviews
- group (01 / 02 / 03)
- classification (TICKET / PASS / TOUR) [system, locked logic]
- pool_member (Yes/No)
- set_member (Yes/No)
- pinned (Yes/No)
- blocked (Yes/No)
- flags (threshold_breach, non_instant, structural_change, dedup_note)

Mode-specific notes:

- All Products: shows provider payload availability (read-only)
- Pool: shows eligibility reason (which thresholds passed/failed)
- Set: shows Set stability state (OK / FLAGGED) and required review reason

Row actions (role-gated; Product Ops)

- Add to Pool (manual override; logged)
- Remove from Pool (logged)
- Promote to Set (auto Pool membership; logged)
- Remove from Set (requires reason; logged)
- Pin / Unpin (Set only; logged)

- Block / Unblock (Pool/Set; logged)
- Open Product Record Detail (Product Ops view)

Stability rule

- Pool membership is recalculated automatically (auto add/remove).
 - Set membership is stable: threshold drops only FLAG items; no auto-removal.
-

3.X.3 Pool (Eligible Candidates)

Pool is a dynamic eligibility layer.

Automatic behavior:

- Products meeting eligibility rules are added automatically
- Products falling below threshold are automatically removed
- Removal does not delete the product from database

Manual behavior:

- Product Ops may manually:
 - Add to Pool (override)
 - Block from Pool
 - Flag as Review Required

If a product is manually added to Set:

→ It is automatically marked as Pool member.

Pool products:

- Do not trigger batch runs
- Do not trigger keyword retrieval
- May be used in fallback contexts

Pool Membership Recalculation

Pool eligibility is recalculated automatically based on threshold rules.

If eligibility drops:

- Product is automatically removed from Pool
- If product belongs to Set → it is flagged, not removed

Set membership always requires manual confirmation.

3.X.4 Set (Curated Collections)

Set contains products selected for:

- City page
- Things to do page
- Attraction page
- Curated Portfolio (B2C)

Rules:

- Only products from Pool may enter Set.
- Manual addition to Set automatically assigns Pool membership.
- Set membership triggers publishing logic.
- If city batch already ran:
→ Single-product AI run must be executed (see Section 3.3.01.16).

Actions available:

- Pin (priority)
- Unpin
- Remove from Set
- Block

- Trigger AI refresh
- Restore RAW (Admin / CMO only)

Automatic removal from Set due to threshold drop is not allowed.

Instead:

- Product is flagged
- Alert generated
- Manual review required

3.X Inventory Change Logging (Extended)

All inventory-related operations must be logged:

- Pool membership change
- Set membership change
- Pin / Unpin
- Block / Unblock
- RAW restore (single)
- RAW restore (mass)
- Single-product AI run
- Threshold override
- Band change impact

Each log entry must include:

- User
- Role
- Timestamp
- ProductCode
- Previous state

- New state
- Trigger type (manual / automatic)

Logs are immutable.

3.3X Product Record Detail View (Product Ops)

This model of UI, bulk actions is based on the detailed model described in [“Product detail view for SEO”](#)

Purpose

This is the operational product view for inventory governance. It is separate from Marketing’s editorial view.

Tabs (Product Ops)

- 1) Overview (read-only summary)
- 2) Inventory (Product Ops actions)
- 3) Versions & Diff (RAW vs AI)
- 4) History / Audit (immutable timeline)

Inventory Tab (Product Ops; Products only)

Shows (read-only):

- Provider identity: productCode, provider, provider_product_id, lastUpdatedAt
- Membership status: RAW / Pool / Set (with timestamps)
- confirmationType: bookingConfirmationSettings.confirmationType
- classification: TICKET / PASS / TOUR (locked)
- group: 01 / 02 / 03
- Archetype Core ID / Variant ID (Phase C output; read-only)
- Flags: threshold_breach, non_instant, structural_change, dedup_note

Actions (role-gated; all logged):

- Add to Pool (manual override)
- Remove from Pool
- Promote to Set (auto Pool membership)
- Remove from Set (requires reason)
- Pin / Unpin (Set only)
- Block / Unblock (Pool/Set)
- Restore RAW (single product only; Admin/Product Ops Lead; confirmation required)

Versions & Diff Tab

Must show:

- RAW payload snapshots (by lastUpdatedAt/hash)
- AI layer snapshots (by ai_version_id)
- Diff viewer:
- RAW → current AI
- AI → Published (if exists)

Restore RAW behavior

- Restore RAW clears the AI layer for this product and creates a required reprocessing marker (execution details defined in Chapter 4 workflow).
- Mass restore across multiple products is Admin-only with double confirmation.

3.3 X Debugging & Explainability

Hard rules

- This tab is **READ-ONLY** for most roles.
- Write actions are limited to **Product Ops Lead + Admin** (only for explicit override fields described below).

- No “silent” overrides: every override requires a **reason** + is written into **Audit Log**.

Purpose

Provide full transparency of **WHY** the system classified, gated, deduplicated, and selected (or rejected) a product.

This is mandatory for premium trust, operational safety, and fast debugging.

This tab **MUST** answer (always, deterministically)

1. Why is this product in **RAW / Pool / Set** (or not)?
 2. Which gates did it pass/fail and with what thresholds?
 3. Which archetype was assigned and why (signals + confidence)?
 4. Was this product deduped? If yes, what cluster and why kept/dropped?
 5. Where is it used on the site (slots/rails/bundles) and why?
 6. What changed since last evaluation (input diffs + decision diffs)?
-

Data shown in the tab (READ-ONLY)

7.1 Identity & Source (READ-ONLY)

Field	Description	Example / Notes
source_provider	Provider of source payload	viator / tiqets / rosotravel / other
source_product_id	Provider primary ID	e.g., Viator productCode
entity_id	Internal stable ID	stable machine key
canonical_url	Public canonical URL	full public URL

primary_city_assignment	Validated primary destination	e.g., “Vienna”
last_evaluated_at	Last evaluation timestamp	ISO datetime
selection_engine_version	Semantic version of logic	e.g., 1.3.2
ruleset_version_id	Active numeric tuning profile	e.g., CFG_2026_03_01

7.2 Destination assignment (critical) (READ-ONLY)

Destination source + validation summary

Field	Value	Source / Computation	Status / Meaning
destinations[]	list	provider payload	show raw list
destinations[].ref	destination reference	provider payload	show per destination
destinations[].primary	primary flag	provider payload	true/false
primary_city_assignment_validated	true/false	computed	must be true for Set

validation_method.meeting_point_city_match	true/false/unknown	meeting point logic	unknown if missing fields
validation_method.pickup_city_match	true/false/unknown	pickup primary logic	unknown if missing fields
decision_state_if_mismatch	label	computed	“REVIEW REQUIRED (PRIMARY CITY)”

Hard rule (shown inline)

- primary_city_assignment must be validated by **meeting point OR primary pickup point** logic.
- If validation fails, product can remain in **RAW** but is **blocked from Set** until resolved.

7.3 Gating & eligibility evaluation (READ-ONLY)

Policy banner (always shown above the table)

- Missing provider field ⇒ value becomes **UNKNOWN** ⇒ **must not influence ranking/selection**.
- Fail-closed applies **only** to hard safety gates (example: **INSTANT required** for Set main stake).

Gate evaluation table (deterministic, complete)

Gate	Gate Block	Gate Name	Scope	Input Fields (normalized)	Threshold / Rule	Severity	Result	Reason Code	Notes
------	------------	-----------	-------	---------------------------	------------------	----------	--------	-------------	-------

ID

(resolved)

G01	City Band	RAW_count(city) band	Pol/Set	raw_count_city	0–299 / 300–599 / 600+	IN FOL	PAS S/FAIL	BAND_ASSIGNED	always shows selected band
-----	-----------	----------------------	---------	----------------	------------------------	--------	------------	---------------	----------------------------

G02	Category Group	Group 01/02/03	Pol/Set	mapped_group_ids, tags/parentTagIds	strictest-category rule	BL OCK	PAS S/FAIL	GROUP_STRICTEST	shows chosen strictest group
-----	----------------	----------------	---------	-------------------------------------	-------------------------	--------	------------	-----------------	------------------------------

G03	Quality	rating / reviews	Pol/Set	rating, reviews_count	thresholds by (band+group)	BL OCK	PAS S/FAIL	QUALITY_FAIL_*	show actual numeric comparison
-----	---------	------------------	---------	-----------------------	----------------------------	--------	------------	----------------	--------------------------------

G04	Confirmation	INSTANT requirement	Set	confirmationType	confirmationType == INSTANT	BL OCK	PAS S/FAIL	NOT_INSTANT_BLOCK	fail-closed for Set
-----	--------------	---------------------	-----	------------------	-----------------------------	--------	------------	-------------------	---------------------

G05	Private	private Tour policy	Pol/Set	privateTour flag	private ruleset (profile)	BL OCK	PAS S/FAIL	PRIVATE_RULE_*	include profile ref
-----	---------	---------------------	---------	------------------	---------------------------	--------	------------	----------------	---------------------

G06	Transport exclusions	Group03 exclusions	Pool/Set	title, tags, mapped_group_ids	exclude patterns for Group03	BL OC K	PAS S/FA IL	TRANSPORT_EXCLUDE D	deterministic pattern ID
G07	Destination validation	primary city validated	Set	primary_city_assignment_validated	must be true	BL OC K	PAS S/FA IL	PRIMARY_CITY_REVIEW	blocks Set when false
G08	Data completeness	required fields present	Pool/Set	required_fields_status	missing ⇒ UNKN OWN (no effect)	W AR N	OK/ WA RN	MISSING_FIELD_*	never blocks unless safety gate

Standardized Result semantics

Result	Meaning
PASS	Gate satisfied
FAIL	Gate not satisfied (may block Pool/Set depending on scope/severity)
WARN	Gate not satisfied but must not block selection (missing field / soft requirement)
OK	Completeness check passed

Reason Code pattern (examples)

Pattern	Example	When used
QUALITY_FAIL_RATING	QUALITY_FAIL_RATING	rating below threshold
QUALITY_FAIL_REVIEWS	QUALITY_FAIL_REVIEWS	reviews below threshold
PRIVATE_RULE_MIN_RATING	PRIVATE_RULE_MIN_RATING	private threshold not met
TRANSPORT_EXCLUDED_PATTERN_01	TRANSPORT_EXCLUDED_PATTERN_01	matched exclusion pattern
MISSING_FIELD_LOGISTICS	MISSING_FIELD_LOGISTICS	missing logistics path

7.4 Membership states (no duplicated records) (READ-ONLY)

Membership flags (shown as a single compact table)

Flag	Value	Meaning	Rule
------	-------	---------	------

is_in_raw	true	record exists in RAW	always true
pool_member	true/false	included in Pool	auto-recalculated
set_member_city	true/false	selected into City Set	Set stability applies
set_member_things_to_do	true/false	selected into Things-to-do Set	Set stability applies
set_member_attraction	true/false	selected into Attraction Set	if applicable
blocked	true/false	blocked from Set/visibility	requires block reason
pinned	true/false	pinned membership	requires pin scope

Inline rules (shown under the table)

- Pool membership is auto-recalculated. Products can enter/leave Pool automatically.
- Set membership never auto-removes; it only flags “**Set stability warning**” when thresholds drop.

7.5 Archetype assignment (core + variant) (READ-ONLY)

Archetype outcome

Field	Value	Notes
-------	-------	-------

archetype_core_id	...	core bucket
archetype_variant_id	...	variant inside core
confidence_score	0–1	numeric
fallback_notes	text	e.g., “UNKNOWN_ROUTE used”

Signals used (explicit list, deterministic)

Signal Group	Signal	Source Field / Logic	Value	Used (Y/N)	Impact Note
Core	linked_attraction_id	linked attraction mapping	...	Y/N	strongest when present
Core	itinerary_points_count	itinerary.itineraryItems[*].*	...	Y/N	supports core separation
Core	vehicle_or_style_keywords	keyword dictionary	matched tokens	Y/N	SOFT if text-based
Variant	pickup_vs_meeting_point	logistics.travelerPickup.*	pickup/meeting	Y/N	variant driver

Variant	skipTheLine	itinerary.skipTheLine	true/false/unknown	Y/N	only if known
Variant	guide_mode	languageGuides[].type	guide/audio/written/unknown	Y/N	normalized

Hard rule (merge transparency)

- If archetype becomes too narrow (**competitive_count < 4 eligible products globally**), system must auto-merge to broader core.
- UI must show event: `merge_event: from X → Y` in History/Audit.

7.6 Dedup & near-duplicate handling (READ-ONLY)

Dedup decision summary

Field	Value	Notes
dedup_cluster_id	...	stable cluster id
decision	KEPT / DROPPED	final
winner_entity_id	...	if DROPPED
winner_reason_code	...	explicit reason

Cluster key factors (explainable)

Factor	Source / Computation	Value	Used (Y/N)
--------	----------------------	-------	------------

attraction_id	linked attraction	...	Y/N
core_id	archetype_core_id	...	Y
pickup_flag	pickup vs meeting	...	Y
format	ticket/tour/workshop etc.	...	Y

Allowed coexistence exceptions (shown as rules list)

- pickup vs meeting point (allowed)
- materially different format (allowed)
Everything is explainable via explicit reason codes.

7.7 Winner/slot exposure (where used) (READ-ONLY)

Exposure matrix (single unified list)

Surface	Container	Slot / Index	Used (Y/N)	Eligibility Reason	Winner Reason	Notes
City Page	Top Picks	1..N	yes/no	ELIGIBLE_* / BLOCKED_* —	WINNER_* / TIEBREAK_* —	include score snapshot
City Page	Rail (name)	index	yes/no	...	...	BV / Premium marker

City Page	Bundle (day role)	day 1/2/ 3	yes/ no	...	...	Anchor/Complement/Flex
Things-to-do	Top Picks	1..N	yes/ no	...	...	
Things-to-do	Rail (name)	index	yes/ no	...	...	
Things-to-do	Bundle (day role)	day 1/2/ 3	yes/ no	...	...	
Attraction on Page	Our Picks	1..2	yes/ no	...	...	best_value / premium
Attraction on Page	Tickets-only	index	yes/ no	...	...	only if ticket logic
Attraction on Page	Private CTA	1	yes/ no	...	...	CTA selection logic
Phase B	curated_catalog_visible	—	yes/ no	...	...	
Phase B	magic_finder_candidate	—	yes/ no	...	...	

Phase B	on_your_request_ only	—	yes/ no	...	...
---------	-----------------------	---	---------	-----	-----

7.8 Change explanation (diffs that matter) (READ-ONLY)

Input diffs (what changed in source/RAW that mattered)

Changed Field / Area	Previous	Current	Detected At	Impact	Event ID
rating	...	...	timestamp	may change quality gate	EVT_*
reviews_count	...	...	timestamp	may change quality gate	EVT_*
confirmationType	...	...	timestamp	Set hard gate	EVT_*
destinations.primary	...	...	timestamp	destination validation	EVT_*
pickupOptionType	...	...	timestamp	variant/dedup	EVT_*

Decision diffs (what changed in outputs since previous evaluation)

Decision Area	Previous	Current	Detected At	Why	Event ID
pool_member	true/false	true/false	timestamp	gate result change	EVT_*

set_member_cit y	true/false	true/false	timestamp	Set rule/stability	EVT_*
archetype_core _id	X	Y	timestamp	signals + merge	EVT_*
dedup decision	KEPT/DROPPE D	KEPT/DROPPE D	timestamp	winner comparison	EVT_*
exposure	used/not used	used/not used	timestamp	slot re-rank	EVT_*

Hard rule (no hidden states)

- Every status change must have an event record. No hidden transitions.

7.9 Allowed actions (role-gated) (WRITE-LIMITED)

Allowed actions table

Action	Allowed Roles	Requires Reason	Audit Event Type	Effects
Block	Admin, Product Ops Lead	YES	EVT_BLOCK	sets blocked=true + scope
Unblock	Admin, Product Ops Lead	YES	EVT_UNBLOCK	clears blocked

Pin	Admin, Product Ops Lead	YES	EVT_PIN	sets pinned=true + pin scope
Unpin	Admin, Product Ops Lead	YES	EVT_UNPIN	clears pin
Manual archetype override	Admin, Product Ops Lead	YES	EVT_ARCHETYPE_OVERRID E	sets override_flag=tru e
Force re-evaluate classification	Admin, Product Ops Lead	NO (optional)	EVT_REEVALUATE	recompute (no rewrite)

No one may (explicitly listed)

- change global rules from this tab (handled in Algorithm Configuration UI)
- bypass hard safety gates silently

3.3X Algorithm Configuration (Product Ops Settings)

Provide safe numeric tuning of Model C without changing core architecture. All changes are versioned, previewed, and audit-logged.

LOCKED (NOT editable)

- RAW → Pool (auto) → Set (stable)

- INSTANT requirement for Set (no exceptions)
- Anti-duplicate rules R1–R6
- 4-layer model (Audience / Intent / Constraints / Value Tier)
- Hidden_gems separation rule (never merged)
- Min candidates per Archetype Core = 4 (global)
- Archetype Core merge order (duration → pickup → POI at city-level only; POI merge forbidden on attraction pages)
- Set stability: Set items are FLAGGED only (no silent auto replacement)
- Ticket classification logic (TICKET / PASS / TOUR) is locked

Editable parameters (bounded; versioned)

A) Pool Eligibility Thresholds (by Band and Group)

Parameter	Default	Min	Max
Group01 min_rating (Band A)	4.5	4.3	4.8
Group01 min_reviews (Band A)	200	50	500
Group01 min_reviews (Band B)	500	100	800
Group01 min_reviews (Band C)	800	150	1200
Group02 min_rating (all)	4.5	4.2	4.8
Group02 min_reviews (Band A)	40	20	150

Group02 min_reviews (Band B)	200	40	300
------------------------------	-----	----	-----

Group02 min_reviews (Band C)	200	40	300
------------------------------	-----	----	-----

Partner private min_rating	4.8	4.6	5.0
----------------------------	-----	-----	-----

Partner private min_reviews	50	20	200
-----------------------------	----	----	-----

B) BQS Weights (within-archetype ordering)

Parameter	Default	Min	Max
BQS_rating_weight	0.60	0.50	0.80
BQS_review_weight	0.40	0.20	0.50

Constraint: rating weight must remain $\geq$ review weight.

C) Language Advantage (ALWAYS ON; multipliers only)

Parameter	Default	Min	Max
language_trigger_min	3	2	5
language_low_ceiling	1	1	2
language_bonus_multiplier	1.20	1.00	1.50

D) Display Limits

Parameter	Default	Min	Max
max_visible_archetypes_city	8	6	10
max_visible_archetypes_cat	6	4	8
variant_coexistence_limit	2	1	2

E) Slot Budgets (counts only; page order is template-governed)

Parameter	Default	Min	Max
city_top_picks_count	6	4	8
city_rails_count	4	3	5
city_top_tickets_count	4	4	6
things_rails_count	6	4	8
things_top_tickets_count	4	4	6
attraction_our_picks_count	2	1	2

F) Bundle Tuning (bundle-only; does not affect winner ranking)

Difficulty weights

Product class	Default	Min	Max
walking_3h_plus	2	1	3
walking_2h	1	0	2
day_trip_any_duration	3	2	4
car_tour_private	2	1	3
boat_cruise_relax	1	0	2
museum_ticket	1	0	2
scenic_relax	0	0	1
tasting_or_workshop	1	0	2
activity_atv_rafting_etc	2	1	3

Audience tuning defaults (v0)

Parameter	first_time_visitor	family_planner	comfort_easy_pace	solo_social	deep_dive	active_adventure
max_daily_difficulty	4	3	3	4	6	5

max_paid_timed_ per_day	2	1	2	2	3	2
allow_two_anchor s_per_day	FALSE	FALSE	FALSE	FALSE	TRUE	FALSE
walking_penalty_ multiplier	1.0	1.4	1.5	1.0	1.0	0.8
car_preference_ multiplier	1.1	1.3	1.4	1.0	1.0	0.8
private_preferenc e_multiplier	1.0	1.3	1.3	1.0	1.1	0.9
pickup_preferenc e_multiplier	1.1	1.3	1.4	1.0	1.0	1.0
day2_large_city_ strategy	CITY_DE EP	CITY_DE EP	CAR_CITY	CITY_ DEEP	CITY_ DEEP	DAY_TRIP
day2_small_city_ strategy	DAY_TRIP	DAY_TRI P	DAY_TRIP	DAY_T RIP	DAY_T RIP	DAY_TRIP

Config safeguards

- Every change creates a new config version (diff + rollback).
- Before saving, UI must run an impact preview for a selected city:
 - Pool in/out counts
 - Set items flagged count

- Archetype cores below 4 candidates
- simulated slot/bundle changes (simulation only)
- High-impact changes require Admin approval (lowering floors or large threshold reductions).

3.X Product Ops: Derived Tags & Coverage (Read-Only Governance Layer)

Purpose

This section defines how “Global Product Tags” are derived, displayed, and governed in Product Operations without duplicating provider data. Tags are used to support Model C classification, archetype assignment, eligibility signals, and bundle feasibility.

Important:

- Full mapping of derived tags to Viator API fields is implemented by engineering during discovery.
- This spec defines tag names, intended usage (HARD vs SOFT), and the Viator fields we have already confirmed from uploaded payload examples.
- Any missing/unknown data MUST result in UNKNOWN coverage state and must not influence selection decisions.

1) Derived Tag Classes (HARD vs SOFT)

HARD tags (structural, API-derived)

- May influence: classification (TICKET/PASS/TOUR), archetype core/variant, constraint signals, bundle feasibility (difficulty/time geometry), INSTANT gating.
- Must be derived only from explicit provider fields (no title/description heuristics).
- If source fields are missing: tag is UNKNOWN and must not affect selection.

SOFT tags (semantic/merchandising)

- Must NOT influence: Pool gates, Set eligibility, exclusion rules.
- May influence only: rail naming, “why this pick” copy, optional last tie-break within already-qualified candidates (if enabled).

2) Derived Tags Snapshot (per Product Record)

Product Record Detail (Product Ops) must display a read-only snapshot:

- derived_tags_hard[] (with KNOWN/UNKNOWN state per tag)
- derived_tags_soft[] (optional; may be disabled)
- derived_tags_coverage_summary:
 - % of HARD tags known (per product)
 - list of missing source fields (by API path)
 - classification confidence (HIGH/MED/LOW)

3) Tag Coverage Report (per City / per Provider)

Product Ops must provide a coverage dashboard:

- HARD tag coverage by city (eligible inventory)
- HARD tag coverage by provider
- Missing-field leaderboard (top missing API paths)
- Impact warnings:
 - “constraints coverage low: X% unknown”
 - “ticket classification confidence low: X products need review”

4) Tag Mapping Registry (Engineering-Owned; Discovery Output)

The system must store a mapping registry:

- tag_name
- tag_class (HARD/SOFT)
- provider (viator/tickets/...)

- source_fields[] (API paths + enum values)
- derivation_rule (deterministic)
- coverage_state rules (KNOWN/UNKNOWN)

Engineering is responsible for finalizing source_fields for all tags during discovery.

3.3X Derived Product Tags — Viator Mapping v0 (Discovery-Driven)

Notes

- Engineering must map **ALL tags** to Viator API fields during discovery.
- This table lists **tag names + Viator field paths** we have already confirmed from uploaded product JSON payloads.
- Any missing source field results in **UNKNOWN**; **UNKNOWN tags must not affect selection.**

Legend

- **HARD** = API-derived (allowed to influence selection/bundles when **KNOWN**)
- **SOFT** = semantic (must not influence gates/Set eligibility)
- **Status**
 - **CONFIRMED** = field observed in uploaded payloads - still need to be verified and mapped by devs
 - **UNKNOWN** = not observed or requires other endpoints; TBD in discovery

A) Format & commerce (HARD)

Tag name	Type	Purpose	Viator source fields (path)	Status
----------	------	---------	--------------------------------	--------

format_private_tour	HARD	private/shared split	itinerary.privateTour	CONFIRMED
format_group_tour	HARD	inverse of private	itinerary.privateTour == false	CONFIRMED
format_ticket_only	HARD	ticket-only classification	languageGuides[].type (no GUIDE) + itinerary.itineraryItems []	CONFIRMED *
format_ticket_plus_guided	HARD	ticket + guide (not Top Tickets)	languageGuides[].type includes GUIDE	CONFIRMED *
format_workshop_class	HARD	class/workshop (Group02)	tags[] (mapped to Group02)	CONFIRMED
format_transfer_transport	HARD	transport (Group03)	tags[] (mapped to Group03)	CONFIRMED
format_day_trip	HARD	day trip vs city short	itinerary.duration.* (bucket) + Group01 mapping	CONFIRMED
format_city_tour	SOFT	semantic label only (do not gate)	title/description dictionary OR provider tag mapping	UNKNOWN
format_cruise_boat_tour	SOFT	semantic label only (do not gate)	title/description dictionary OR provider tag mapping	UNKNOWN

format_activity_experience	SOFT	semantic label only (do not gate)	title/description dictionary OR provider tag mapping	UNKNOWN
----------------------------	------	-----------------------------------	--	---------

- Ticket-only vs tour-like logic is **locked in Model C** and must use provider fields only.
 CONFIRMED fields for this logic: `itinerary.itineraryItems[ ]` exists;
`languageGuides` exists either at top-level or per option.

B) Time & duration (HARD where possible)

Tag name	Type	Purpose	Viator source fields (path)	Status
duration_short_0_2h	HARD	duration bucket	itinerary.duration.fixedDurationInMinutes	CONFIRMED
duration_half_day_2_5h	HARD	duration bucket	itinerary.duration.fixedDurationInMinutes	CONFIRMED
duration_full_day_5_10h	HARD	duration bucket	itinerary.duration.fixedDurationInMinutes	CONFIRMED
duration_multi_day	HARD	multi-day / long duration	itinerary.duration.variableDuration*	UNKNOWN
time_morning	HARD	start-time bucket (needs availability)	availability/start-times endpoint (TBD)	UNKNOWN

time_evening	HARD	start-time bucket (needs availability)	availability/start-times endpoint (TBD)	UNKNOWN
--------------	------	--	---	---------

C) Pickup & logistics (HARD where possible)

Tag name	Type	Purpose	Viator source fields (path)	Status
pickup_meeting_point_only	HARD	meeting-only format	logistics.travelerPickup.pickupOptionType	CONFIRMED
pickup_hotel_availability	HARD	hotel pickup availability	logistics.travelerPickup.pickupOptionType (enum TBD)	UNKNOWN
pickup_airport	HARD	airport pickup	logistics.travelerPickup.locations[] (needs location types)	UNKNOWN
pickup_cruise_port	HARD	cruise pickup	logistics.travelerPickup.locations[] (needs location types)	UNKNOWN
pickup_train_station	HARD	station pickup	logistics.travelerPickup.locations[] (needs location types)	UNKNOWN
confirmation_instant	HARD	Set gate	bookingConfirmationSettings.confirmationType == INSTANT	CONFIRMED

cancellation_free	HARD	free cancellation signal	cancellationPolicy.type + refundEligibility[]	CONFIRMED
non_refundable	HARD	non-refundable signal	cancellationPolicy.type == ALL_SALES_FINAL	CONFIRMED

D) Group size & crowd proxies (HARD/TBD)

Tag name	Type	Purpose	Viator source fields (path)	Status
group_small_le_12	HARD	small group preference	itinerary.maxTravelersInShared Tour	CONFIRMED
group_medium_13_25	HARD	medium group	itinerary.maxTravelersInShared Tour	CONFIRMED
group_large_ge_26	HARD	large group	itinerary.maxTravelersInShared Tour	CONFIRMED
crowd_low	SOFT	crowd heuristic (do not gate)	TBD (no reliable API field confirmed)	UNKNOWN
crowd_medium	SOFT	crowd heuristic (do not gate)	TBD	UNKNOWN

crowd_high	SOFT	crowd heuristic (do not gate)	TBD	UNKNOWN
------------	------	----------------------------------	-----	---------

E) Physical effort & accessibility (HARD when explicit, otherwise UNKNOWN)

Confirmed source pattern

- `additionalInfo[]` contains objects: `{ type, description }`
- Confirmed types seen in uploaded payloads: `PHYSICAL_EASY`, `WHEELCHAIR_ACCESSIBLE`, `STROLLER_ACCESSIBLE`

Tag name	Type	Purpose	Viator source fields (path)	Status
pace_easy	HARD	low effort	<code>additionalInfo[].type == PHYSICAL_EASY</code>	CONFIRMED
wheelchair_accessible	HARD	wheelchair support	<code>additionalInfo[].type == WHEELCHAIR_ACCESSIBLE</code>	CONFIRMED
stroller_friendly	HARD	stroller support	<code>additionalInfo[].type == STROLLER_ACCESSIBLE</code>	CONFIRMED
pace_standard	HARD	effort level	<code>additionalInfo[]</code> (no confirmed type)	UNKNOWN
pace_fast	HARD	effort level	<code>additionalInfo[]</code> (no confirmed type)	UNKNOWN

walking_low	HARD	walking load	additionalInfo[] (no confirmed type)	UNKNOWN
walking_moderate	HARD	walking load	additionalInfo[] (no confirmed type)	UNKNOWN
walking_high	HARD	walking load	additionalInfo[] (no confirmed type)	UNKNOWN
stairs_low	HARD	stairs load	additionalInfo[] (no confirmed type)	UNKNOWN
stairs_moderate	HARD	stairs load	additionalInfo[] (no confirmed type)	UNKNOWN
stairs_high	HARD	stairs load	additionalInfo[] (no confirmed type)	UNKNOWN

F) Access & inclusions (HARD where explicit; others TBD)

Tag name	Type	Purpose	Viator source fields (path)	Status
skip_line_claimed	HARD	skip-the-line feature	itinerary.skipTheLine	CONFIRMED
guide_included	HARD	guided tour signal	languageGuides[].type includes GUIDE	CONFIRMED*

audio_guide_available	HARD	audio guide signal	languageGuides[].type includes AUDIO/WRITTEN	CONFIRMED*
reserved_entry	HARD	reserved entry	TBD (needs ticket/availability fields)	UNKNOWN
timed_entry	HARD	timed entry	TBD (availability or ticket fields)	UNKNOWN
lift_access	SOFT	lift availability (copy/tie-break)	additionalInfo[].description (text) TBD	UNKNOWN
stairs_only	SOFT	stairs-only (copy/tie-break)	additionalInfo[].description (text) TBD	UNKNOWN
meals_included	SOFT	meals in inclusions	inclusions[] / description text (TBD parsing)	UNKNOWN

- `languageGuides` are present either at top-level (`languageGuides[ ]`) or inside `productOptions[ ].languageGuides[ ]` depending on the product. Discovery must normalize both.

G) Value tier (product-side; deterministic, not a provider field)

Value tier is computed by Model C from within-archetype context.
No direct Viator field mapping exists. This is intentional.

Tag name	Type	Purpose	Viator source fields	Status
----------	------	---------	----------------------	--------

best_value	HARD	computed tier	computed from BQS + fit	N/A
premium	HARD	computed tier	computed from fit + signals	N/A
balanced_cor e	HARD	computed tier	computed	N/A
niche	HARD	computed tier	computed	N/A

3.X Single Product Reprocessing (Governance Layer)

This section defines when and how a single product AI reprocessing may be triggered from the Product Operations panel.

This does not redefine the end-to-end workflow described in Chapter 4.

It defines governance triggers only.

Trigger Scenarios

A. Manual Add to Set

If a product is manually added to Set and the city batch has already been executed,
→ the system must trigger a single-product AI run.

B. Manual Force Refresh

Product Ops may trigger a single-product refresh.

This regenerates only product-level AI content.

C. RAW Restore

If RAW version is restored,

→ the product must undergo single-product reprocessing.

D. Provider Structural Change (Flag Only)

If provider payload changes product structure (e.g. itinerary shape, privateTour flag, duration model),

→ system flags product for review.

Automatic reprocessing is not triggered here.

Eligibility impact is evaluated separately.

Scope of Single Product Run

Single-product run may:

- ✓ Regenerate AI-transformed description
- ✓ Regenerate highlights
- ✓ Regenerate FAQ block (if applicable)
- ✓ Refresh structured data for this product
- ✓ Update product embedding
- ✓ Recalculate archetype classification (Phase C)

Single-product run must NOT:

- ✗ Rebuild city batch
- ✗ Recalculate slot allocation

- ✗ Trigger near-page creation
- ✗ Modify city keyword ownership
- ✗ Modify other products in Set
- ✗ Change city band classification

Stability Rule

Single-product reprocessing does not automatically alter Set ranking.

Ranking and slot logic is handled by the Phase C selection engine during full city evaluation.

All single-product runs must generate:

- run_id
- triggered_by
- timestamp
- version_before
- version_after
- embedding_version (exact information are defined in chapter 4.7)

3.3X Internal Products - need to edit questionnaire of adding own products

Product Bookings Reviews General administration Admin only

Debug info:
current step: 1
product_id:
country_id:
city_id:

The form contains errors. Please correct them to continue.

What is the customer-facing title of your product?
The tour title field is required.

Please insert your text in English

80 characters left

Create a product reference code (optional)

Please insert your text in English

20 characters left

Save tour

Purpose

Internal (Rosotravel-created) products require a dedicated UI area to ensure Model C metadata parity and safe portfolio selection.

This is a UI requirement, not only a migration note.

Hard rule:

Internal products **MUST NOT** be managed using legacy “city section tags” as the primary classification mechanism.

Internal products must use the “Model C Classification (Internal Product)” questionnaire to be eligible for Set.

Navigation (Product Ops / Product module)

Add a highlighted entry:

- Internal Products

Internal Products must provide:

1) Internal Product List (filterable)

- migration_state: NOT_STARTED / IN_PROGRESS / COMPLETED / BLOCKED

- set_eligibility_status: ELIGIBLE / BLOCKED

- blocking_reason_codes (if blocked)

- quick filters: “Blocked from Set”, “Missing required fields”, “Needs review”

2) Internal Product Edit (existing product wizard) — required change

- Replace legacy “Tags – sections on city page” with:

“Model C Classification (Internal Product)”

- All new Model C decision fields must live ONLY in that section (single place).
- Legacy tag checkboxes must be removed from the UI (deprecated).

3) Read-only “Decision Snapshot” (same section)

Show read-only signals from other sections (languages, pickup, duration, availability, cancellation, guide type, etc.) to reduce navigation burden.

4) Explicit Set blocking behavior

If migration_state != COMPLETED:

- internal product is blocked from Set (main stack)
- blocking is explicit and visible in Product Ops and in Product Detail

Reference for full migration plan

All rollout steps, migration queue operations, and bulk processes are specified in:
Chapter 5 — Migration & Rollout Plan (Internal Products → Model C Parity)

Implementation note:

This Chapter 3 section defines the UI surface that MUST exist.

Chapter 5 defines HOW migration is executed operationally.

Internal Product Edit Wizard

1) MODIFY existing menu section (single place for all new fields)

TARGET SECTION (existing):

“Tags – sections on city page”

RENAME TO:

“Model C Classification (Internal Product)”

Hard requirement:

All NEW Model C decision fields must be added ONLY in this section.

No new decision fields should be added to other tabs.

2) REMOVE legacy tags UI (checkboxes) — no longer used

REMOVE from UI (deprecated):

- Bestsellers
- Top Attractions
- Walking Tours & City Tours
- One Day Tour
- Tasting Tours & Nightlife
- For Families
- World War II Tours

Migration note:

These legacy tags may remain in DB as deprecated_legacy_tags (read-only) for historical reference,

but must not be editable and must not influence any decision logic.

3) ADD: "Model C Minimal Questionnaire" (editable)

URGNET: internal and external products model C need to be aligned for details refer to point 4.8 of this scope especially (3.3-3.7 below points).

Inside "Model C Classification (Internal Product)", add ONE questionnaire panel - all of below queries should be added to 1 menu option

3.1 Declared Product Class (required)

Field: declared_product_class

Options:

- TOUR
- ACTIVITY
- TICKET
- CITY_PASS_CITY_CARD
- WORKSHOP
- TRANSFER

Auto-derived (read-only):

- category_group (computed, not a user field)
- if declared_product_class in {TOUR, ACTIVITY} → Group 01 (Tours & Activities)
- if declared_product_class in {TICKET, CITY_PASS_CITY_CARD, WORKSHOP} → Group 02
- if declared_product_class == TRANSFER → Group 03 (Transfers)

3.2 Decision Rail Role (required)

Field: decision_rail_role

Purpose:

A stable "portfolio placement role" used by Model C to place internal products into the correct rails/bundles.

Options (global allowlist):

- Top Attraction Tour (guided experience; not a ticket)
- City Highlights / Orientation
- Museum & Culture (Ticket-first)
- Food & Tastings
- Night Experience
- Day Trip / Outside City
- Activity / Adrenaline
- Transfer / Transport

Clarifications:

- "Top Attraction Tour" is for TOUR/ACTIVITY only.
- "Museum & Culture (Ticket-first)" means the portfolio prefers ticket/pass as the primary surface. Guided tours (if any) are treated as premium/upgrade options and must not contaminate Top Tickets rules.

Validation rules:

- If declared_product_class in {TICKET, CITY_PASS_CITY_CARD} and decision_rail_role == Top Attraction Tour → validation error.
- If declared_product_class == TRANSFER → decision_rail_role must be Transfer / Transport.

3.3 Value Tier (required, with auto behavior)

Field: value_tier

Options:

- best_value_for_money
- balanced_core
- premium
- niche

Auto rule (internal products):

- If product is Group Tour (existing field in Facilities: Group tour): value_tier is forced to best_value_for_money (read-only).
- If product is Private tour: value_tier is editable (default = premium).

Definition:

- balanced_core = “core curated default” (not positioned as cheap and not positioned as premium comfort).

3.4 Physical Difficulty (required, simple like Viator)

Field: physical_level

Options:

- LOW
- MODERATE
- HARD

System mapping (read-only internal):

physical_level → internal bundle feasibility weight used only for bundle pacing/feasibility.

3.5 Confirmation Type (read-only for internal)

Field: confirmation_type

Value:

- INSTANT (read-only)

3.6 Audience + Intents (minimal required)

Field: primary_audience (required; single select)

- **Refer to point 4.8 of this scope**

Field: intents (required; choose 1–2)

Options:

- **Refer to point 4.8 of this scope**

Validation:

- more than 2 intents → error.

3.7 Accessibility (optional; only keep 2)

Field: accessibility_flags (optional; multi-select)

Options:

- Refer to point 4.8 of this scope

Derived automatically (no manual fields here):

- hotel_pickup_preferred derived from pickup mode and premium/comfort profile
- free_cancellation_preferred derived from cancellation policy and value tier
- low_walking_required derived from physical_level=LOW (where applicable)

4) ADD: "Model C Decision Snapshot" (read-only, same section)

Below the questionnaire show a compact read-only snapshot of already-existing product fields:

- Start city/country (Location)
- Linked attractions / POIs (Location)
- Pickup mode + meeting point (Meeting point and pickup)
- Group/Private (Facilities)
- Guide type (Facilities)
- Tour vehicle (Facilities)
- Skip-the-line flag (Facilities)
- Languages offered (Option setup)
- Duration minutes + cut-off (Option setup)
- Availability schedule summary (Availability)
- Cancellation policy (Cancellation Policy)

5) ADD: Migration State (read-only banner + blocking)

Show:

"Model C Decision Metadata: COMPLETE / INCOMPLETE"

Rules:

- If any required questionnaire field is missing → INCOMPLETE and product is blocked from Set.
- Blocking is explicit and surfaced in Product Ops and Explainability.

6) Logging / Governance (hard)

All questionnaire edits are audit-logged:

- actor, role, timestamp
- before/after values
- reason (required)

7) OTA vs Internal Product Differences (Important) — Model C Parity Contract

Purpose

Internal (Rosotravel-created) products must participate in Model C selection under the same governance rules as feed products.

They must be classifiable, gated, deduplicated, ranked, and bundled using the same decision engine.

Hard rules

1) Internal products must use:

“Model C Classification (Internal Product)” (single section) to provide minimal decision metadata.

2) Legacy “city section tags” UI is deprecated and must not be used for decision logic.

3) Internal products are blocked from Set (main stack) unless decision metadata is COMPLETE.

4) Internal products are not special-cased winners.

Migration reference

The operational migration plan (queues, rollout, reporting, bulk ops) is specified in:
Chapter 5 — Internal Products → Model C Parity Migration

Acceptance Criteria — Internal Products → Model C Classification (Model C v1)

A. Navigation & Visibility

AC-01: Backend must expose a highlighted navigation entry “Internal Products” inside Product Ops / Product module.

AC-02: “Internal Products” list must show: migration_state, set_eligibility_status, blocking_reason_codes, and filters (Blocked from Set / Missing required fields / Needs review).

B. Single-section requirement (no scattered fields)

AC-03: All new Model C decision inputs for internal products must exist ONLY inside the renamed section:

“Model C Classification (Internal Product)”.

No new decision inputs may be added to other tabs.

C. Legacy tag removal

AC-04: The legacy checkbox block (“Bestsellers”, “Top Attractions”, “Walking Tours & City Tours”, “One Day Tour”, “Tasting Tours & Nightlife”, “For Families”, “World War II Tours”) must be removed from UI.

AC-05: If legacy tags exist in DB, they must be read-only and must not influence any decision logic.

D. Questionnaire fields and validation

AC-06: The questionnaire must include required fields:

- declared_product_class (TOUR / ACTIVITY / TICKET / CITY_PASS_CITY_CARD / WORKSHOP / TRANSFER)
- decision_rail_role (from allowlist)
- value_tier (best_value_for_money / balanced_core / premium / niche)
- physical_level (LOW / MODERATE / HARD)
- confirmation_type (INSTANT, read-only)

- primary_audience (single select)
- intents (1–2)

Optional:

- accessibility_flags (wheelchair_accessible, stroller_accessible)

AC-07: category_group must be auto-derived (read-only) from declared_product_class:

- TOUR/ACTIVITY → Group 01 (Tours & Activities)
- TICKET/CITY_PASS_CITY_CARD/WORKSHOP → Group 02
- TRANSFER → Group 03 (Transfers)

AC-08: Validation errors must block saving completion when:

- declared_product_class in {TICKET, CITY_PASS_CITY_CARD} AND decision_rail_role == “Top Attraction Tour”
- declared_product_class == TRANSFER AND decision_rail_role != “Transfer / Transport”
- intents selection count > 2
- any required field is missing

AC-09: value_tier auto rule must apply:

- If Facilities indicates “Group tour”, value_tier is forced to best_value_for_money (read-only).
- If Facilities indicates “Private tour”, value_tier remains editable (default premium).

E. Decision Snapshot (read-only)

AC-10: The same section must render a read-only “Decision Snapshot” showing fields sourced from other tabs:

- Start city/country and linked attractions (Location)
- Pickup mode + meeting point (Meeting point and pickup)
- Group/Private, guide type, tour vehicle, skip-the-line (Facilities)
- Languages, duration minutes, cut-off time, max group size (Option setup / Tour options)
- Availability schedule summary (Availability)
- Cancellation policy (Cancellation Policy)

F. Migration state + Set blocking

AC-11: Every internal product must have migration_state: NOT_STARTED / IN_PROGRESS / COMPLETED / BLOCKED.

AC-12: If any required questionnaire field is missing, migration_state must become BLOCKED and the product must be blocked from Set eligibility.

AC-13: Blocking must be explicit in Product Ops list and in Product detail (reason codes + missing fields list).

G. Logging / audit

AC-14: Every change to any questionnaire field must be audit-logged with: actor, role, timestamp, old_value, new_value, and reason (required).

AC-15: System must support “who changed what” timeline on the product (immutable log).

H. Set safety rule (instant)

AC-16: confirmation_type for internal products is INSTANT and read-only in v1.

If future MANUAL is enabled, it must block Set eligibility and allow only Phase B exposure.

I. Explainability link

AC-17: Product Ops and Product Detail must expose a link to Explainability (reason codes, eligibility status, archetype assignment) once that module is implemented.

3.3X Inserting media by guides and admin “Recent Experience & Human Face”

This subsection should include two separate blocks:

Block 1 — Guide / Host Anchor

Purpose:

manage the human anchor used on product pages.

Fields / actions:

- select guide mode:
 - exact guide
 - guide team
 - curated by Rosotravel
- assign guide_id or guide_team_id
- upload/select guide anchor image (if not from guide profile)
- short credential line
- optional quote
- preview product-page placement

Block 2 — Snapshot Media & Recent Proof

Purpose:

attach manual media to a specific product for snapshot use.

Fields / actions:

- upload photo / video
- source_type = admin_snapshot
- link to product_id (required)
- optional city_id / attraction_id
- optional snapshot_date
- optional caption / note
- moderation status
- snapshot preview

Why this is needed

Even if guide-uploaded media is not available in MVP, Rosotravel admins must be able to:

- enrich the Recent Experience layer,
- attach trusted media to specific products,
- and maintain higher-quality proof blocks for priority tours.

43) Chapter 3 — Guides UI requirement

Add or expand a dedicated Guides / Experts management area with:

- guide profile image
- quote
- specialties
- assigned cities
- assigned products (optional)
- active status
- team membership

This is the source of truth for Guide / Host Anchor when exact per-product guide is not available.

44) Chapter 3 — review media linkage requirement

For customer reviews, the admin/UI layer must support:

- media moderation
- review-to-product linkage visibility
- source labeling
- automatic eligibility for snapshot clustering

This should exist either:

- in the reviews moderation panel,
- or
- in the product detail panel under “Recent Experience & Human Face”

45) Chapter 3 insertion guidance (paste-ready note)

INSERT in Chapter 3 under the Product detail / Product Ops UI section:

A new subsection:

“Recent Experience & Human Face”

with:

- Guide / Host Anchor management
- Snapshot media upload and linkage
- Snapshot preview
- Review media moderation visibility

Also INSERT in the Guides / Experts UI section:

- guide profile fields required by the Human-Face layer

3.4 UI main panel: Static Pages (blog, expert, about us) (not part of marketing Panel)

Owner: SEO Team (single governance owner; no separate Legal Admin)

Scope: all low-volume, high-visibility pages linked from footer/top navigation that are not product pages and not programmatic near-pages.

0. Purpose (Why this UI exists)

This UI is the authoritative place to create, maintain, translate, validate, and publish static pages used for brand trust, policies, support, and key discovery hubs.

It exists to ensure:

- every footer-linked page is always complete, consistent, and published in EN
- translations are safe, traceable, and non-destructive (never overwriting manual work)
- missing/invalid content is surfaced via persistent errors (not dismissible)
- system-generated pages (Sitemap / AI Sitemap / API Feed) are configured and monitored, not edited as free text

1. Pages in scope

- Add access to all of the below landing pages, it can be created as sub-options but should be easy to navigate.

1.1 Brand / Company

- About Rosotravel
- Who We Are
- Mission & Values
- How We Work
- Sustainability
- Careers
- Press & Media

1.2 Discovery / Content hubs (static + system listing)

- Destinations

i. Experiences

j. Travel Guides

`author_id` (stable internal ID; SSOT for attribution)
`author_display_name`
`author_profile_url` (e.g., `/experts/julia-rossi/`)
`author_image_url`
`author_short_bio` (1–2 sentences)
`author_job_title`
`author_sameAs[]` (LinkedIn, X, Wikipedia, etc.)
`author_affiliation_org_id` (if not Rosotravel)

k. Experts

`author_id` (stable internal ID; SSOT for attribution)
`author_display_name`
`author_profile_url` (e.g., `/experts/julia-rossi/`)
`author_image_url`
`author_short_bio` (1–2 sentences)
`author_job_title`
`author_sameAs[]` (LinkedIn, X, Wikipedia, etc.)
`author_affiliation_org_id` (if not Rosotravel)

l. Creators

`author_id` (stable internal ID; SSOT for attribution)
`author_display_name`
`author_profile_url` (e.g., `/experts/julia-rossi/`)
`author_image_url`
`author_short_bio` (1–2 sentences)
`author_job_title`
`author_sameAs[]` (LinkedIn, X, Wikipedia, etc.)
`author_affiliation_org_id` (if not Rosotravel)

m. World Attraction Map

1.3 Business / Partnerships

n. Travel Agents / DMC

o. Affiliate Program

p. Content Creator Program

1.4 Support / Operations

- q. Contact Us
- r. Help Center
- s. Manage Booking

1.5 Policies / Trust

- t. Refund Policy
- u. Safety Standards
- v. Terms & Conditions
- w. Privacy Policy
- x. Cookies

1.6 SEO / System outputs

- y. Sitemap
- z. AI Sitemap
- aa. API Feed

Hard rule: pages referenced in footer must exist and be published in EN, otherwise the system shows a RED Blocking Error.

2. Roles & permissions (hard)

2.1 Roles

- SEO Editor: edit EN content, trigger translations, publish pages, lock blocks in non-EN languages
- SEO Lead / Admin: all SEO Editor permissions + unlock locked blocks + override publish blocks with reason + change indexing config for static pages

2.2 Hard restrictions

- AI cannot invent legal claims or policy meaning.
 - On Policy pages, AI is allowed only for translation + formatting suggestions (approval required).
 - System pages (Sitemap / AI Sitemap / API Feed) are not edited as free text (configuration only).
-

3. Content model (Blocks)

Each static page is composed of Content Blocks. Blocks are edited and validated independently.

3.1 Standard blocks (default)

- H1 (required for most pages)
- Snippet / Intro (optional, recommended; 180–260 chars)
- Main Body (required for most pages)
- FAQ (optional; allowed mainly for Help/Trust pages)
- CTAs (optional; e.g., “Contact us”, “Join program”, “Explore map”)
- Meta Title (required)
- Meta Description (required)
- OpenGraph Image (recommended)
- Indexing Profile (system-controlled)

3.2 Page type block requirements (hard)

Table: Required blocks by page group

Page group	H1	Snippet	Main Body	Meta Title	Meta Description	FAQ	Notes
Brand / Company	Required	Recommended	Required	Required	Required	Optional	Trust + brand consistency
Discovery hubs	Required	Recommended	Optional/Required (per page)	Required	Required	Optional	Has system-driven listings below copy
Partnerships	Required	Recommended	Required	Required	Required	Optional	Contains program CTA blocks
Support	Required	Recommended	Required	Required	Required	Recommended	Help Center may use FAQ

Policies / Trust	Required	Optional	Required	Required	Required	Optional	Strong audit + translation safety
SEO / System outputs	Optional	Optional	N/A	Required	Required	N/A	Config + preview only

Hard rule: If a required block is missing → Publishing blocked (RED).

4. Ownership per block (Source of truth)

4.1 Ownership types

- Manual (SEO Team): authoring and final approval
- AI-assisted: translation + rewrite suggestions, never auto-publish
- System: computed/generated outputs (Sitemap, AI Sitemap, API feed content)

4.2 Ownership rules (hard)

- EN is canonical for all pages.
- Non-EN languages are derived from EN by default via translation.
- Manual edits in a non-EN language can be locked per block, preventing overwrite.

5. Language & translation model (deterministic)

5.1 Canonical language

- English (EN) is the source language for all static pages.

5.2 Per-block per-language state machine

Each block in each language has exactly one state:

- AUTO — automatically translated from EN
- EDITED — manually edited (not locked yet)
- LOCKED — manually finalized; translation overwrite forbidden
- OUTDATED — EN changed after this language version was generated/edited
- OUTDATED (LOCKED) — locked block is now outdated; requires manual review (no overwrite)

5.3 Deterministic update rules (hard)

When an EN block changes:

- all AUTO translations → become OUTDATED

- all LOCKED translations → become OUTDATED (LOCKED)
- no translation is overwritten automatically

5.4 Allowed translation actions

- Translate missing languages: generates AUTO blocks for languages that have no content
- Re-translate OUTDATED AUTO blocks: allowed only for blocks in state OUTDATED and not locked
- Never overwrite LOCKED blocks (hard)

5.5 Manual edit and lock behavior

- editing a translation changes state to EDITED
- clicking Lock block changes state to LOCKED
- Unlock block is role-gated (SEO Lead/Admin) and requires a reason

6. Completion & Health (mandatory and persistent)

Static pages are low in count and high impact → must have a strict, visible completeness model.

6.1 Where Completion appears

- Dashboard List View: per page completion and issue counts
- Page Editor: persistent top status bar (always visible)

6.2 Completion dimensions

Each page has computed scores:

1. Required blocks complete (EN)
2. SEO readiness
 - meta title present and within configured limits
 - meta description present and within configured limits
 - canonical and indexing profile valid
3. Translation coverage
 - % languages complete
 - count of OUTDATED blocks

- count of OUTDATED (LOCKED) blocks

4. Link integrity (recommended)

- internal links valid
- footer reference integrity (page exists & published)

6.3 Status levels (hard)

- GREEN (Complete): all required blocks filled; no blocking issues
- YELLOW (Incomplete): non-blocking issues exist (e.g., missing OG image, OUTDATED translations)
- RED (Blocked): missing required blocks, invalid meta rules, or footer integrity violation

6.4 Blocking rules (hard)

Publishing is blocked if:

- required H1 missing (where required)
- meta title missing
- meta description missing
- main body missing (where required)
- page is referenced in footer but EN is not published
- indexing profile violates policy (e.g., API Feed incorrectly set to indexable if policy forbids)

7. Notifications & Visibility (errors cannot be ignored)

7.1 Notification types

- RED: Blocking — must be fixed to publish
- ORANGE: Action required — typically OUTDATED translations or failed generation
- YELLOW: Warning — recommended improvements

7.2 Persistent Issue Banner (hard requirement)

Acknowledging a notification does not clear the underlying problem.

A persistent banner remains visible at the top of:

- Static Pages Dashboard
- Each Page Editor

Banner shows:

- count of RED / ORANGE / YELLOW
- “View Issues” opens an Issues Drawer with:
 - issue type
 - affected blocks
 - affected languages
 - direct jump links to fix location
 - “Mark as reviewed” (does not remove flags; only records acknowledgment)

Hard rule: Issues disappear only when fixed, never by dismissal.

8. UI layout and flows

8.1 Static Pages Dashboard (default view)

Default view lists all static pages.

Each row shows:

- Page name
- Page group
- Canonical URL
- Indexing profile (Indexable / Noindex / System-generated)
- EN Completion (GREEN/YELLOW/RED)
- Translation coverage (e.g., “9/12 complete”)
- Issue counters (RED/ORANGE/YELLOW)
- Last edit (timestamp + user)
- Actions: Open / Translate / Publish / Unpublish (role-gated)

Bulk selection:

- Select row checkboxes + “Select all”
- Bulk actions apply only to selected pages

Bulk actions available

- Translate missing languages

- Re-translate OUTDATED AUTO blocks
- Export issues report
- Publish selected (only if none are RED-blocked)

8.2 Page Editor (single page)

Editor contains:

- Persistent Completion/Issue banner at top
- Language selector (EN default)
- Block list with per-block state indicators (AUTO/EDITED/LOCKED/OUTDATED)
- Preview panel (optional)

Per block actions:

- Edit
- Save
- Lock / Unlock (role-gated)
- Translate / Re-translate (only for AUTO + OUTDATED)
- View diff (recommended)

9. Page-type specifics (hard rules)

9.1 Policy pages (Refund / Safety / T&C / Privacy / Cookies)

- Main body required in EN
- Translation allowed, but AI restricted to translation + formatting suggestions
- Mandatory audit log for every change
- Recommended: “Last updated” timestamp displayed (manual)

9.2 Discovery hubs (Destinations / Experiences / Guides / Experts / Creators / World Attraction Map)

- Copy blocks (H1/snippet/intro) are editable + translatable
- Listings below are system-driven; editor can configure:
 - section enable/disable
 - order mode (system default, manual pinning allowed)

- Translation affects only copy blocks, not system entity data

9.3 SEO/System pages (Sitemap / AI Sitemap / API Feed)

- No free-text editing of generated content
 - UI provides:
 - configuration fields
 - last generation timestamp
 - preview (read-only)
 - validation status + errors
 - Completion/Health covers configuration and generation success
-

10. Audit log (mandatory)

Every change must be recorded:

- user
- timestamp
- page + block + language
- old vs new value hash (or full diff)
- lock/unlock events + reasons
- publish/unpublish events
- translation events (source EN version + target language)

Hard rule: audit log is immutable.

11. Indexing and SEO safety (static pages)

Each page has an Indexing Profile:

- Indexable
- Noindex
- System-generated (rules-driven)

Hard rules

- Policies: indexable unless explicitly configured otherwise

- System pages: follow global SEO policy (e.g., AI Sitemap may be indexable or not; API Feed typically not indexable unless explicitly desired)
- Changing indexing profile is role-gated (SEO Lead/Admin) and requires a reason

3.5 SEO governance: SEO rules, Sitemaps, index/noindex, robots, eligibility

01. SEO general rules Indexation:

0. Core Principle (Most Important Rule)

Rosotravel **MAY index supplier-sourced content ONLY** if it has been materially transformed by AI into Rosotravel-owned copy.

Material transformation means:

- rewritten wording,
- different sentence structure,
- different emphasis and ordering,
- no reuse of supplier phrasing,
- output owned by Rosotravel.

If content is **not rewritten by AI**, it is treated as **supplier content** and **must not be indexed**.

1. Absolute Prohibitions (Never Index, Never Rewrite)

The following elements are **NEVER indexable**, **NEVER rewritten**, and **NEVER transformed into SEO content**, regardless of source:

Images sourced from Viator

- Supplier photos (Viator)
- User-generated photos (UGC)

- Review photos
Images may be displayed to users, but:
- never indexed,
- never included in image sitemaps,
- never reused on other landing pages.

✗ **Reviews sourced from Viator**

- Individual review text
- Review authors
- Review dates
- Review photos
- Reviews must never be rewritten, paraphrased, summarized verbatim, or exposed in crawlable HTML
- can not generating average ratings in global SEO
- building "customer reviews" blocks on other subpages
- search engine autosuggestions.

✗ **viatorUniqueContent sourced from Viator**

- Must not be indexed.
- Must not be rewritten.
- Must not be used as AI input.
- Preferably not ingested at all.
If ingested for UX, it must never appear in page source (HTML or inline JS).

✗ **Full itineraries (step-by-step) sourced from Viator**

- Supplier itinerary text
- Route descriptions
- Stop-by-stop schedules

1.2 Internal prohibitions Runtime vs Indexable Layers Clarification

To prevent SEO fragmentation:

The following UI components are runtime-only layers and must not generate indexable URLs:

- Audience segments (family, premium, less crowded, etc.)
- Intent filters
- Tour-language filters
- Magic Finder
- Fallback pools
- Availability-based alternative rendering

These layers:

- Do not create canonical variants
- Do not generate hreflang entries
- Do not produce separate sitemap entries
- Do not alter URL structure

Indexable Pages Are Limited To:

- City Page
- Things to Do Page
- Attraction Page
- Approved Near Pages

City Band Changes

Changes in inventory band (0–299 / 300–599 / 600+) affect slot allocation only.

They must not alter URL structure or indexing logic.

Language Clarification.

There are two separate systems:

1. Site language (hreflang controlled)
2. Tour language (runtime filter only)

Tour language selection does not create new landing pages.

Canonical rules must enforce this separation.

2. Rendering Safety Rule (Viator Compliance)

For **all non-indexable supplier content**:

Protected content must **never appear directly in the page source** (HTML or inline JavaScript).

Allowed display method:

- external JavaScript endpoint
- endpoint blocked in `robots.txt`

`noindex` alone is **NOT sufficient** if the content is visible in source.

3. Indexable vs Non-Indexable Entities (Hard Rules)

Indexable entities (`index`, `follow`)

Only if **published** and **SEO-enabled**:

- Product pages
- City pages (≥1 published product)
- Attraction pages (SEO-enabled only)
- Destination pages (countries / regions with ≥1 product)

- Things-to-do pages (city-level aggregators)
 - Expert profile pages
 - Blog articles (human-written, explicitly published)
 - Aggregate review components:
 - rating summary
 - AI sentiment summary
 - Reviews from own customers (first-party only)
-

Non-indexable entities (`noindex`, `follow` unless stated)

- Draft products
- Draft attractions
- Near-pages (initial state)
- Dynamic filters (all parameterized URLs)
- Pagination URLs
- Sorting URLs
- Internal search results
- Tag pages (unless explicitly whitelisted)
- QA / preview / internal tools

Supplier-related (always non-indexable):

- Raw review content from Viator
- Supplier photos / UGC photos
- Full itineraries
- `viatorUniqueContent`

None of the above may appear in sitemaps.

4. Product Page Indexation Rules

Indexable elements

ONLY if rewritten by AI (Rosotravel-owned copy):

- H1
- Meta title
- Meta description
- Product overview / description
- Highlights
- FAQ (AI-generated)
- “Why travelers love this tour” (AI sentiment summary)
- Guide / expert section
- Trip DNA attributes
- Pricing & availability blocks
- Related tours module
- Factual blocks (duration, opening hours, meeting point, cancellation policy, availability) are system-owned and supplier-derived. They may be rendered and indexed as structured factual data but **MUST NEVER** be rewritten, paraphrased, summarized, or stylistically altered by AI

Non-indexable elements (always)

- Raw customer reviews
- Individual review texts
- Review metadata (author, date)

- Review photos
 - Supplier photos
 - Product gallery images from Viator
 - Full itinerary (step-by-step)
 - Supplier itinerary text
-

Special Rule: Itinerary

- Full itinerary **must be visible to users** but **never indexable**
- Must not appear in crawlable HTML
- Must not be in sitemap

Allowed indexable substitute:

- `itinerary_summary` (AI-generated, 3–4 sentences)
- Describes:
 - scope
 - pace
 - suitability
 - main landmarks
- Must NOT describe steps, stops, or schedule details
- Must NOT reproduce supplier wording
- Sample what means itinerary on Viator

Overview
What's Included
Meeting and Pickup
Itinerary
Additional Info
Cancellation

[Open in Google Maps](#) >

Goddess Diana Fountain (listed as Constance Fund Fountain of Diana on Google Maps) in Green Park. Enter Green Park next to the Ritz Hotel; the fountain is a few steps down the path on the right.

[Open in Google Maps](#) >

Conclusion Point: Westminster Abbey

Start time · 9:00 am

Itinerary

- Green Park**

Begin your experience in Green Park, where your expert guide introduces London's royal history and sets the scene for the tour.

10 minutes • Admission Ticket Free
- Buckingham Palace**

Stroll past Spencer House, the Green Park Gates, and Canada Gate while your guide shares stories of royal connections and historic architecture, before discovering Buckingham Palace, learning about the Br ... [Read more](#)

From PLN

Standard price

Date
Wed, Jan 7

Free cancel
experience

Reserve Now
while stayin

Book ahead
On average,

5. Attraction Page Indexation Rules

Indexable (ONLY if AI-rewritten)

- Attraction name
- Intro description
- Historical / factual overview
- Visiting tips
- “How to visit”
- AI-generated FAQs
- Related tours
- Nearby attractions

Non-indexable

- Supplier attraction descriptions
- viatorUniqueContent

- Raw reviews
- UGC photos
- Supplier photos

Attraction variants

- Minor variants are NOT separate pages
 - One canonical attraction page only
 - Variants appear as subsections
-

6. City, Region, Country Pages

City pages

- Indexable when ≥ 1 product is published
- Navigational & broad commercial intent only
- No transactional keyword ownership

Region / country pages

- Auto-published when ≥ 1 product exists
 - Minimal SEO by default
 - Full SEO Auto requires manual trigger
 - No embeddings or clustering by default
-

7. Near-Pages Rules

Two sources exist:

- **Immediate Near-Pages** (created immediately from predefined templates for controlled intent coverage)

- **GSC Near-Pages** (created every 5 weeks from real GSC queries)

All near-pages are created as **noindex,follow** by default and **are excluded from sitemaps**.

Near-pages become indexable only when they reach **Index-Eligible** state:

- eligibility is evaluated by **hard gates** (ownership, similarity, intent allowlist) plus **demand signals** (GSC performance or validated volume).

Promotion to index can be:

- **automatic** (default, when gates pass - [near page spawn](#)), or
- **manual** (SEO override for strategic topics).

8. Reviews Rules (Strict)

OTA / Viator reviews

- Must never be indexed
- Must never be rewritten or paraphrased
- Must never appear in page source

Allowed indexable outputs:

- Aggregate rating (numeric)
- AI sentiment summary (signals-only, no text reuse)
- Guide-level reputation summary

Own customer reviews

- Only first-party Rosotravel reviews may be indexed

9. Translation & Indexation Timing

- English is the source language
 - Translations generated post-publish
 - Translation is indexable only if:
 - EN version is indexable
 - valid `hreflang`
 - self-canonical language URL
-

10. Final Enforcement Rules

- No page enters any sitemap unless indexable
 - No page becomes indexable without a published parent
 - Keyword ownership must be locked before indexation
 - AI sitemap + AI summary feeds required for every published entity
 - Any violation defaults to `noindex`, `follow`
 - Supplier compliance overrides SEO in all conflicts
-

11. Viator Compliance Summary (Explicit)

- Supplier text may be displayed to users **only if not crawlable**
- Supplier content may be indexed **only after AI rewriting**
- Reviews, photos, itineraries, viatorUniqueContent are **never indexable**
- AI sentiment summaries are allowed
- Raw supplier content must never appear in page source

Viator requirements for sourced content -

<https://docs.viator.com/partner-api/technical/#section/Key-concepts/Itineraries>

- We have a special permission for marketing for AI changed content, we can index product only if the content is changed by AI
- We still can not index: user reviews, user photos, product photos, itineraries (it is a special itinerary block)

a. Sitemaps rules

01. Sitemap Rules (Standard Sitemaps)

On every publish / unpublish event, update:

- sitemap.xml (index only)
 - sitemap-products.xml
 - sitemap-attractions.xml
 - sitemap-destinations.xml (countries / regions / cities)
 - sitemap-things-to-do.xml
 - sitemap-blog.xml
 - sitemap-experts.xml

Sitemap inclusion rules:

- Only indexable URLs
- Canonical URLs only
- No draft, no near-pages, no filters

Each entry must include:

- `<loc>` canonical URL
- `<lastmod>` from `last_modified`
- Optional `<changefreq>`
- Optional `<priority>`

02. AI Sitemap & LLM Feeds

AI Sitemap

- Path: `/ai-sitemap/`
- Format: JSON or XML
- Includes all indexable entities + selected non-indexable discovery nodes
- AI Sitemap may include non-indexable discovery nodes only if they are explicitly flagged as `discovery_only: true`.
- Fields:
 - URL
 - `entity_type`
 - `primary_keyword`
 - `geo` (if applicable)
 - `isPartOf` / `hasPart` / `nearby relations`

- language

AI Summary Feed

- Path: `/ai-summary.json`
- One record per entity
- Fields:
 - short_summary (3–5 sentences)
 - key highlights
 - top FAQs
 - entity identifiers (city_id, attraction_id, product_id)
 - geo coordinates (if applicable)

These feeds are refreshed on publish, update, and translation.

G.07 Dynamic Filters & Non-Indexable URLs

G.07.1 Filter parameter handling

The site will use URL parameters for filters (duration, language, price range, time of day, rating, etc.).

Rules:

- Filtered URLs must be crawlable only partially (e.g. some core combinations) or not at all, depending on SEO strategy.
- By default, all filtered URLs must have canonical to the base page, unless explicitly whitelisted.
- CMS must allow an SEO admin to mark certain filter combinations as “indexable landing pages” in rare cases (e.g. `/things-to-do/?family-friendly=true`), with separate canonical rules.

G.04.3 meta robots rules

- Default for all published commercial pages is `index, follow`.
- Noindex must be set automatically for:
 - Very thin content (below minimum word count threshold, if flagged).
 - Experimental A/B test variants if they should not be indexed.
 - Duplicate or deprecated pages preserved only for UX.
- Non-indexable filter pages (with many query params) must explicitly carry `noindex, follow` if they can be accessed directly.

04. Review & Rating Schema Gate (AggregateRating / Review)

Ratings and reviews are high-risk schema fields and must follow strict consistency and compliance rules.

Hard rules:

1. Emit **aggregateRating** only when the rating is allowed for that page type and is consistent with UI.
2. **ratingValue** and **reviewCount** must come from a declared source (internal reviews, allowed partner source, or approved external source).
3. Do not emit **review** objects unless policy and licensing allow displaying review text and attribution.
4. If only aggregated values are allowed, emit **only aggregateRating + reviewCount** (no review text).
5. Any mismatch between UI and schema for rating/count blocks must block publishing (or force noindex + alert).

DoD:

- Rating/review schema is emitted only when allowed and always matches UI.
- No unlicensed review text is emitted in schema.

SEO Governance & indexation control:

5.4 SEO Rules and metrics

This section is safety, growth control, and indexation governance on Rosotravel.

It defines mandatory, deterministic rules governing **keyword allocation, ownership, intent routing, canonical logic, index/noindex decisions, near-page lifecycle, AI rewrite limits, migration memory, multilingual SEO, seasonality handling, crawl budget protection, and URL deprecation.**

These rules apply to **AI-generated, human-generated, and hybrid content** and are **enforced automatically by the publishing layer**. Any violation is treated as a **critical defect** and must block publishing or trigger rollback.

For every URL, before generation, publishing, or indexation, the system MUST store:

- `page_type`
(country, region, city, attraction, product, category, package, near-page, blog, filter, experiment)
- `language`
- `source` (AI / human / hybrid)
- `intent_type` (SEMrush: informational, navigational, commercial, transactional)
- `primary_keyword` (string or null)
- `owned_keywords` (array, approved variants only)
- `keyword_owner` (true / false)
- `owner_url` (winning URL for the intent cluster)
- `semrush_volume` (monthly)
- `semrush_intent`
- `keyword_embedding` (vector)
- `entity_flag` (true / false – Wikidata or Google Places confirmed)
- `seasonal_flag` (true / false)
- `index_state` (index / noindex)
- `canonical_target` (URL)
- `hreflang_cluster_id`
- `created_at` / `updated_at`
- `last_gsc_validation_at` (date or null)
- `gsc_impressions_28d`
- `gsc_clicks_28d`

If ANY field is missing → page defaults to:

`noindex, follow + canonical to owner/parent + not eligible for promotion.`

G.07.1 Intent Classification & Validation (SEMrush + GSC)

Intent Source of Truth

- **SEMrush is the authoritative source for intent labels.**
- The system **MUST NOT** invent intent types.

GSC Validation Layer (Mandatory)

- **Google Search Console is the authoritative source for intent validation, query class separation, and promotion eligibility.**
- GSC does NOT override SEMrush intent labels.
- **GSC does control index eligibility, near-page graduation, and distinct query confirmation.**

SEMrush defines *what the keyword is*.

GSC validates *how Google actually treats it*.

G.07.1.1 GSC Proof Definitions

GSC is the only valid source of “proof” that Google treats a topic as a distinct query class. The system **MUST** compute and store proof using only **EN GSC**, in rolling windows:

Required stored fields per URL (EN):

- `gsc_impressions_28d`, `gsc_clicks_28d`, `gsc_ctr_28d`, `gsc_avg_position_28d`
- `gsc_top_queries_28d[ ]` (top 50 queries with impressions + position)
- `gsc_query_cluster_hash` (stable hash of normalized top queries)

Normalization rules (queries):

- lowercase
- remove stopwords
- strip brand tokens unless the query is navigational

- normalize plurals and punctuation

Distinct Query Class Proof (DQCP)

A page is considered to have “distinct query class proof” **ONLY IF ALL are true:**

1. `gsc_impressions_28d >= 100`
2. `unique_queries_28d >= 10` (within top 50)
3. **Overlap with parent owner query set is limited:**
`overlap_ratio <= 0.60` where $\text{overlap_ratio} = \frac{\text{shared_queries}}{\min(\text{query_count_page}, \text{query_count_owner})}$
4. The page has **at least 3 unique head queries** that are not present on the owner URL:
head query = query with `impressions_28d >= 10` OR `clicks_28d >= 1`
5. The page has been live for at least `min_age_days_for_gsc_proof` days.

CONFIG (must be editable in CMS config):

- `min_age_days_for_gsc_proof = 35`
- `dqcp_min_impressions_28d = 100`
- `dqcp_min_unique_queries_28d = 10`
- `dqcp_max_overlap_ratio = 0.60`
- `dqcp_min_head_queries = 3`

If DQCP is not met → the system **MUST** treat the page as **not proven**, regardless of SEMrush volume.

G.07.2 Page Type → Allowed Intent Ownership (Hard Allowlist)

Page Type	Allowed Intent Ownership
Product	Transactional only
Attraction / POI	Informational, Navigational
Near-page	Informational, Commercial
City / Region / Country	Navigational, Commercial
Category / Things-to-do	Commercial, Navigational
Package / Bundle	Commercial only (flagship only)
Blog / Guides	Informational only (human-edited only)
Filters / UI variants	None (never indexed)

If a keyword's SEMrush intent is not allowed for the page type → assignment is blocked.

G.07.3 Hierarchical Keyword Routing (Deterministic)

- **Transactional keywords** → Product pages only (1:1 fit required)
- **Commercial keywords** → City / Category / Near-page
- **Informational keywords** → Attraction / Near-page / Blog
- **Navigational keywords** → City / Region / Attraction

Attraction pages may contain CTAs but **may never own transactional keywords**.

G.07.4 Deterministic Product Fit Test (1:1 Enforcement)

`product_fit = true` ONLY if all are true:

1. Exactly one strong entity anchor detected
2. Booking modifier present (*book, tickets, skip-the-line, guided tour, private tour, entry*)
OR intent = Transactional
3. One product cluster exceeds similarity threshold AND
no second product is within **0.03 similarity** of the top match

If `product_fit = false`, the keyword may NEVER be owned by a Product page.

G.07.5 Keyword Ownership & Keyword Map (Single Source of Truth)

Each product:

- 1 primary transactional keyword
- up to 2 secondary transactional keywords

Locked keywords:

- cannot be reused by any other page
- cannot be targeted by near-pages
- cannot be overridden by AI

`keyword_map` is mandatory, audited, and immutable without human approval.

G.07.6 Similarity Thresholds

All similarity thresholds (keyword similarity and page similarity) are defined in **D.05 [Similarity Thresholds](#)** and MUST be used by all pipelines (keyword discovery, ownership locking, near-page candidacy, duplicate page prevention).

G.07.7 Weekly Keyword Re-Sync (Growth Without Chaos)

- Run once per week
 - No page creation from variants
 - New URLs only from **new topic + allowed intent + no ownership conflict**
-

G.07.8 Variant Band Safety (No Rewrites)

Allowed:

- meta description only
- FAQ answer clarification only

Forbidden:

- H1
 - SEO body
 - intent
 - canonical
 - URL
 - internal linking
-

G.07.9 Volume Thresholds (Index Gates – EN Only)

Page Type	Index Eligibility
Product	≥ 20 + product_fit + bookable
Attraction / POI	≥ 50 OR entity_flag
City / Region	≥ 300
Category	≥ 100 + distinct GSC query class
Near-page	≥ 30 (but always noindex on creation)
Blog / Guide	≥ 50 + human-edited

Below threshold → **noindex,follow**.

G.07.10 Entity Override Rule

If **entity_flag = true**, volume thresholds may be bypassed.

Entity override **never bypasses**:

- intent allowlists
 - ownership
 - near-page noindex-on-start
 - canonical hierarchy
-

G.07.11 Seasonality Handling (Production-Grade)

Seasonal detection sources:

1. Keyword token dictionary
2. SEMrush / GSC trend curve
3. Structured availability window

Seasonal Pre-Index Window:

- Index starts **90 days before season**

- Ends **30 days after season**

Applies to:

- products
 - attractions
 - categories
 - city pages
-

G.07.12 Canonical & Indexation (Best-Practice Model)

Language Canonical (Correct SEO Standard)

- **Every language page uses self-canonical**
- Canonical NEVER crosses languages
- Hreflang connects variants
- Ownership & promotion decisions are EN-driven only

This allows:

- local ranking
 - proper hreflang behavior
 - cost-controlled SEO governance
-

Page-Type Canonical Rules

Country / Region / City

- canonical: self
- index: index, follow

Category

- self ONLY if GSC proves distinct intent
- else → canonical to city

Attraction

- self ONLY if independent demand
- else → canonical to city

Product

- self ONLY
- NEVER canonical to city or attraction

Package

- self ONLY if flagship + ownership
- else → canonical to city

Near-pages

- creation: canonical to parent
- graduation: self

Blog

- self ONLY if human-edited

Filters / Params

- canonical to base clean URL
- always noindex

Countries and Regions may exist as published pages, but they are **not part of city batching** and they MUST NOT generate near-pages.

Country (Lite Page)

Allowed content scope:

- short intro (300–600 words)
- 8–12 general FAQs
- internal links to ACTIVE cities only (top 10–30)

Forbidden:

- near-pages
- programmatic “things-to-do” subpages under country
- long-tail expansion beyond hub intent

Default SEO state:

- `index_state = noindex, follow` until manually promoted
- `canonical = self`
- ownership limited to broad navigational/commercial queries only

Region (Lite Page)

Definition: administrative region (e.g., province/voivodeship).

Allowed content scope:

- short intro (250–500 words)
- 6–10 general FAQs
- links to ACTIVE cities in the region

Forbidden:

- near-pages

- region-level “things-to-do” expansion unless explicitly approved as flagship

Default SEO state:

- `index_state = noindex, follow` until manually promoted
- `canonical = self`

Promotion eligibility (optional, deterministic):

- `active_cities_in_region >= 3`
- AND (`gsc_impressions_28d >= 300` OR `semrush_volume >= 300` on the assigned primary keyword)

CONFIG

- `country_index_promotion_manual_only = true`
- `region_index_promotion_manual_only = true`

G.07.13 Winner Rule & Duplicate Resolution

Winner selected by:

1. intent clarity
2. EN GSC performance
3. ownership timestamp
4. structural priority

Losers:

- `canonical + noindex`
- OR 301 redirect if permanently deprecated

Detection without resolution is forbidden.

G.07.14 Near-Page Lifecycle (Hard Safety)

- Max **1 new near-page per city per week**
 - Start: noindex + canonical to parent
 - Graduate ONLY after:
 - EN GSC distinct query class
 - ownership lock
 - 5 weeks minimum age
-

G.07.15 AI Rewrite Safeguards

- cooldown: **45 days**
- max: **2 rewrites / quarter**
- rewrites only if:
 - **≥ 1,000 impressions / 28 days**

Allowed:

- meta description
- FAQ answers

Forbidden:

- H1
- SEO body
- intent
- canonical
- URL
- internal linking

G.07.16 Crawl Budget Protection

- Global index cap enforced
- If exceeded → new pages enter noindex queue
- Promotion requires full re-validation

Additionally:

- Max **20 new indexable URLs / city / week**
-

G.07.17 Redirect & Deprecation Policy (Mandatory)

- Permanent winner → **301**
 - Temporary suppression → canonical + noindex
 - Removed without successor → **410**
 - Redirects maintained minimum **18 months**
-

G.07.18 SEO Memory (Migration Safety)

Mandatory snapshot before any migration:

- URL → queries
 - URL → intent
 - URL → ownership
 - URL → performance
 - URL → canonical & hreflang
-

G.07.19 Authority & Link Timing

Link building allowed ONLY after:

- core nodes stable
 - ownership validated
 - near-pages proven safe
-

G.07.20 Explicitly Forbidden

The system must NEVER:

- index near-pages on creation
 - canonicalize products upward
 - allow translated pages to own intent
 - index filters or parameters
 - rewrite core SEO body automatically
 - create pages from minor variants
-

G.07.21 Strategic Outcome

- Transactional intent is owned exclusively by products
- Discovery intent expands safely via controlled near-pages
- Long-tail is captured without crawl or cannibalization risk
- Seasonality is handled deterministically
- Multilingual SEO works with zero ownership chaos
- The site scales weekly without structural debt

3.6 Semrush Rules: option in UI Marketing panel

6.1 Weekly Sync

Only refresh tracked keywords.

No new keywords discovered during weekly sync.

Weekly refresh updates:

- Search Volume (monthly avg)
- Keyword Difficulty
- CPC
- SERP Features
- Intent Classification
- 12-month trend (bundled, cheap)

Stored in:

`semrush_keywords_weekly`

Important constraints

- Do NOT request historical database
 - Do NOT fetch multi-year history
 - Use only 12-month trend included in keyword overview response
 - No “keyword_history” endpoints allowed
-

6.2 Monthly Keyword Discovery (Expensive — Controlled Budget)

Runs **once per month** on a limited set.

Sources:

- SEMrush Keyword Magic (limited)

- SEMrush Related Keywords
- SEMrush People Also Ask
- GSC new queries (daily, cost-free)

Discovery capped by:

- max monthly keyword budget
- max cost per run
- max number of new topics per entity

Stored in:

`keyword_discovery_monthly`
`gsc_query_log`

Discovery Strict Rules

- Do NOT re-fetch keywords that exist
- Do NOT fetch full databases
- Only fetch suggestions for selected parent keywords
- Only run for top-demand cities (configurable)

6.2A New City Acceleration Runs (Week-2 and Week-4, Early Learning, Cost-Safe)

Purpose: capture early long-tail signals for newly published cities and update existing owners **without creating new URLs**.

Definition

- “Week-2” means: run when `days_since_city_publish >= 14`
- “Week-4” means: run when `days_since_city_publish >= 28`

Eligibility window

Applies only to newly published City pages in the first:

- `new_city_acceleration_window_days` (default **35**)

Runs

- Run A: at `t = 14 days` after city publish
- Run B: at `t = 28 days` after city publish

Data sources (strict order)

1. GSC queries for that City URL (free)
2. SEMrush Batch Keyword Overview for those GSC queries (cheap verification only)

Allowed outputs (hard)

- attach queries as secondary keywords to existing owners
- expand FAQ lists on existing owners
- expand “mini-sections” on existing owners (no intent shift, no H1/body rewrite)
- flag near-page candidates as `topic_candidate` only (do NOT publish; do NOT index)

Forbidden outputs (hard)

- creating or indexing near-pages
- creating any new indexable URLs
- running expensive SEMrush discovery endpoints (related/questions/fullsearch)

Strict limits

- `new_city_acceleration_gsc_query_limit` (default **25**) per run
- SEMrush endpoints allowed: `batch_keyword_overview` only

CONFIG

- `new_city_acceleration_window_days = 35`

- `new_city_acceleration_run_offsets_days = [14, 28]`
 - `new_city_acceleration_gsc_query_limit = 25`
 - `new_city_acceleration_semrush_allowlist = ["batch_keyword_overview"]`
-

6.2B Ongoing City Refresh Loop (Every 5 Weeks, Continuous, Cost-Safe)

Purpose: continuously refresh city-related owners (City, Category/Things-to-do, Attractions) based on real demand signals and controlled SEMrush usage.

Cadence (hard)

For every ACTIVE City, run a refresh loop:

- every `city_refresh_interval_days` (default **35** days)

This is NOT one-off. It repeats indefinitely while the city is ACTIVE.

Inputs

1. GSC new queries since last refresh (free)
2. SEMrush verification for selected candidates only (unit-capped)

Allowed outputs

- attach new variant keywords to the existing owner (no new URL)
- add/refresh FAQs and mini-sections (no intent shift)
- queue near-page drafts ONLY if near-page eligibility gates and lifecycle rules pass (including 5-week learning window + DQCP)

Near-page rule alignment

Near-pages may only enter “draft creation” after:

- the city has been live for at least `nearpage_min_age_days` (default **35**)
AND
- the candidate passes cannibalization + similarity + volume/GSC gates

CONFIG

- `city_refresh_interval_days = 35`
- `city_refresh_gsc_query_limit = [FILL]` (recommended 50–200 depending on city size)
- `city_refresh_semrush_units_cap = [FILL]`
- `city_refresh_top_demand_only = [FILL true/false]`
- `city_refresh_allowlist[ ] = [FILL]`

Hard cost rule

If `city_refresh_semrush_units_estimate > city_refresh_semrush_units_cap` → abort SEMrush calls and execute GSC-only outputs.

6.2C Scheduling Rule (Deterministic Triggering)

All runs MUST be scheduled deterministically per City using `city_published_at` and `last_city_refresh_run_at`.

New City schedule (exact):

- First acceleration refresh: run when `days_since_city_publish = 14` (± 1 day scheduler tolerance).
- Second acceleration refresh: run when `days_since_city_publish = 28` (± 1 day scheduler tolerance).
- After acceleration completes, the city enters the ongoing refresh loop: run every 35 days counted from the timestamp of the last completed refresh run (`last_city_refresh_run_at`).

Clarification (hard):

Acceleration runs may surface near-page candidates, but they MUST be stored only as `topic_candidate` and MUST NOT create or index new near-page URLs until the near-page lifecycle gates are met (minimum age + GSC proof + cannibalization check).

Required stored fields (per City):

- city_published_at
- last_acceleration_run_14_at
- last_acceleration_run_28_at
- last_city_refresh_run_at
- last_gsc_ingested_at

This removes the unclear parenthesis and states the timeline exactly: 14d → 28d → then every 35d from the last refresh.

6.3 Forbidden SEMrush Actions

Developer must NEVER call:

- `keyword_history/*`
- `keyword_historical_trends/*`
- `domain_organic_historical/*`
- `database_history/*`

These calls = **high-cost historical datasets** → strictly forbidden.

6.4 SEMrush main settings and rules

The system MUST budget SEMrush usage in **units**, not “number of keywords”.

Known unit costs (current endpoints we allow):

- Batch Keyword Overview / Keyword Overview: **10 units / keyword line**
- Broad Match (phrase_fullsearch): **20 units / keyword line**
- Related Keywords: **40 units / keyword line**
- Phrase Questions: **40 units / keyword line**

Hard rule

Any “historical” mode or historical endpoint is treated as **~5× cost** and is forbidden.

6.5 Allowed Endpoint Allowlist + Hard Display Limits

The developer **MUST** enforce an allowlist and per-call limits.

Allowed endpoints (only these):

- `batch_keyword_overview` (default verification)
- `keyword_overview` (single keyword verification)
- `phrase_fullsearch` (broad match)
- `related_keywords`
- `phrase_questions`

Forbidden endpoints (hard-ban):

- any keyword history endpoints
- any domain historical endpoints
- any database history endpoints
(keep your existing forbidden list; this section adds enforcement)

Display limits (mandatory caps, per call)

- `phrase_fullsearch_display_limit = 100`
- `related_keywords_display_limit = 50`
- `phrase_questions_display_limit = 50`

If a request exceeds `display_limit` → truncate response and log a warning.

6.6 Trend Handling Without History (12-Month Trend Only)

The system **MUST NOT** fetch historical datasets.

Seasonality/trend is computed using:

- **12-month trend values returned in current Keyword Overview responses**, saved as snapshots.

Snapshot rule

On every successful `keyword_overview` / `batch_keyword_overview` call: store:

- `trend_12m_values[]` (12 monthly points)
- `trend_snapshot_date`

This allows internal trend analytics without SEMrush historical endpoints.

6.7 Unit Caps (Abort Rules)

Discovery **MUST** be capped by **units**, and must abort safely.

CONFIG (must be editable)

- `max_units_per_run` = **200,000**
- `max_units_per_month` = **4,000,000**
- `max_monthly_keyword_budget` = **50,000** (stored keywords after dedupe; not “lines fetched”)
- `max_new_topics_per_entity` = **20**
- `top_demand_cities_only` = **true**
- `top_demand_cities_allowlist_max_size` = **300**

Abort behavior:

If `run_units_estimate` > `max_units_per_run` → abort before calling SEMrush.

3.65 Media Library & Rights Panel: option in UI Marketing panel

Purpose

Single operational UI for media governance and licensing safety at scale (including Gooru content). The system must store all assets (owned vs imported) but only “Processed + Approved” assets may be used on indexable pages.

Marketing MUST be able to:

- upload/import assets into a media_asset catalog with provenance fields:
 - source_type: owned / supplier_import / licensed / user_contributed (future)
 - license_type, rights_confirmation, author/credit, capture_date (if available)
 - entity associations (city_id / attraction_id / product_id / article_id)
- run processing pipeline (or view processing status):
 - virus scan, EXIF extract (if present), conversion to WebP, renditions, compression,
 - alt-text draft generation, safe filename normalization
- approve / reject assets for SEO usage (approval gate)
- search/filter assets by city/entity/type/approval state/tags
- view usage references (which pages/articles use an asset)
- bulk operations: upload, process, approve, deprecate

Marketing MUST NOT be able to:

- mark unprocessed assets as usable on indexable pages
- bypass approval gating for blog/article assets
- delete assets that are in use (must deprecate with replacement workflow)

Hard rule (usage gating)

Only assets with usage_state = approved_processed may be selectable in:

- Blog/Gooru article editor (image picker),
- any indexable landing pages requiring editorial media.

Partners Panel (Shortcut / Operational)

Purpose

Provide a CMS/backoffice operational entry point to Partner Programs without mixing it with Bookings → Discount Codes.

This panel is a shortcut that references the full functional specification in Chapter VII.

Marketing MUST be able to:

- access Partners → Applications / Affiliates / Travel Agents / Gooru Guides
- view partner assets (coupon + deeplinks), statements, and high-level performance counters (no customer PII)
- manage Gooru Topic Pool and review queue (SEO workflow)

Marketing MUST NOT be able to:

- modify core checkout logic
- access supplier-only data or customer PII

3.7 Alerts & Inbox

01. UI Main Panel add: Inbox (alerts)

- Add in a visible place near the logo, if a new alert occurs you can see a red circle on it
- The alert circle is visible until notification will not be clicked and read
- You can mark the notification as unread even if clicked
- Low scores (GA4)
- Processing Bag
- Sampling

02. Persistent top-of-page banner (must exist)

For any entity page in admin UI (City / Attraction / Product), show a persistent banner at the top until resolved:

- Red (Blocking): prevents indexation / sitemap inclusion
- Yellow (Warning): publishable but requires review
- Green: OK

Banner must include:

- Error code (e.g., RAW_CONTENT_MISSING, FACTUAL_MISSING, OFFER_MISSING, RELATIONSHIP_MISSING, AI_FAILED, MEDIA_REVIEW_REQUIRED)
- Affected entity + version domain (raw_version / offer_version / realtime_offer_version)
- “What is blocked” (Indexing / Sitemap / AI feed / Publish)
- Recommended action (Regenerate allowed blocks / Fix mapping / Retry ingestion)

2.3 Error Panel (always reachable, never hidden)

Admin UI must include a permanent “Errors & Flags” drawer/panel showing:

- All active errors for the entity
- First detected timestamp

- Last attempted remediation
- Current status (Open / In progress / Resolved)
- Links to:
 - Pipeline actions (regenerate)
 - Content & Structure view (manual fix / lock)
 - Relationship editor / mapping (if applicable)

Content + Schema Validation Alerts (Editor Facing)

- **BLOCKER:** Missing Title/H1/Meta on an indexable page
- **BLOCKER:** Schema emitted but module not rendered (hidden schema)
- **BLOCKER:** Offer invalid on Product/Package page (price/currency/availability missing)
- **WARNING:** HowTo eligible but not present
- **WARNING:** TOC expected but not rendered (Article templates)

Each alert must show: URL, page type, failing rule, recommended action, and “who can fix” (Marketing/SEO/IT).

3) How to scheme alert

Add alerts:

- “HowTo blocked: steps invalid / missing fields”
- “HowTo blocked: module not rendered but schema emitted”
- “HowTo blocked: contradictions with Offer/Availability”
Each alert must show: URL, page type, failing rule, and fix action.

3) Missing or Broken RAW Content Fields (PRODUCT only)

3.6 UI notification

- RAW_CONTENT_MISSING

- Severity:
 - Blocking if title missing and no safe fallback

Warning otherwise

4) Missing or Inconsistent FACTUAL Data (CITY/ATTRACTION)

4.4 UI notification

- FACTUAL_MISSING
- Severity:
 - Blocking if coordinates missing
 - Warning otherwise

5) Missing OFFER for UI

5.4 UI notification

- OFFER_MISSING (Blocking)

6) Missing REALTIME Offer Fields (Availability/Pricing) — PRODUCT realtime domain - UI notifications

6.5 UI notification

- REALTIME_OFFER_MISSING
- Severity:
 - Blocking if both price and availability are unusable
 - Warning if only start_times missing

7) Missing or Broken RELATIONSHIP Fields (Hierarchy / Mapping) — ALL entity types - UI

7.4 UI notification

- RELATIONSHIP_MISSING (Blocking)
Must link directly to:

- Relationship/mapping view
- Dedupe conflict resolution view (if collision)

9) AI Generation Failures (Any entity where AI blocks exist) - UI

9.3 UI notification

- AI_GENERATION_FAILED (Warning or Blocking depending on required AI-owned blocks)

10) API Rate Limits / Outages (System-level) - UI

Non-blocking banner in CMS: “External API degraded — using cached data”.

11) Product problem alert

- Alert types:
 - Rating below eligibility threshold
 - Reviews count below threshold
 - Product removed from provider
 - Product in Set no longer eligible
 - New eligible product detected in city
 - City RAW_count band change
 - Availability failure for pinned hero product
 - Manual override event (Pool or Set)
 - RAW restore event
- Single-product AI run triggered
 - Product removed from Pool due to threshold drop
 - New product added to Pool (meets thresholds)
 - Set item violates Pool eligibility (flag only; manual review required)
 - Set item is non-instant (confirmationType != INSTANT) (blocking flag for Set)
 - Structural product change detected (classification-risk flag)
 - City band change detected (flag; does not auto-change Set)
 - Manual override event (Pool/Set membership change)
 - RAW restore event
 - Algorithm config version change (with approval chain)

Alerts must include:

- - timestamp
- - city
- - productCode
- - severity
- - responsible role
- - recommended action

Each alert must include:

- Timestamp
- ProductCode
- City
- Severity level
- Responsible role

Alerts must be stored and logged.

Structural Product Change Flag

b. Structural change to product:

If provider payload indicates a structural change that may affect product classification (e.g. change in privateTour flag, itinerary shape, duration type, category group),

the system must:

- Flag the product
- Notify Product Operations
- Re-evaluate Pool eligibility automatically
- Flag Set membership (manual review required)

Set membership is not automatically removed.

Manual review is required to maintain stability.

Alerts UI Semantic layer

1) FAQ Collision Alerts

If marketer attempts to add a FAQ question that is ≥ 0.92 similar to an existing one:

- show collision alert

block save until manually overridden

Notification per role

Editor Alerts (Inbox)

- Low CTR
- Schema errors
- FAQ duplication
- New near-page suggestion
- Canonical suggestion
- Page marked as dead
- A/B candidate ready

Editor Alerts (Inbox) — add blog-specific items

- Blog pending SEO review (Gooru submission; published as noindex)
- Blog freshness due (Gooru) — refresh required within 120 days to remain commission-eligible
- Blog freshness expired (Gooru) — commission attribution disabled until updated
- Blog optimization recommended — triggered by low CTR / impression thresholds (same mechanics as other pages)

DevOps Alerts

- Qdrant latency high
- SEMrush quota exceeded
- GPT rate limits
- Feed import failures
- Validation script errors

CMS “Optimization Inbox”

Central queue showing:

- Pages triggered by thresholds
- Pages needing optimization
- “Dead Pages”
- A/B candidates
- Near-page candidates
- FAQ duplication warnings
- Schema error items
- Blog Review Queue:
 - articles published_noindex awaiting SEO approval
 - articles flagged for refresh (freshness due/expired)
 - articles flagged for on-page improvements (CTR, schema warnings, FAQ duplication, etc.)

Columns (reuse existing):

- URL, page_type=blog, primary_keyword, CTR/impressions (30 days), status flag, recommended action, action buttons

Columns:

- URL

- page_type
- primary_keyword
- CTR / impressions (30 days)
- status flag
- recommended action
- action buttons (AI regenerate, send to editor, ignore 30 days)

CMS “Performance Tab” (per entity)

Visible on Product / Attraction / City page.

Includes:

SEO Metrics

- Impressions (30–90 days)
- Clicks
- CTR
- Position
- Cannibalization risk

UX Metrics

- Engagement time
- Bounce rate
- Scroll depth
- Conversion rate

Ownership

- Primary keyword
- Secondary keywords
- Locked keywords
- FAQ questions owned
- FAQ duplication alerts

Technical Status

- Schema validation
- Alternate / hreflang status
- Indexable / noindex flag

Actions

- Regenerate meta
- Regenerate intro
- Generate new FAQ

FAQ & LONG-TAIL — ROLLOUT PLAN

FAQ: ship early with base pages (enables rich snippets).

Long-tail: after base indexation (when CTR > 1.5%), driven by PAA, Search Console, and intent analysis (e.g., “for families”, “skip-the-line”, “in winter”).

- Lock/unlock fields
- Force re-run

Entity-level (Product / City / Attraction panels)

- View GSC metrics (impressions, clicks, CTR, position)
- View GA4 metrics (engagement, bounce, conversions)

- View semantic similarity warnings
- View keyword ownership + FAQ ownership
- View optimization flags + reason
- **Clear intent/CTAs (Book Now / Check Availability).**
- **Monitor indexation (GSC, crawl logs).**

SEO Control Center (Dashboard)

1 CMS HOME / SEO CONTROL CENTER (DASHBOARD)

A. Index Health

- Total indexed URLs
- Total noindex URLs
- Index cap (limit)
- % utilization
- Trend 7d / 28d

B. Risk Alerts (RED / ORANGE / GREEN)

- Ownership conflicts
- Canonical violations
- Near-pages awaiting approval
- Rewrite cooldown violations
- Crawl budget risk

👉 click on the alert → leads to the panel and URL

C. Weekly Changes

- New pages generated
- Pages promoted to index
- Pages demoted
- Rewrites executed
- Redirects added

Alerts & Blocking Panel

Purpose

Immediate visibility of governance violations.

System MUST auto-flag:

- index rule violations
- ownership conflicts
- canonical contradictions
- rewrite limit breaches
- near-page rule violations

Marketing MUST be able to:

- block publishing
- rollback last batch
- escalate to IT

UX

List of violations:

- Severity
- Rule broken
- Affected URLs
- Auto-fix possible? YES/NO

Buttons:

- Block publishing
- Rollback last batch
- Escalate to IT

Seasonality Panel

Seasonality & Calendar Panel

Purpose

Control seasonal visibility without manual hacks.

Marketing **MUST** be able to:

- view detected seasonal keywords and pages
- see season window logic
- manually confirm or reject seasonality flag
- adjust season window dates (within ± 14 days)
- approve seasonal pre-indexing

Marketing cannot disable seasonality safeguards globally.

UX Elements

- Calendar view (months)
- Seasonal keywords list
- Affected URLs

Actions:

- Confirm season window
- Adjust window (± 14 days)
- Approve pre-indexing

Badge:

SEASONAL INDEX ACTIVE

Schema & Relations Consistency Alerts

The CMS Inbox must surface blocking issues:

- Schema invalid / missing required schema blocks (by page type)
- Relations mismatch (schema claims ItemList/nearby/breadcrumbs that UI does not render)

- Offer mismatch (price/currency/availability differs between schema and CTA). Each alert must show: URL, page type, missing fields, and the exact publish gate that failed

Crawl Budget & Index Growth Panel

Purpose

Avoid index bloat during scale-up.

Marketing **MUST** see:

- total indexable URLs
- index cap
- per-city index growth
- noindex queue size

Marketing **MAY**:

- prioritize pages in noindex queue
- pause index growth per city or page type

Marketing cannot raise global caps without IT approval.

UX

Graphs:

- Indexed URLs over time
- Per city growth
- Noindex queue size

Controls:

- Pause index growth (city / page type)
- Reorder noindex queue (priority)

Read-only:

- Global index cap (IT-owned)

Redirect & Deprecation Panel

Redirect & Deprecation Panel (Migration Safety)

Purpose

Control permanent SEO decisions.

Marketing **MUST** be able to:

- approve 301 redirects
- approve 410 removals
- see redirect chains
- see redirect age and expiration
- see traffic loss/gain after redirect

No redirect may be deleted before 18 months.

UX Table

- Old URL
- Target URL
- Type (301 / 410)
- Age (days)
- Traffic before / after

Actions:

- Approve redirect
- Mark as permanent
- View chain

No delete button < 18 months.

SEO Memory & Migration Panel

SEO Memory & Migration Panel (Read-Only, Mandatory)

Purpose

Ensure no migration wipes SEO history.

Marketing **MUST** be able to view:

- URL → query mappings
- ownership history
- performance snapshots

- canonical history
- hreflang history

This panel is read-only.
Deletion is forbidden.

UX

Search by URL → show:

- historical queries
- ownership timeline
- performance snapshots
- canonical & hreflang history

Export button (CSV / JSON).

3.75 UI main panel: Sales

3.76 Partners Panel: option panel of sales

Purpose

Provide a CMS/backoffice operational entry point to Partner Programs without mixing it with Bookings → Discount Codes.

This panel is a shortcut that references the full functional specification in Chapter VII.

Marketing MUST be able to:

- access Partners → Applications / Affiliates / Travel Agents / Gooru Guides
- view partner assets (coupon + deeplinks), statements, and high-level performance counters (no customer PII)
- manage Gooru Topic Pool and review queue (SEO workflow)

Marketing MUST NOT be able to:

- modify core checkout logic
- access supplier-only data or customer PII

3.8 Roles & Permissions

01. Security, Access Control & Roles

Purpose of the Security Layer

The security & role-model layer ensures that:

- Only authorized users can trigger AI writes, keyword ownership changes, schema regeneration, or page publishing.
- Sensitive operations (keyword locking, factual overrides, AI unlocks, near-page approvals) are protected from junior users.
- All changes (AI, human, system, external API) are audit-logged and recoverable.
- Data ownership logic (keywords, FAQs, schema, dynamic near-pages) is not unintentionally broken by non-technical editors.
- Developer and DevOps actions are fully separated from content team actions.

Additional Security / NFR Controls (Partners + Tracking)

Deeplink open-redirect prevention (hard)

- The deeplink redirect route MUST NOT accept arbitrary URLs.
- /go/{deeplink_id} MUST:
 - 1) look up stored target_url server-side,
 - 2) validate it is a Rosotravel canonical URL (allowlist),
 - 3) set attribution cookie,
 - 4) redirect only to the stored validated target_url.

- If validation fails → return 404 (fail-closed) and log an abuse event.

Rate limits (hard)

- Rate limit deeplink creation (Deeplink Builder) per partner_id + IP (burst + daily caps).
- Rate limit topic claim operations per gooru_guide_id + IP (burst + daily caps).
- Rate limit coupon validation attempts to prevent brute-force (already required in Partners security).

Audited settings and locks (hard)

- Any change to partner settings, approval status, commissions/discounts, and any keyword/topic lock/unlock MUST be audit-logged:
 - actor_id, role, timestamp, old_value, new_value, reason, affected entities/URLs.
- Audit records must be immutable and exportable for compliance/reconciliation.

This layer governs BOTH:

1. **CMS Permissions** (what users can do in the dashboard)
2. **API-Level Permissions** (which endpoints can be called by whom)
3. **AI & Pipeline Permissions** (who can trigger AI or override it)

02. Role Categories (4-Level Permission Model)

We define **4 core roles** and **1 optional high-security role**.

1. Admin (Full System Owner)

Has unrestricted control.

Can:

- ✓ Modify anything in CMS
- ✓ Change all AI locks
- ✓ Rebuild schema for entire batches
- ✓ Unlock keyword ownership
- ✓ Override keyword collisions
- ✓ Trigger or stop pipeline runs
- ✓ Manage integrations (Qdrant keys, SEMrush keys, API credentials)

- ✓ Access all logs and dashboards
- ✓ Delete/restore entities
- ✓ Manage users & roles
- ✓ Perform mass RAW restore (with double confirmation)
- ✓ Override inventory governance locks
- ✓ Force removal from Set without eligibility check

Cannot:

- ✗ Nothing — Admin equals full system root.
-

2. SEO Lead / Head of Content

Has full authority over all SEO-critical components.

Can:

- ✓ Approve or reject AI drafts
- ✓ Trigger AI rewriting on any page
- ✓ Lock/unlock individual fields
- ✓ Override keyword ownership
- ✓ Approve dynamic near-pages
- ✓ Merge/redirect cannibalizing pages
- ✓ Approve blog topic pool
- ✓ Edit schema (metadata sections only)
- ✓ Manage FAQ ownership conflicts
- ✓ Access SEO dashboards
- ✓ Access keyword & FAQ logs
- ✓ Revert versions
- ✓ Trigger AI regeneration for specific sections
- ✓ Access GSC & GA4 deep metrics per URL
- ✓ View Set membership (read-only)
- ✓ Request product reclassification review (cannot enforce change)

Cannot:

- ✗ Change system-level settings (API keys, Qdrant config, embeddings model)
 - ✗ Modify database structure
 - ✗ Modify pipeline logic or queues
-

3. Content Editor / Marketer

Primary content creators.

Can:

- ✓ Manually edit any content section (except factual fields)
- ✓ Mark fields as “manual override”
- ✓ Request AI rewrite (requires SEO Lead approval if configured)
- ✓ Add FAQs manually
- ✓ Add blog posts
- ✓ Adjust images & galleries
- ✓ View SEO dashboards (limited)
- ✓ Resolve basic schema warnings (non-technical)

Cannot:

- ✗ Unlock keyword ownership
 - ✗ Change primary keywords
 - ✗ Trigger full AI pipeline
 - ✗ Edit factual fields (Google Places / Wikidata overrides)
 - ✗ Approve near-pages
 - ✗ Edit schema JSON-LD directly
 - ✗ Change URL slugs
-

4. Junior Editor

Restricted content handler.

Can:

- ✓ Edit body text (long description, highlights)
- ✓ Upload images
- ✓ Add internal notes
- ✓ Suggest FAQs (need approval)
- ✓ Access simplified performance view (CTR, impressions simplified)

Cannot:

- ✗ Trigger ANY AI regeneration
 - ✗ Change meta, H1, schema
 - ✗ Modify slugs
 - ✗ Create near-pages
 - ✗ Own or modify keywords
 - ✗ Override factual data
 - ✗ Approve drafts
-

5. Developer / CTO Role (Optional Advanced Security Role)

This role is separate from Admin.

Can:

- ✓ Manage Qdrant collections
 - ✓ Rebuild embeddings
 - ✓ View and flush pipeline queues
 - ✓ Diagnose API integrations
 - ✓ See raw logs (AI, Qdrant, SEMrush, GSC)
 - ✓ Run migration scripts
 - ✓ Reconfigure pipeline thresholds
 - ✓ Force-run near-page generation
 - ✓ Create new entity types
-
- ✓ View inventory layer transitions (RAW / Pool / Set) logs

Cannot:

- ✗ Modify content
- ✗ Trigger publishing
- ✗ Modify keyword ownership
- ✗ Edit product descriptions

6. The Product Operations Lead is responsible for:

- Managing RAW / Pool / Set transitions
- Manual promotion of products to Pool
- Manual promotion of products to Set
- Pin / Unpin logic within Set
- Blocking products from Pool or Set
- Reviewing threshold violations
- Monitoring band changes (0–299 / 300–599 / 600+)
- Triggering single-product AI run when required (after Set insertion)
- Reviewing inventory-related alerts
- Approving fallback pool usage
- Managing city-level product distribution

Can:

- ✓ View full RAW provider payload
- ✓ Compare RAW vs AI-transformed version
- ✓ Manually add product to Pool (override with log)
- ✓ Manually remove product from Pool
- ✓ Manually add product to Set
- ✓ Pin / Unpin product in Set
- ✓ Trigger single-product AI run (does not modify SEO ownership)
- ✓ View product version history
- ✓ Restore RAW version (single product only)
- ✓ Access inventory alerts dashboard
- ✓ Flag product for SEO review
- ✓ View city band classification
- ✓ View archetype classification (read-only until Phase C defined)

Cannot:

- ✗ Modify system-level thresholds
- ✗ Modify API keys or integrations
- ✗ Override keyword ownership
- ✗ Approve near-pages
- ✗ Modify schema JSON-LD
- ✗ Change URL slugs
- ✗ Run full city batch pipelines
- ✗ Mass restore RAW versions (Admin only)
- ✗ Modify embeddings or Qdrant configuration

Additional Governance Rules:

- Manual Add to Set automatically assigns Pool membership.
- If city batch has already been processed, single-product AI run must be triggered.
- All inventory actions must generate immutable audit logs.
- RAW restore action must log previous AI version snapshot.

7. External Partner Roles (Portal-only; not CMS)

Principle (hard)

External partners (Affiliates, Travel Agents, Gooru Guides) **MUST NOT** access the Marketing CMS Panel or any SEO Governance UIs. They must authenticate into a separate “Partners Panel” (Partner Portal) with strictly limited capabilities.

Rationale:

- Prevent accidental SEO governance changes (keyword ownership, indexing, schema, near-pages).
- Reduce attack surface and protect business-critical controls.
- Provide a partner-friendly UX for profile, links, and payouts without exposing CMS complexity.

Shared restrictions (hard, applies to ALL partner roles)

Partners **MUST NOT** be able to:

- ✗ view any customer PII (name, email, phone, address)
- ✗ view supplier IDs/endpoints/raw payloads
- ✗ change keyword ownership, intent routing, indexing rules, sitemaps, canonicals, hreflang
- ✗ edit schema JSON-LD
- ✗ change URL slugs

- ✗ trigger global AI pipeline runs
- ✗ access raw logs, API keys, system caps, or strategy settings

Partners MAY be allowed to:

- ✓ manage their own public profile fields (display name, bio, photo, social links)
- ✓ manage attribution assets within their program (deeplinks, coupon code view)
- ✓ view their own statements, balances, and transactions (non-PII)
- ✓ upload invoices (affiliate/gooru) or view savings (travel agent)

Role: Partner (Affiliate)

Can (Partner Portal only):

- ✓ view affiliate assets: coupon code (read-only unless admin changes), primary deeplink, campaign deeplinks
- ✓ create/revoke campaign deeplinks via Deeplink Builder (target allowlist enforced)
- ✓ edit own profile: name/display name, bio, social links, website
- ✓ view balances + statements + transaction list (non-PII)
- ✓ upload invoice and manage payout details

Cannot:

- ✗ access any blog/topic pools
- ✗ create/edit Rosotravel pages or content

Role: Partner (Travel Agent)

Can (Partner Portal only):

- ✓ edit own company profile: display name, contact email, website, social links
- ✓ view assigned discount rates (OWN vs PARTNER)
- ✓ view order list + savings summary (non-PII)

Cannot:

- ✗ create coupon codes or deeplinks

- ✗ access affiliate/gooru ledgers or statements

Role: Partner (Gooru Guide)

Can (Partner Portal only):

- ✓ edit own guide profile: display name, bio, languages, social links, headshot
- ✓ view city-scoped Topic Pool (informational only), claim topics, release topics (TTL rules)
- ✓ create/edit/submit articles assigned to the guide
- ✓ use Media Library picker (approved assets only)
- ✓ view earnings, statements, and article-level performance (no PII)

Cannot:

- ✗ approve indexation (noindex→index), cannot override SEO review decisions
- ✗ change primary keywords or ownership outside the claim flow
- ✗ access other guides' drafts or topics

03. CMS Permission Matrix (Full Table)

Legend

- **A** – Admin
- **S** – SEO Lead
- **C** – Content Editor
- **J** – Junior Editor
- **D** – Developer/CTO

CMS Action	A	S	C	J	D
Edit page content	✓	✓	✓	✓	✗

AI rewrite (full page)	✓	✓	ask S	✗	✗
AI rewrite (section)	✓	✓	ask S	✗	✗
Lock/unlock AI fields	✓	✓	✗	✗	✗
Change primary keyword	✓	✓	✗	✗	✗
Modify keyword ownership	✓	✓	✗	✗	✗
Approve near-pages	✓	✓	✗	✗	✗
Trigger pipeline job	✓	✓	✗	✗	✓
View SEO dashboard	✓	✓	✓	limited	✓
Edit schema fields (minor)	✓	✓	limited	✗	✓
Edit factual fields	✓	S-only	✗	✗	✗
Add blog articles	✓	✓	✓	✓	✗
Change slugs / URL	✓	✓	✗	✗	✗
Restore old versions	✓	✓	ask S	✗	✓
Manage users & roles	✓	✗	✗	✗	✗
Manage API integrations	✓	✗	✗	✗	✓
Access raw logs	✓	✓	✗	✗	✓

G.07.CMS.2 Roles & Permissions (Explicit)

Role	Capabilities
Marketing	Approve, lock, promote, demote, request changes
SEO Lead	All marketing rights + conflict resolution
IT	System rules, caps, logic
AI	Execute only

Editor Content only, no SEO governance

Marketing **never codes rules**,
AI **never decides ownership**,
IT **never decides intent**.

IV. System Contract & Data/AI Architecture

1. Workflow table

Step	Owner	Tooling	Trigger	What it does	Output
1. Ingest raw feed	appinventive	Viator API → DB	Schedule (hourly/daily)	Store raw vN JSON (draft_raw)	/raw/vN/{id}.json
2. Versioned diff	Appinventiv BE	Diff engine (hash per section)	On new vN	Compare vN vs active_version; produce change_map	change_map.json
3. Routing	Appinventiv BE	Policy rules	After diff	Classify: minor / medium / major	regen_queue tasks
4. Programmatic gen	Appinventiv AI	OpenAI + LangChain	medium/major OR first publish	Create structured content (meta, H1–H3, summary, FAQ, tags, links, alt)	/drafts/{slug}.json
5. Validation	Appinventiv BE	Schema linter + business rules	After gen	Check consistency (fields), schema.org validity	validated_draft.json
6. Uniqueness (EN only)	Appinventiv + SEO QA	Copyscape (EN, ai_modified)	First publish + medium/major re-gen	External duplicate check; if fail → human fix → recheck	uniqueness_report.json
7. Human QA	Rosotravel Marketing	Internal panel	Post-validation	Spot-check, edits, approve	Approved draft
8. Publish	Appinventiv FE	Publishing service + sitemap builder	On approve	Set published status, lock URL, set index/canonical/hreflang flags, update	Live URL